



香港科技大學  
THE HONG KONG  
UNIVERSITY OF SCIENCE  
AND TECHNOLOGY

COMP 4901B  
Large Language Models

# Recurrent Neural Networks, Transformers

Junxian He

Sep 12, 2025

# Recap: Autoregressive Language Models

$$\begin{aligned} p(\text{the, mouse, ate, the, cheese}) &= p(\text{the}) \\ &\quad p(\text{mouse} \mid \text{the}) \\ &\quad p(\text{ate} \mid \text{the, mouse}) \\ &\quad p(\text{the} \mid \text{the, mouse, ate}) \\ &\quad p(\text{cheese} \mid \text{the, mouse, ate, the}). \end{aligned}$$

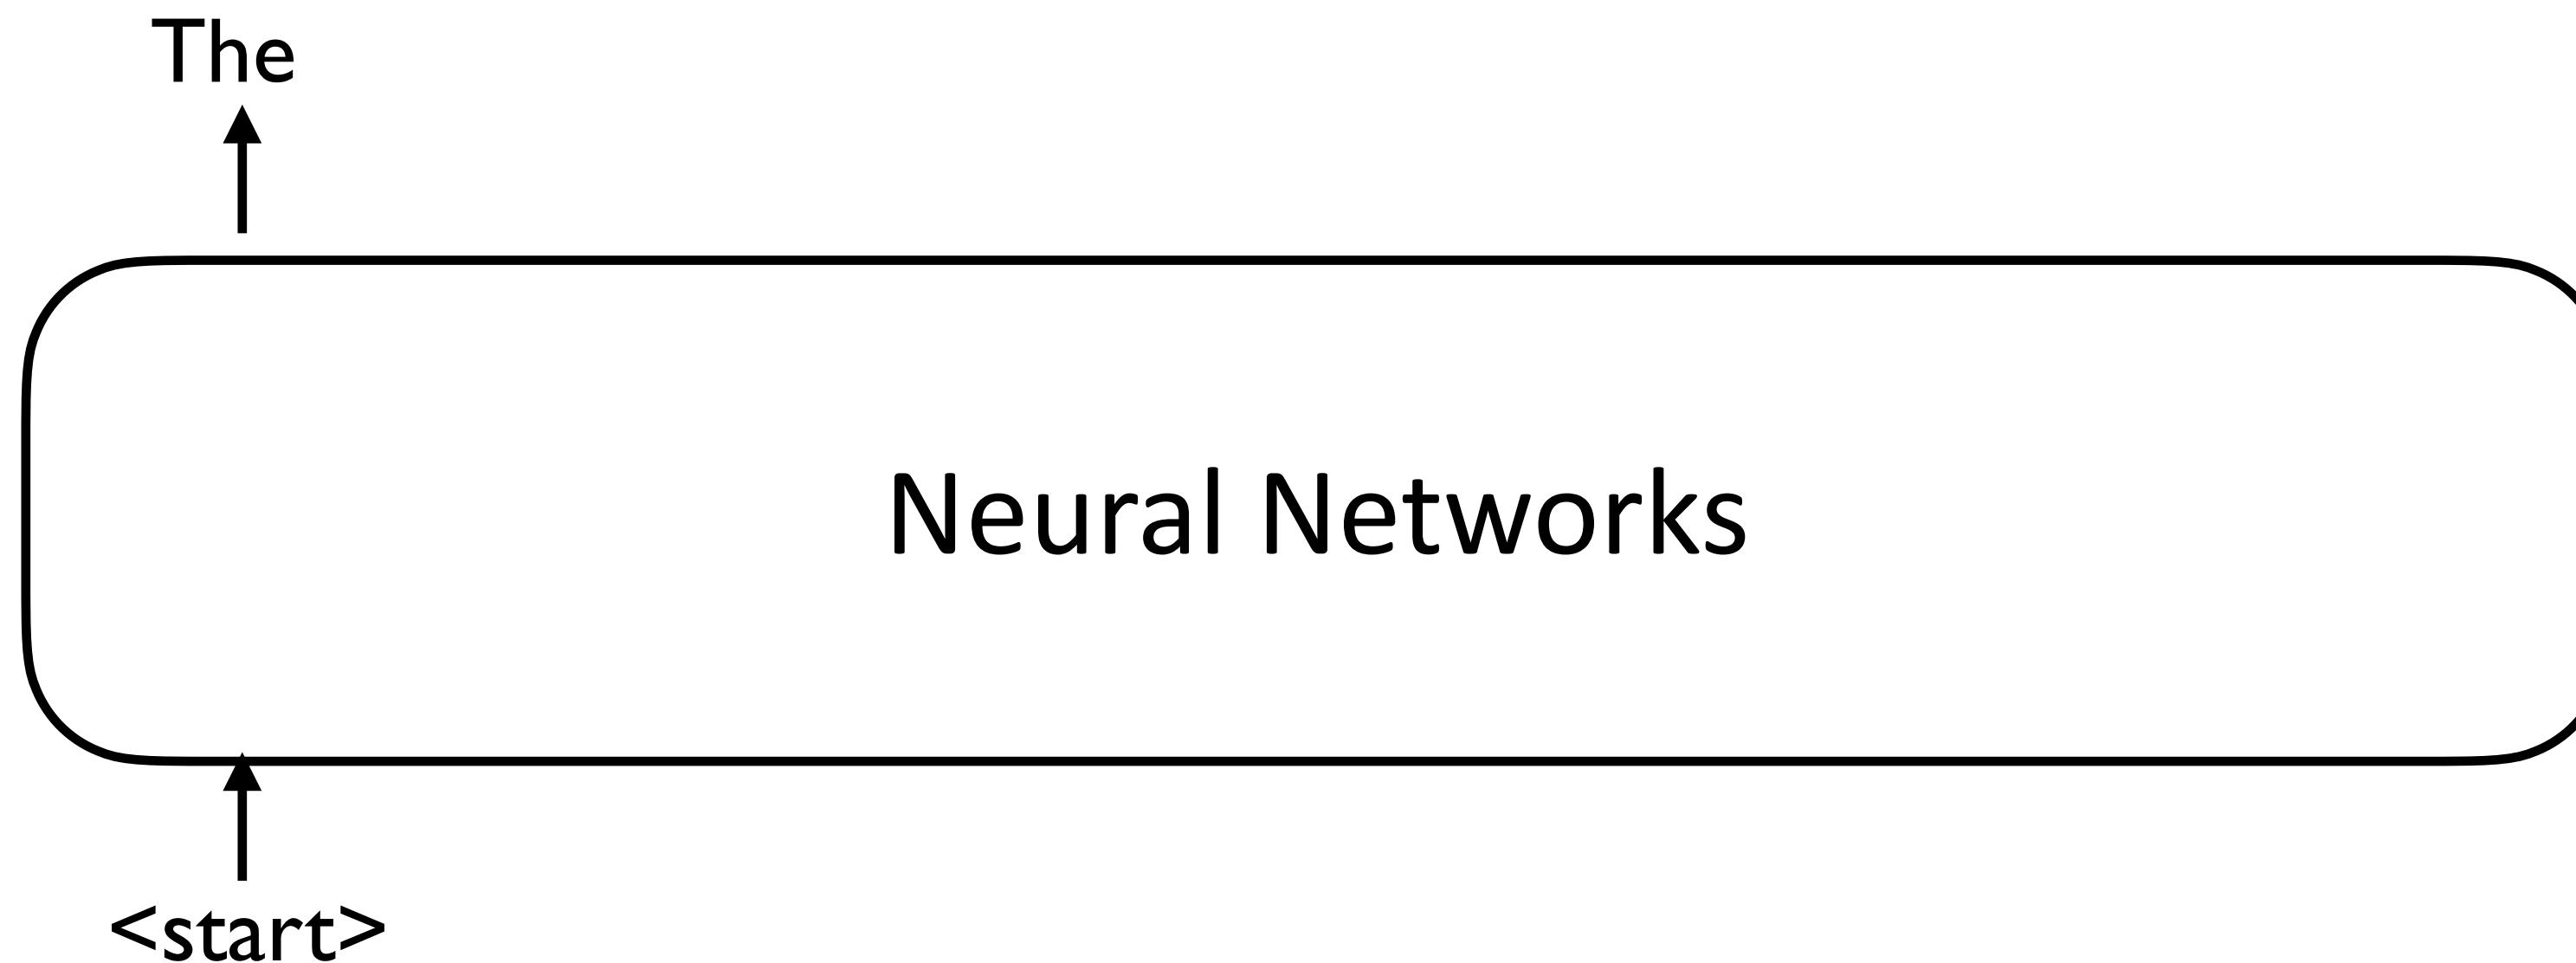
$$p(x_1, x_2, \dots, x_I) = \prod_{i=1}^I p(x_i \mid x_{1:i-1})$$

Next Word      Context

# Recap: Neural Language Models

Neural language models are typically autoregressive

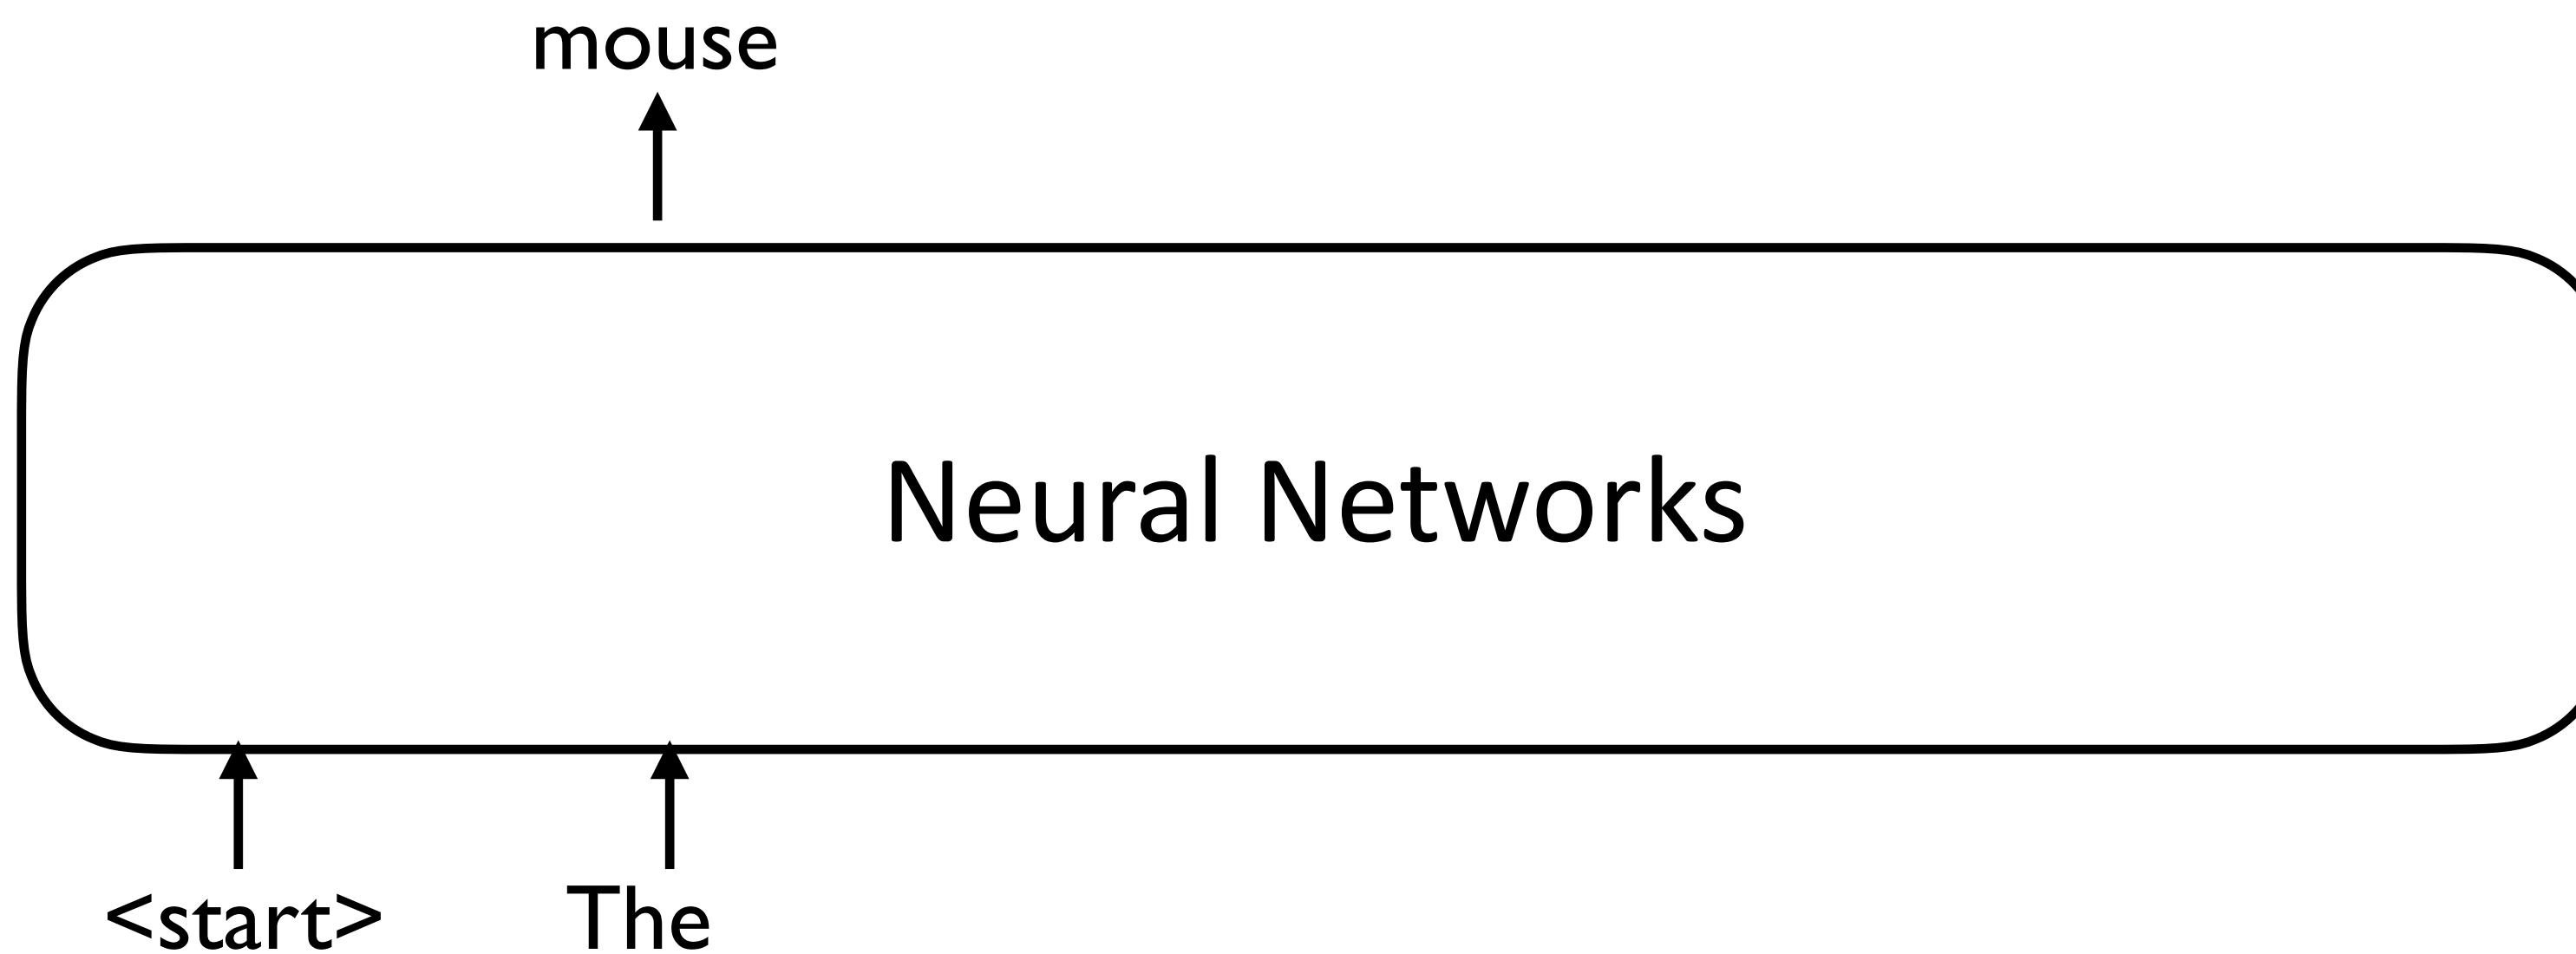
Data: “The mouse ate the cheese.”



# Recap: Neural Language Models

Neural language models are typically autoregressive

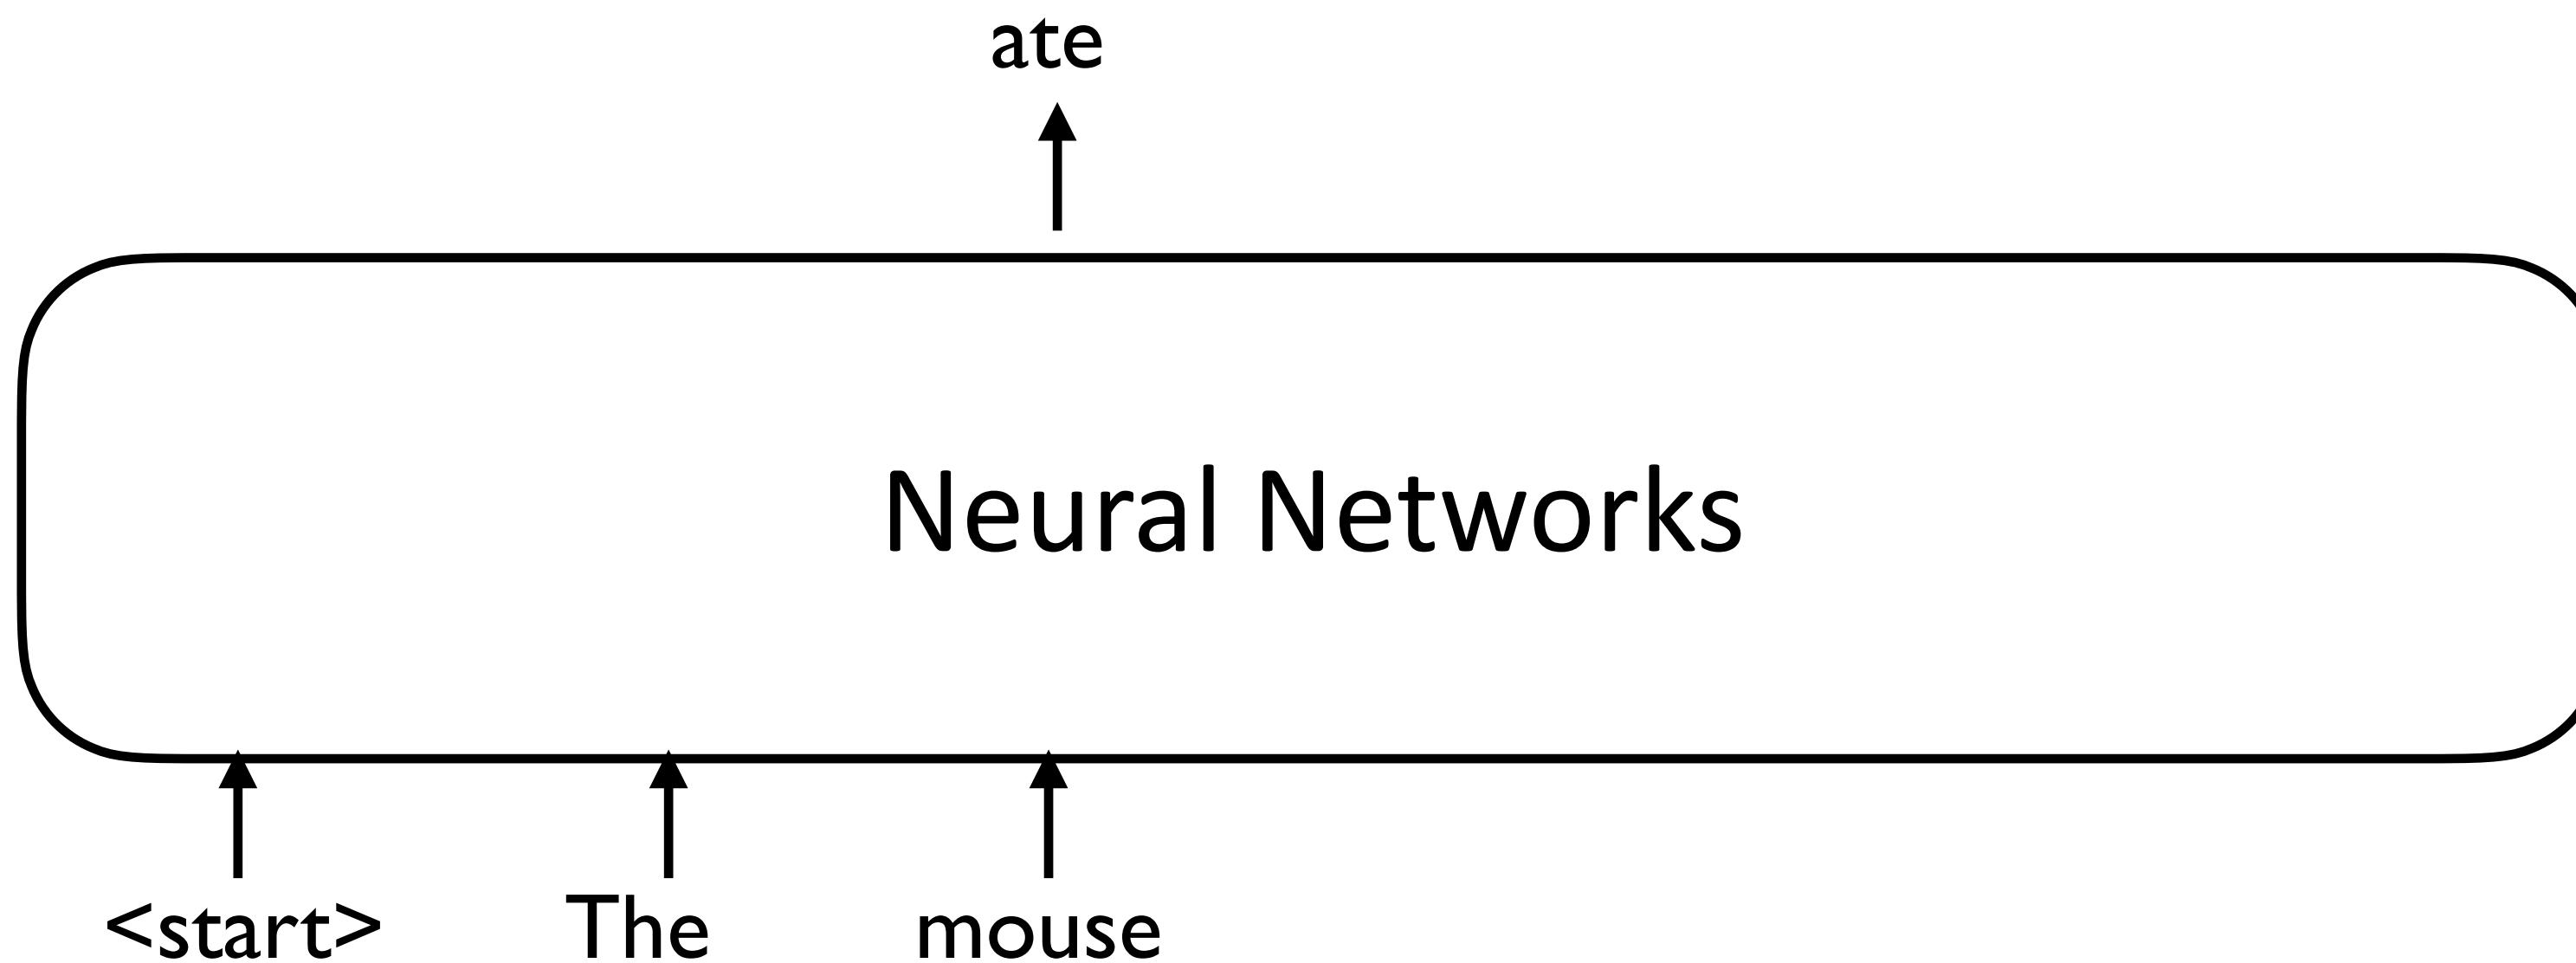
Data: “The mouse ate the cheese.”



# Recap: Neural Language Models

Neural language models are typically autoregressive

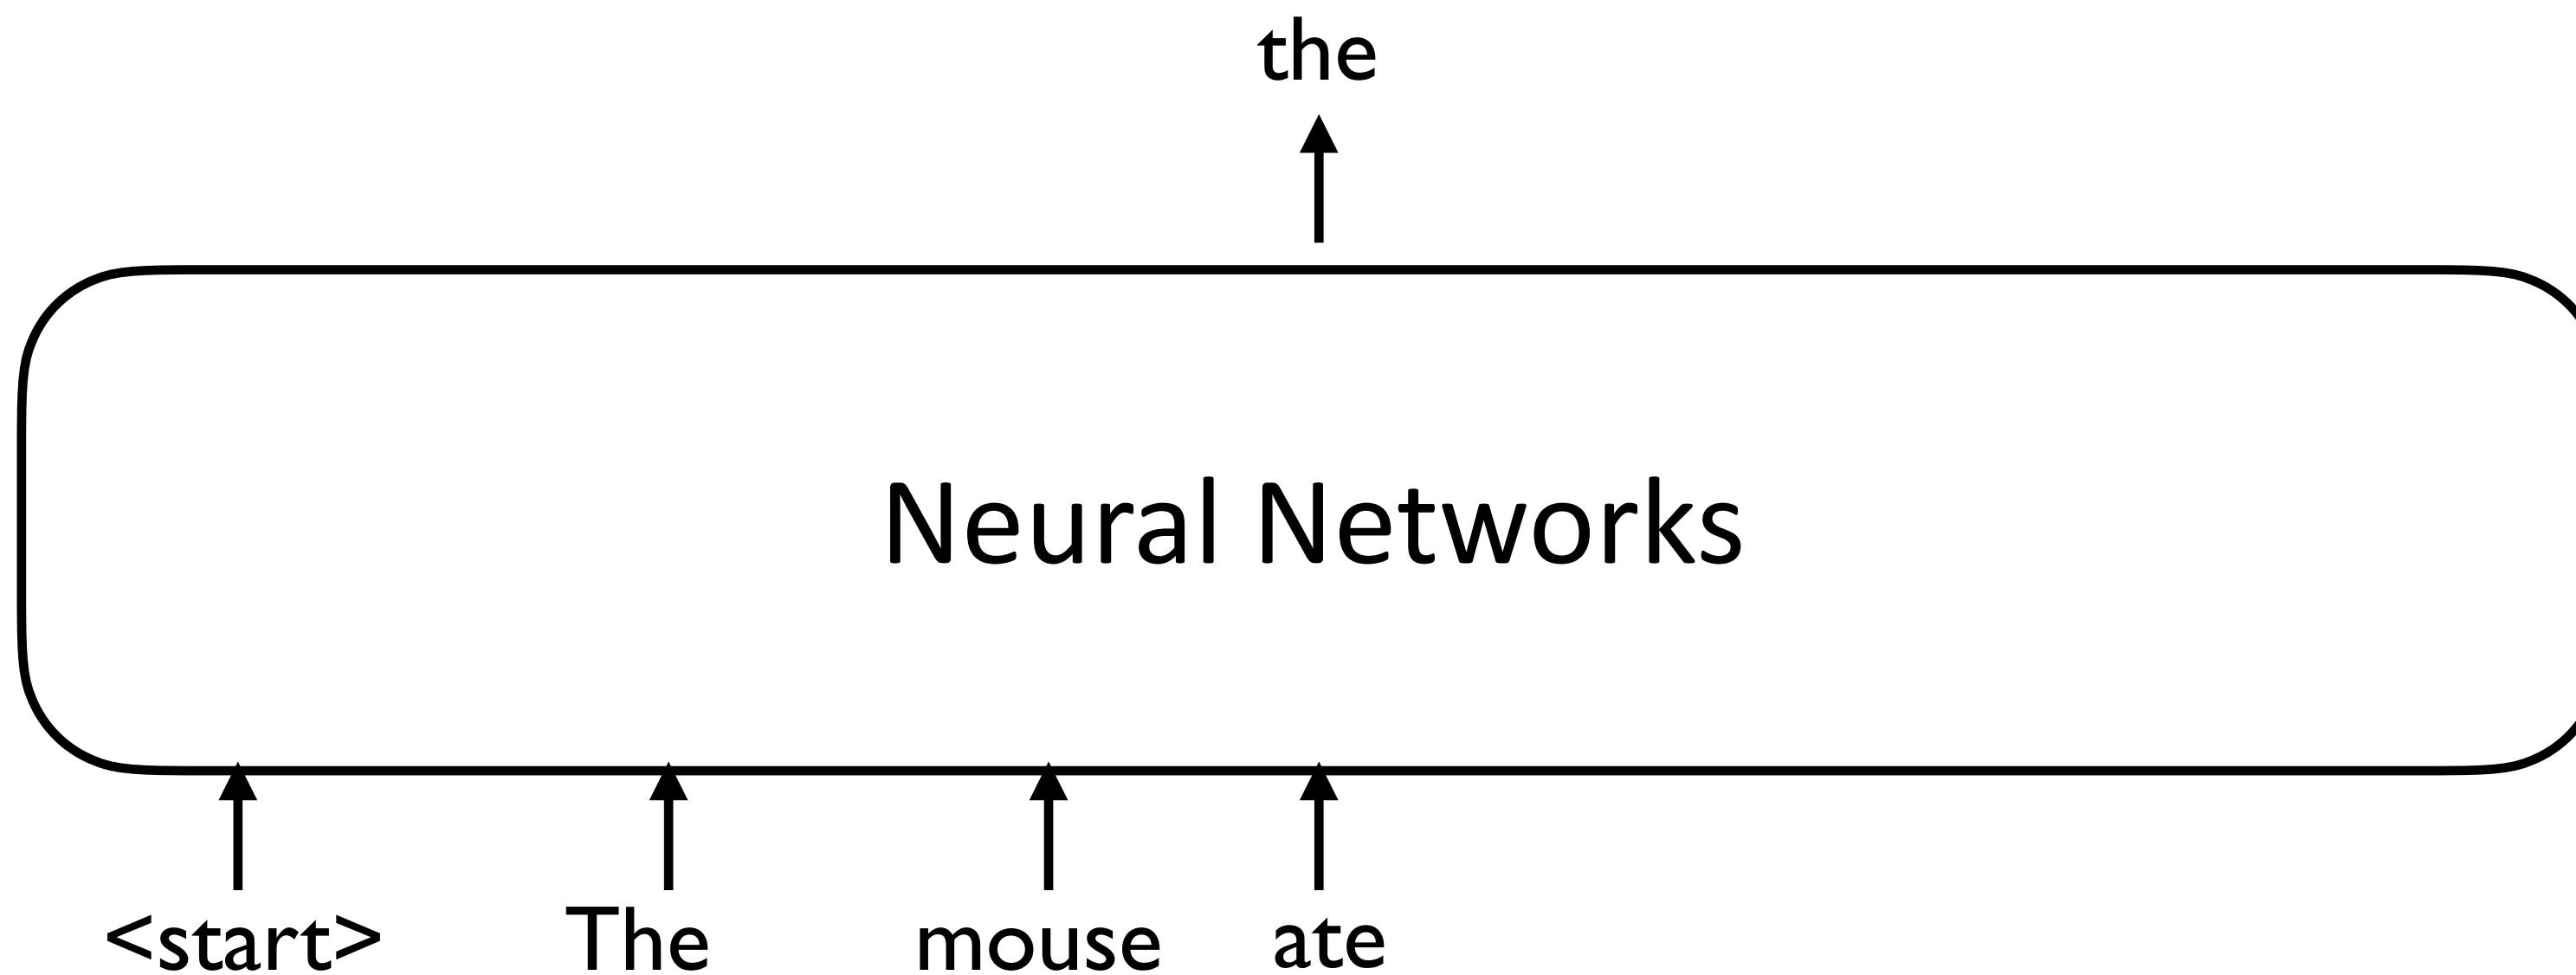
Data: “The mouse ate the cheese.”



# Recap: Neural Language Models

Neural language models are typically autoregressive

Data: “The mouse ate the cheese.”

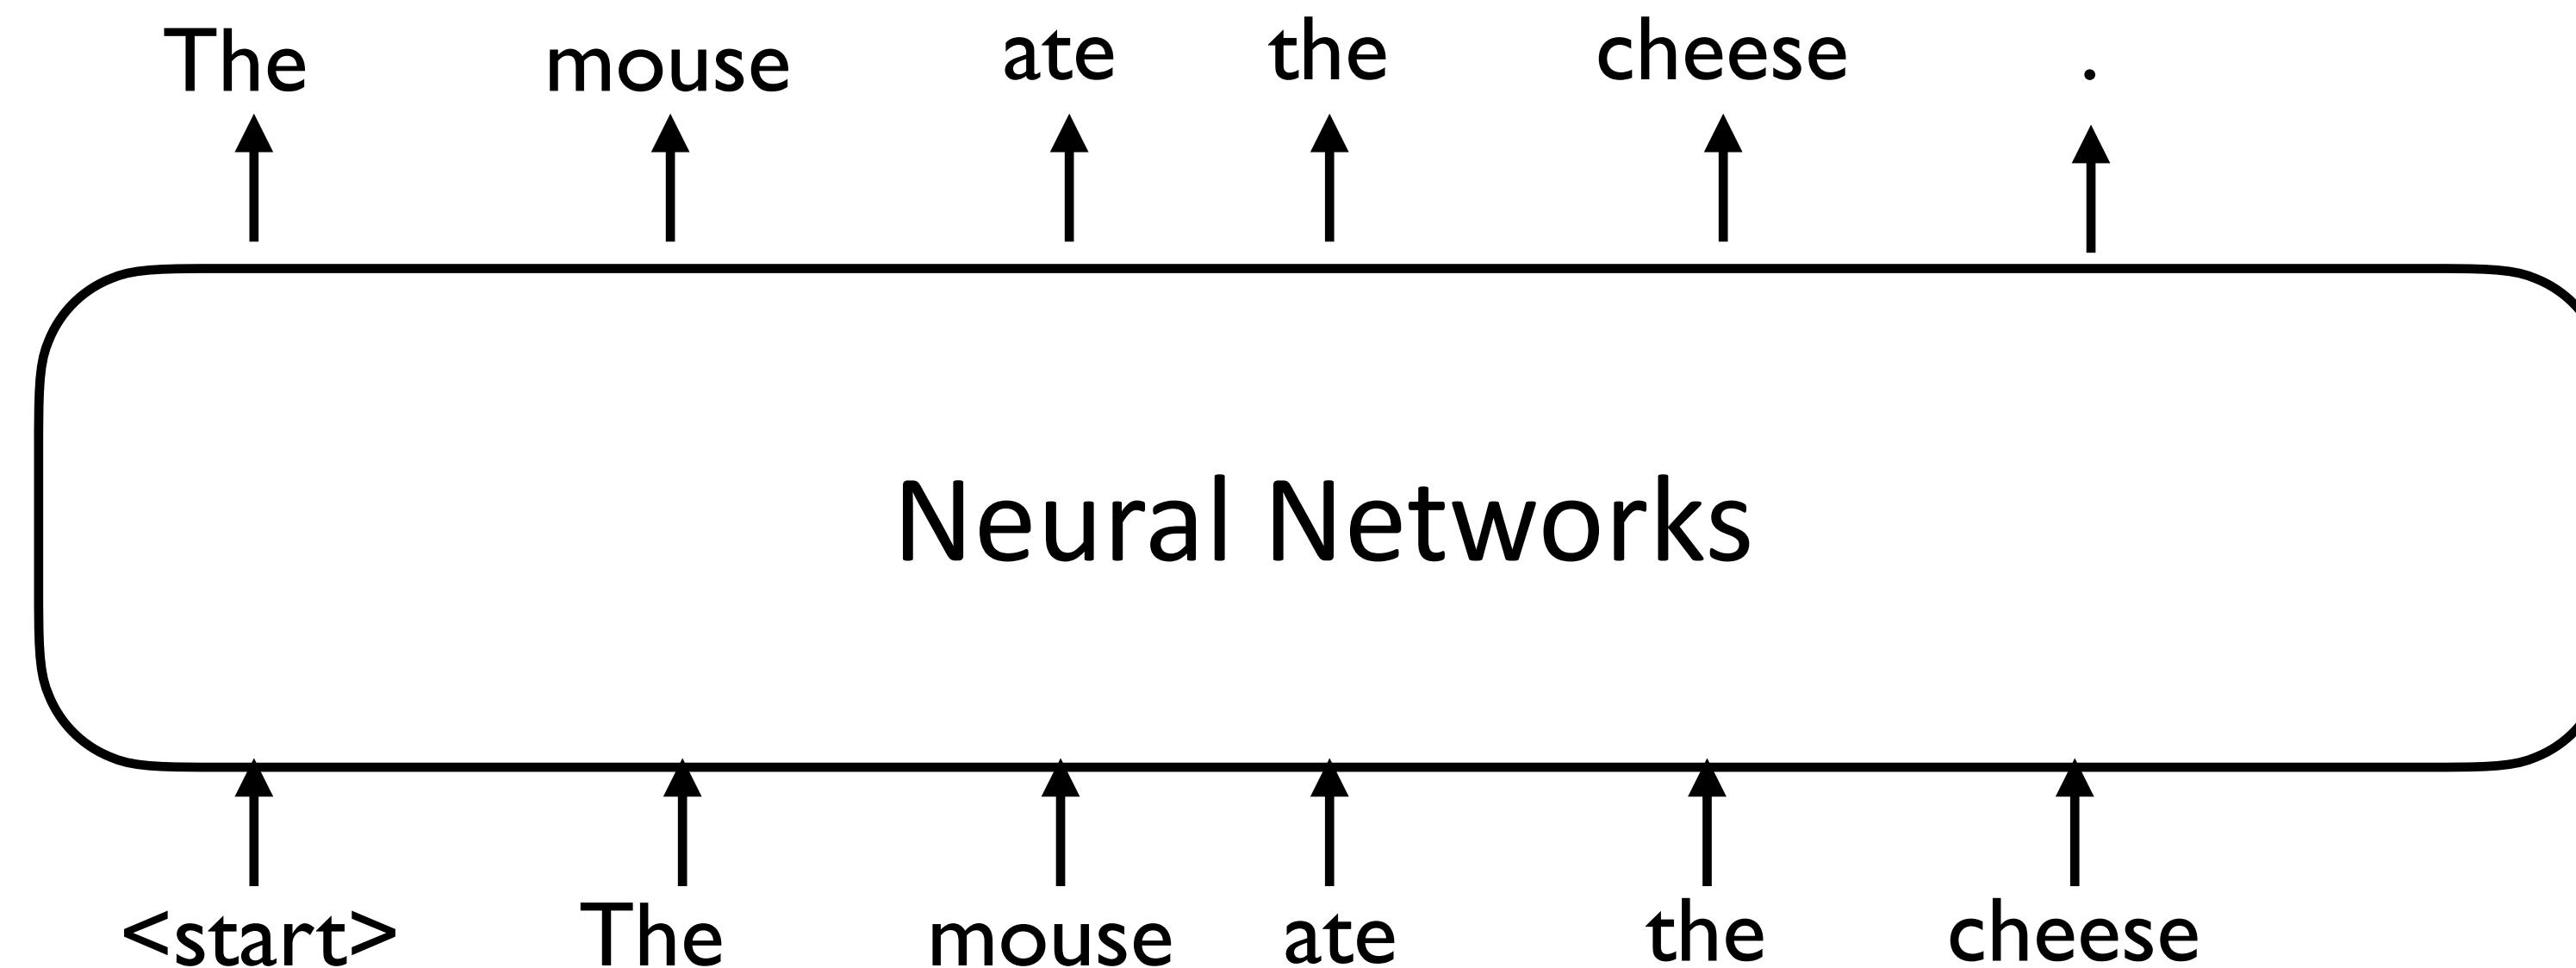


We can compute the loss on every token in parallel

# Recap: Neural Language Models

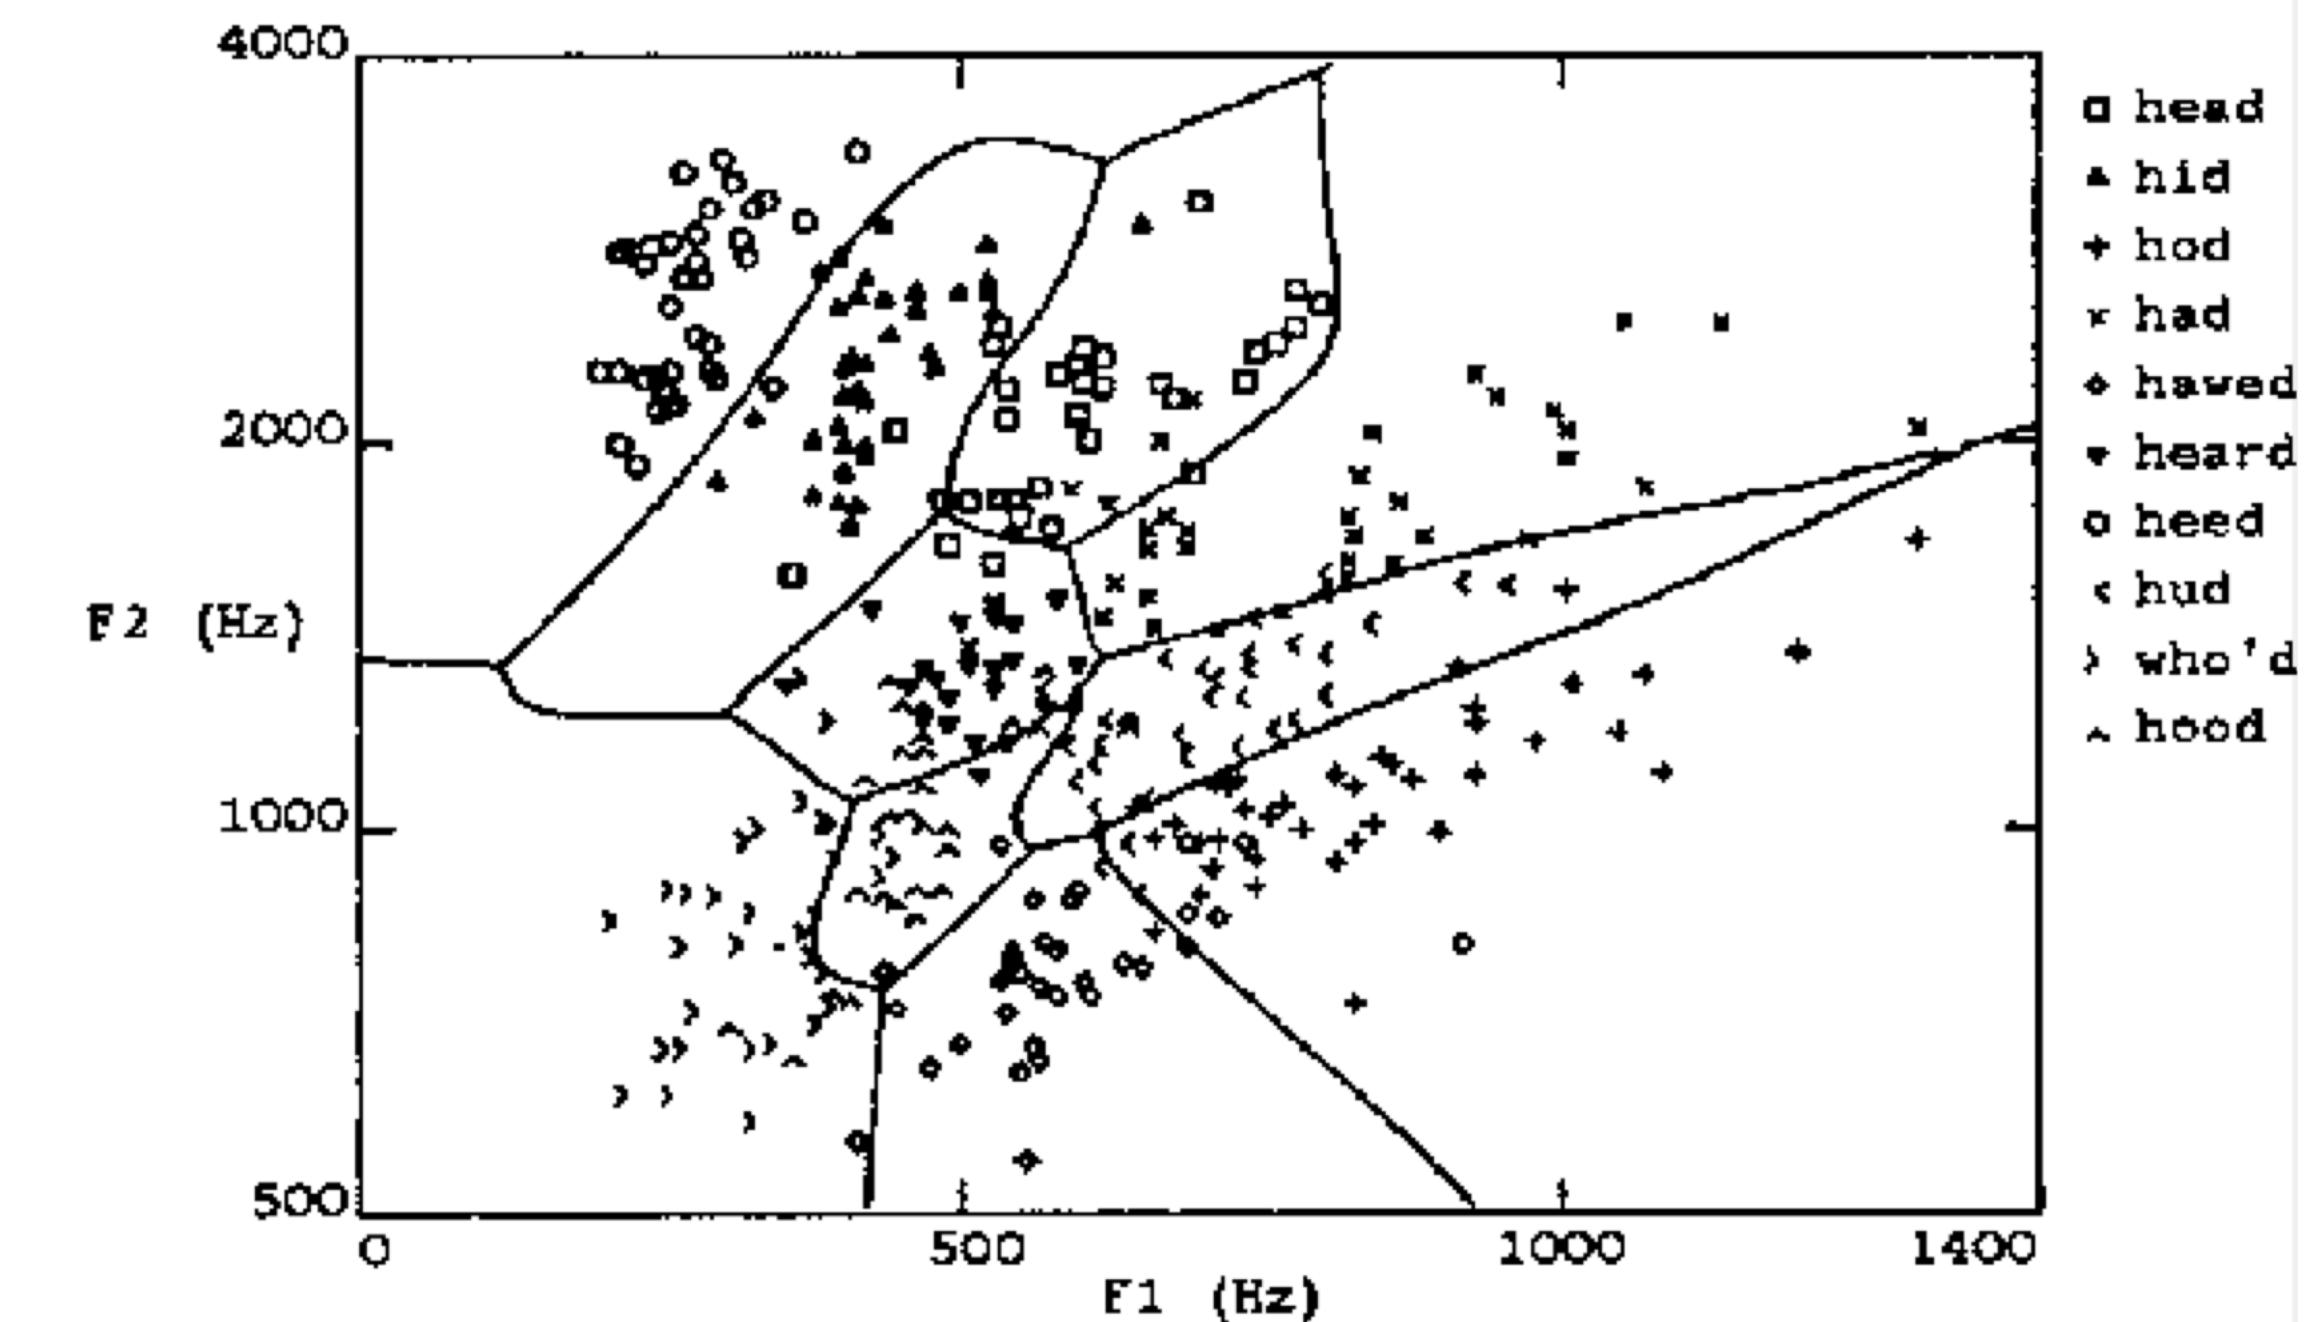
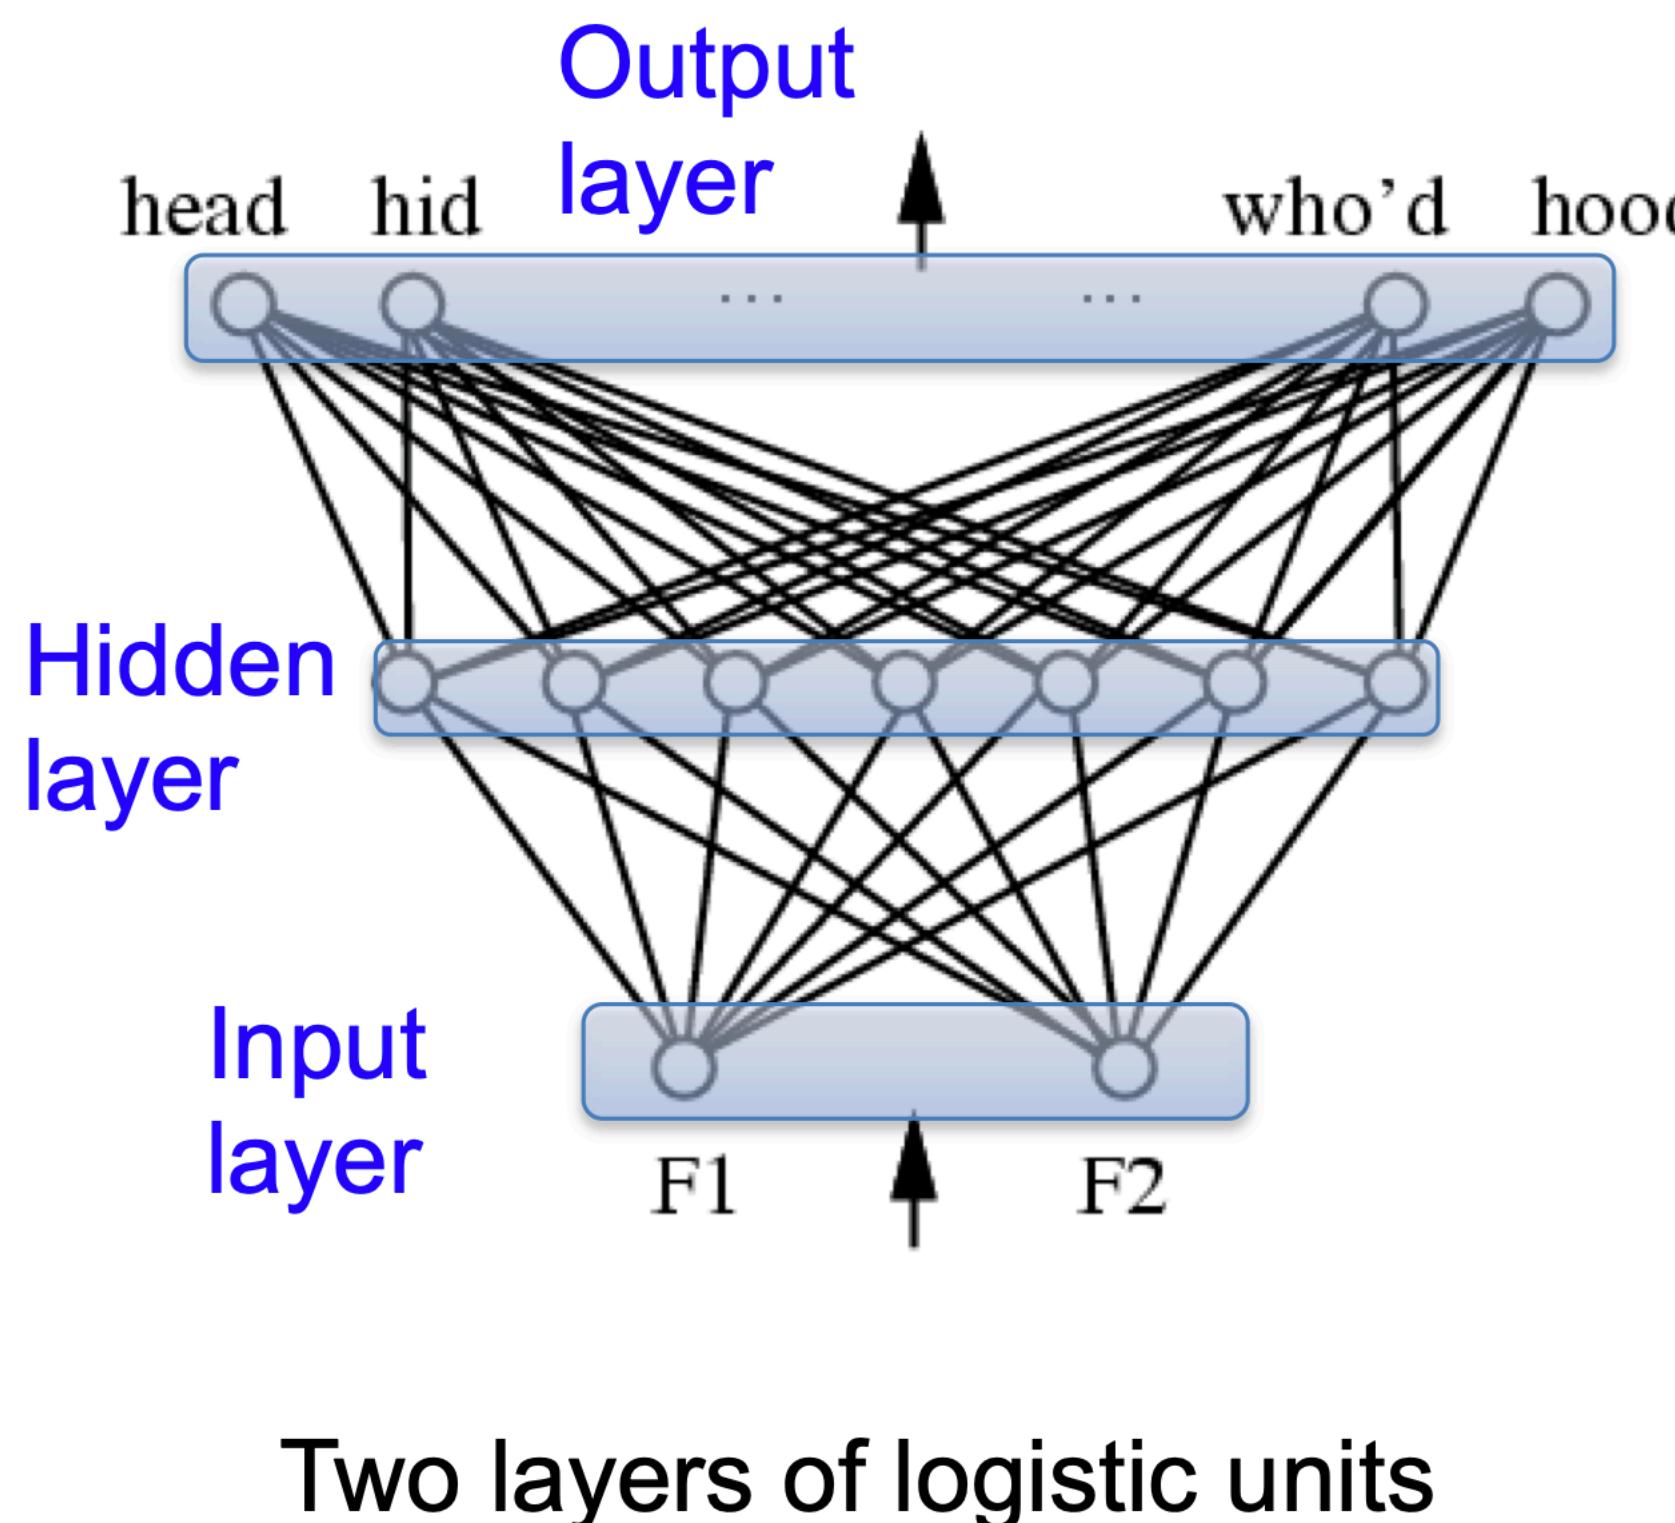
Neural language models are typically autoregressive

Data: “The mouse ate the cheese.”



Each prediction only sees the inputs on its left

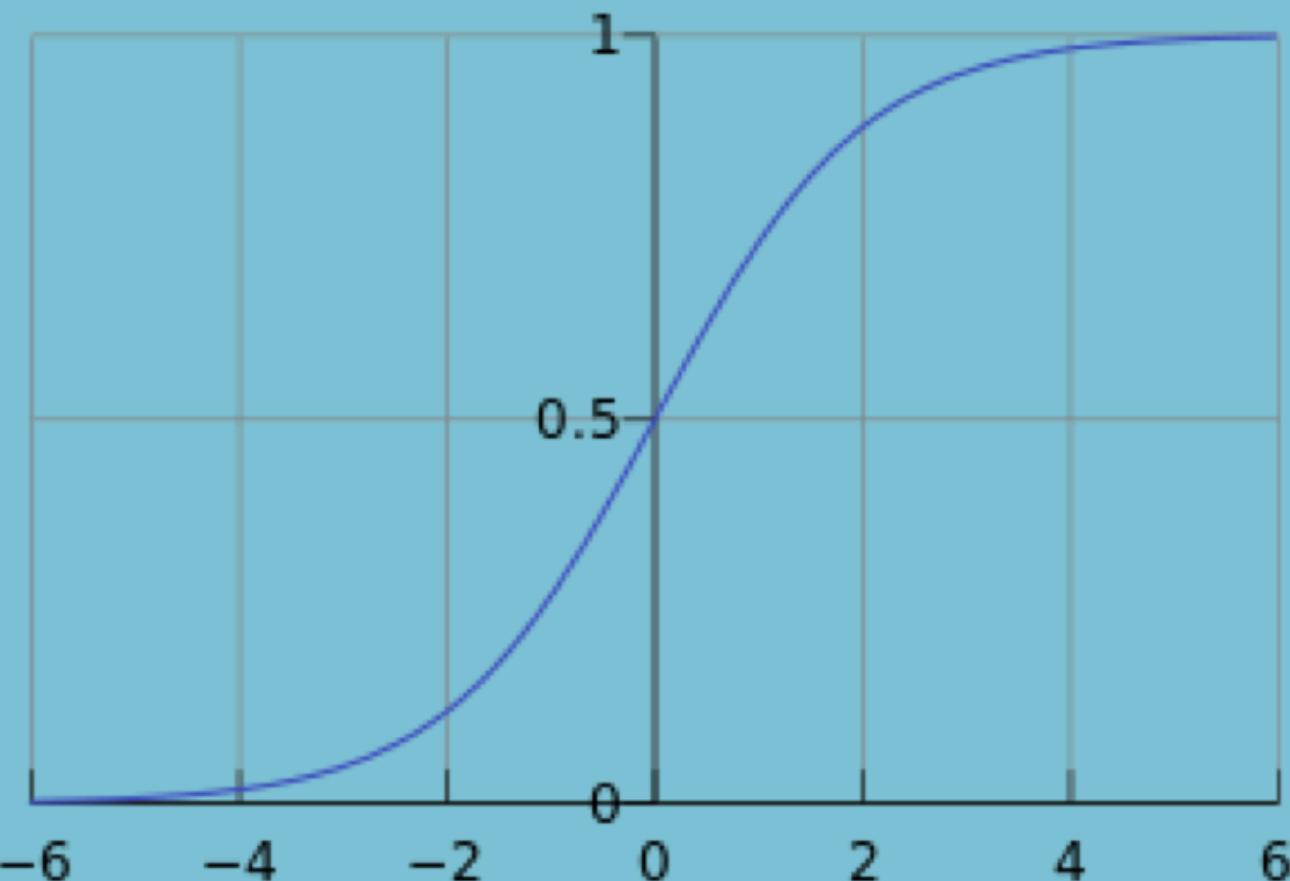
# Recap: Multilayer Networks of Sigmoid Units



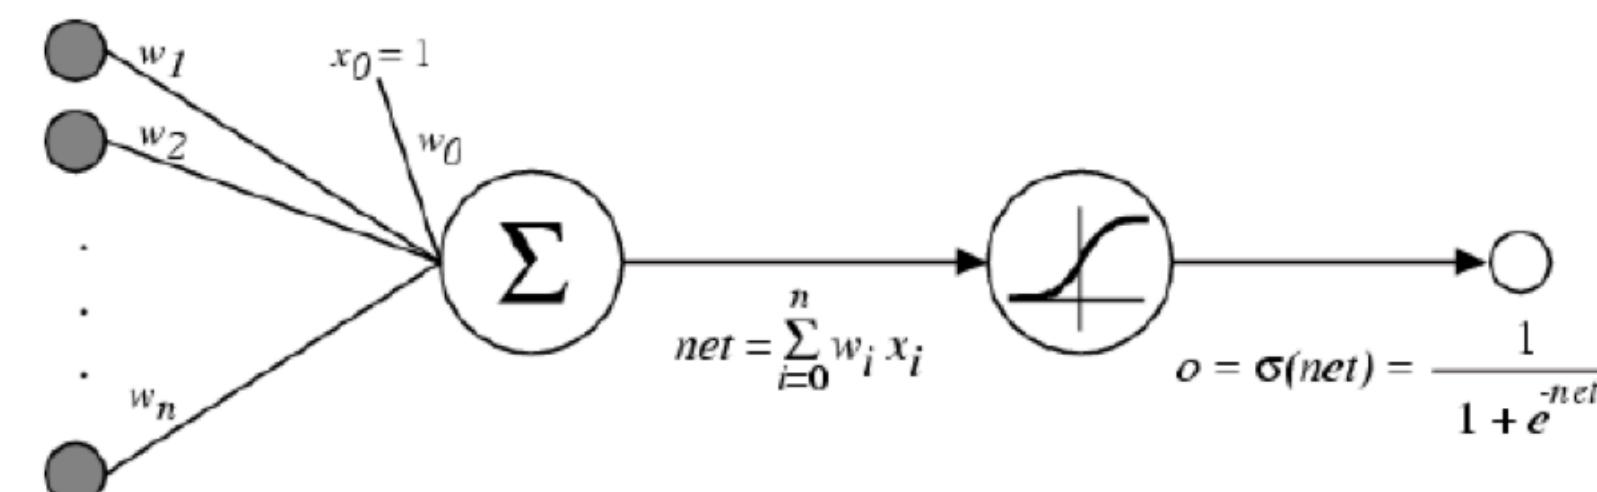
# Activation Functions

Sigmoid / Logistic Function

$$\text{logistic}(u) = \frac{1}{1 + e^{-u}}$$

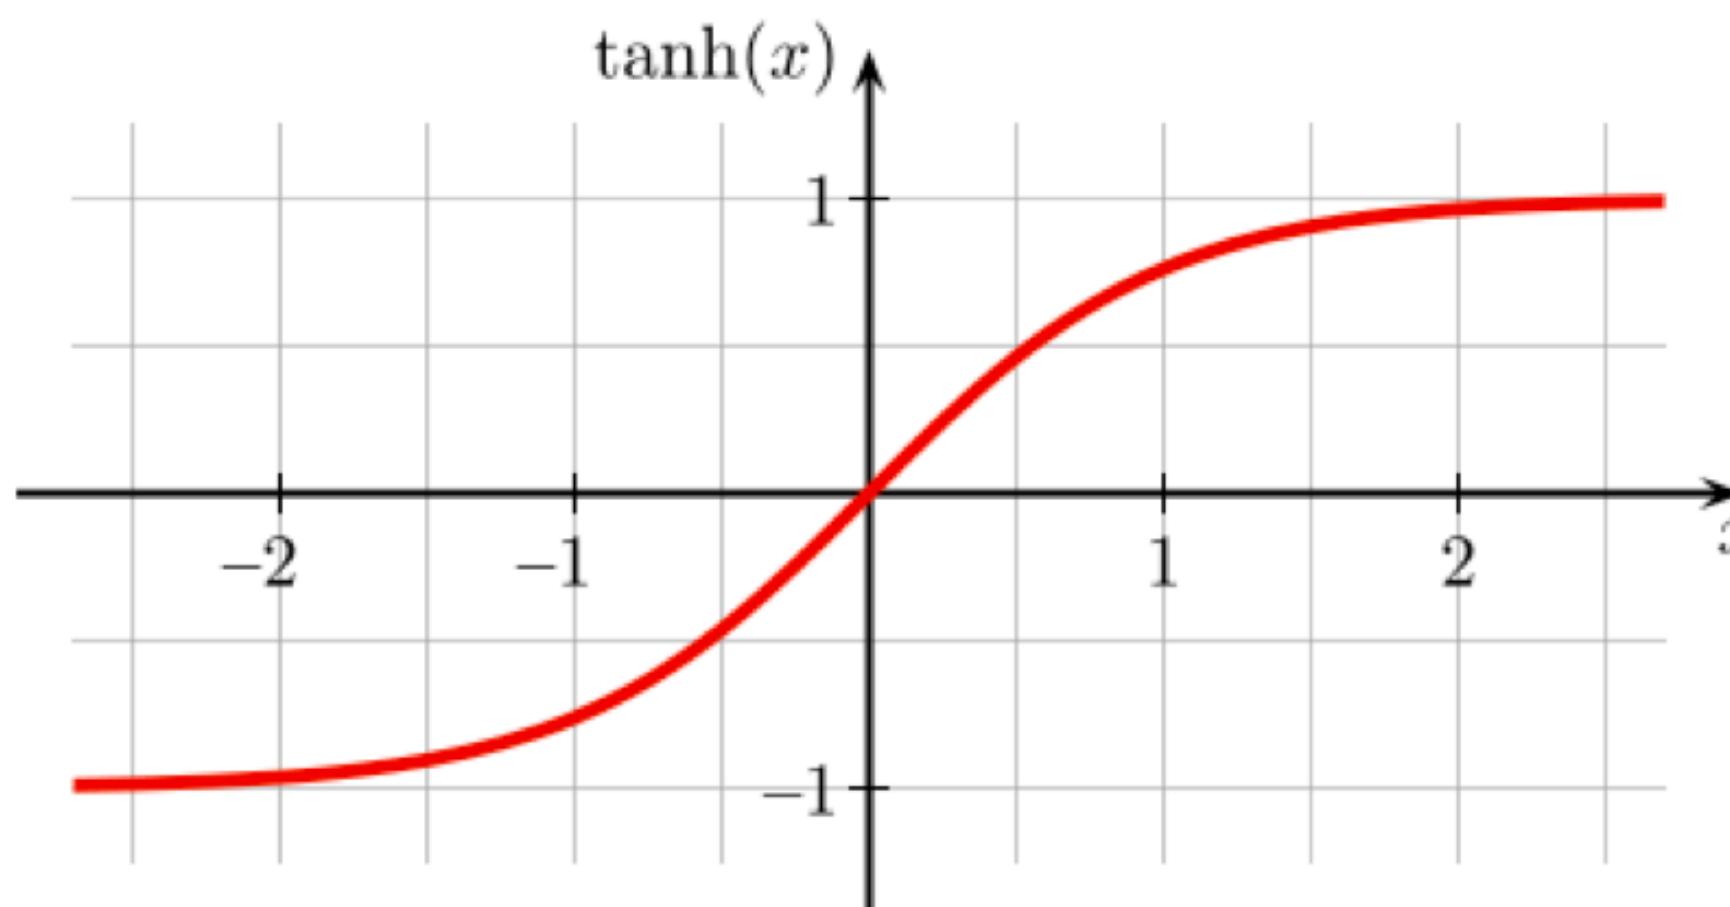


So far, we've assumed that the activation function (nonlinearity) is always the sigmoid function...



# Tanh

- A new change: modifying the nonlinearity
  - The logistic is not widely used in modern ANNs



Alternate 1:  
 $\tanh$

Like logistic function but  
shifted to range [-1, +1]

# Activation Function

## Understanding the difficulty of training deep feedforward neural networks

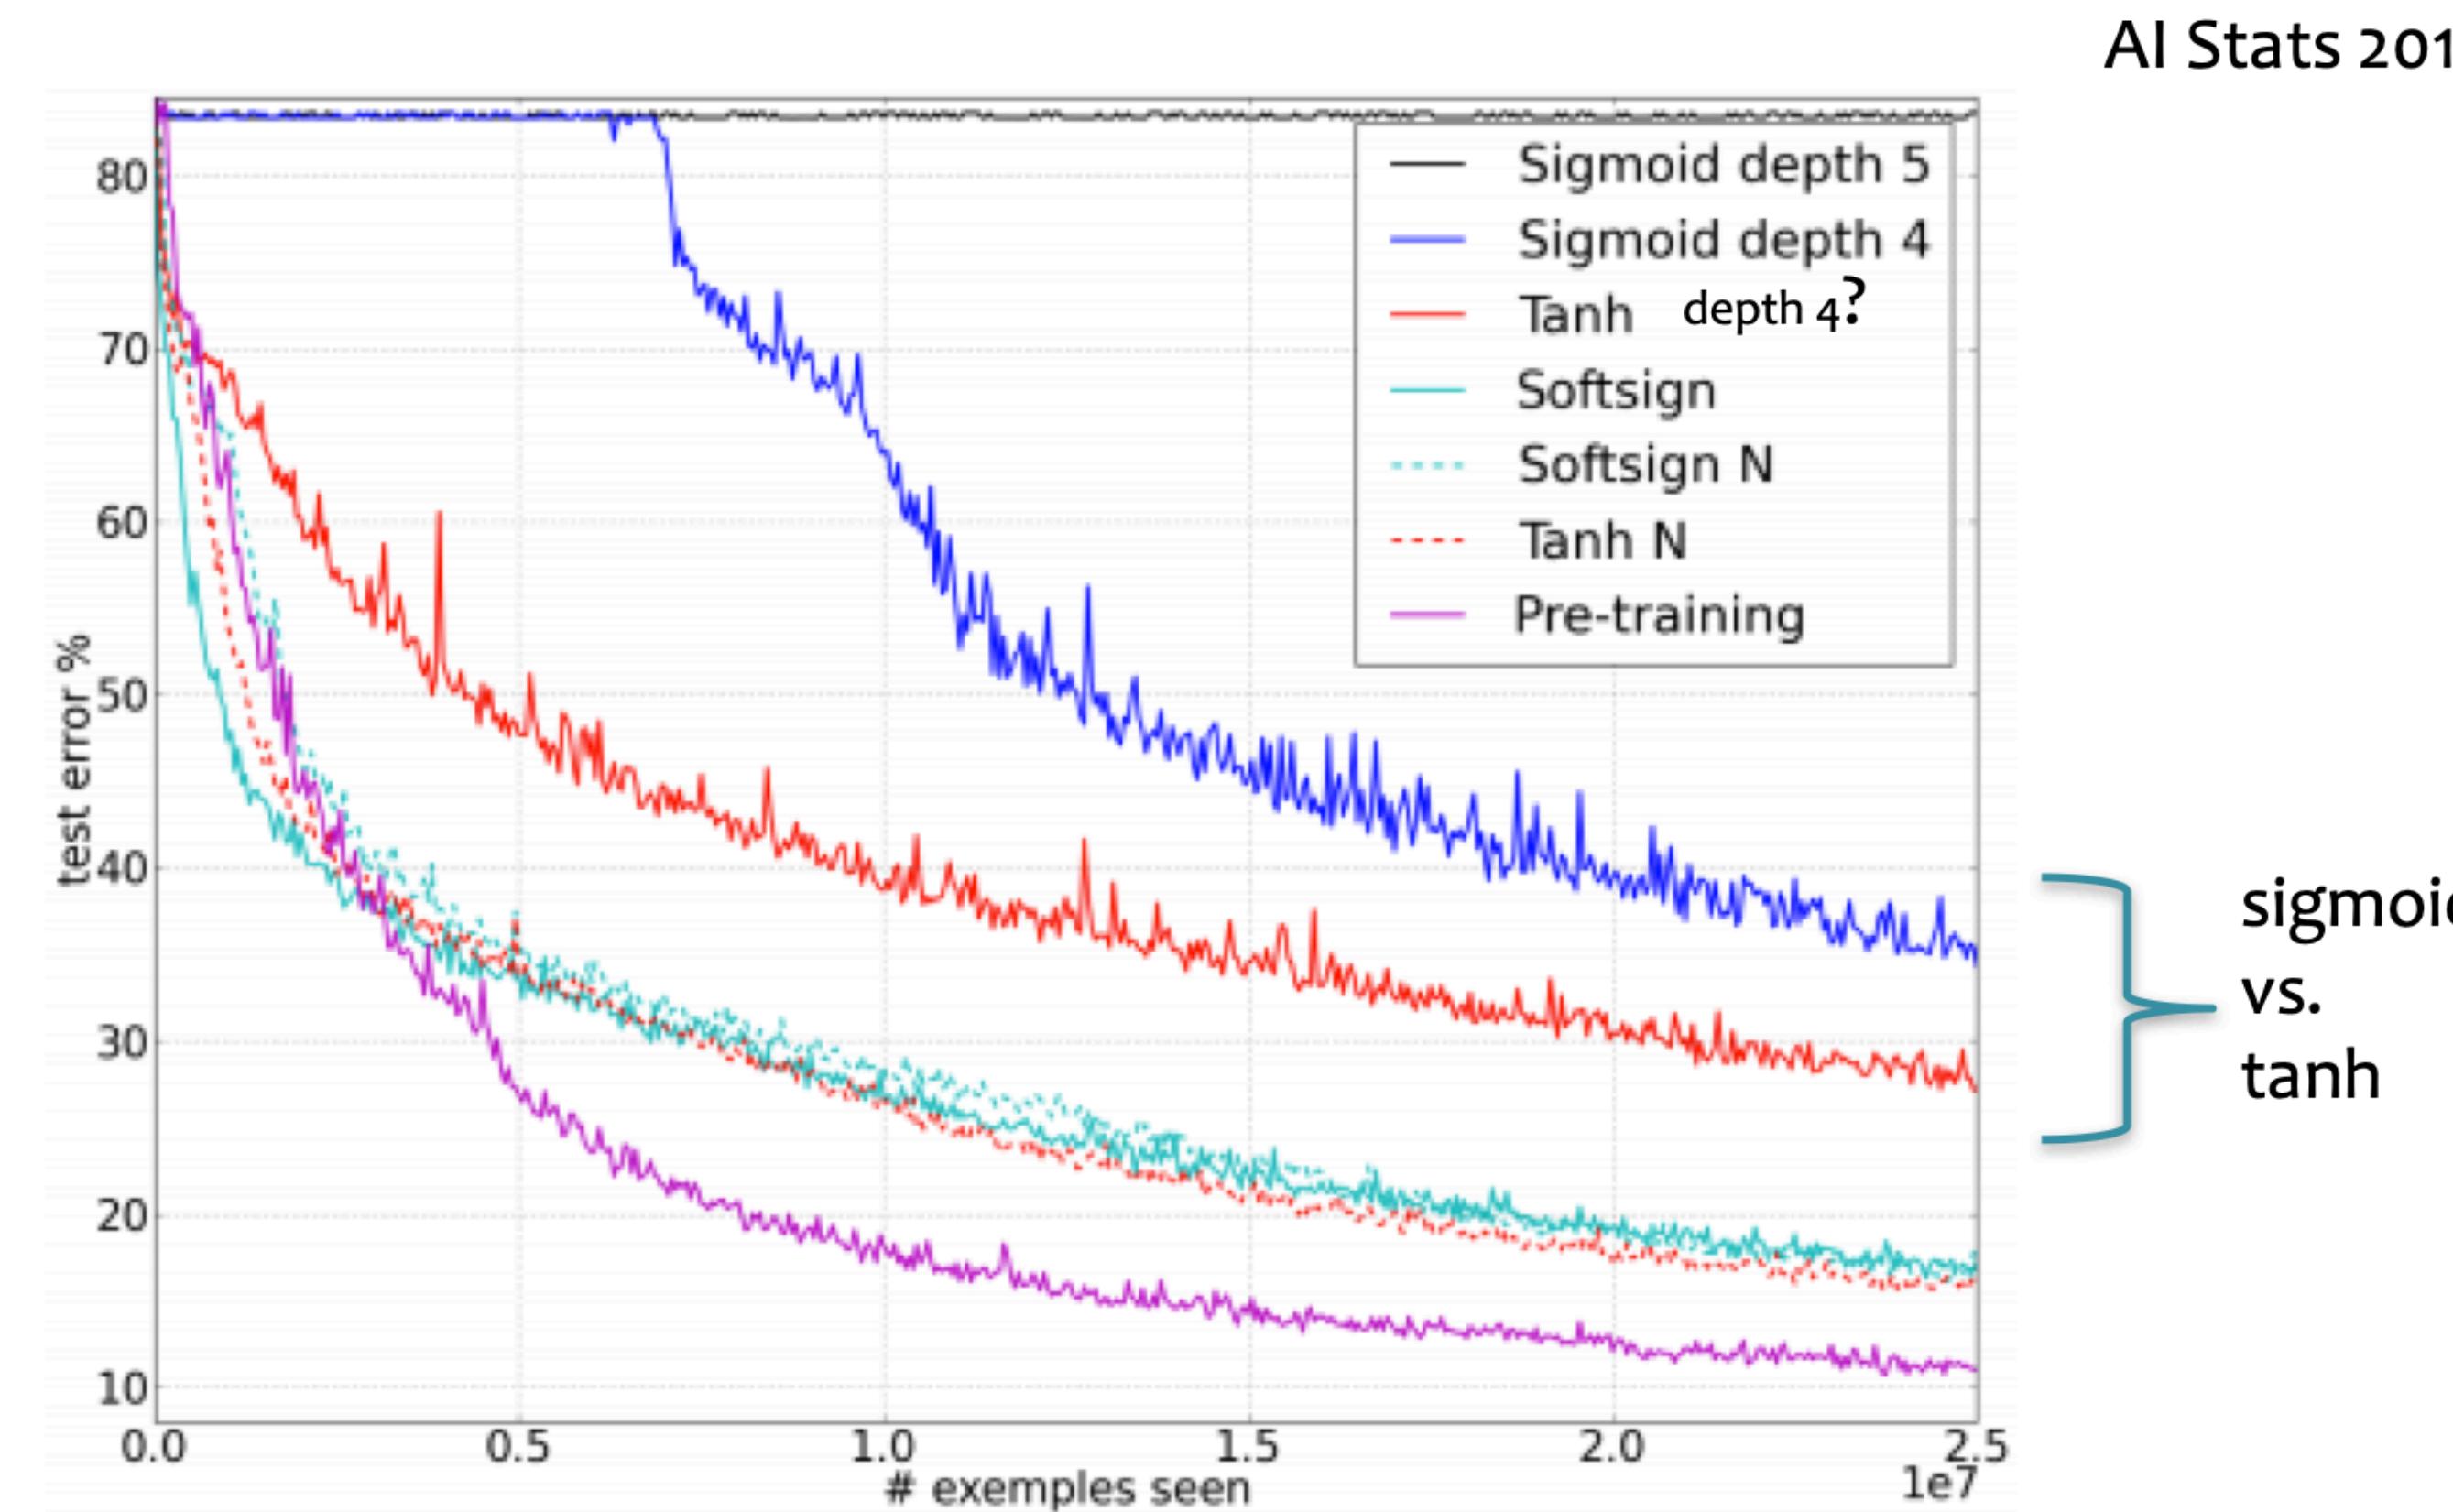
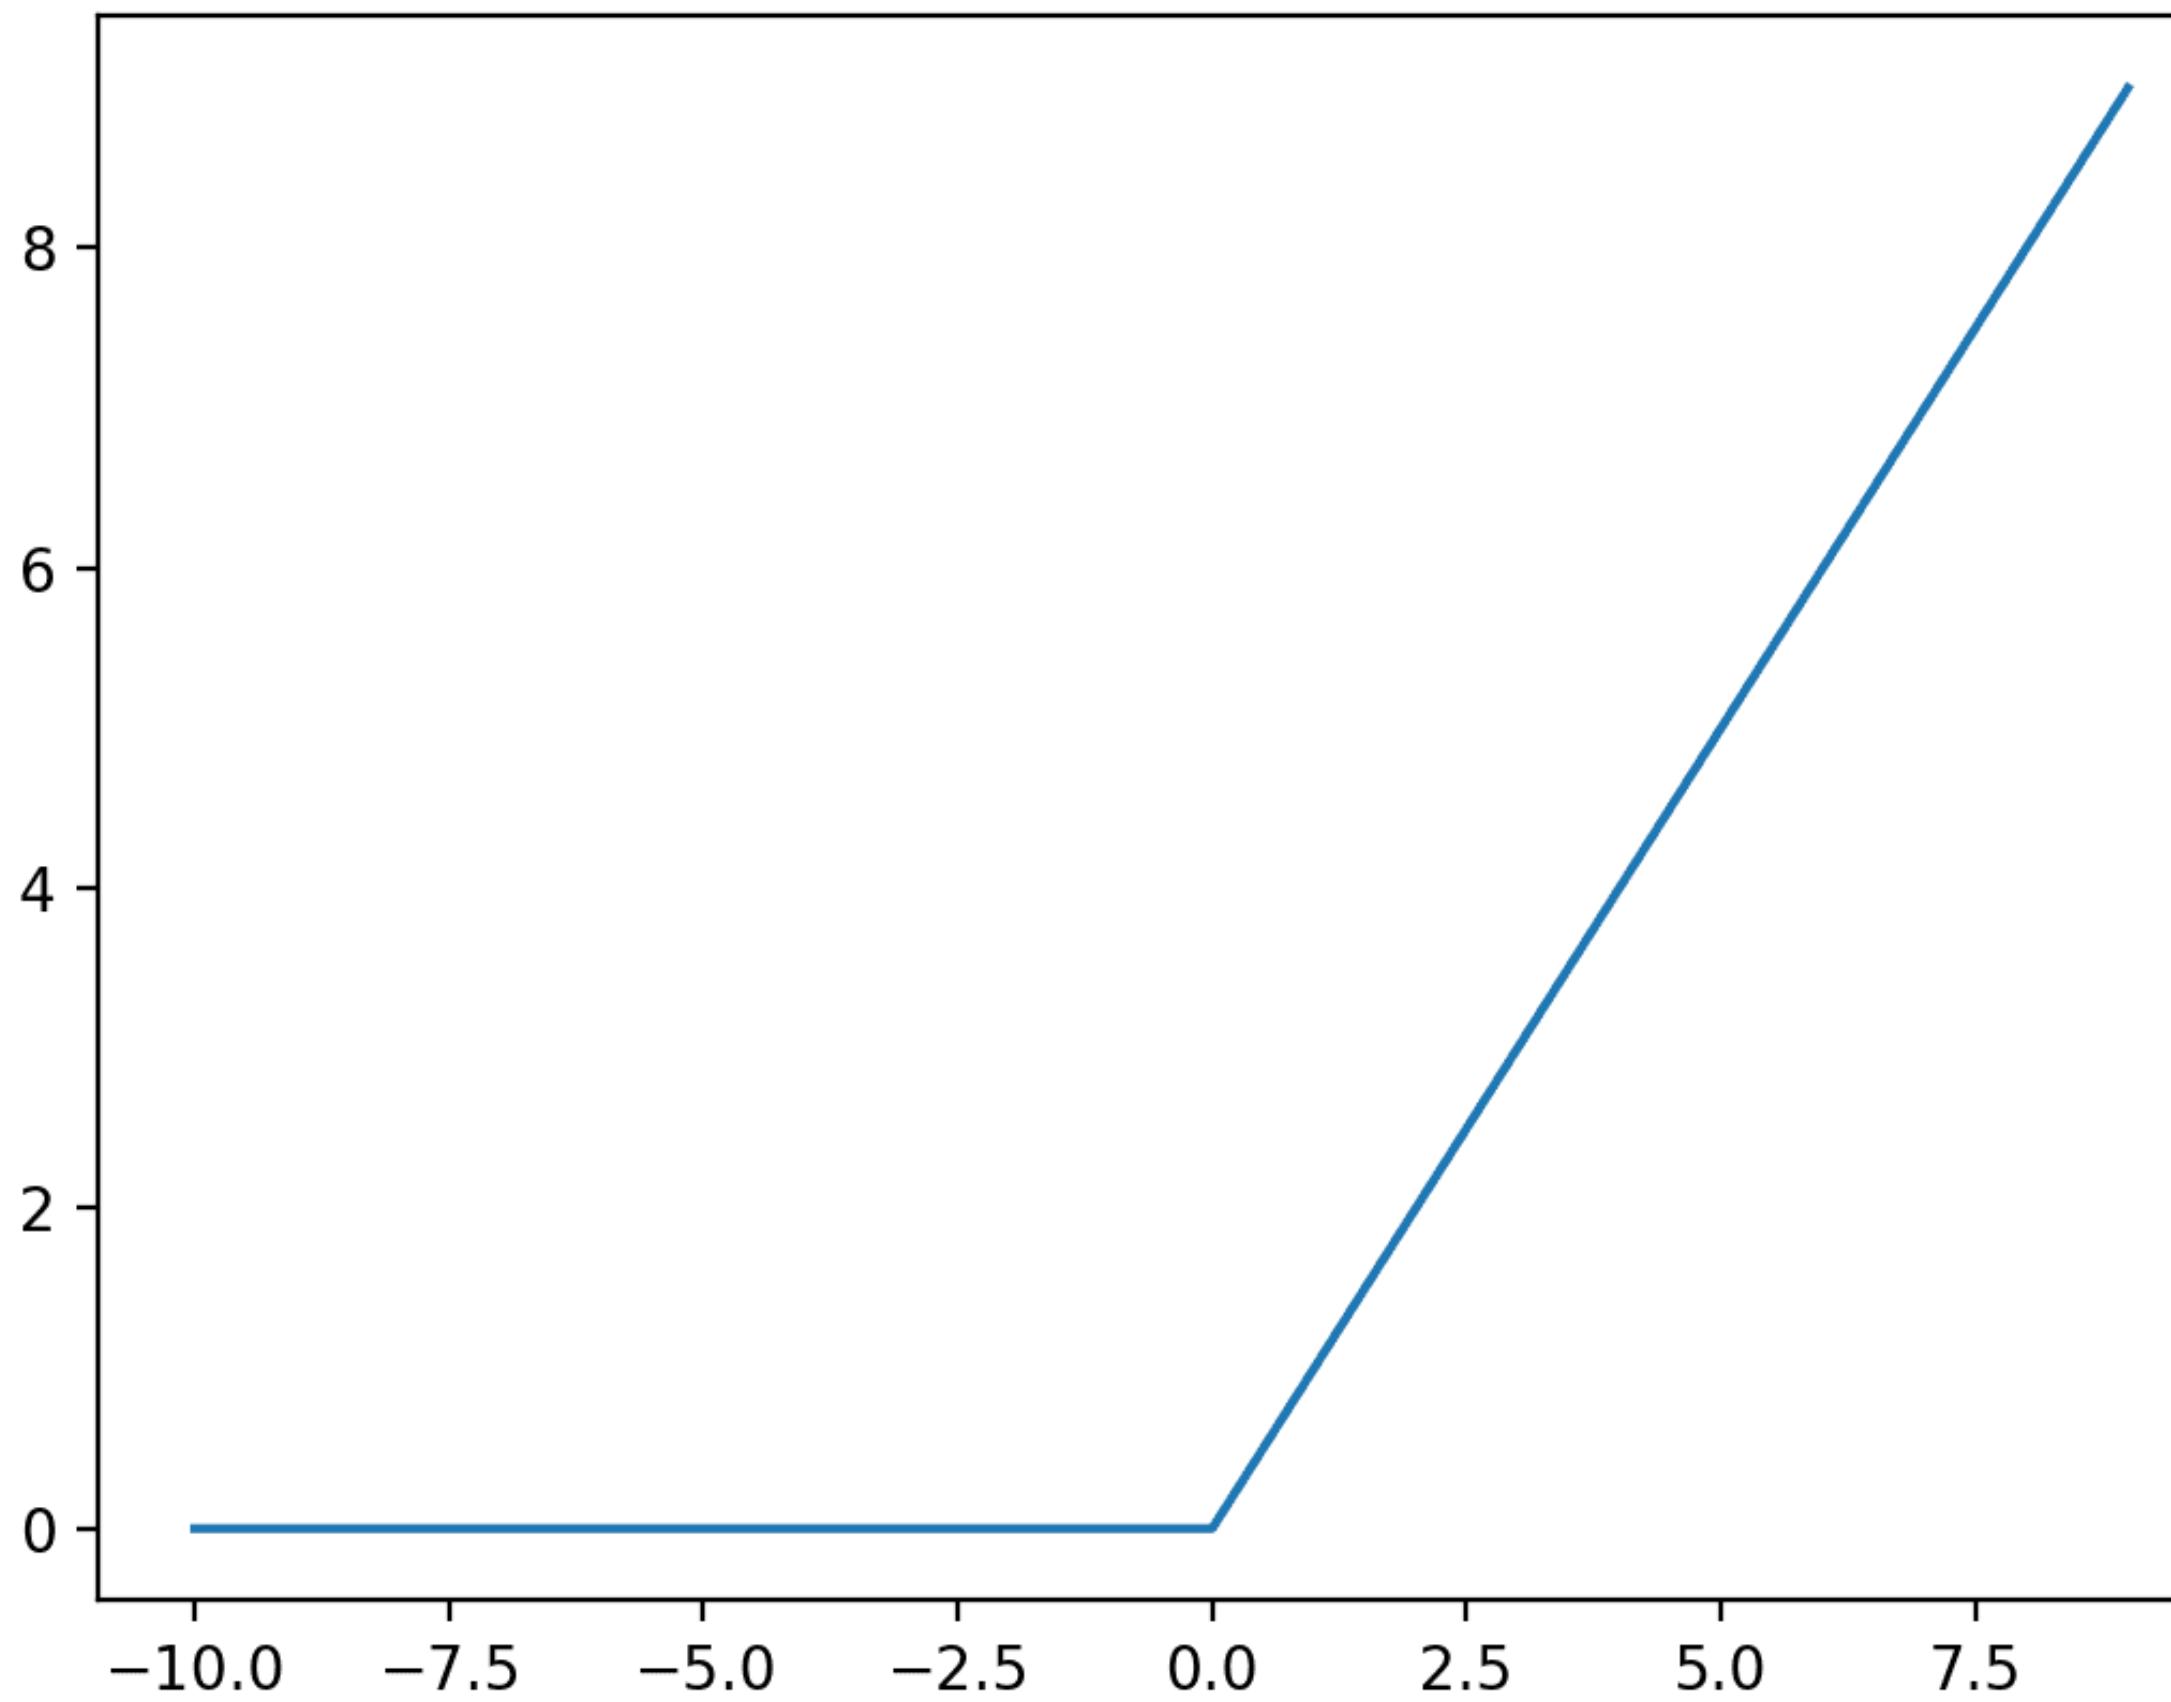


Figure from Glorot & Bentio (2010)

# ReLU



# Other Activation Functions

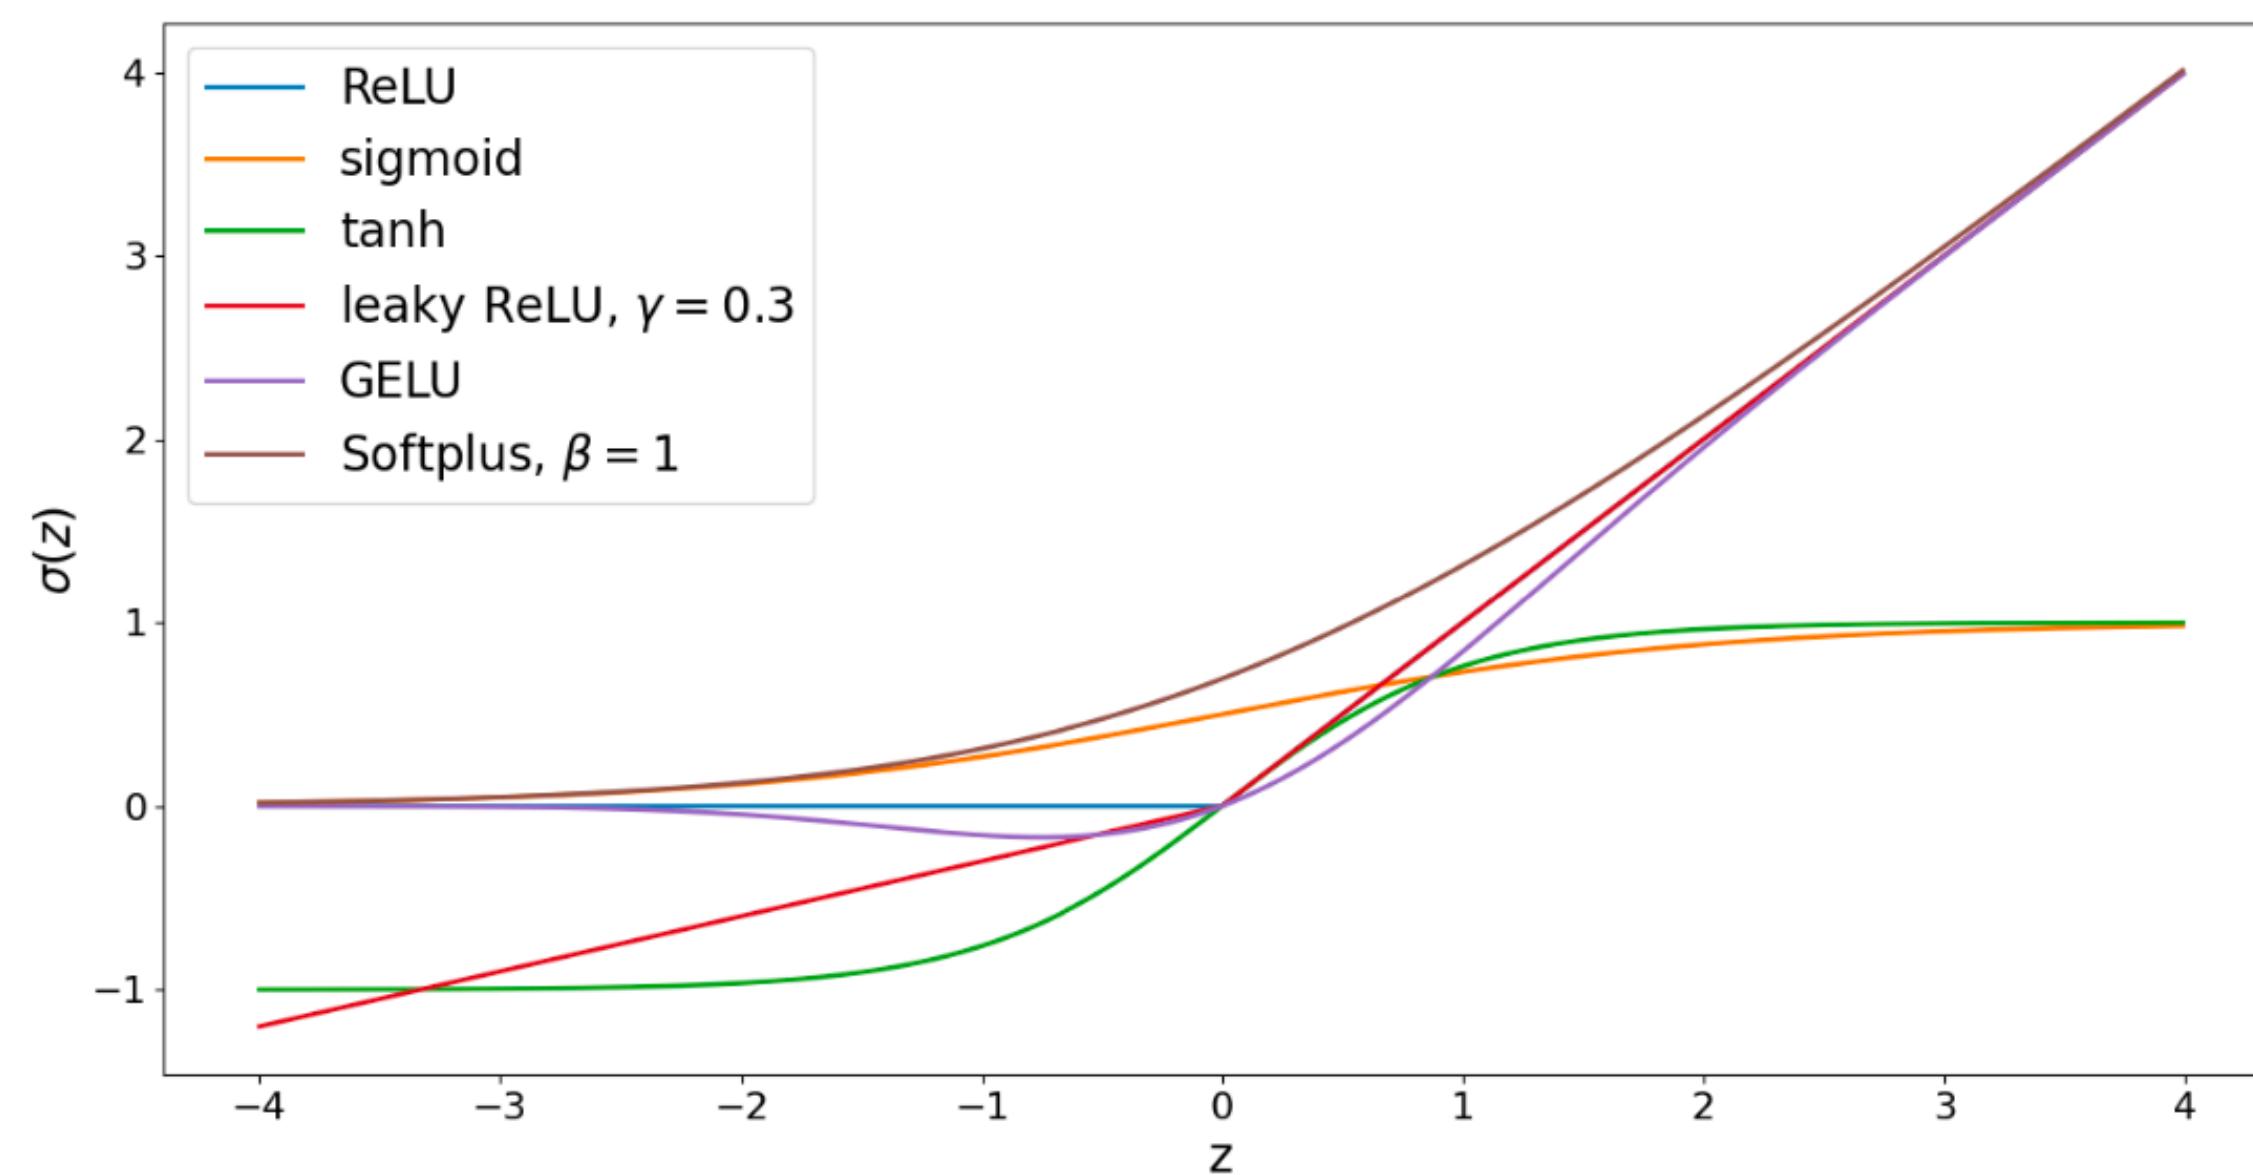
$$\sigma(z) = \frac{1}{1 + e^{-z}} \quad (\text{sigmoid})$$

$$\sigma(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}} \quad (\tanh)$$

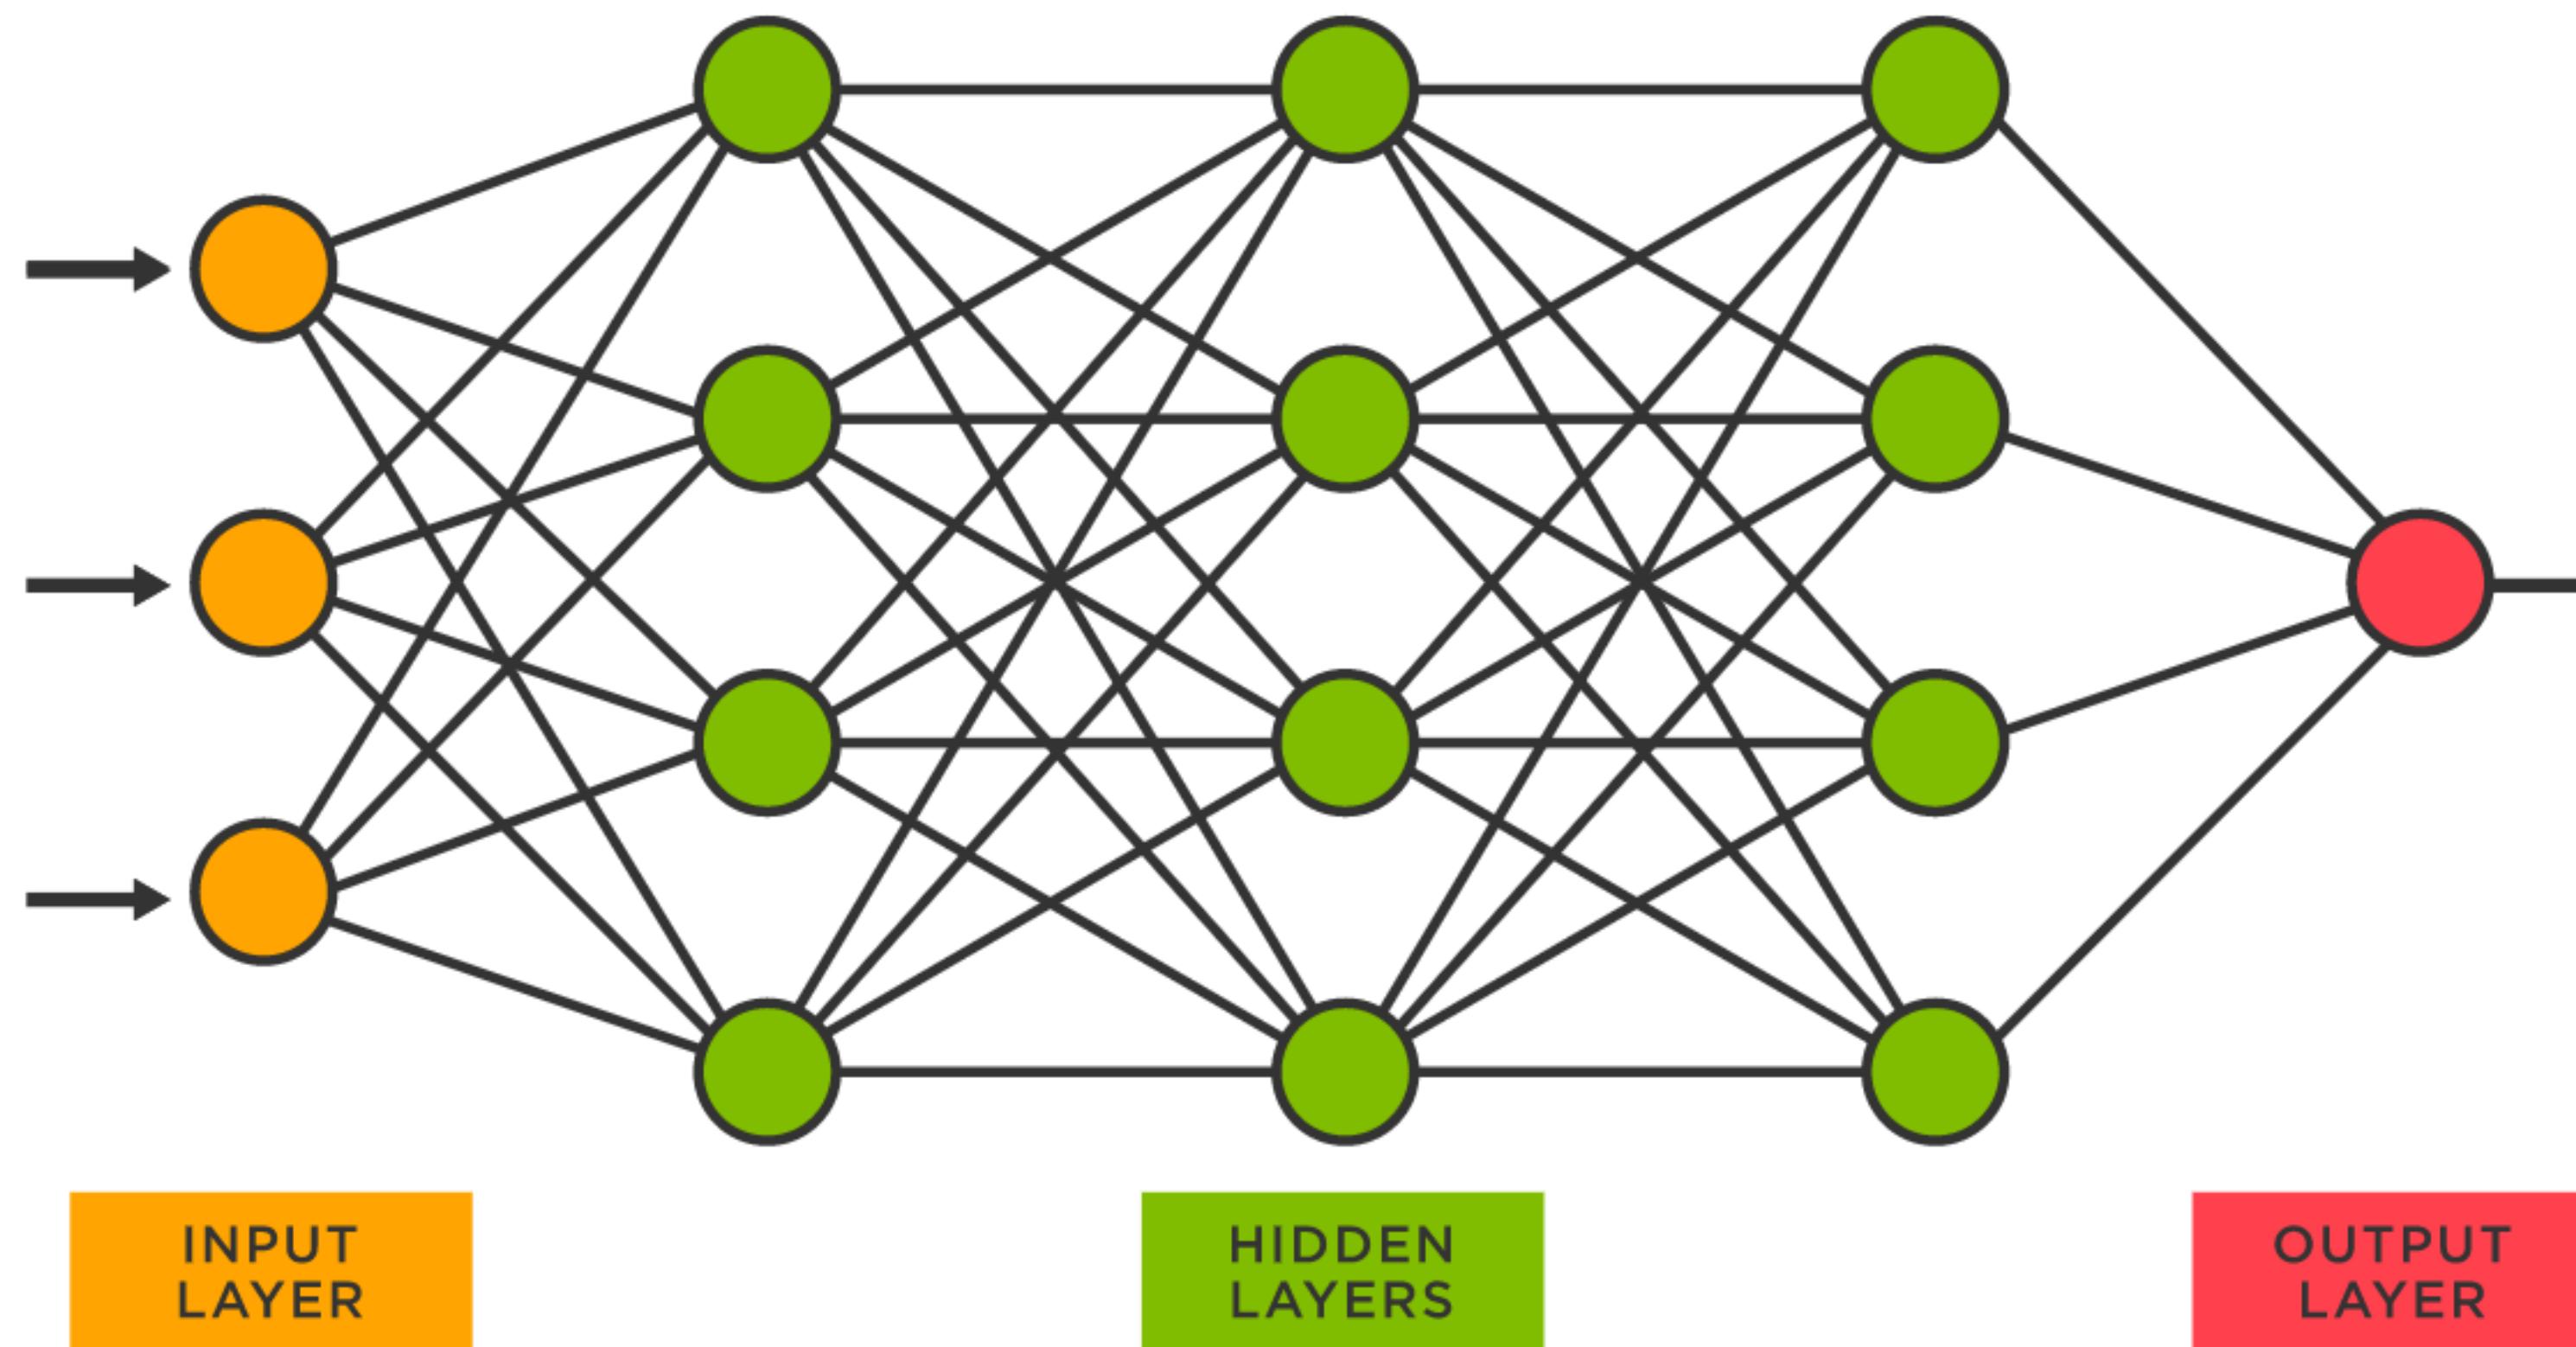
$$\sigma(z) = \max\{z, \gamma z\}, \gamma \in (0, 1) \quad (\text{leaky ReLU})$$

$$\sigma(z) = \frac{z}{2} \left[ 1 + \operatorname{erf}\left(\frac{z}{\sqrt{2}}\right) \right] \quad (\text{GELU})$$

$$\sigma(z) = \frac{1}{\beta} \log(1 + \exp(\beta z)), \beta > 0 \quad (\text{Softplus})$$



# Multilayer Perceptron Neural Networks (MLP)



# Residual Connection

We want deeper and deeper NNs, but going deep is difficult

- Troubles accumulate w/ more layers
- Signals get distorted when propagated
- in forward and backward passes

Commonly used techniques to train “Deep” NNs:

Weight initialization

Normalization modules

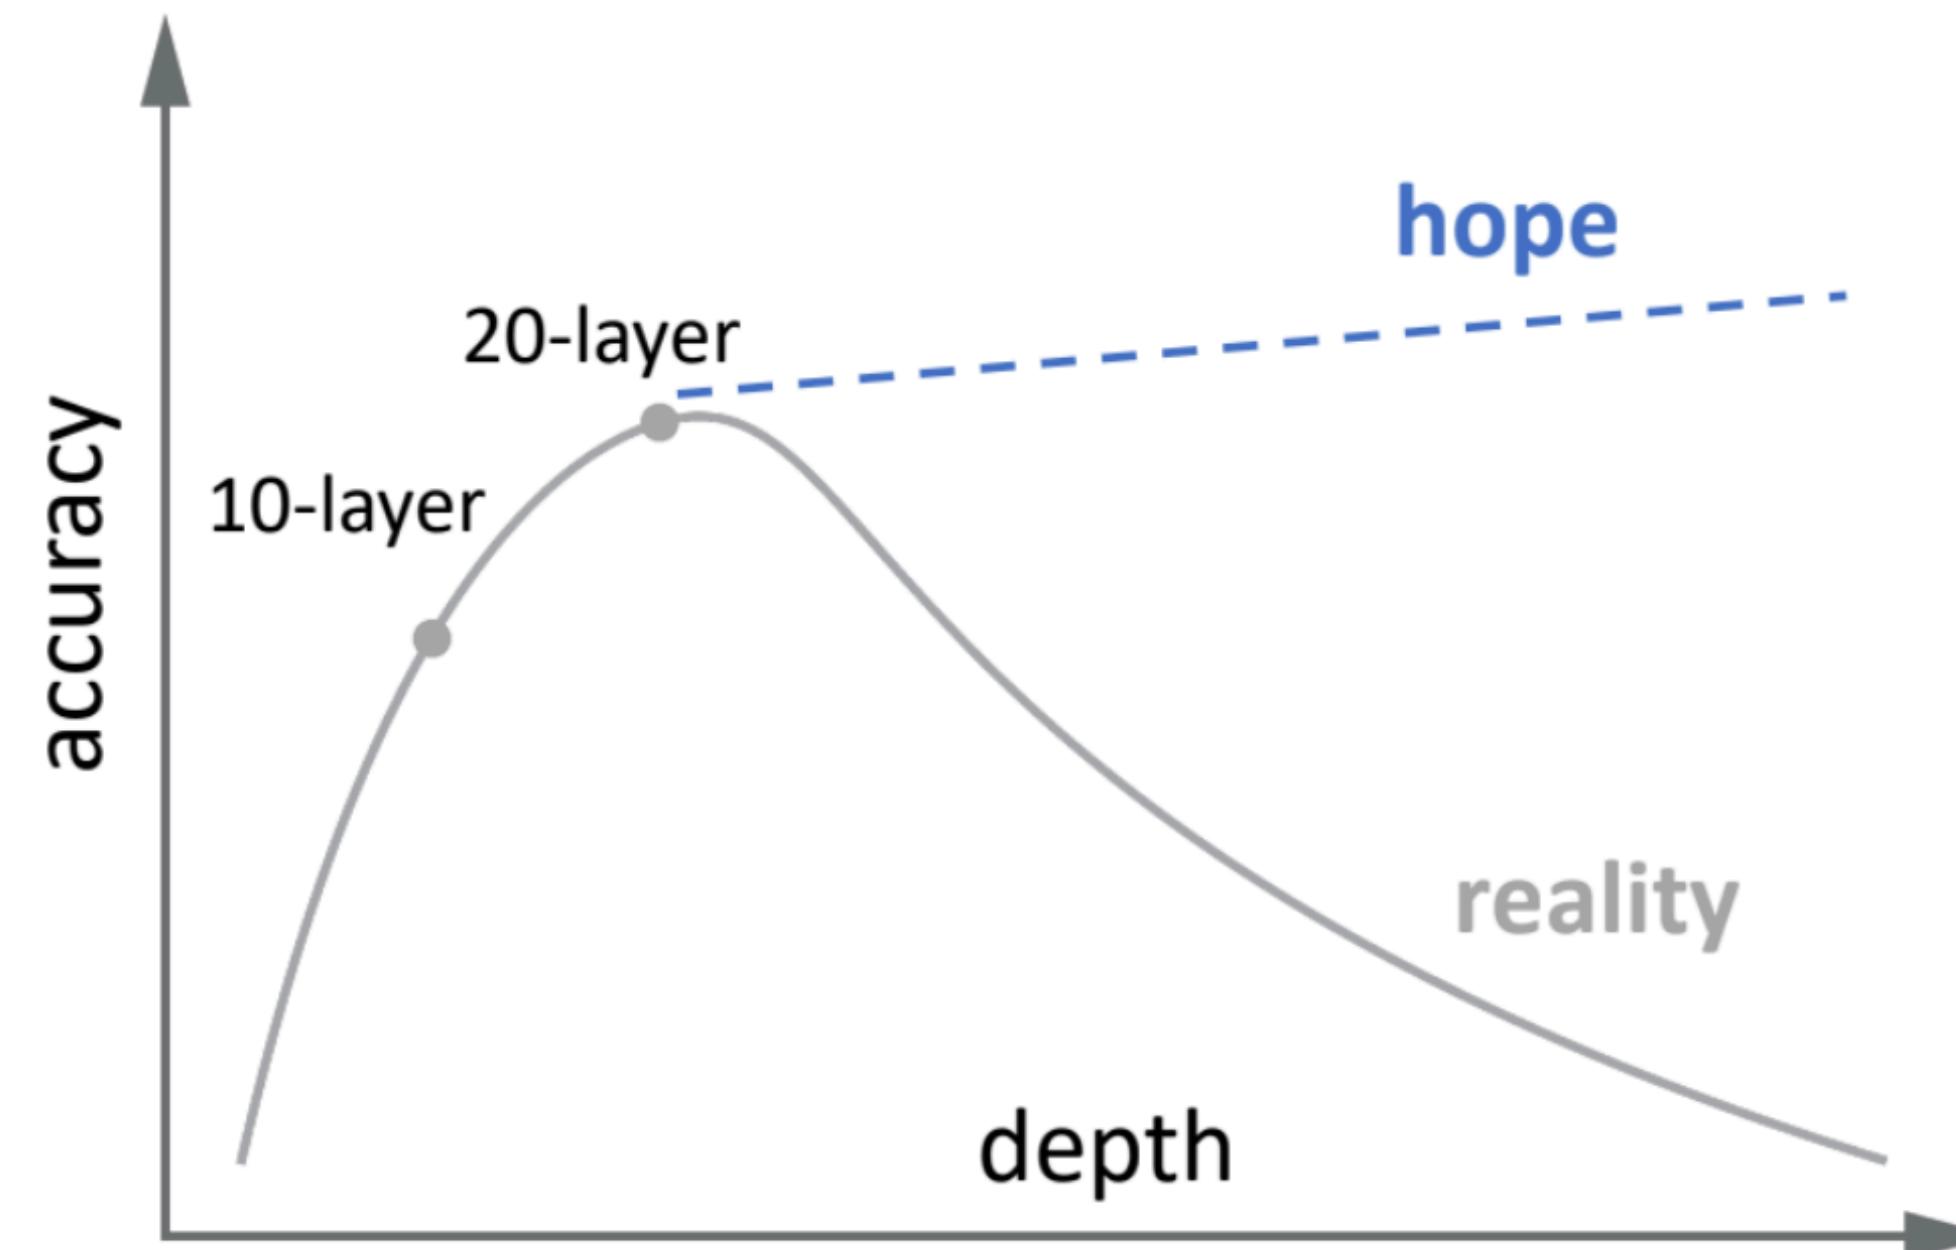
Deep residual learning



# The Degradation Problem

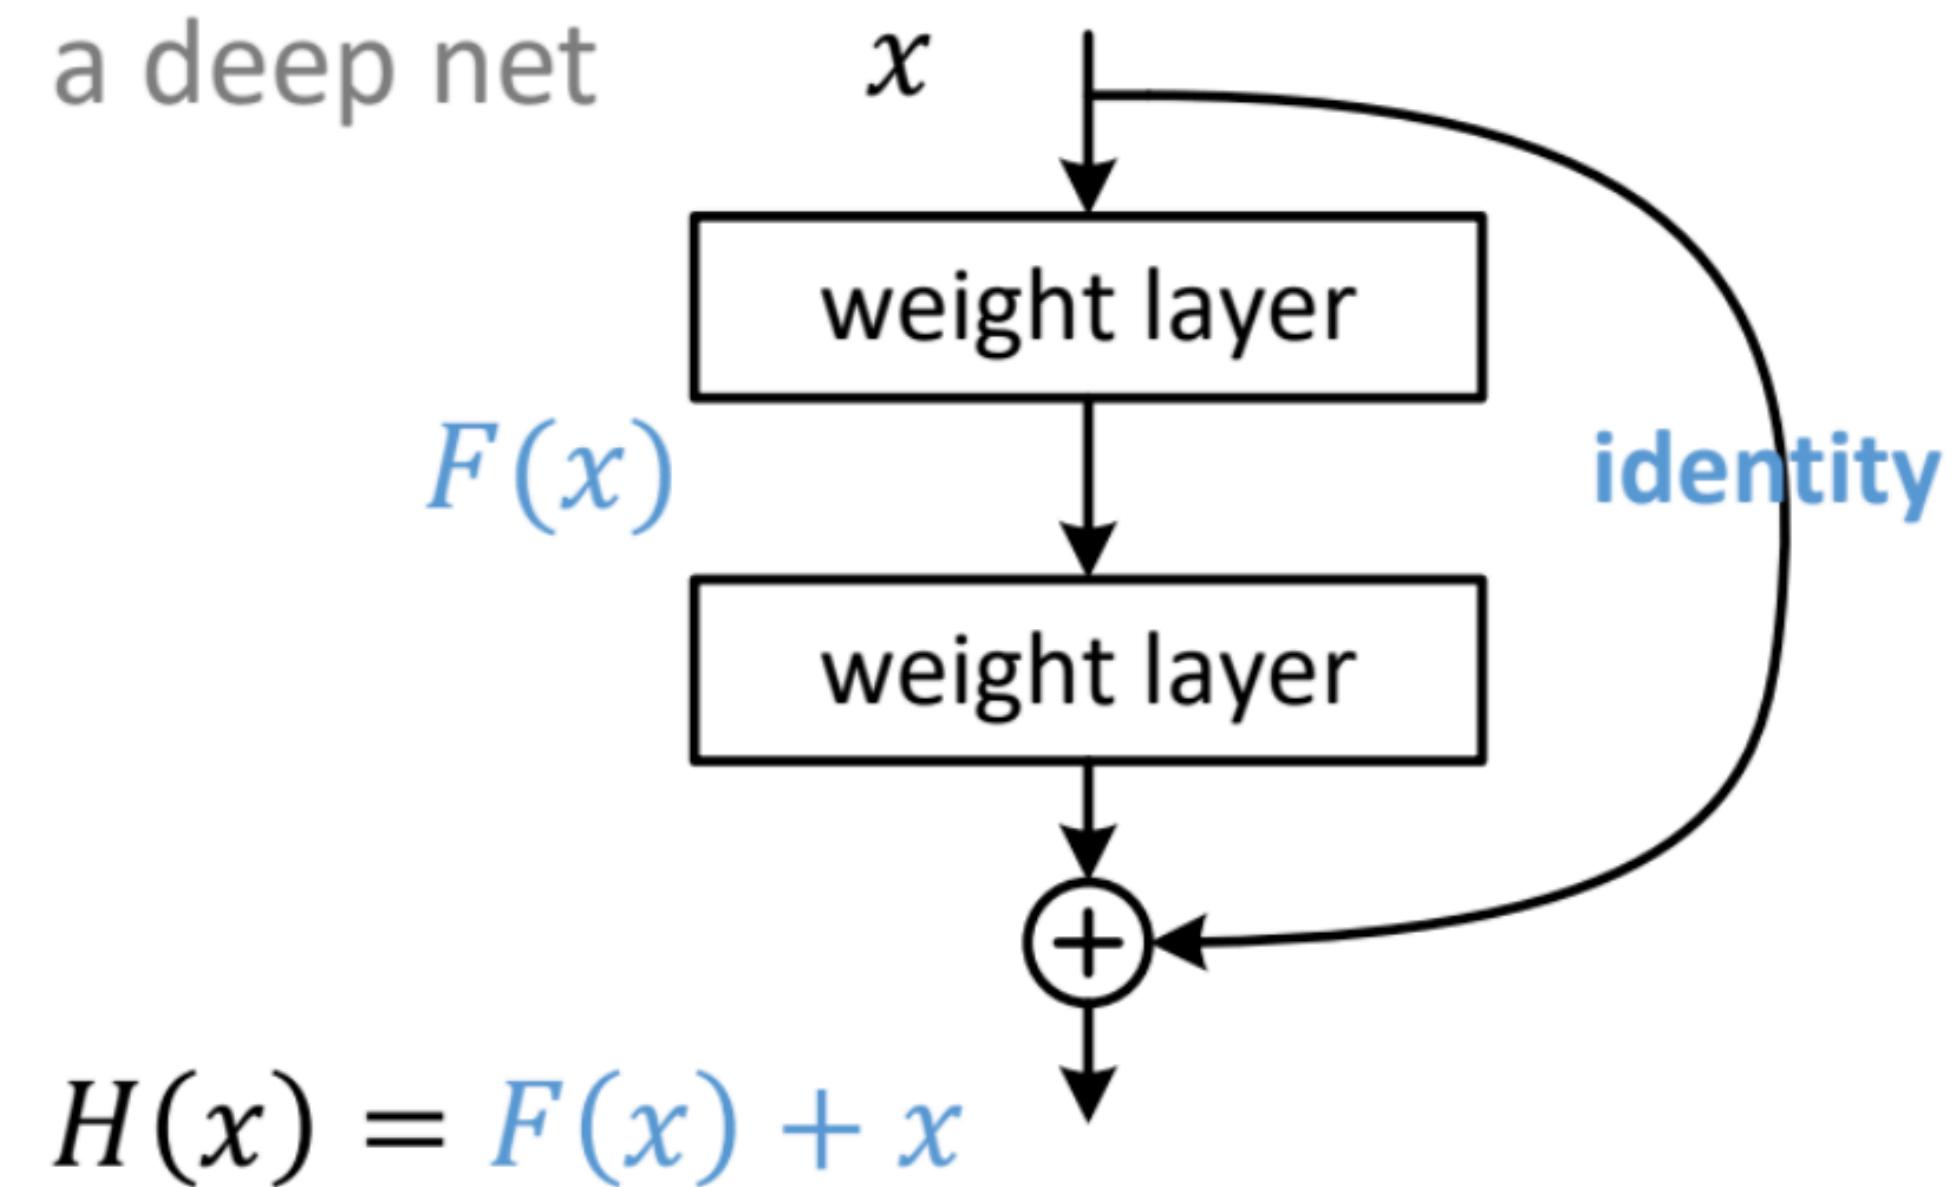
- Good init + norm enable training deeper models
- Simply stacking more layers?

- Degrade after  $\sim 20$  layers
- Not overfitting
- Difficult to train



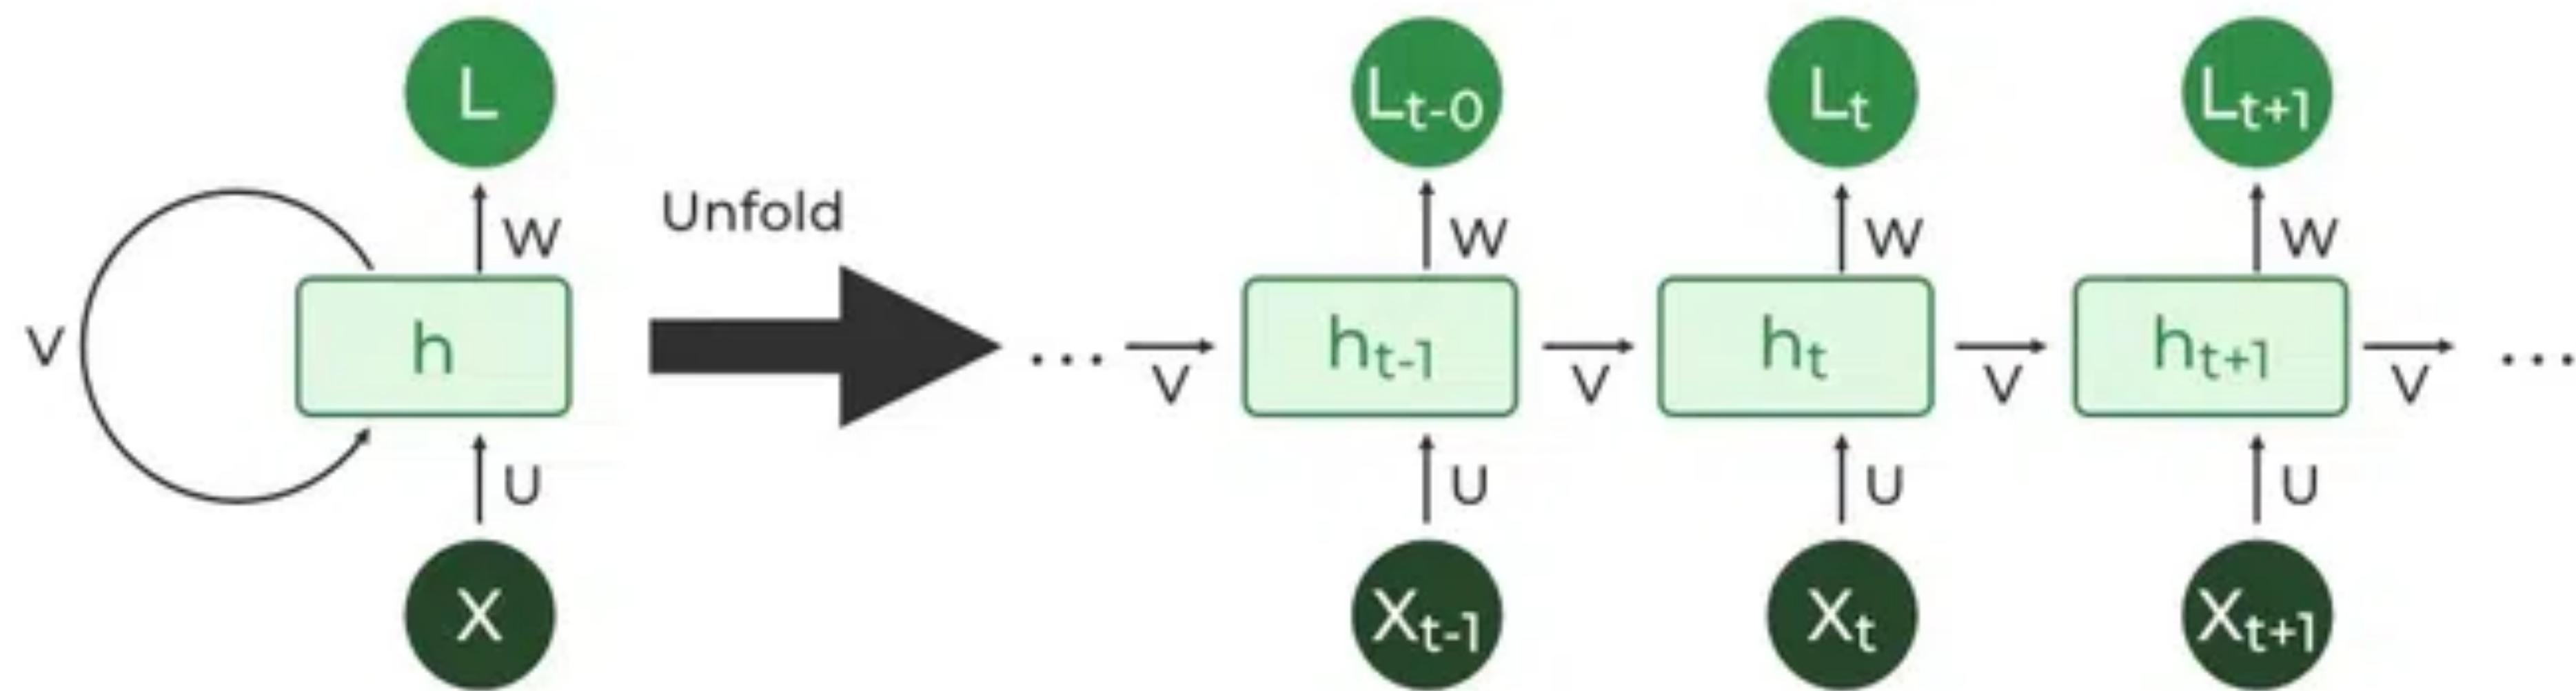
# Deep Residual Learning

a subnet in  
a deep net



**MLP network is hard to handle  
sequence data with varying length**

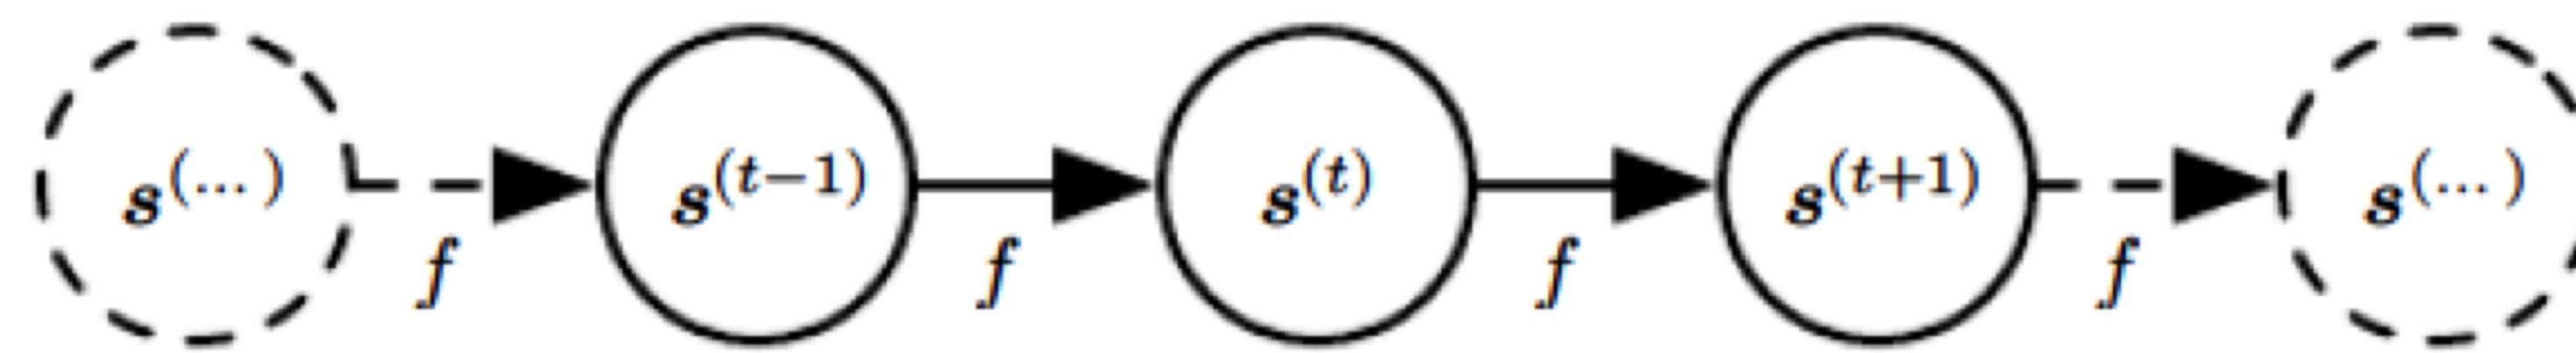
# Recurrent Neural Networks (RNNs)



# Recurrent Neural Networks

- Dates back to (Rumelhart *et al.*, 1986)
- A family of neural networks for handling sequential data, which involves variable length inputs or outputs
- Especially, for natural language processing (NLP)

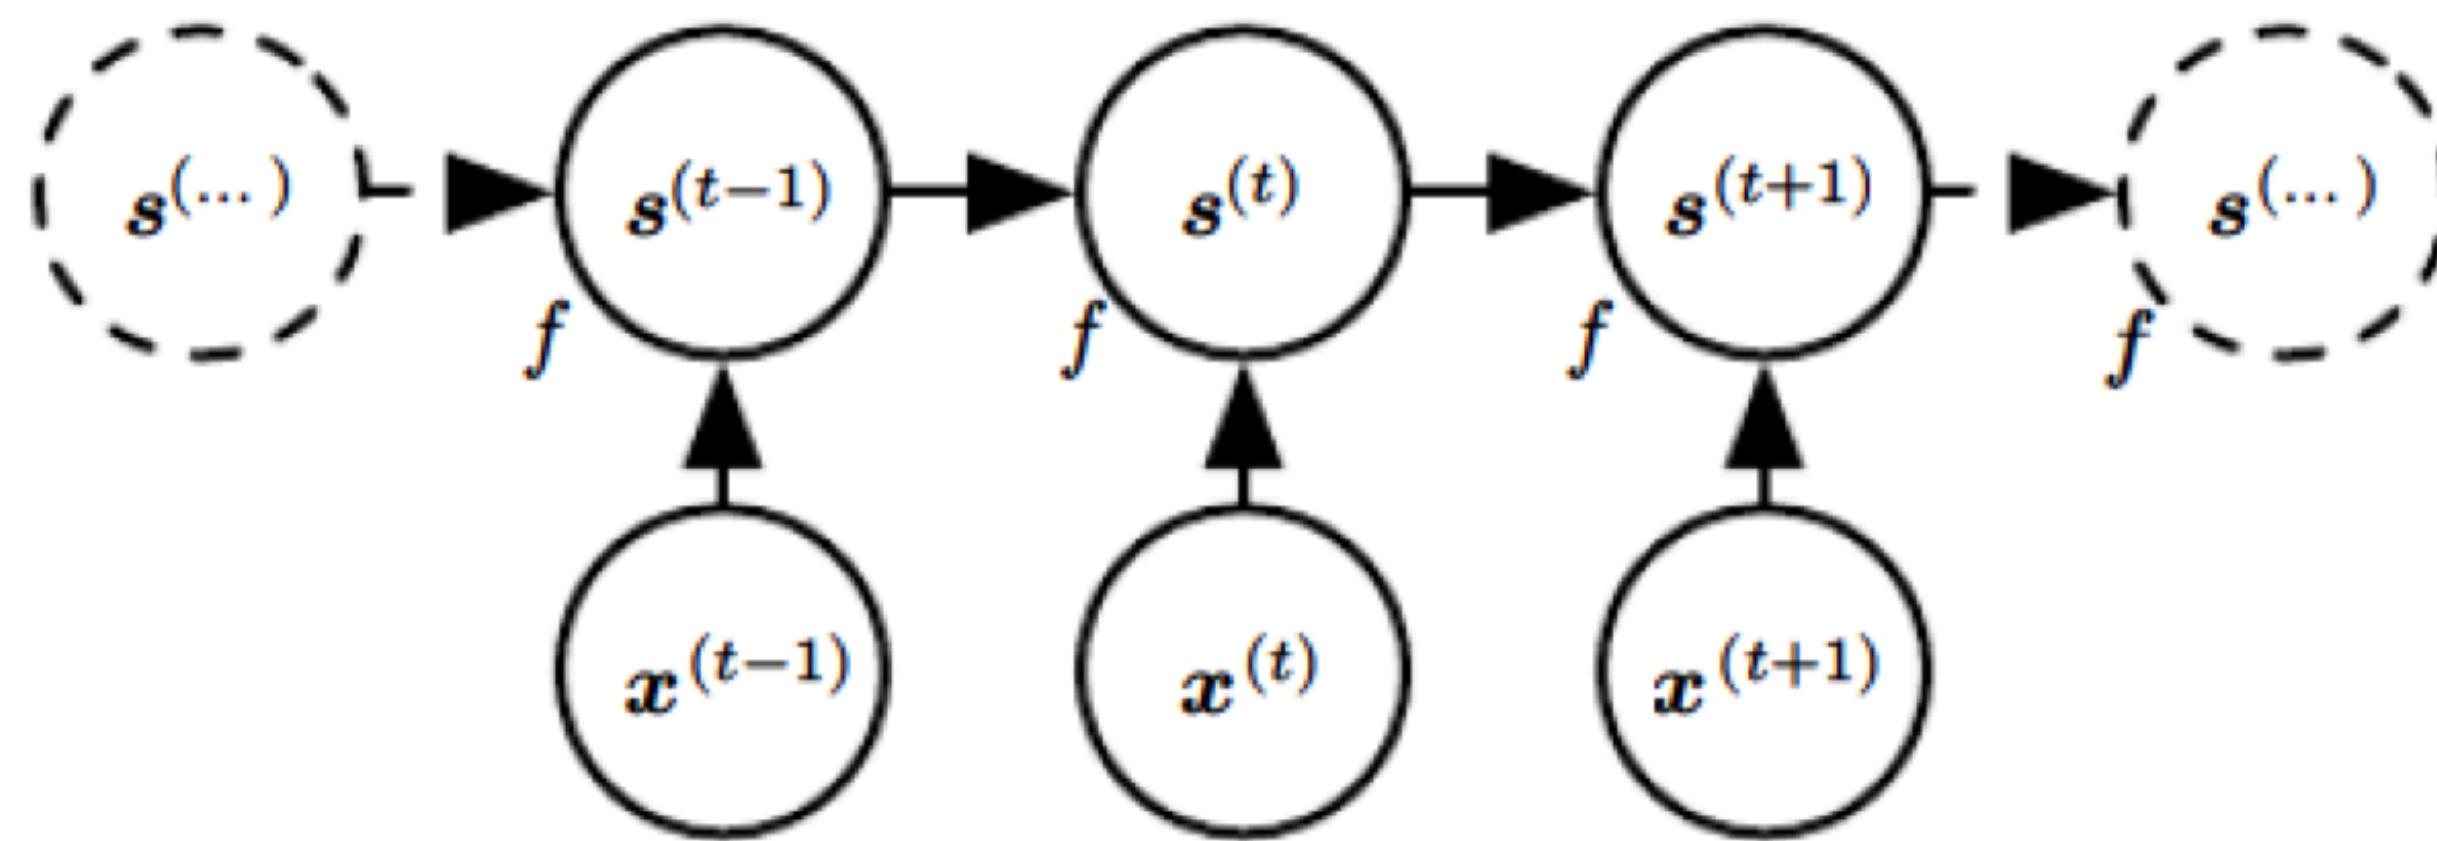
# Computation Graph



$$s^{(t+1)} = f(s^{(t)}; \theta)$$

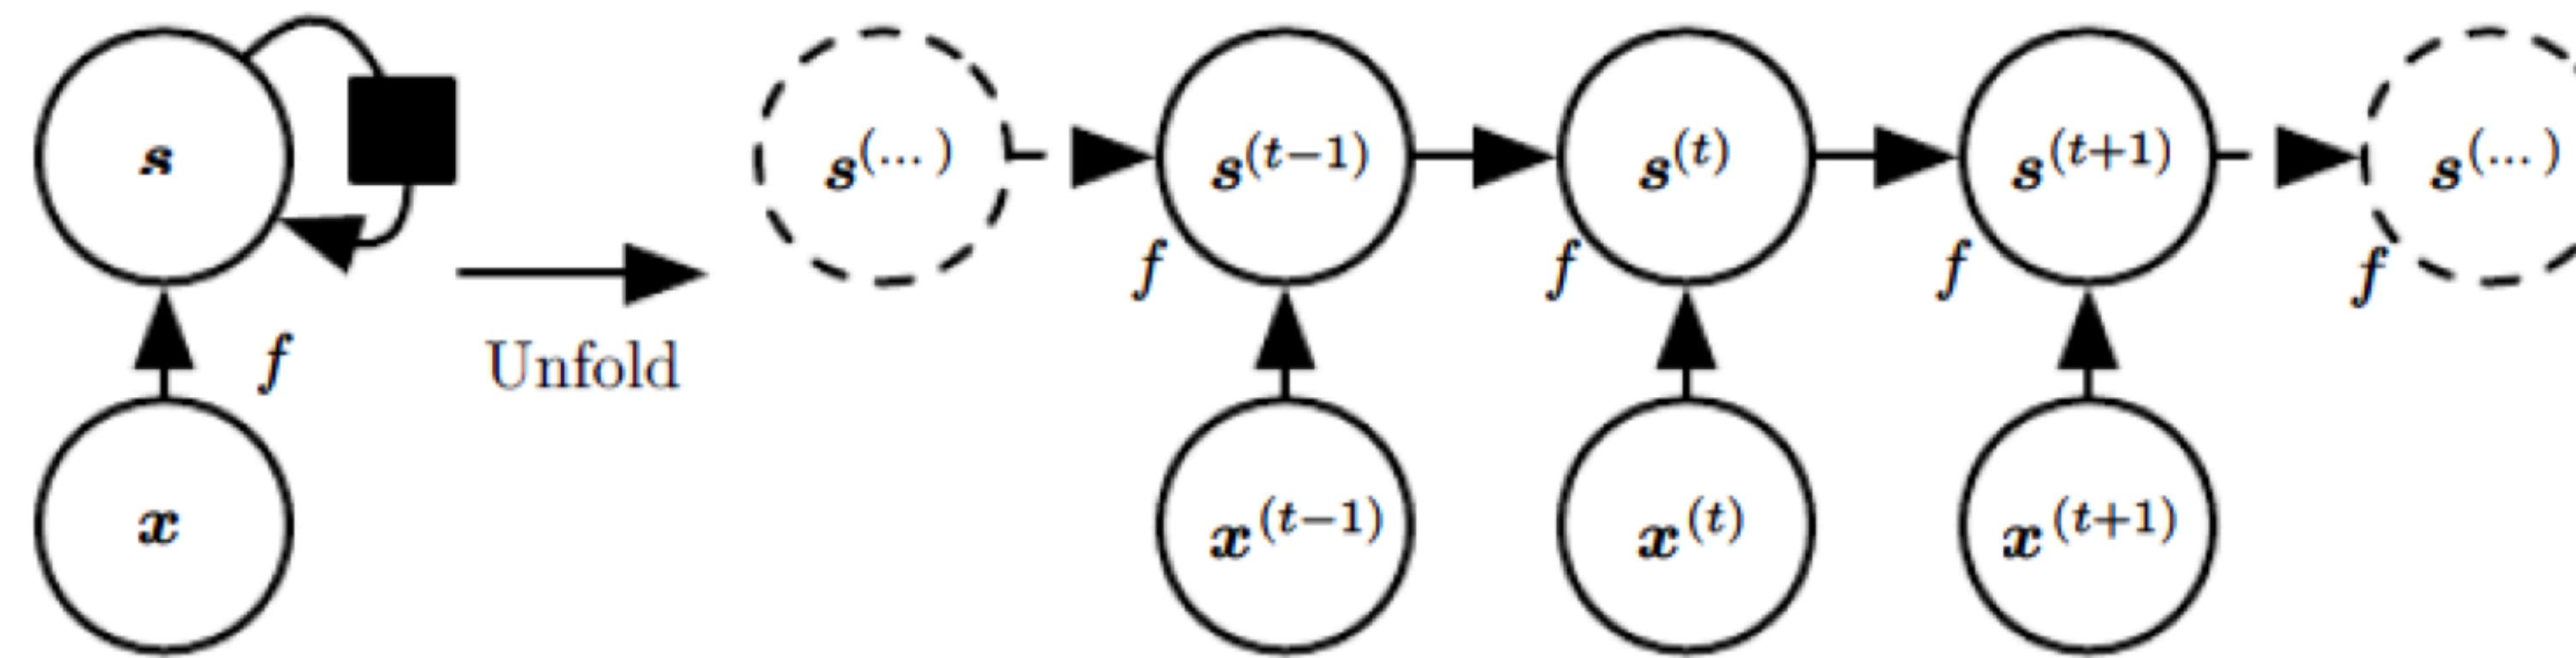
Figure from *Deep Learning*,  
Goodfellow, Bengio and Courville

# Computation Graph



$$s^{(t+1)} = f(s^{(t)}, x^{(t+1)}; \theta)$$

# Compact view



$$s^{(t+1)} = f(s^{(t)}, x^{(t+1)}; \theta)$$

Key: the same  $f$  and  $\theta$   
for all time steps

# Recurrent Neural Networks

- Use **the same** computational function and parameters across different time steps of the sequence
- Each time step: takes the input entry **and the previous hidden state** to compute the output entry
- Loss: typically computed every time step

# Recurrent Neural Networks

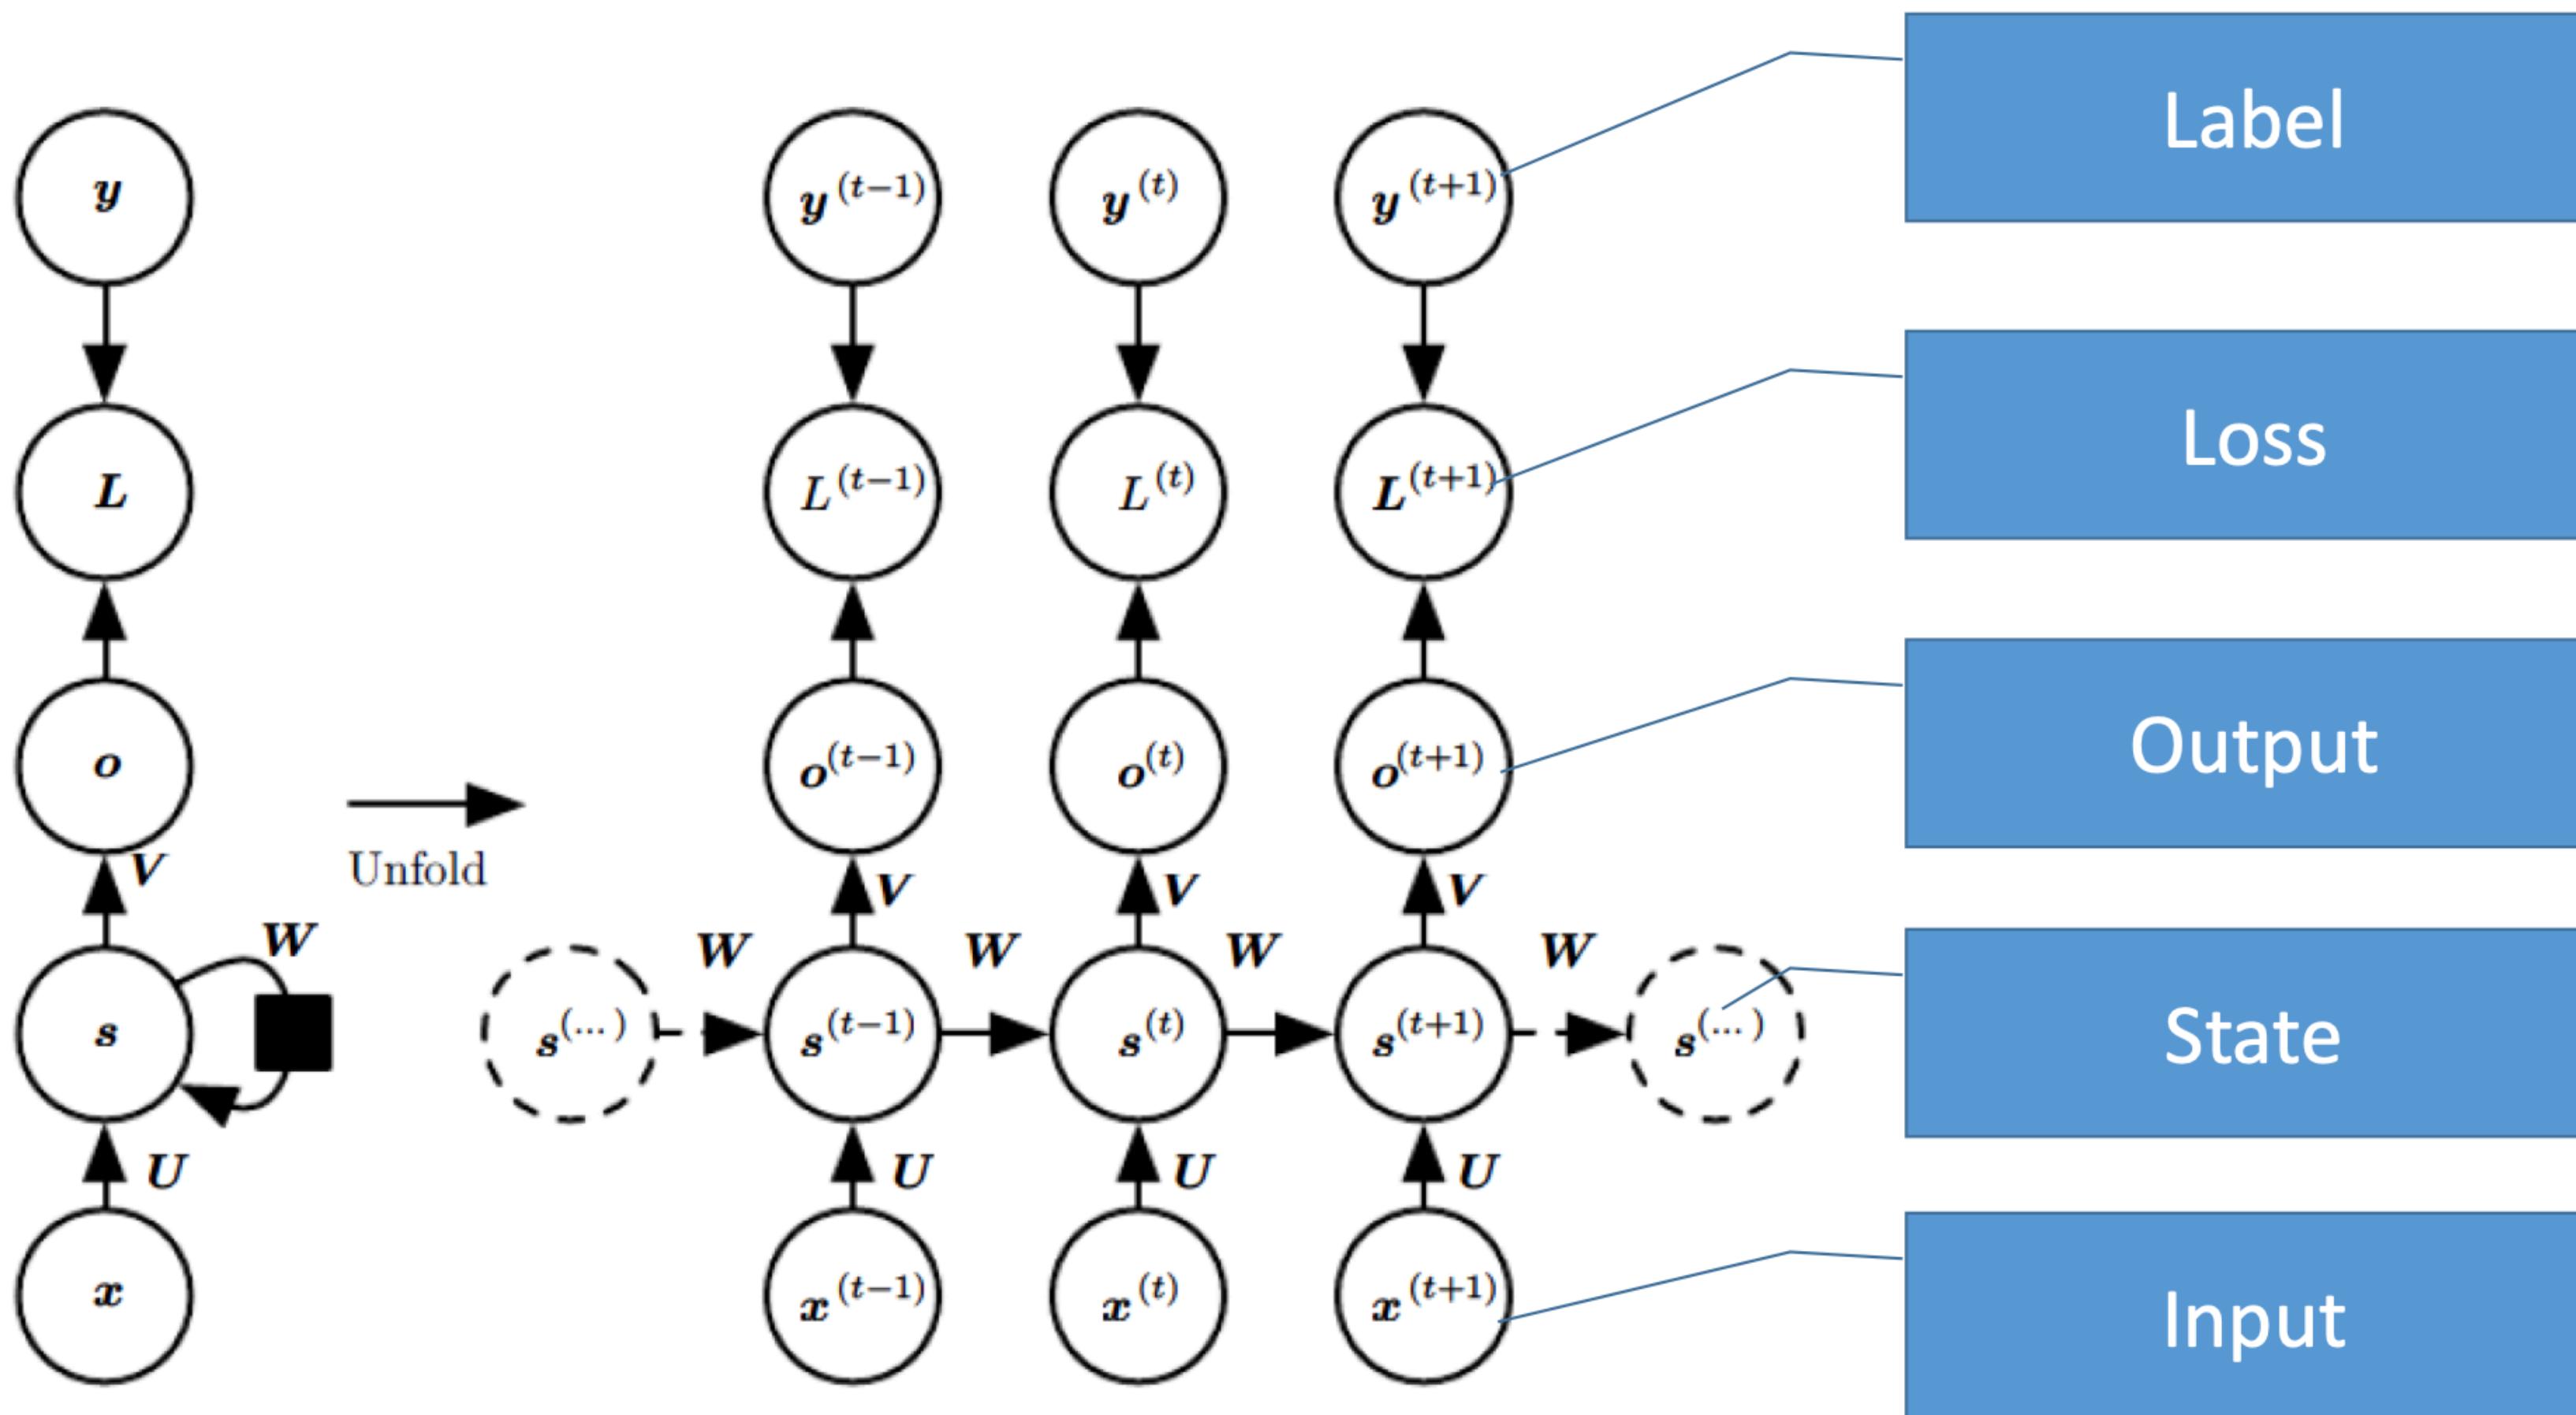


Figure from *Deep Learning*, by Goodfellow, Bengio and Courville

# Recurrent Neural Networks

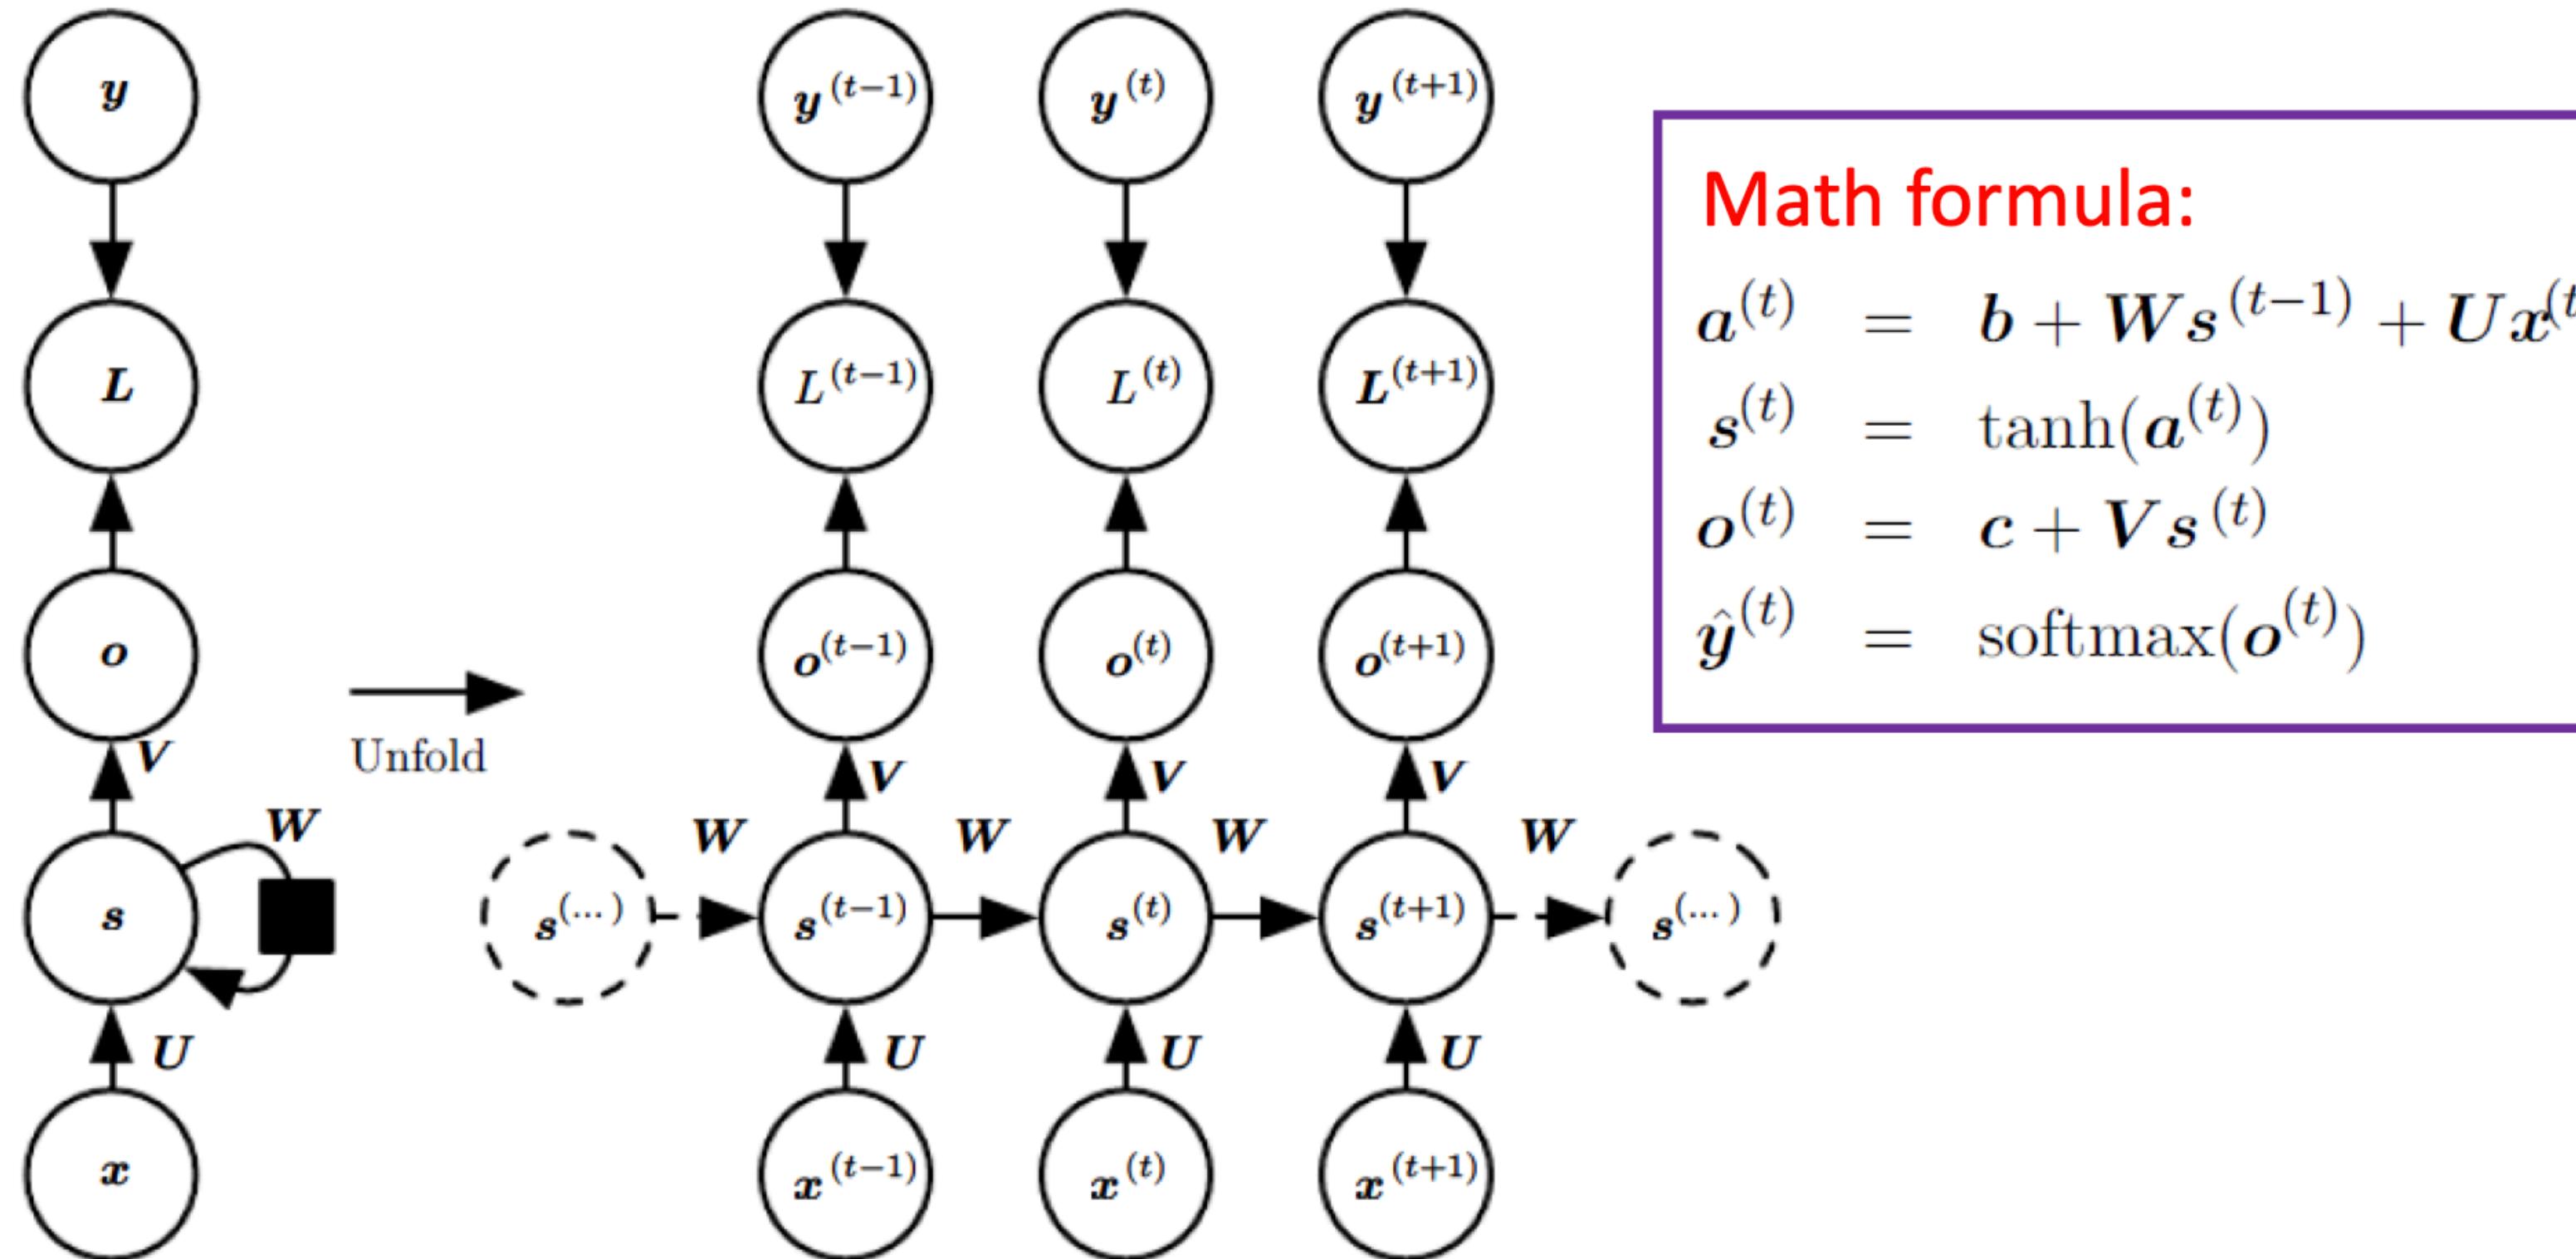
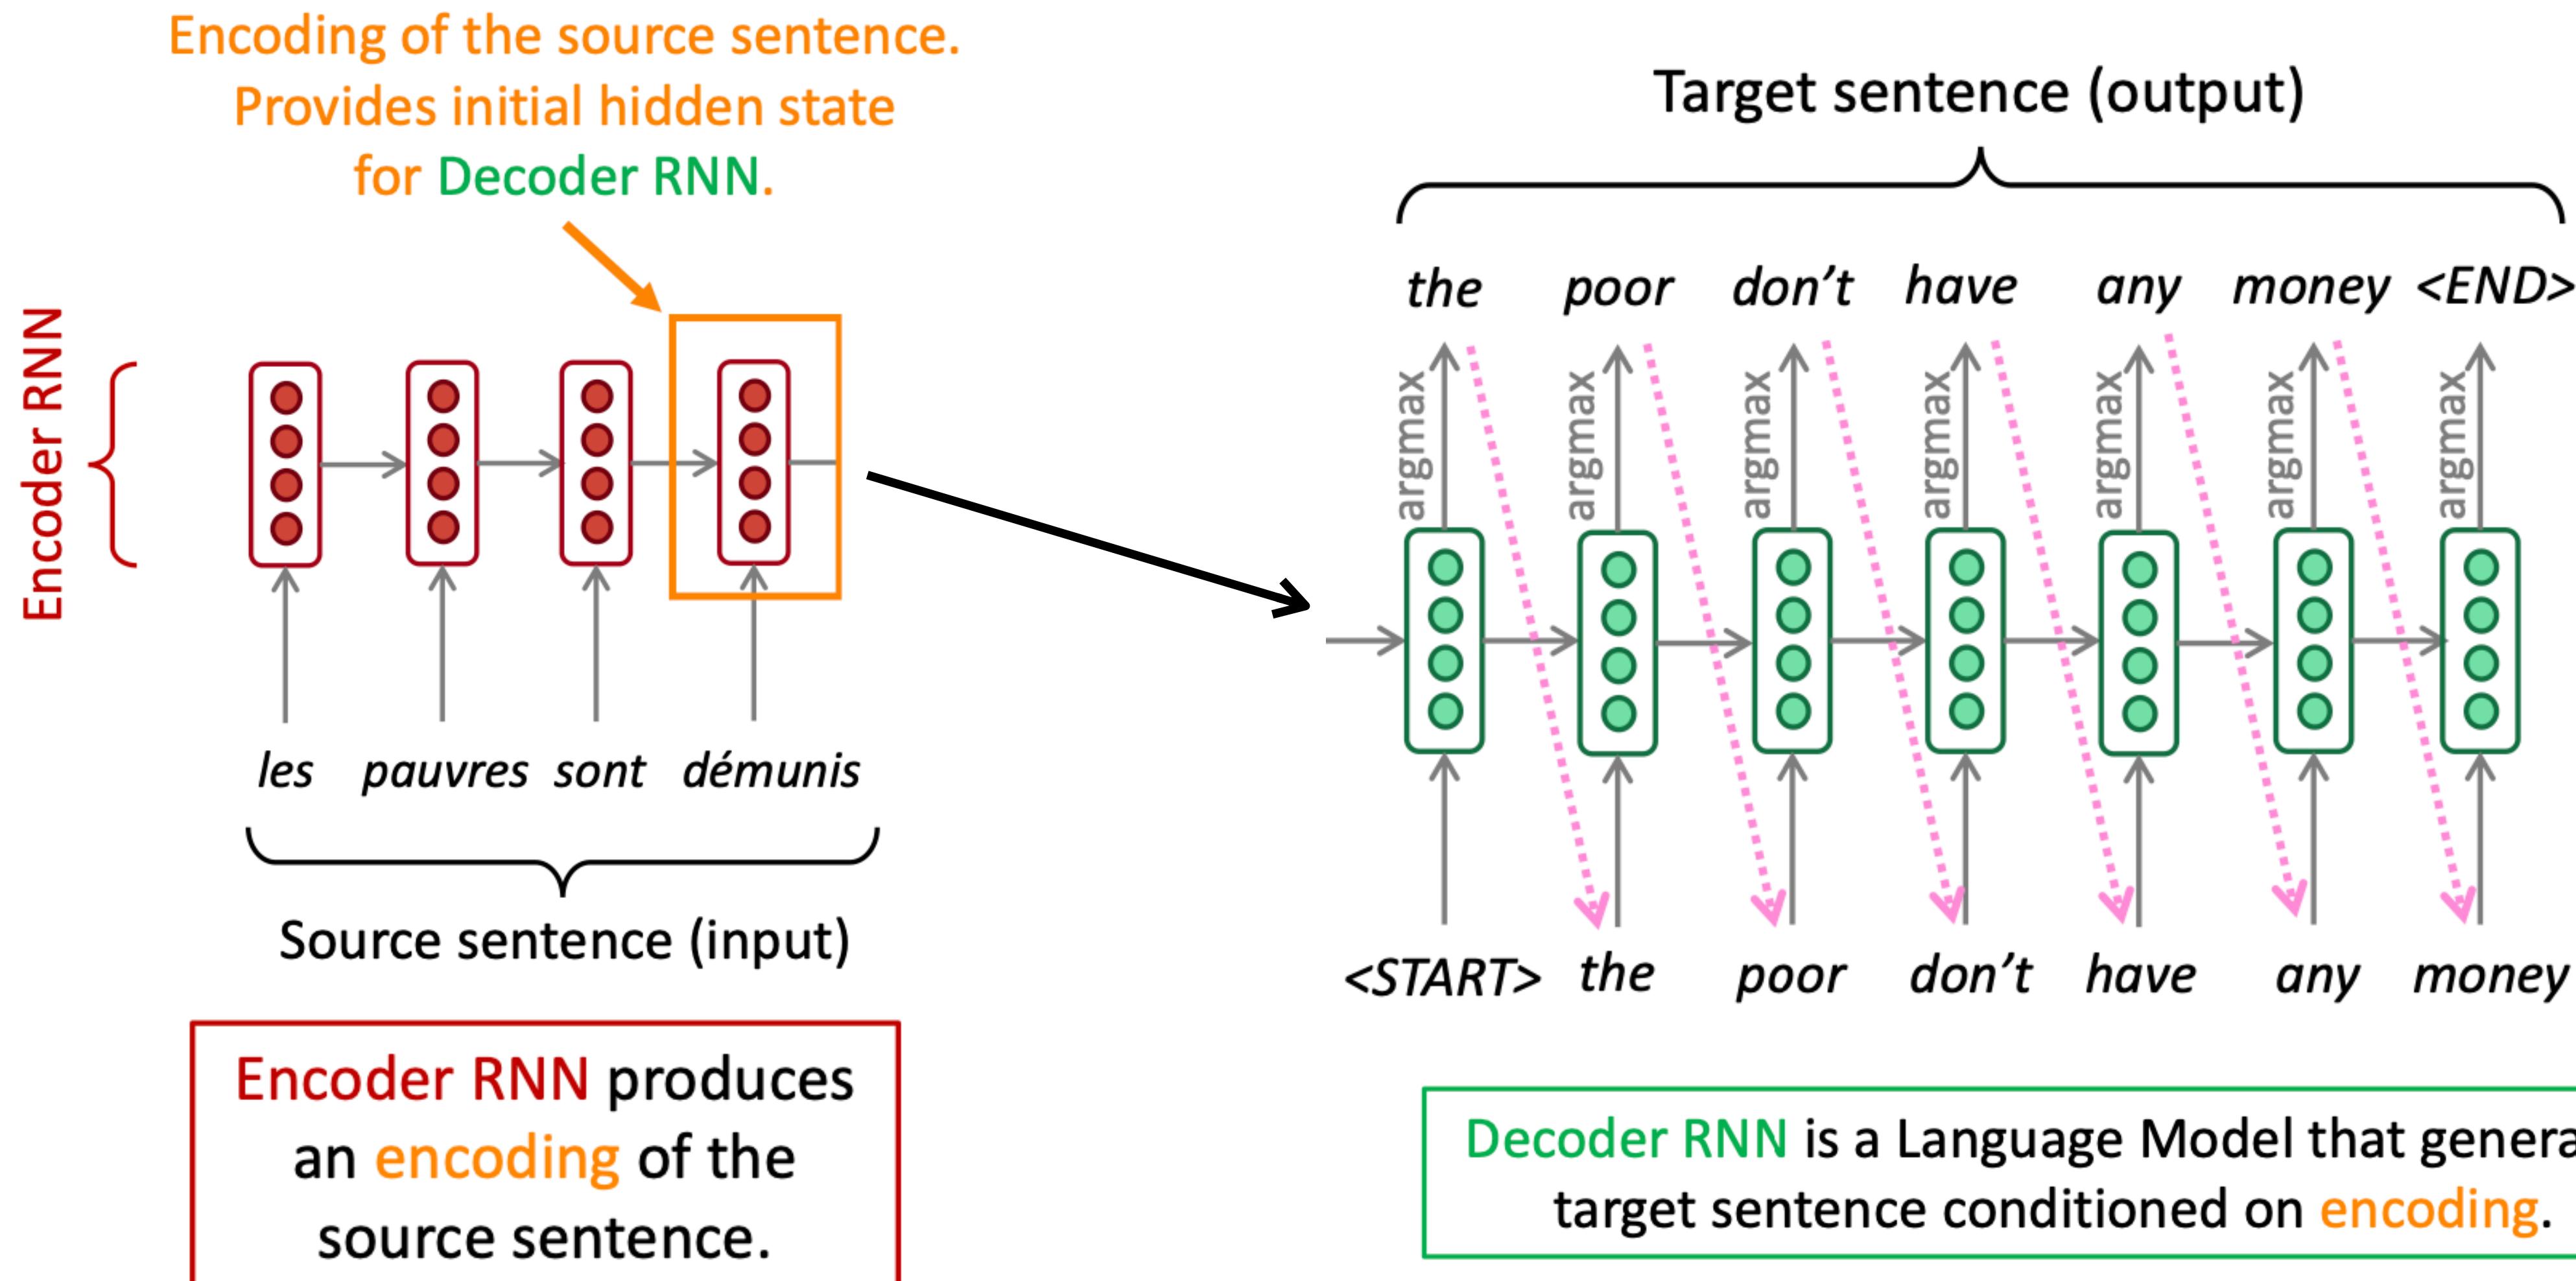


Figure from *Deep Learning*,  
Goodfellow, Bengio and Courville

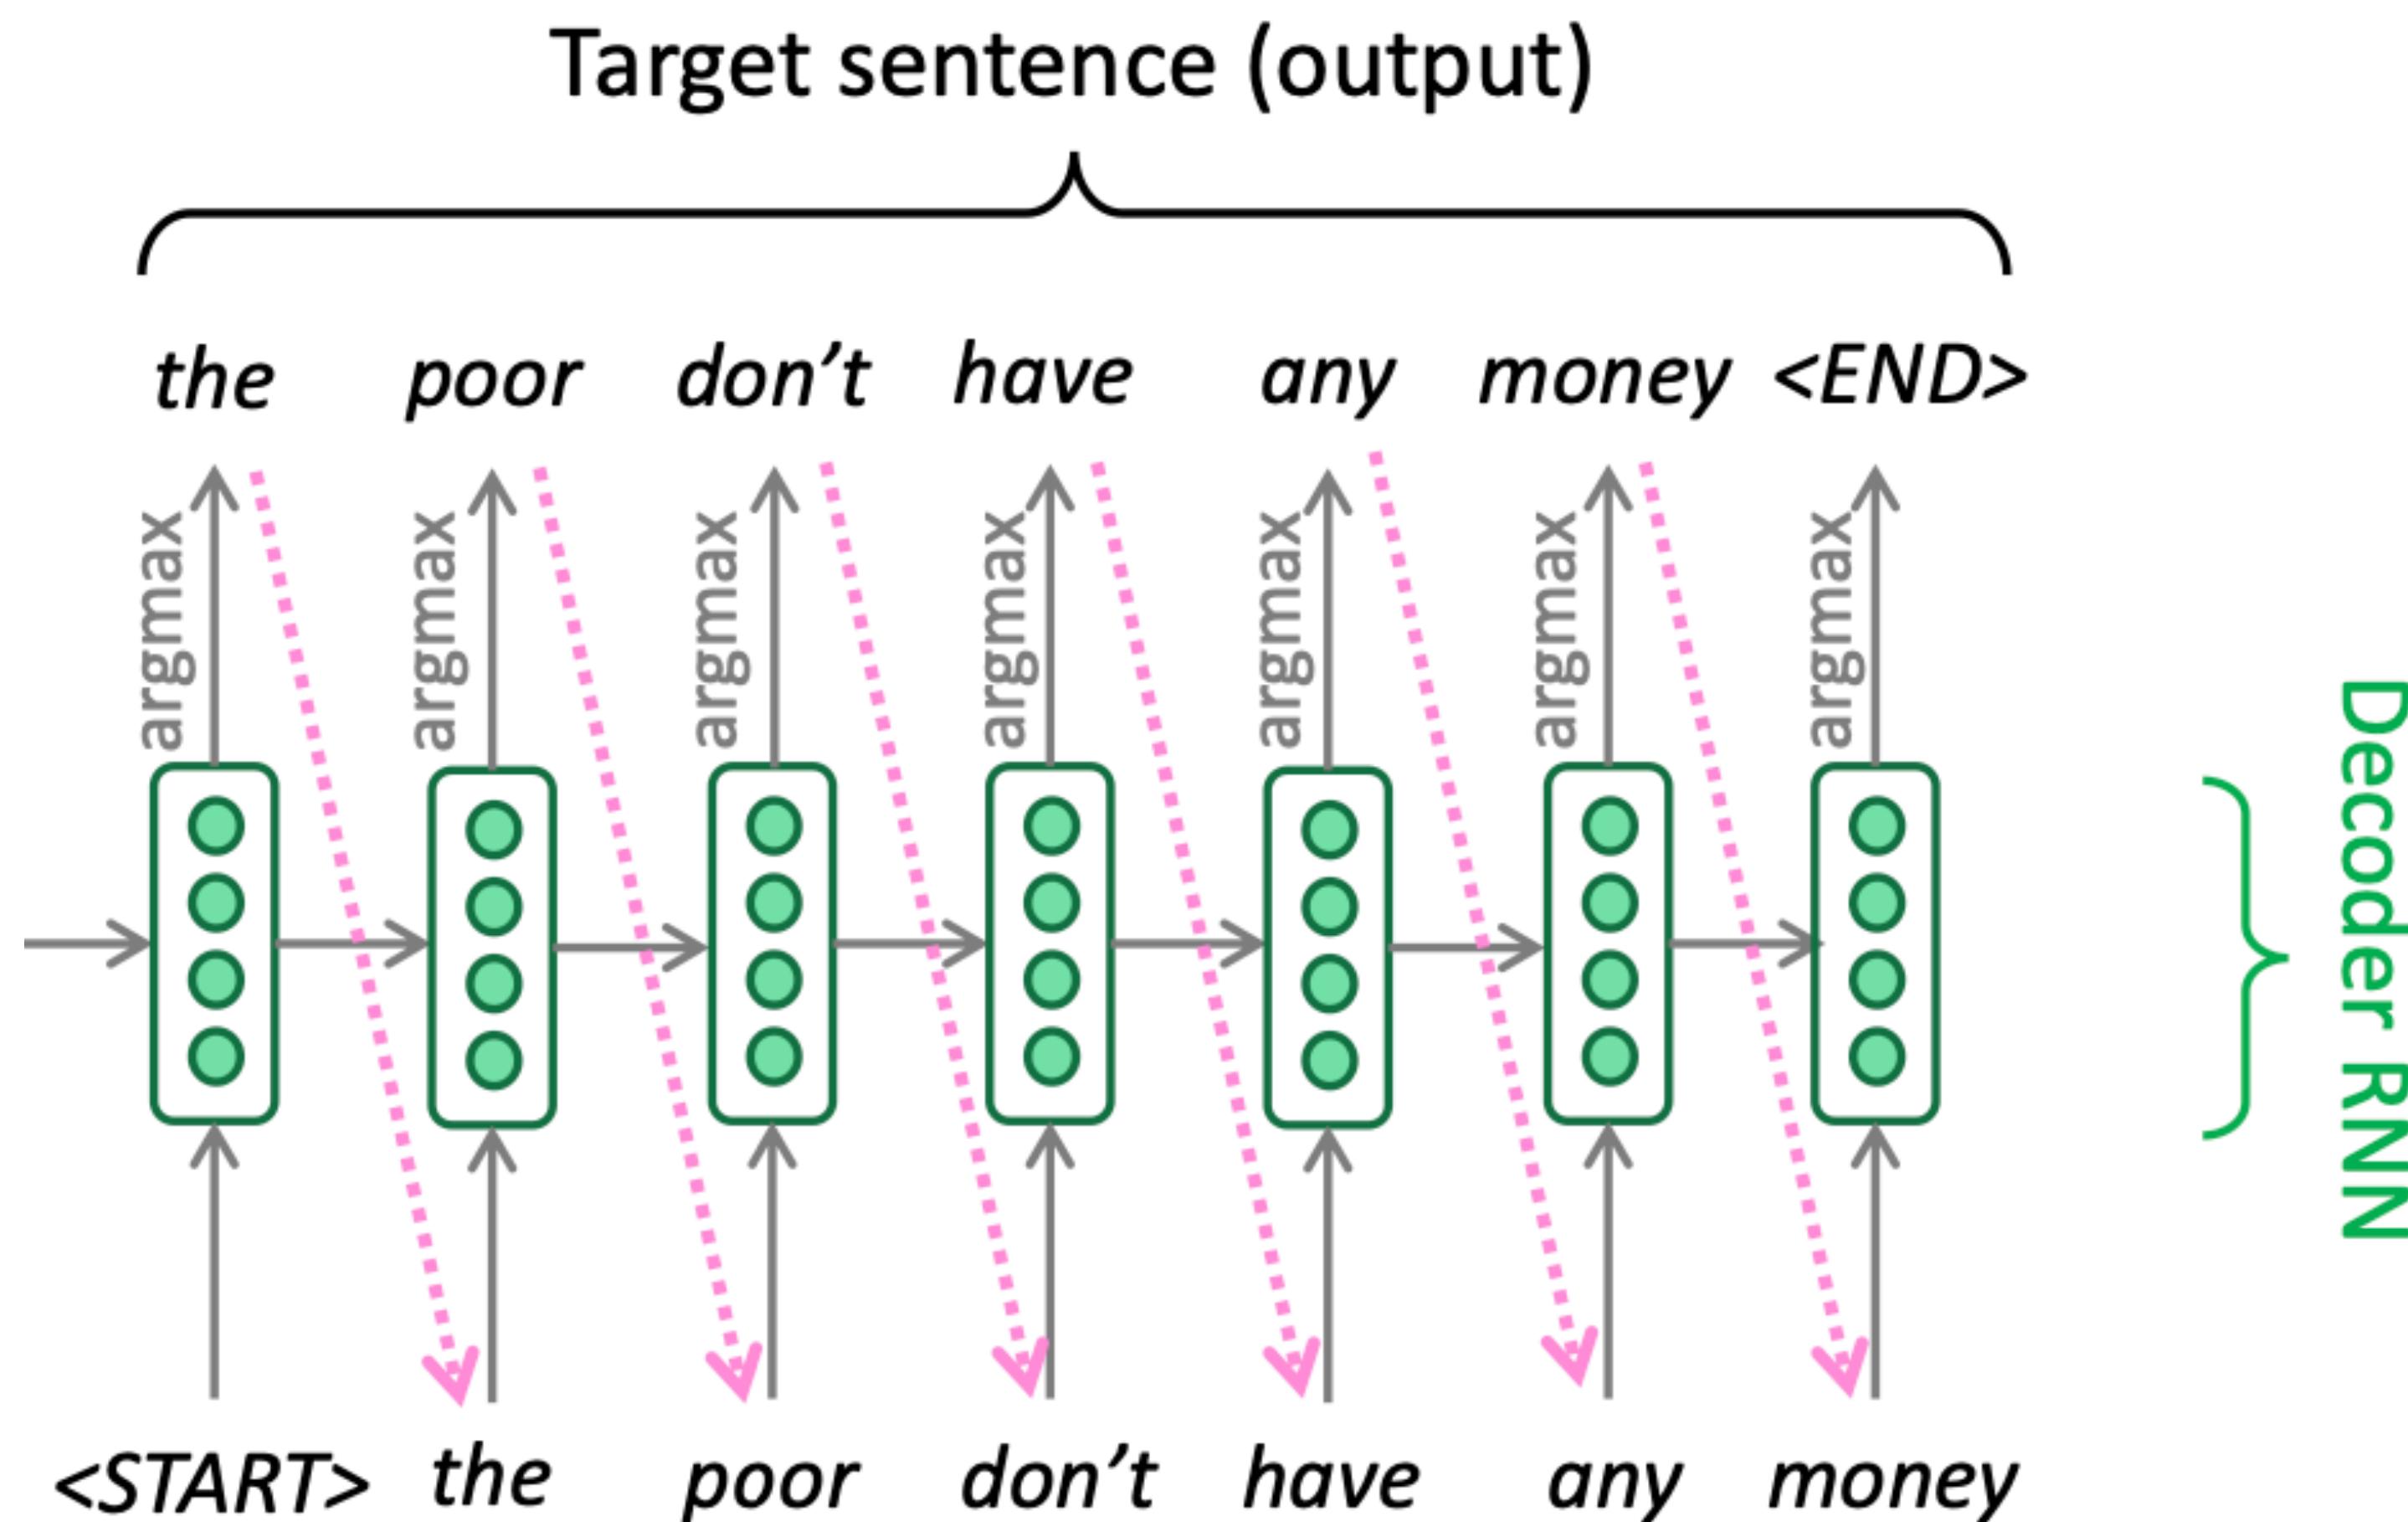
There are many variants of RNNs since the functional form to compute  $s^{(t)}$  can vary, e.g., LSTM

# Sequence-to-Sequence Learning

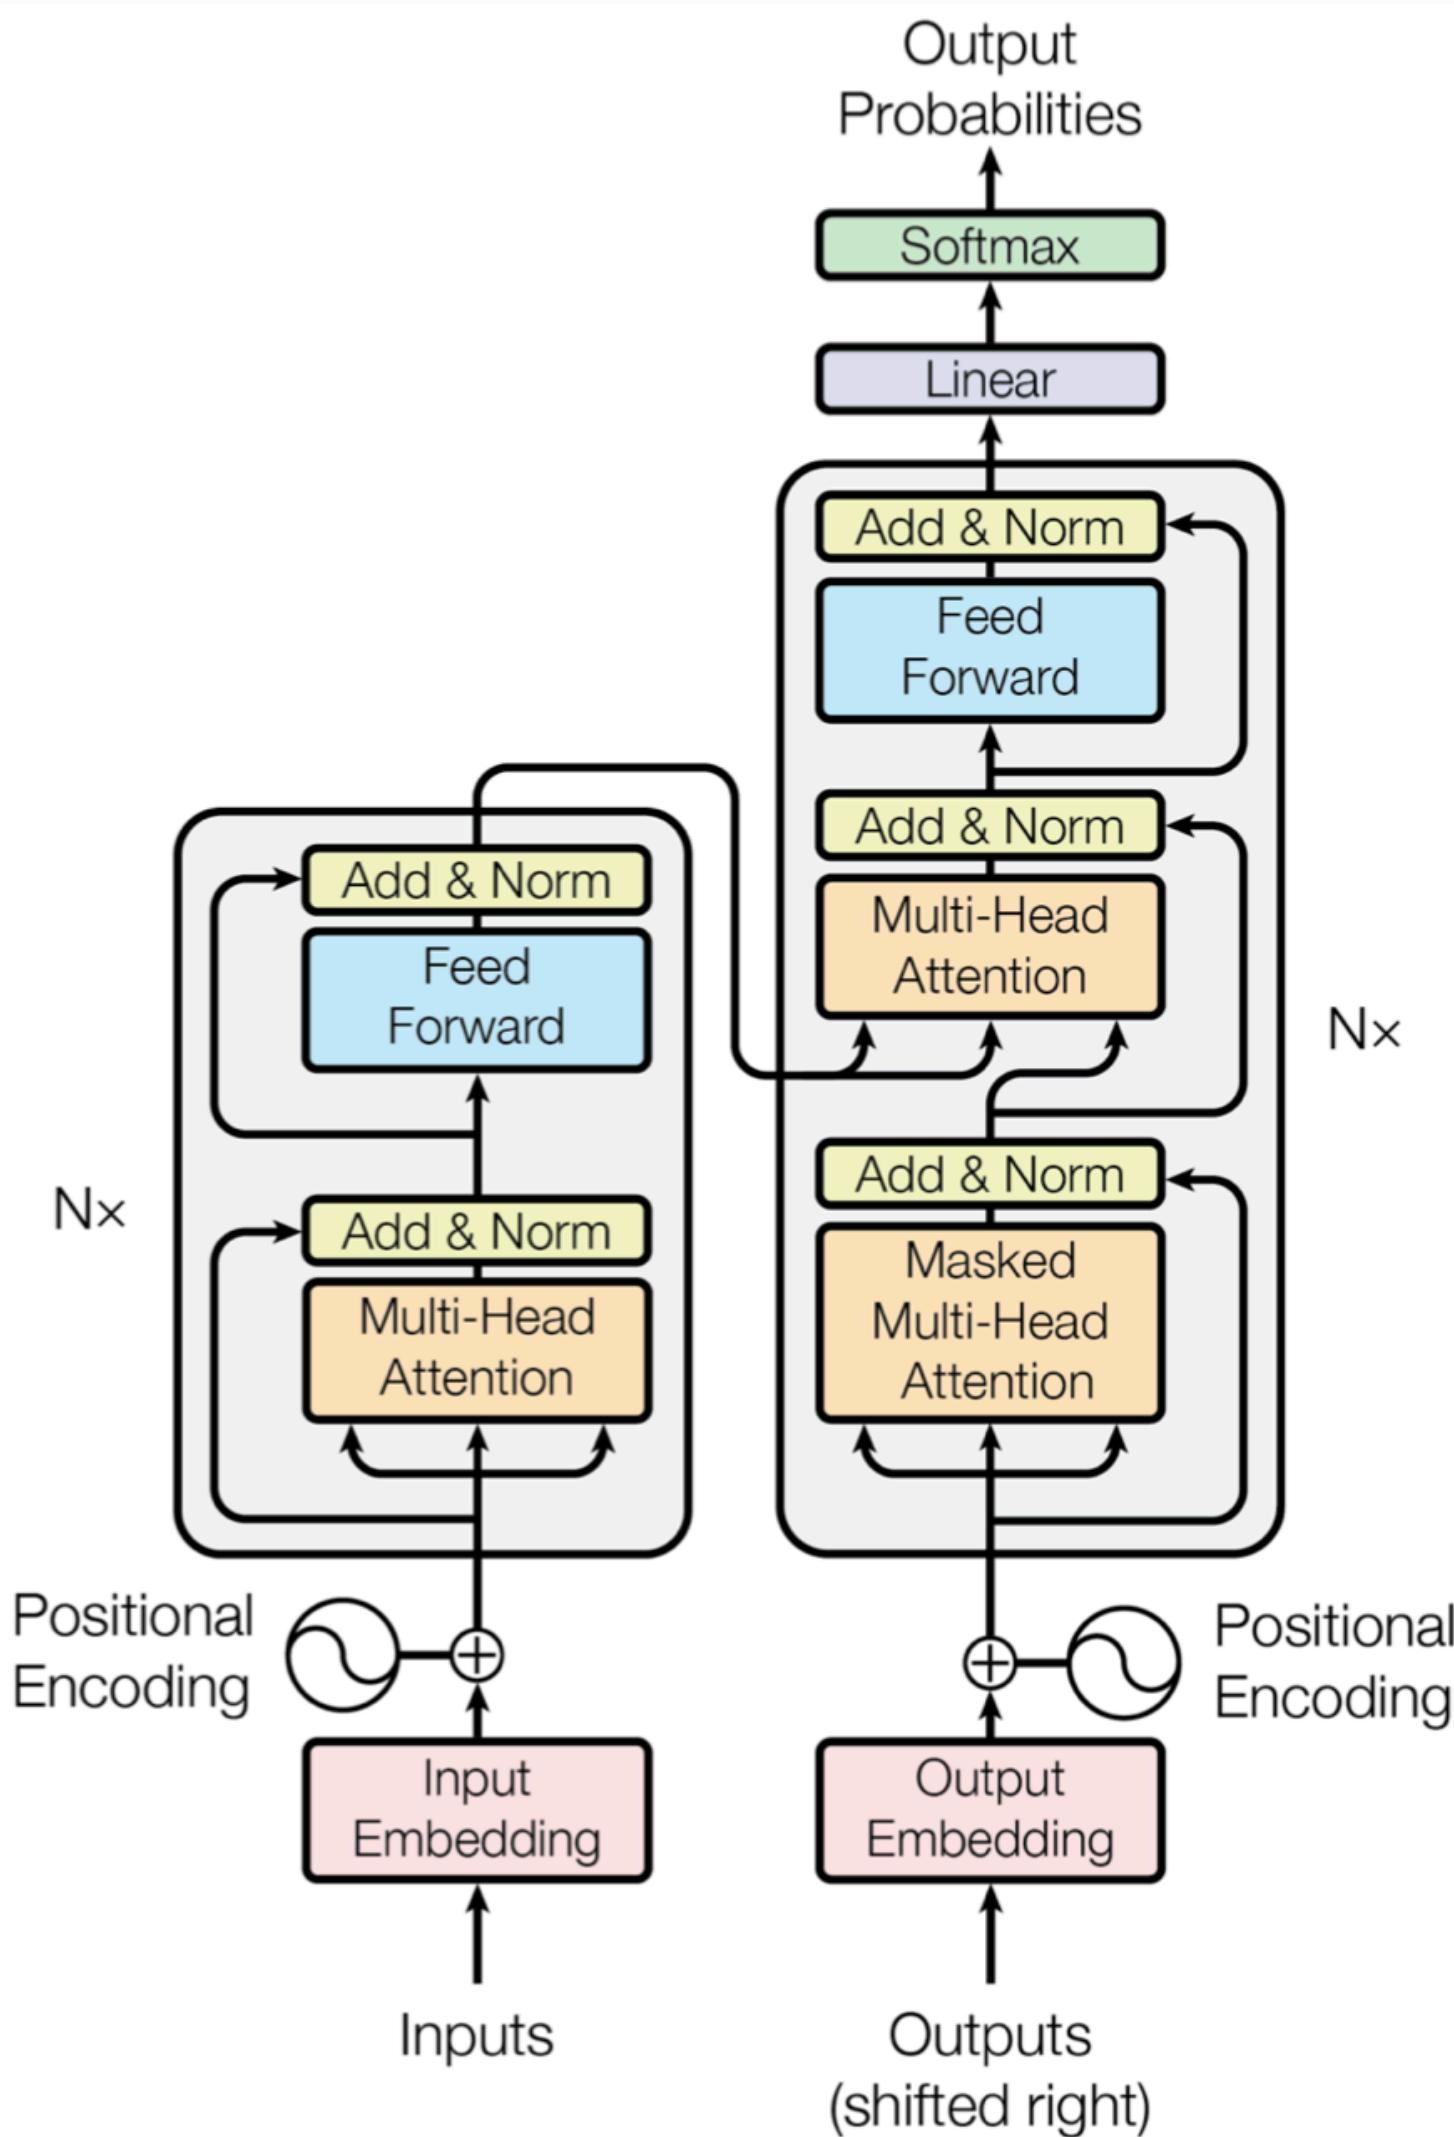
## Example of Neural Machine Translation



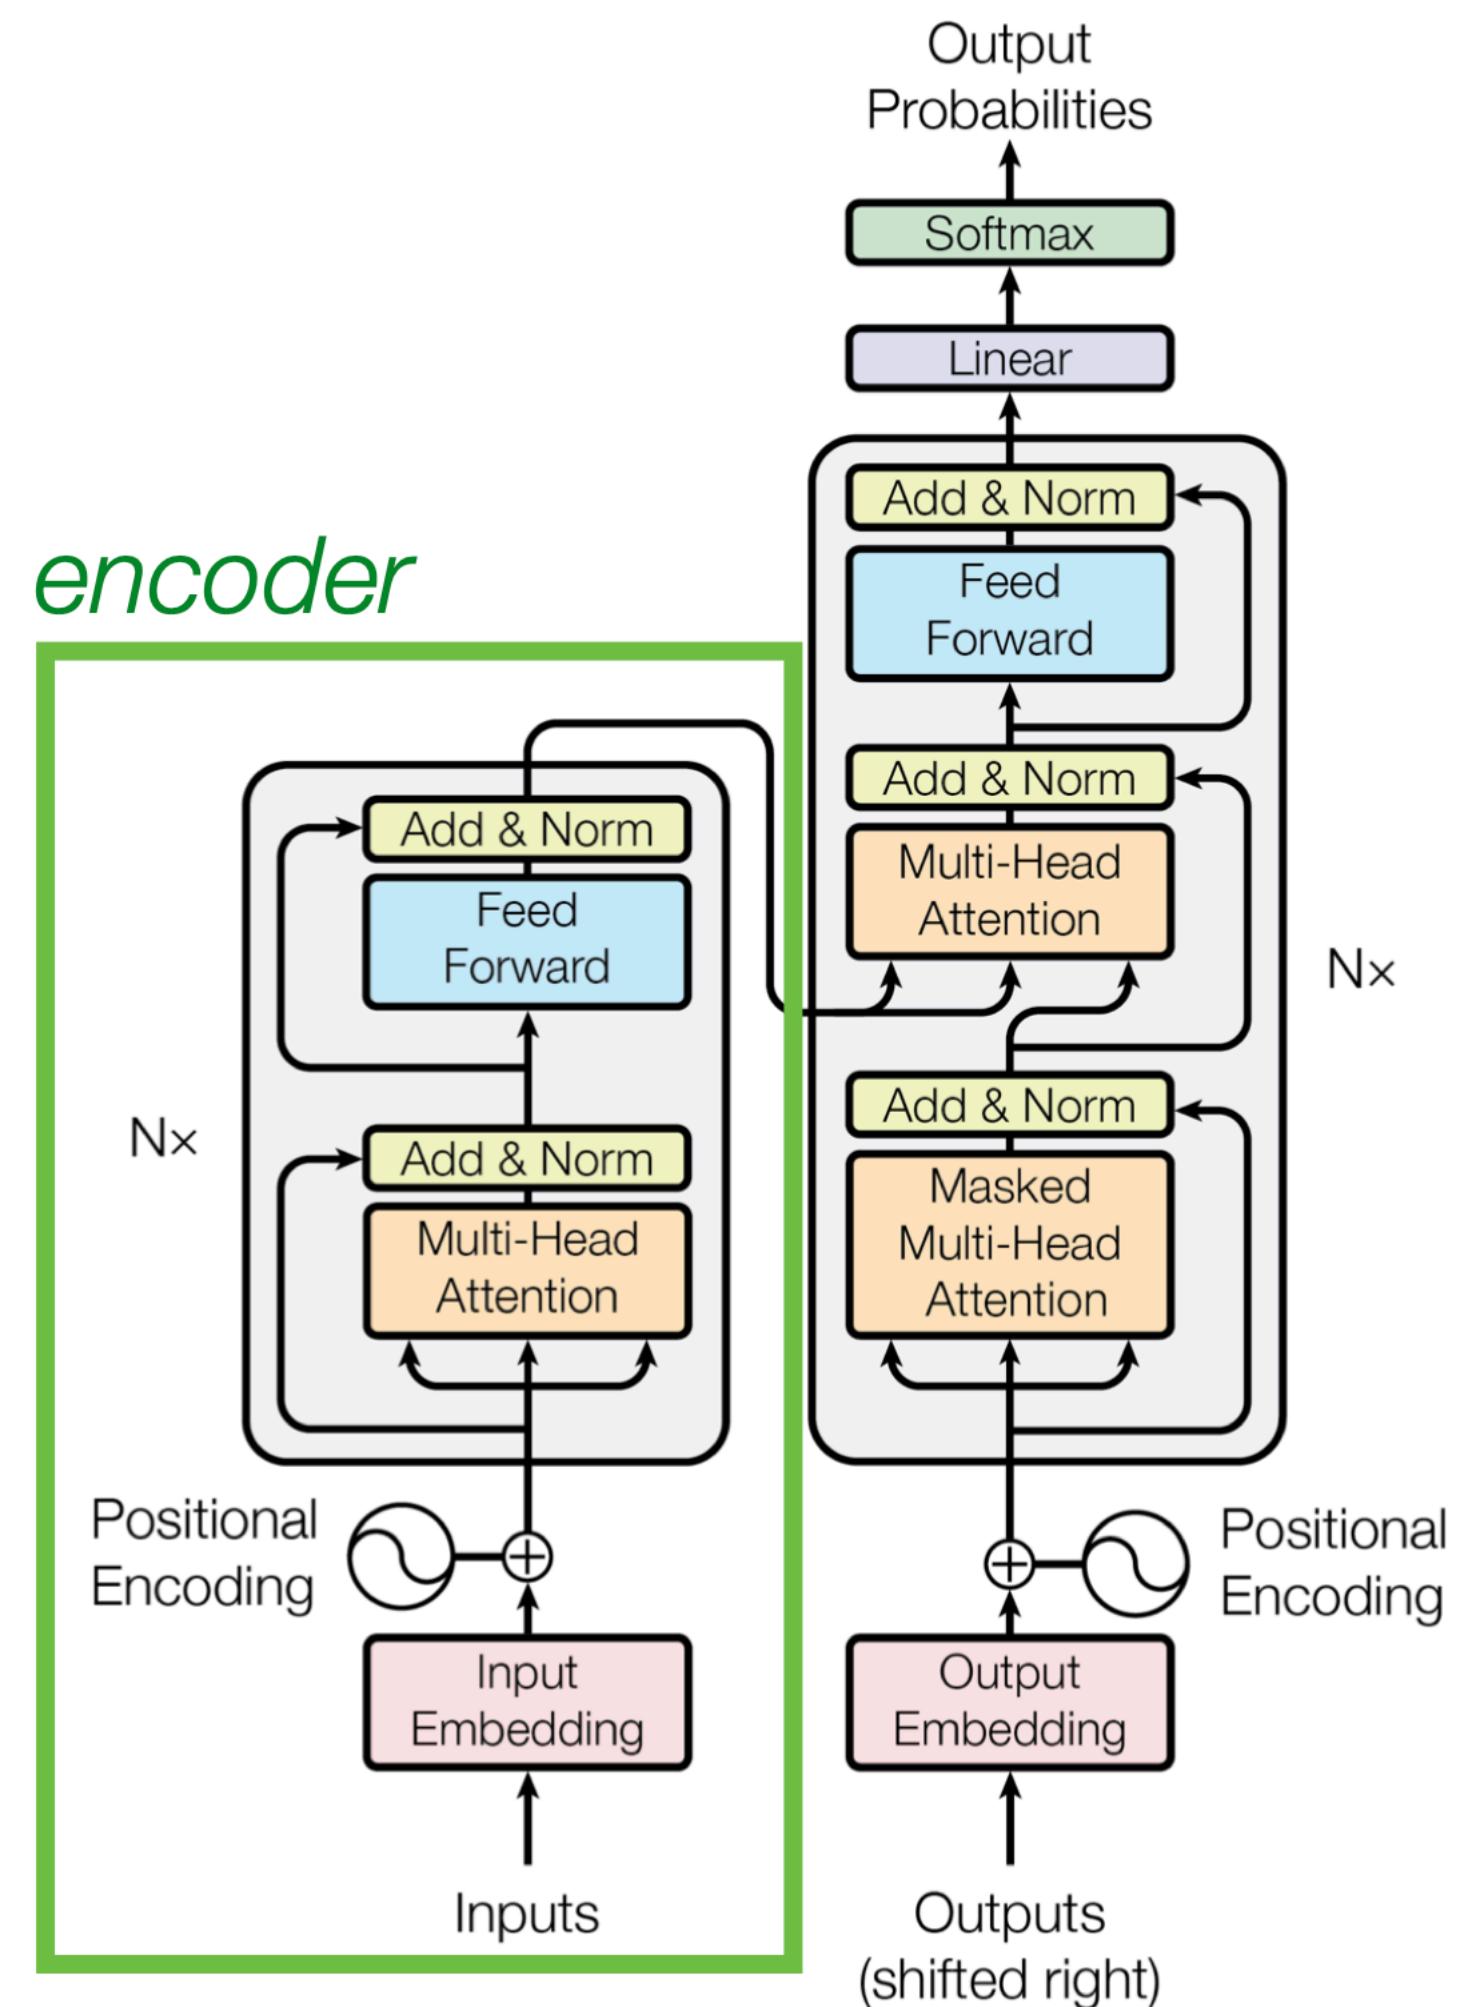
# RNN Language Model



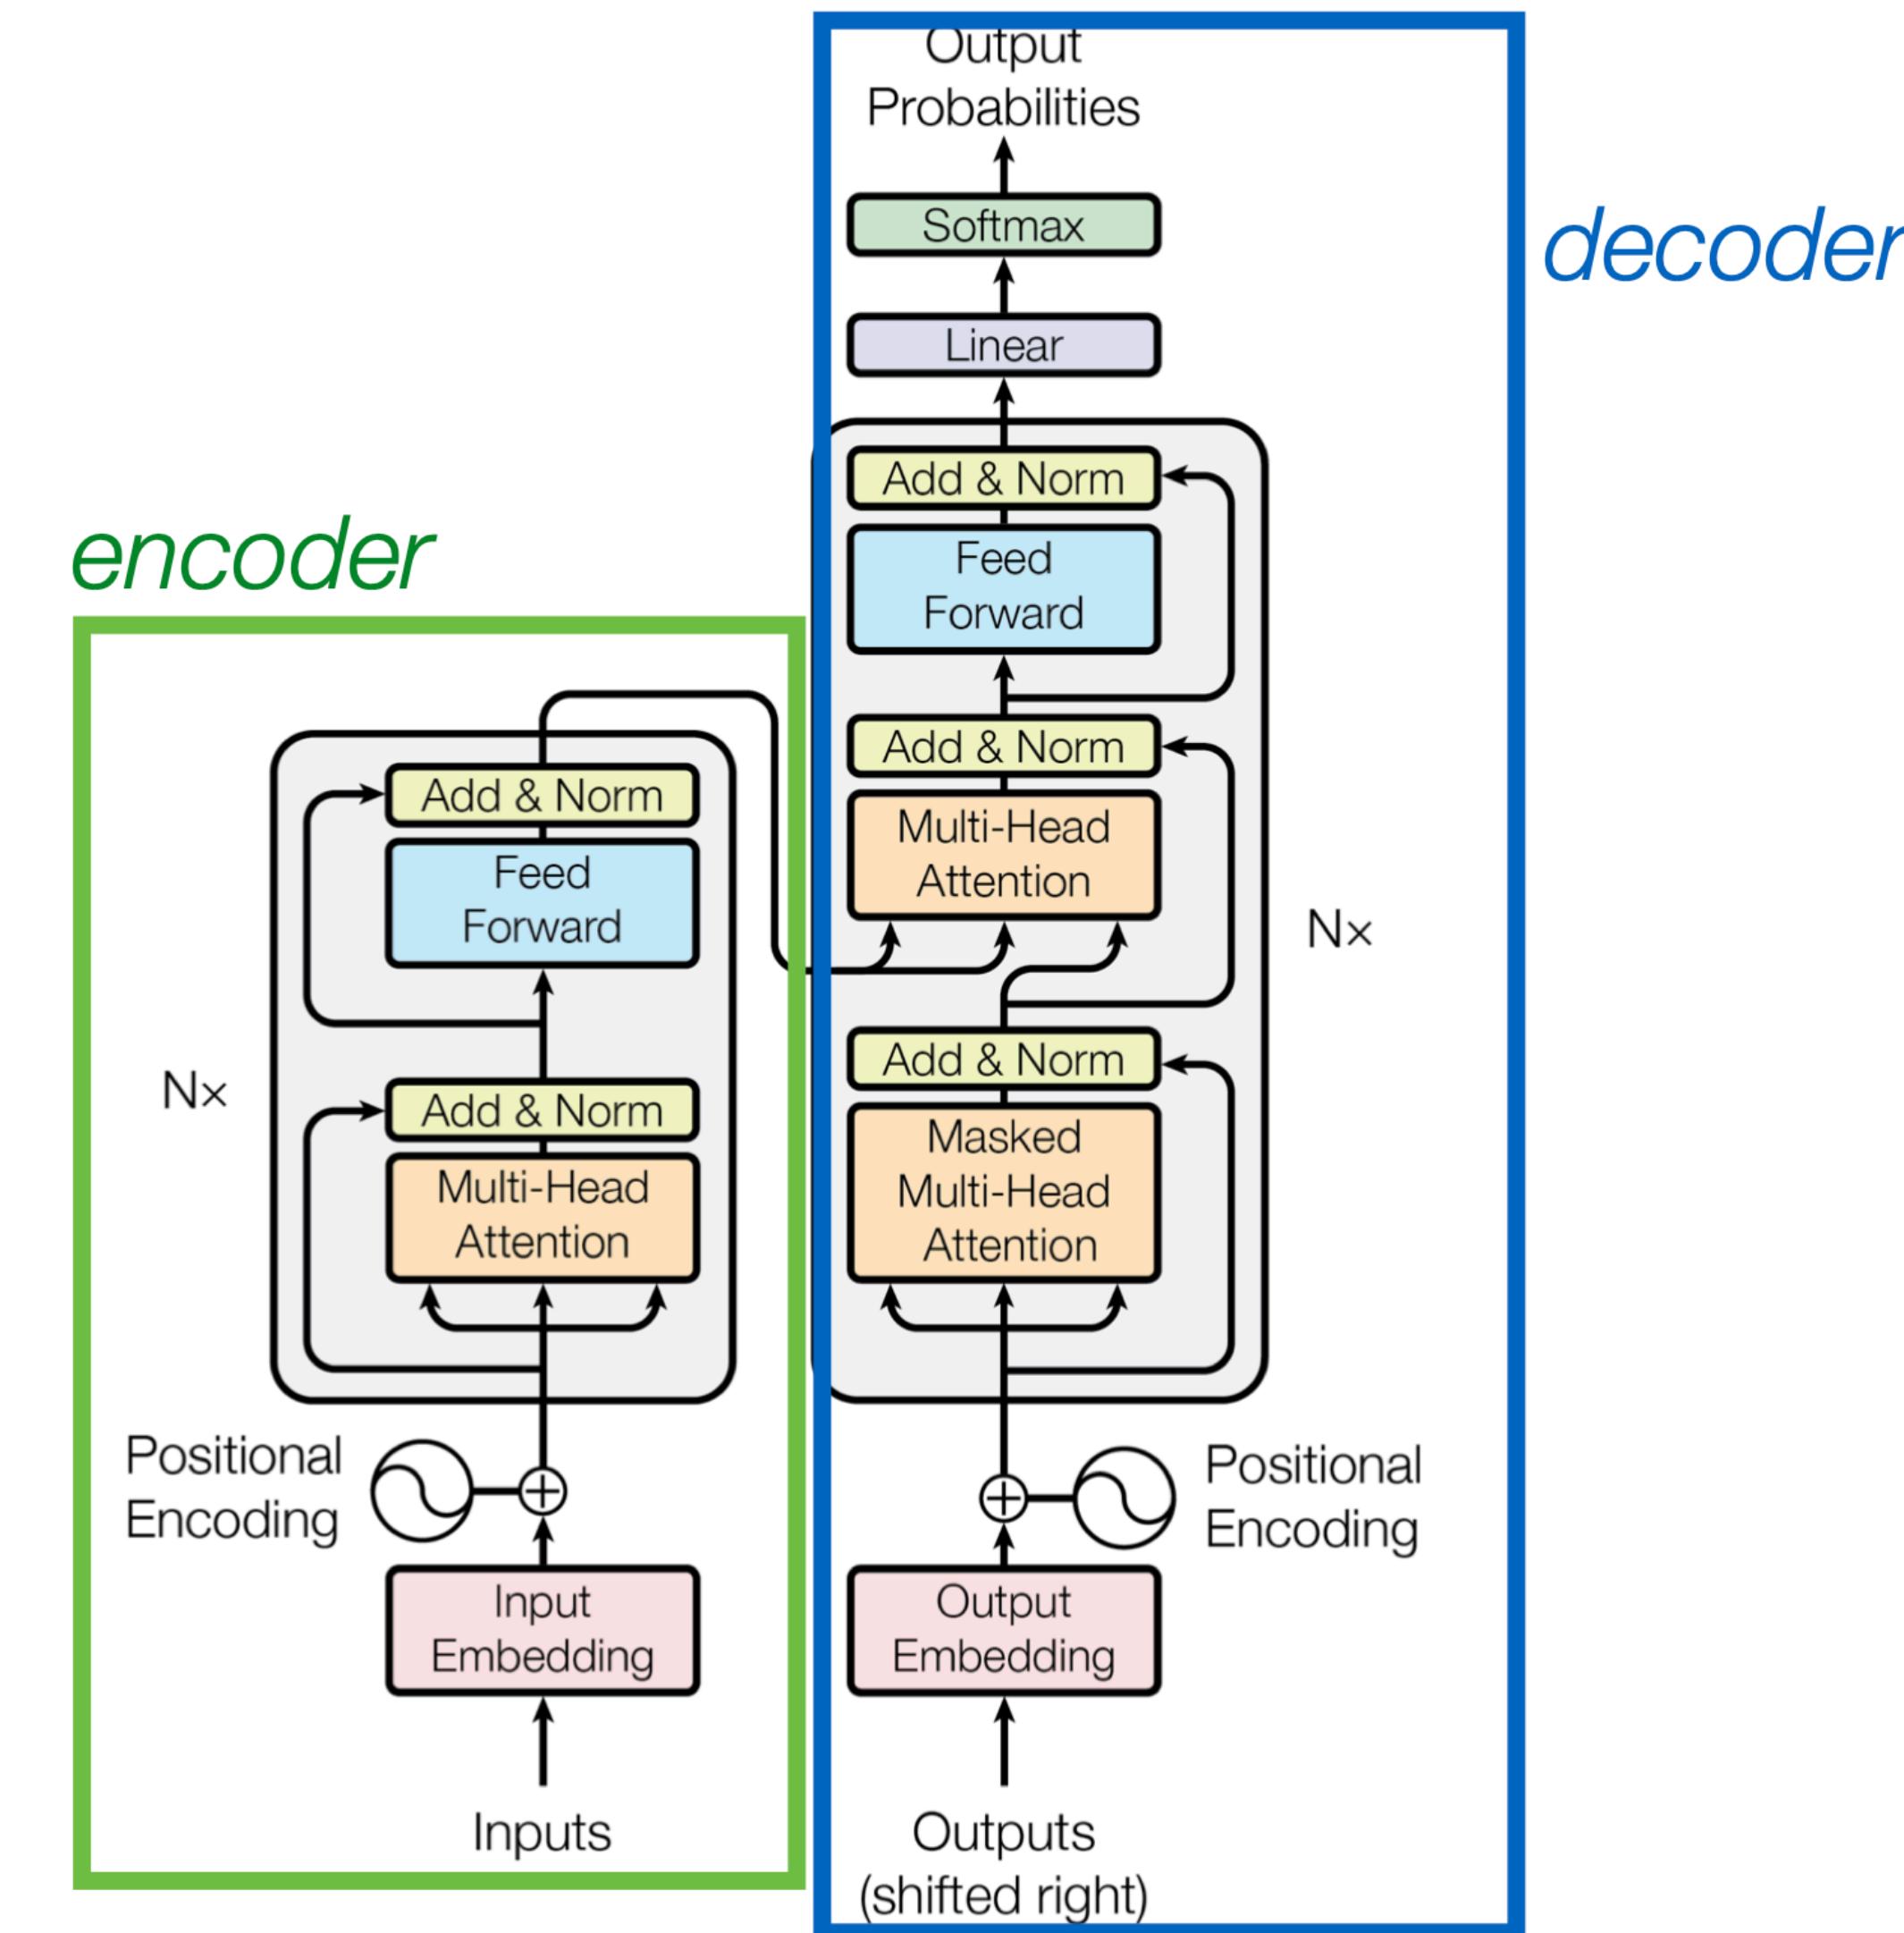
# Transformer



# Encoder

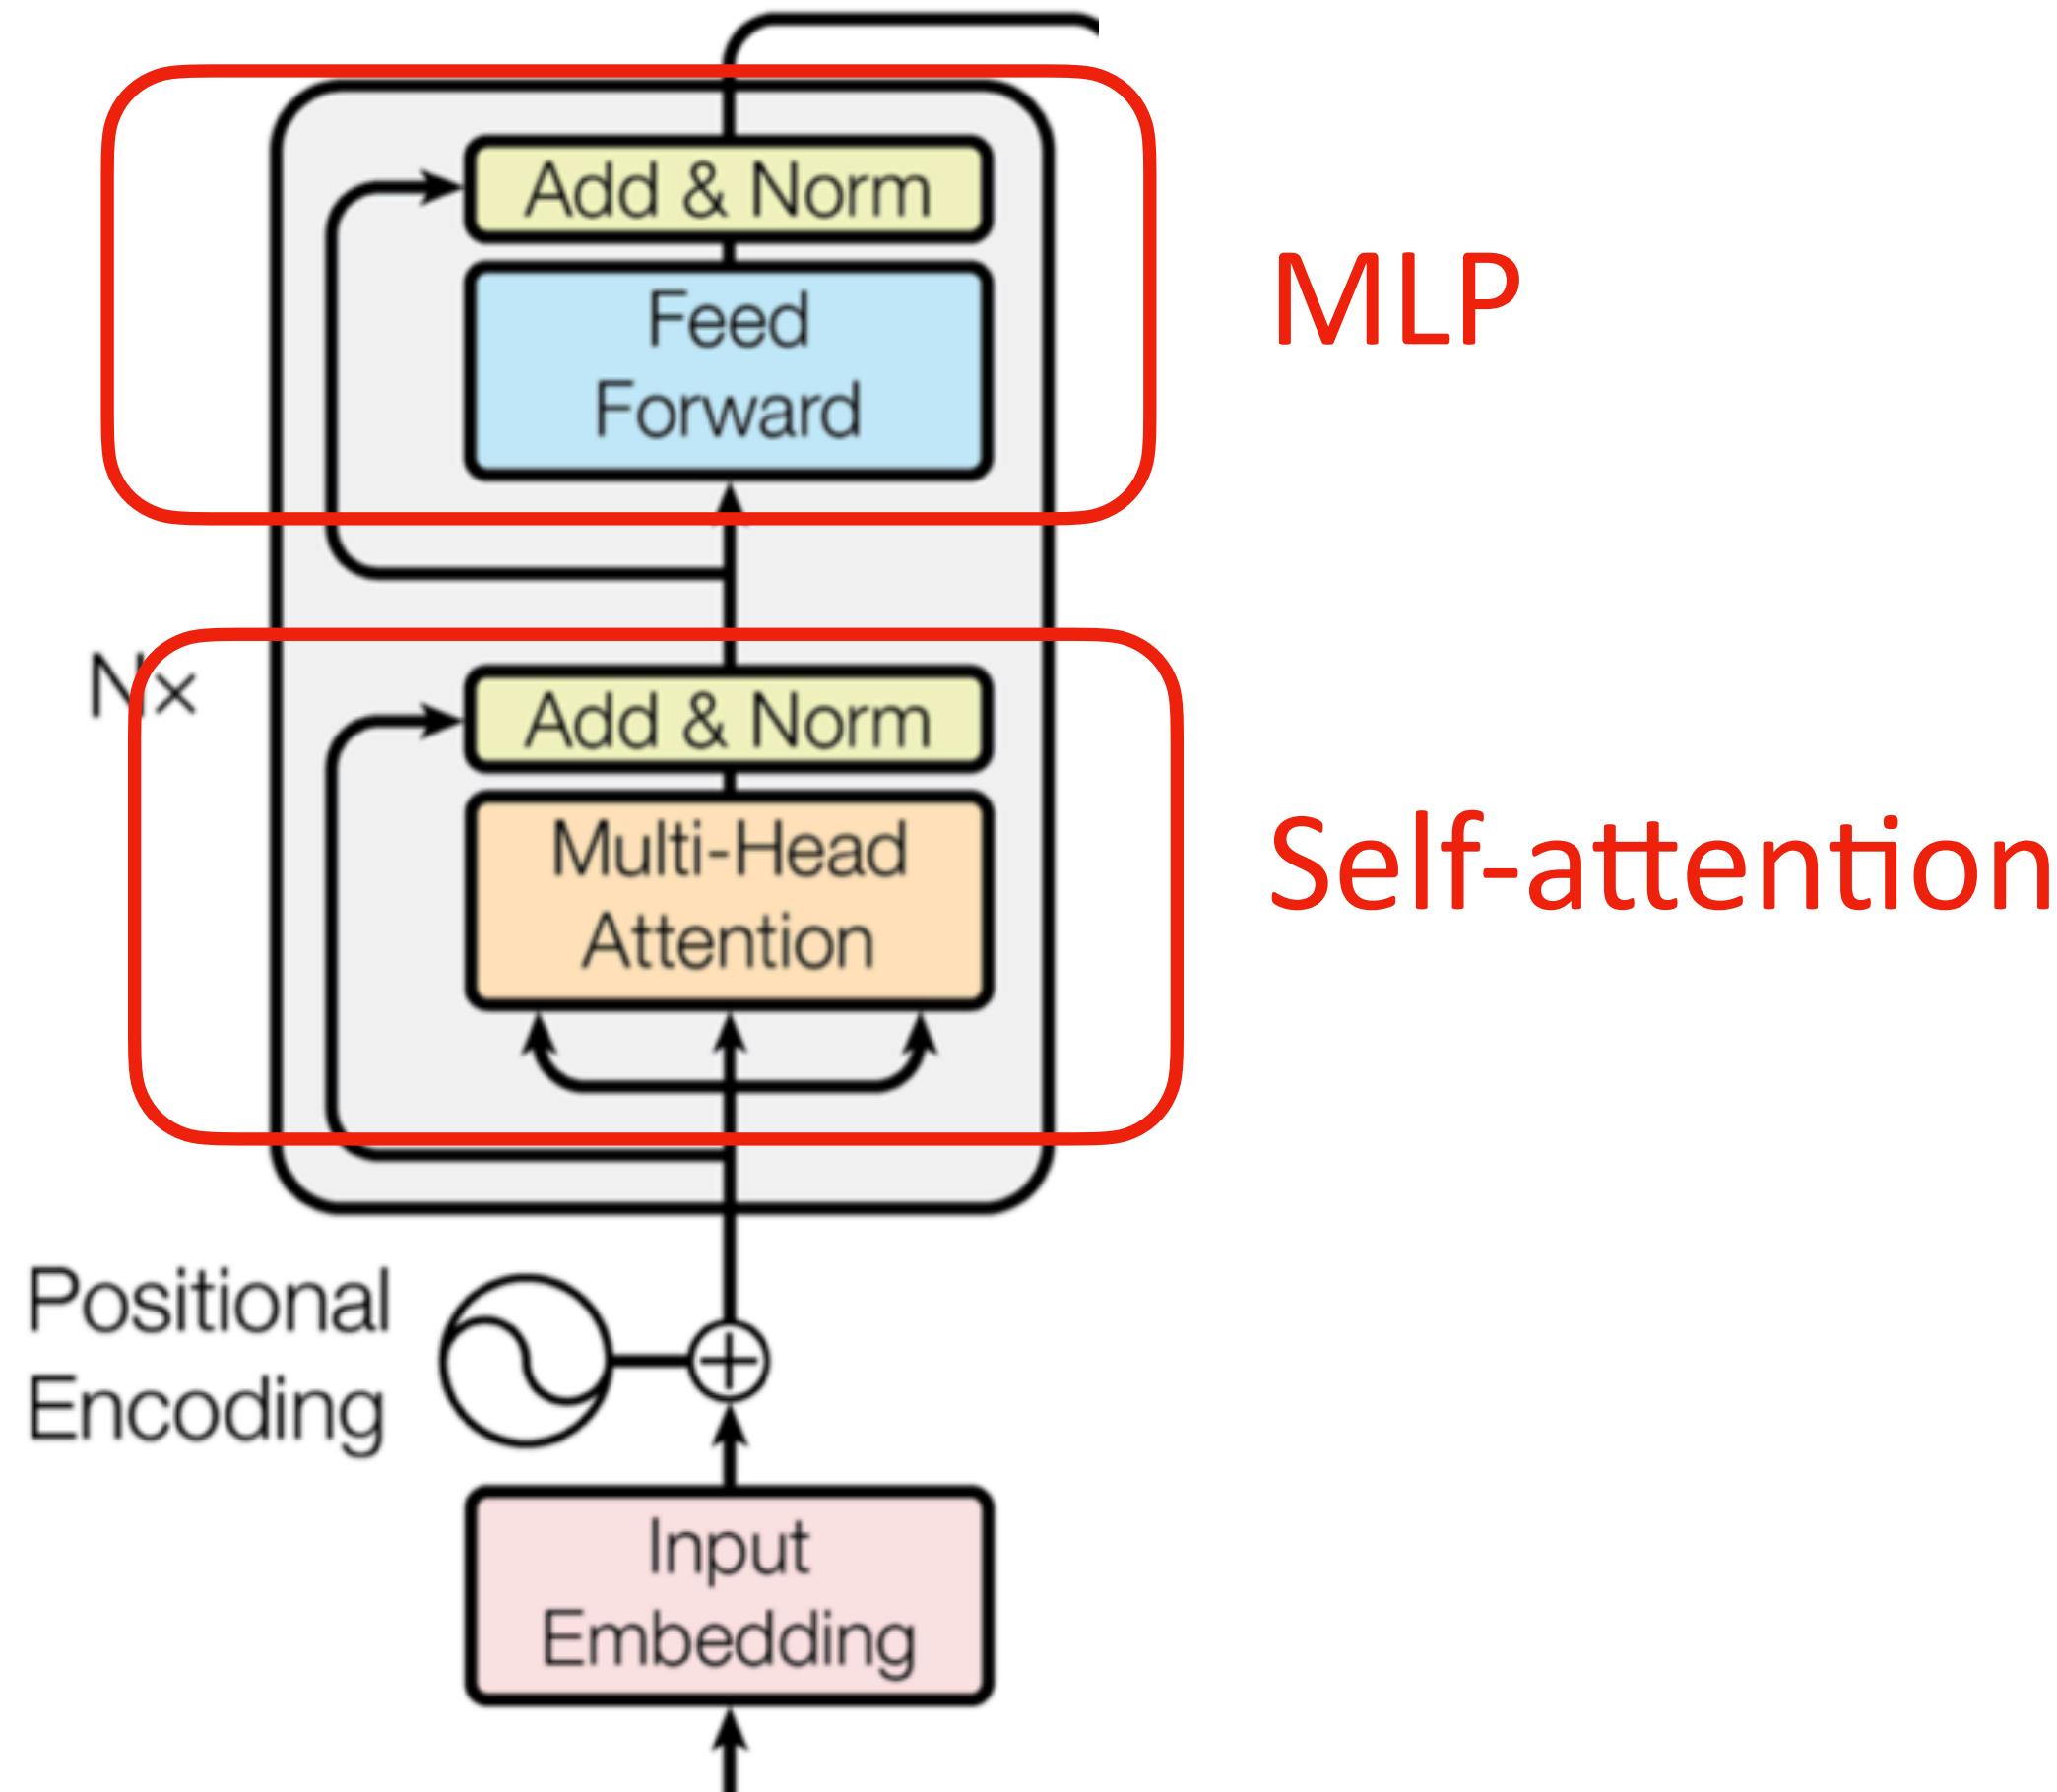


# Decoder



# Transformer Encoder

Residual  
connection

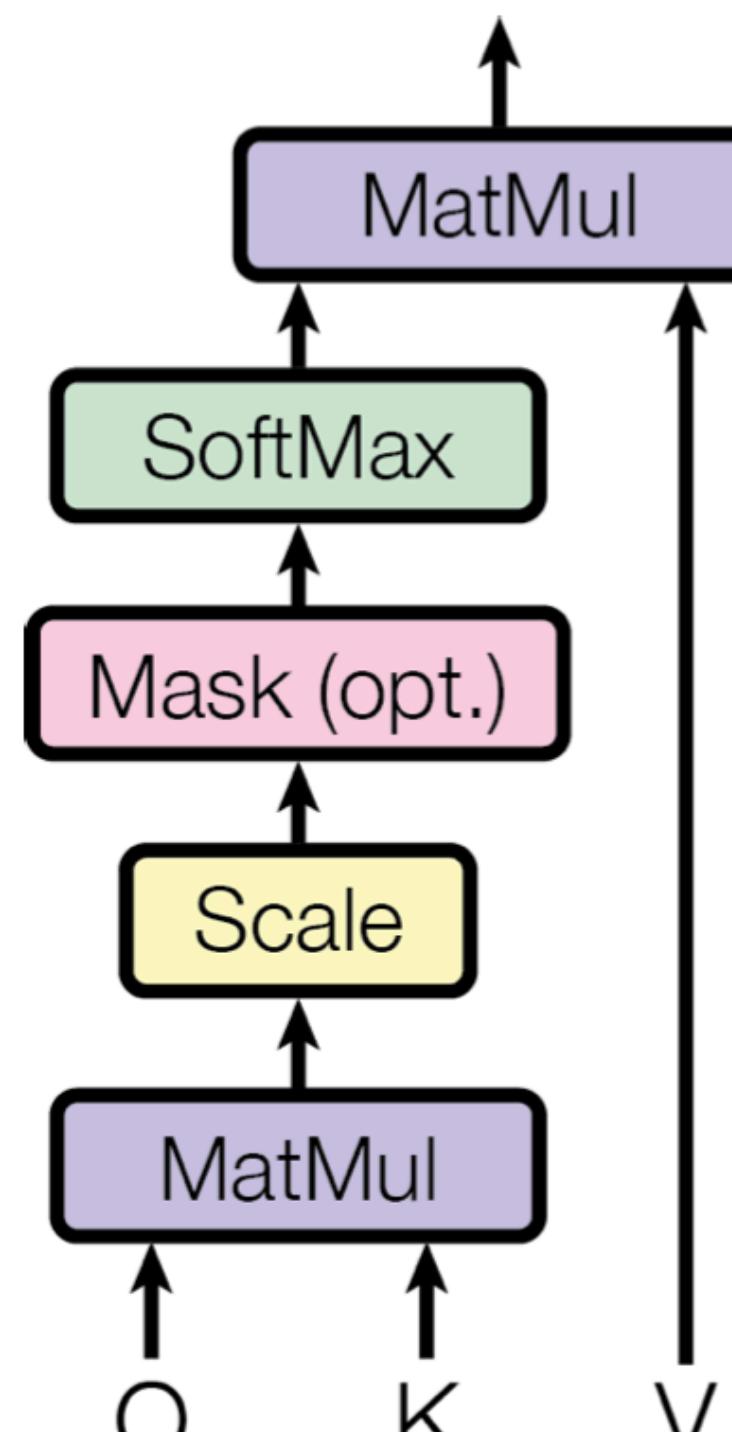


# What is Attention

Scaled Dot-Product Attention

$$Q \in R^{n \times d} \quad K \in R^{m \times d} \quad V \in R^{m \times d}$$

We have n queries, m (key, value) pairs



$$\text{Attention weight} = \text{softmax}(QK^T)$$

Dot-products grow large in magnitude

$$\text{Scaled Attention weight} = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right) \text{ Shape is mxn}$$

Attention weight represents the strength to “attend” values V

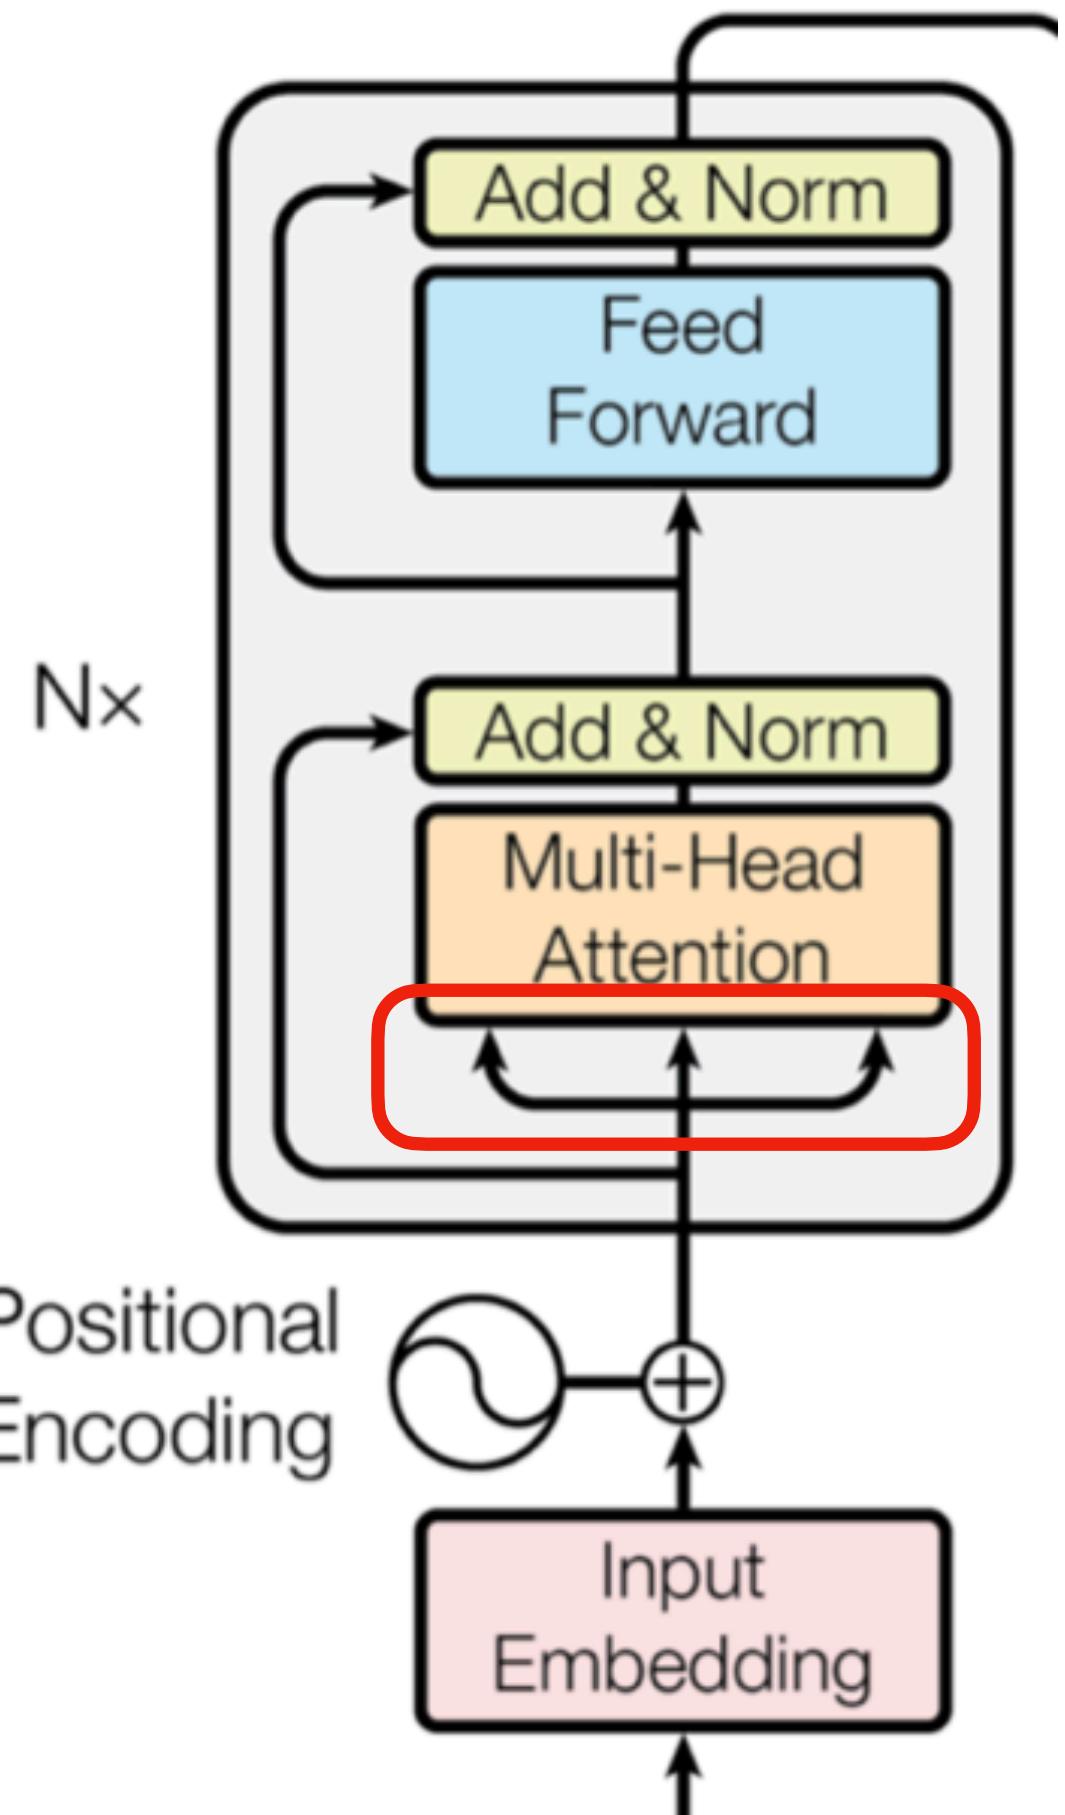
$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

Q: Query

K: key

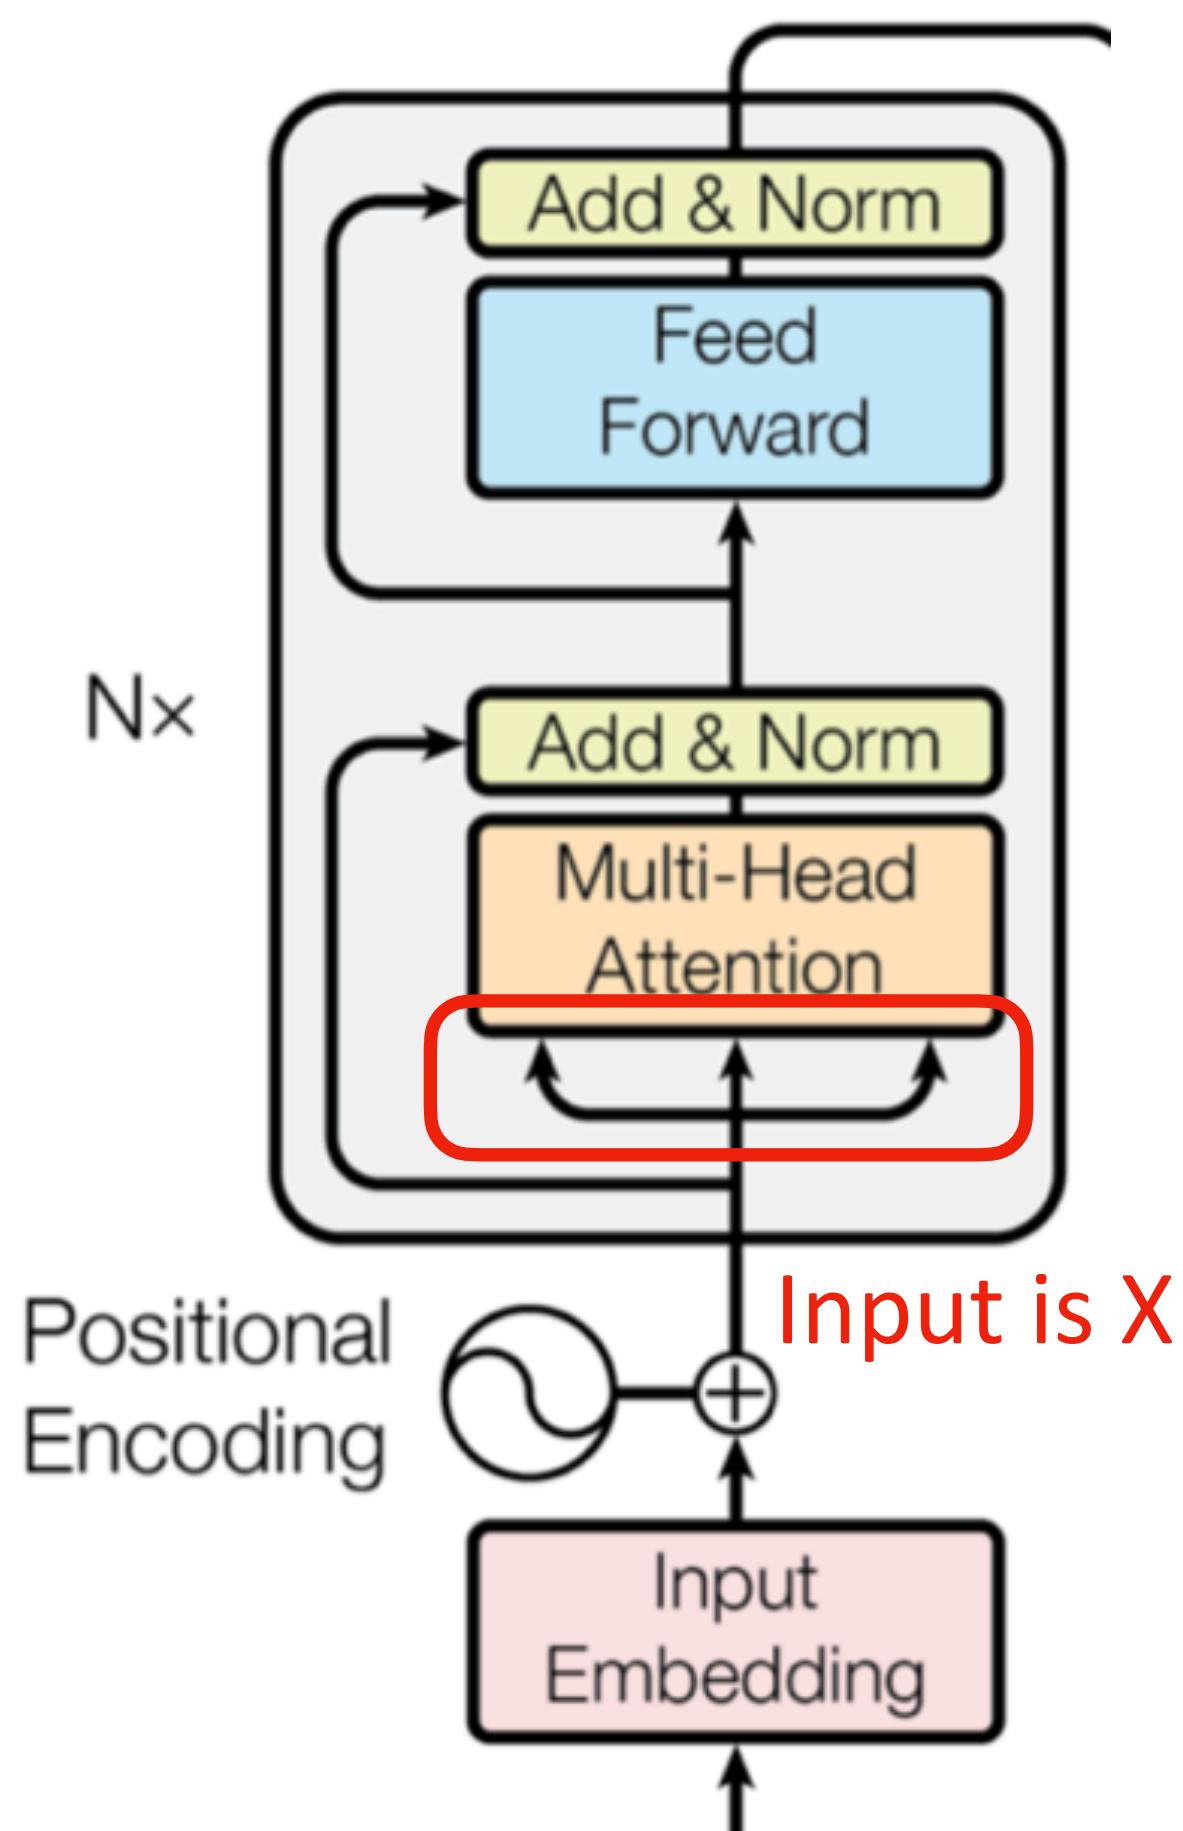
V: value

# Q, K, V



What are Q, K, V in the transformer

# Self-Attention



35

$$X \times W^Q = Q$$

$$X \times W^K = K$$

$$X \times W^V = V$$

Query, key, and value are from the same input, thus it is called “self”-attention

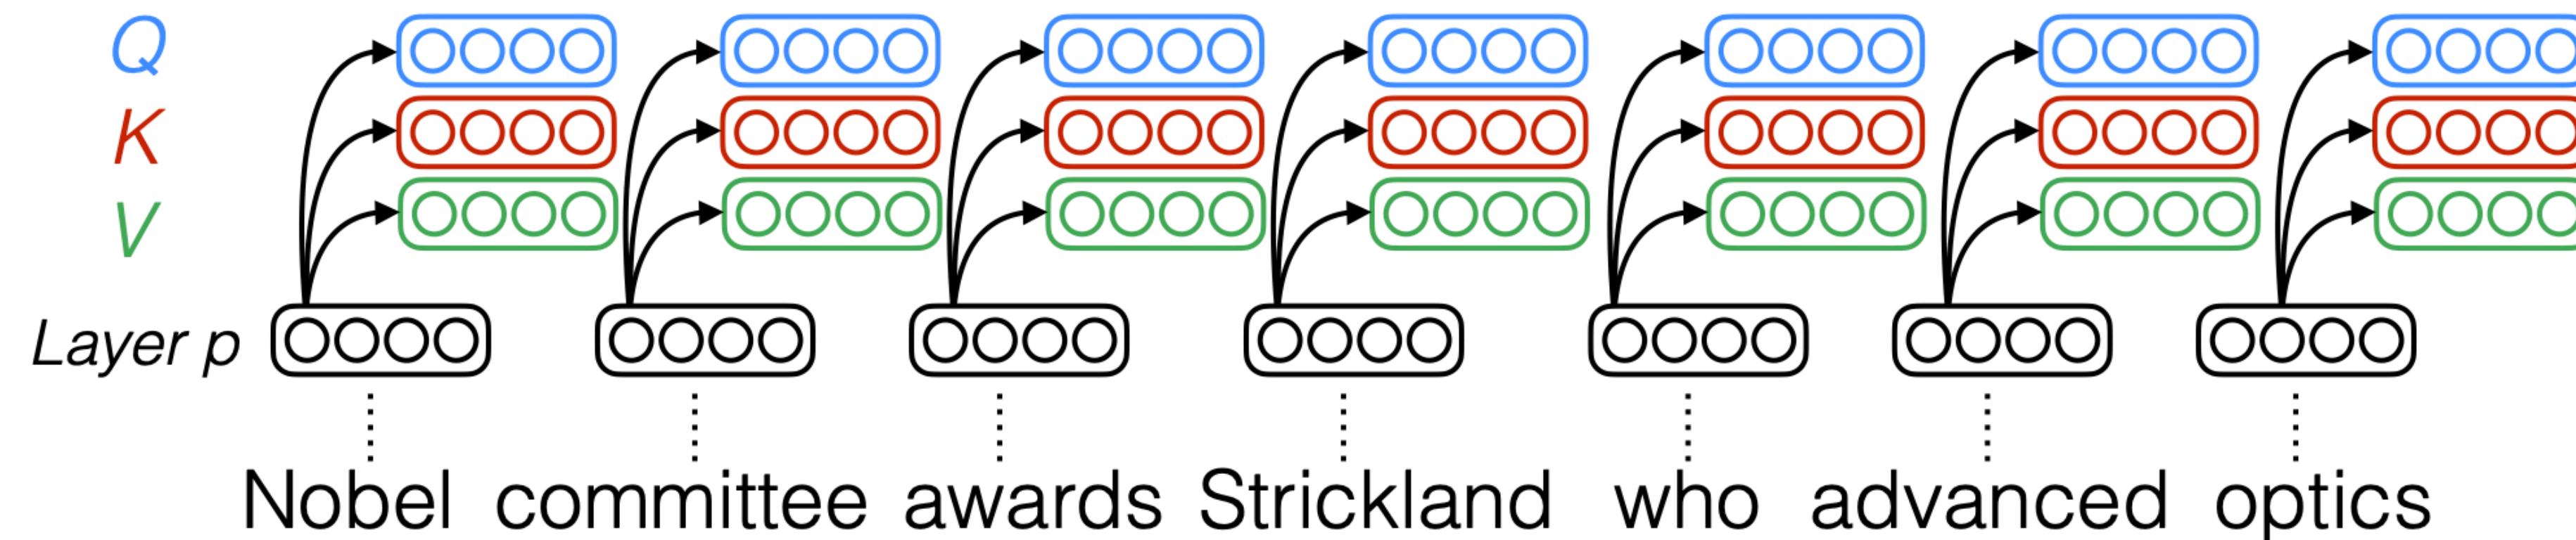
$$\text{softmax} \left( \frac{Q \times K^T}{\sqrt{d_k}} \right) \times V = Z$$

35

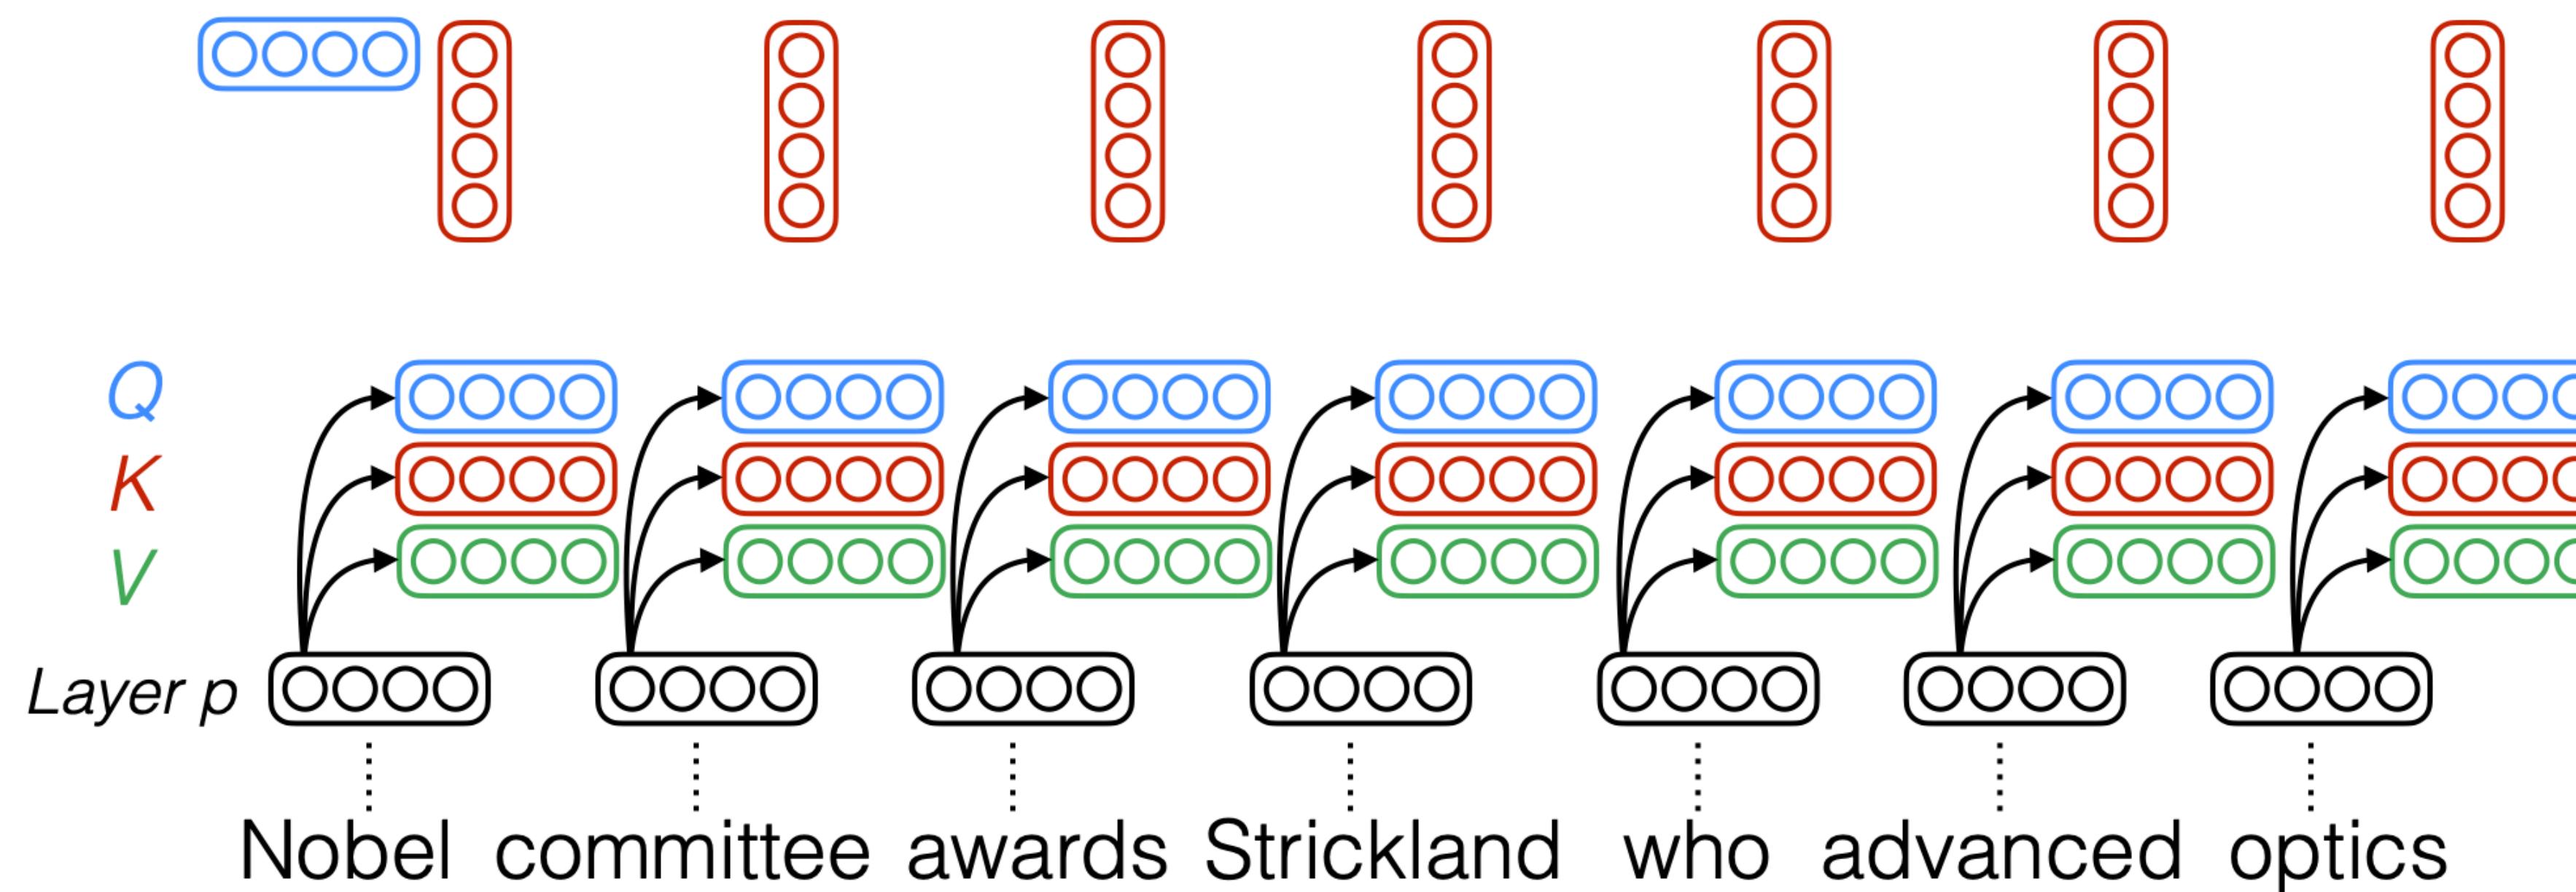
Jay Alammar. The Illustrated Transformer.

# Self-Attention

At each step, the attention computation attends to all steps in the input example

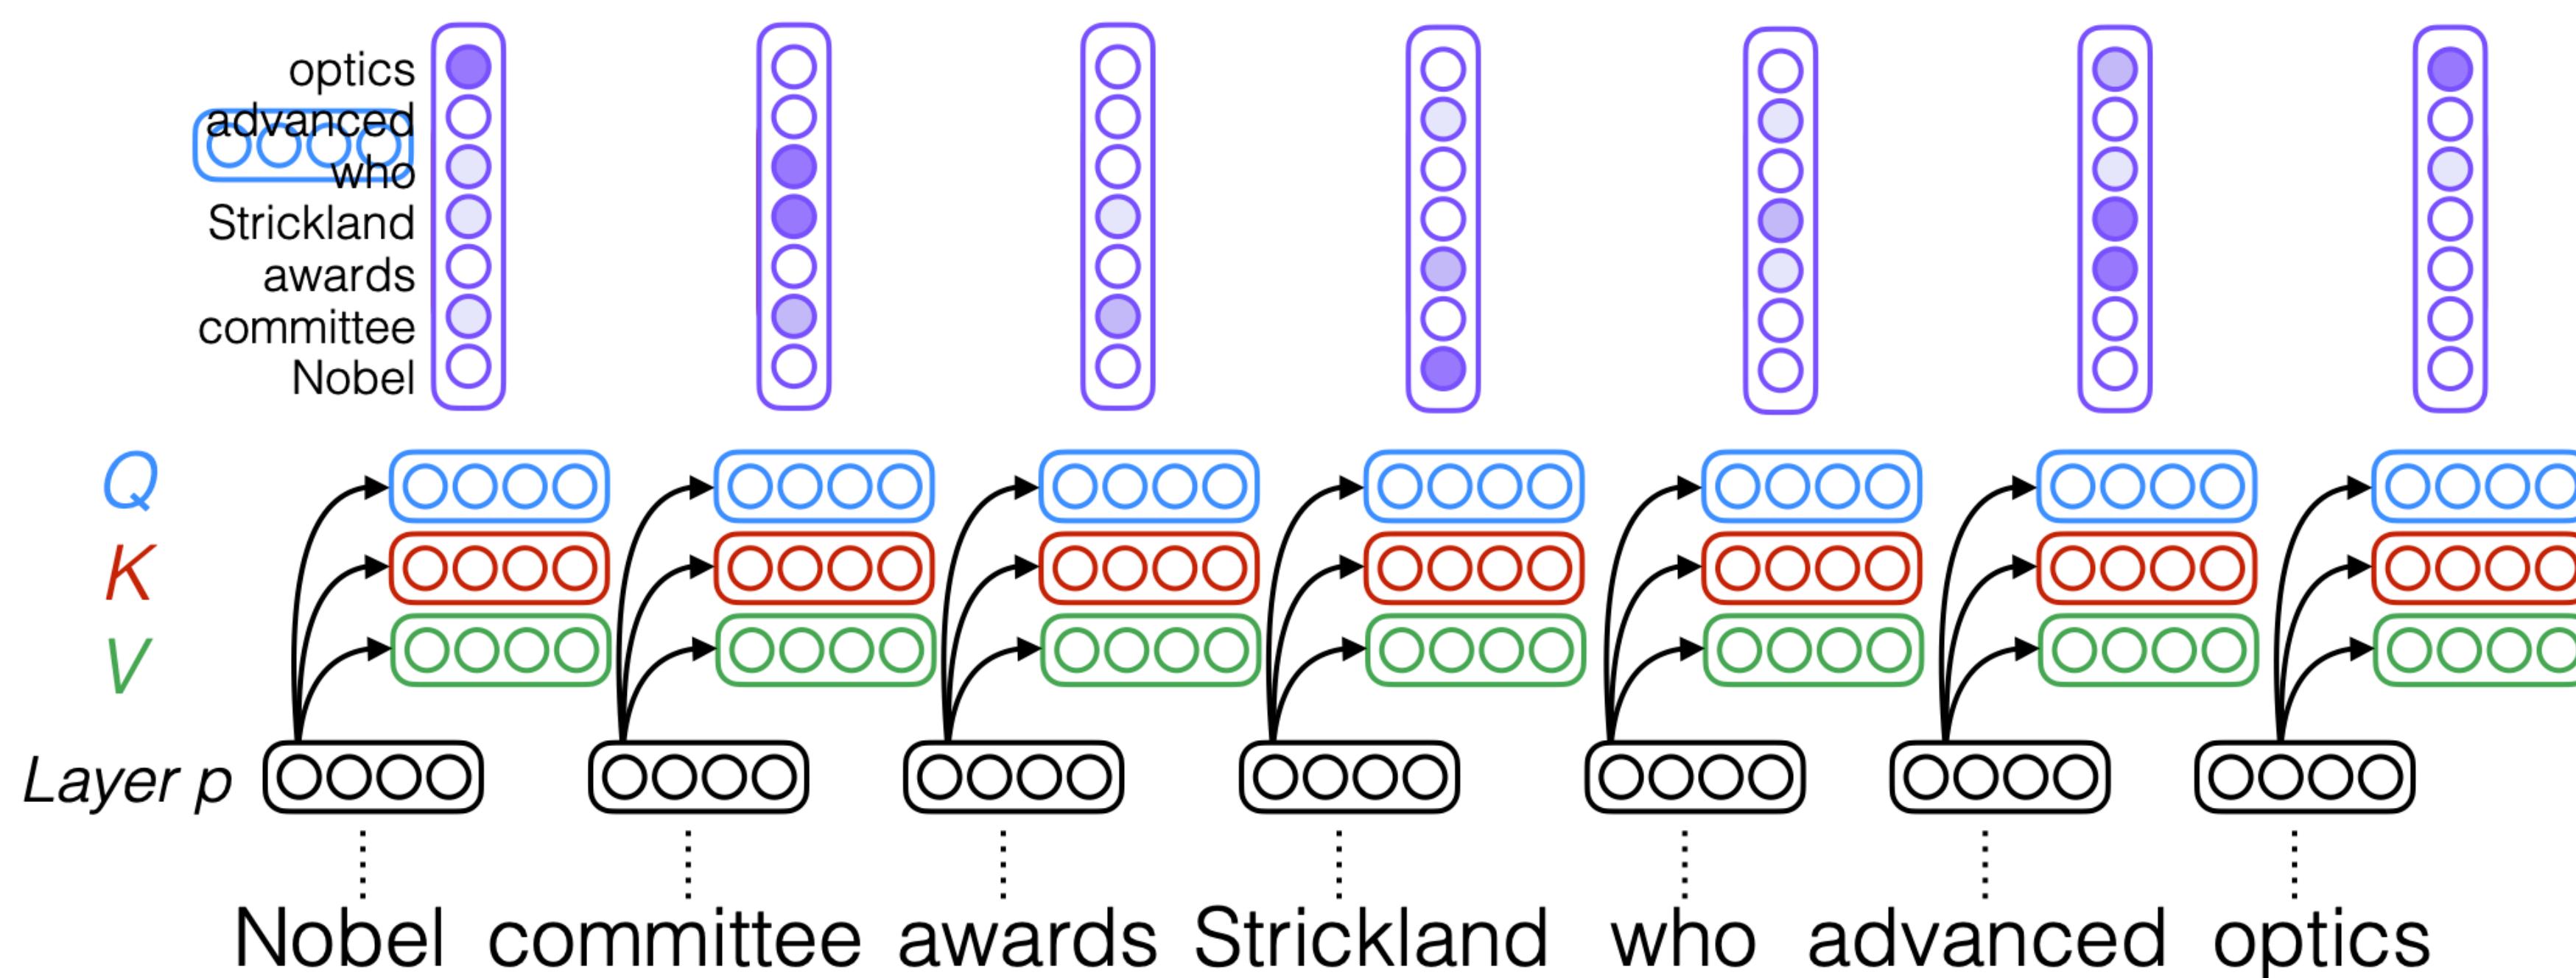


# Self-Attention

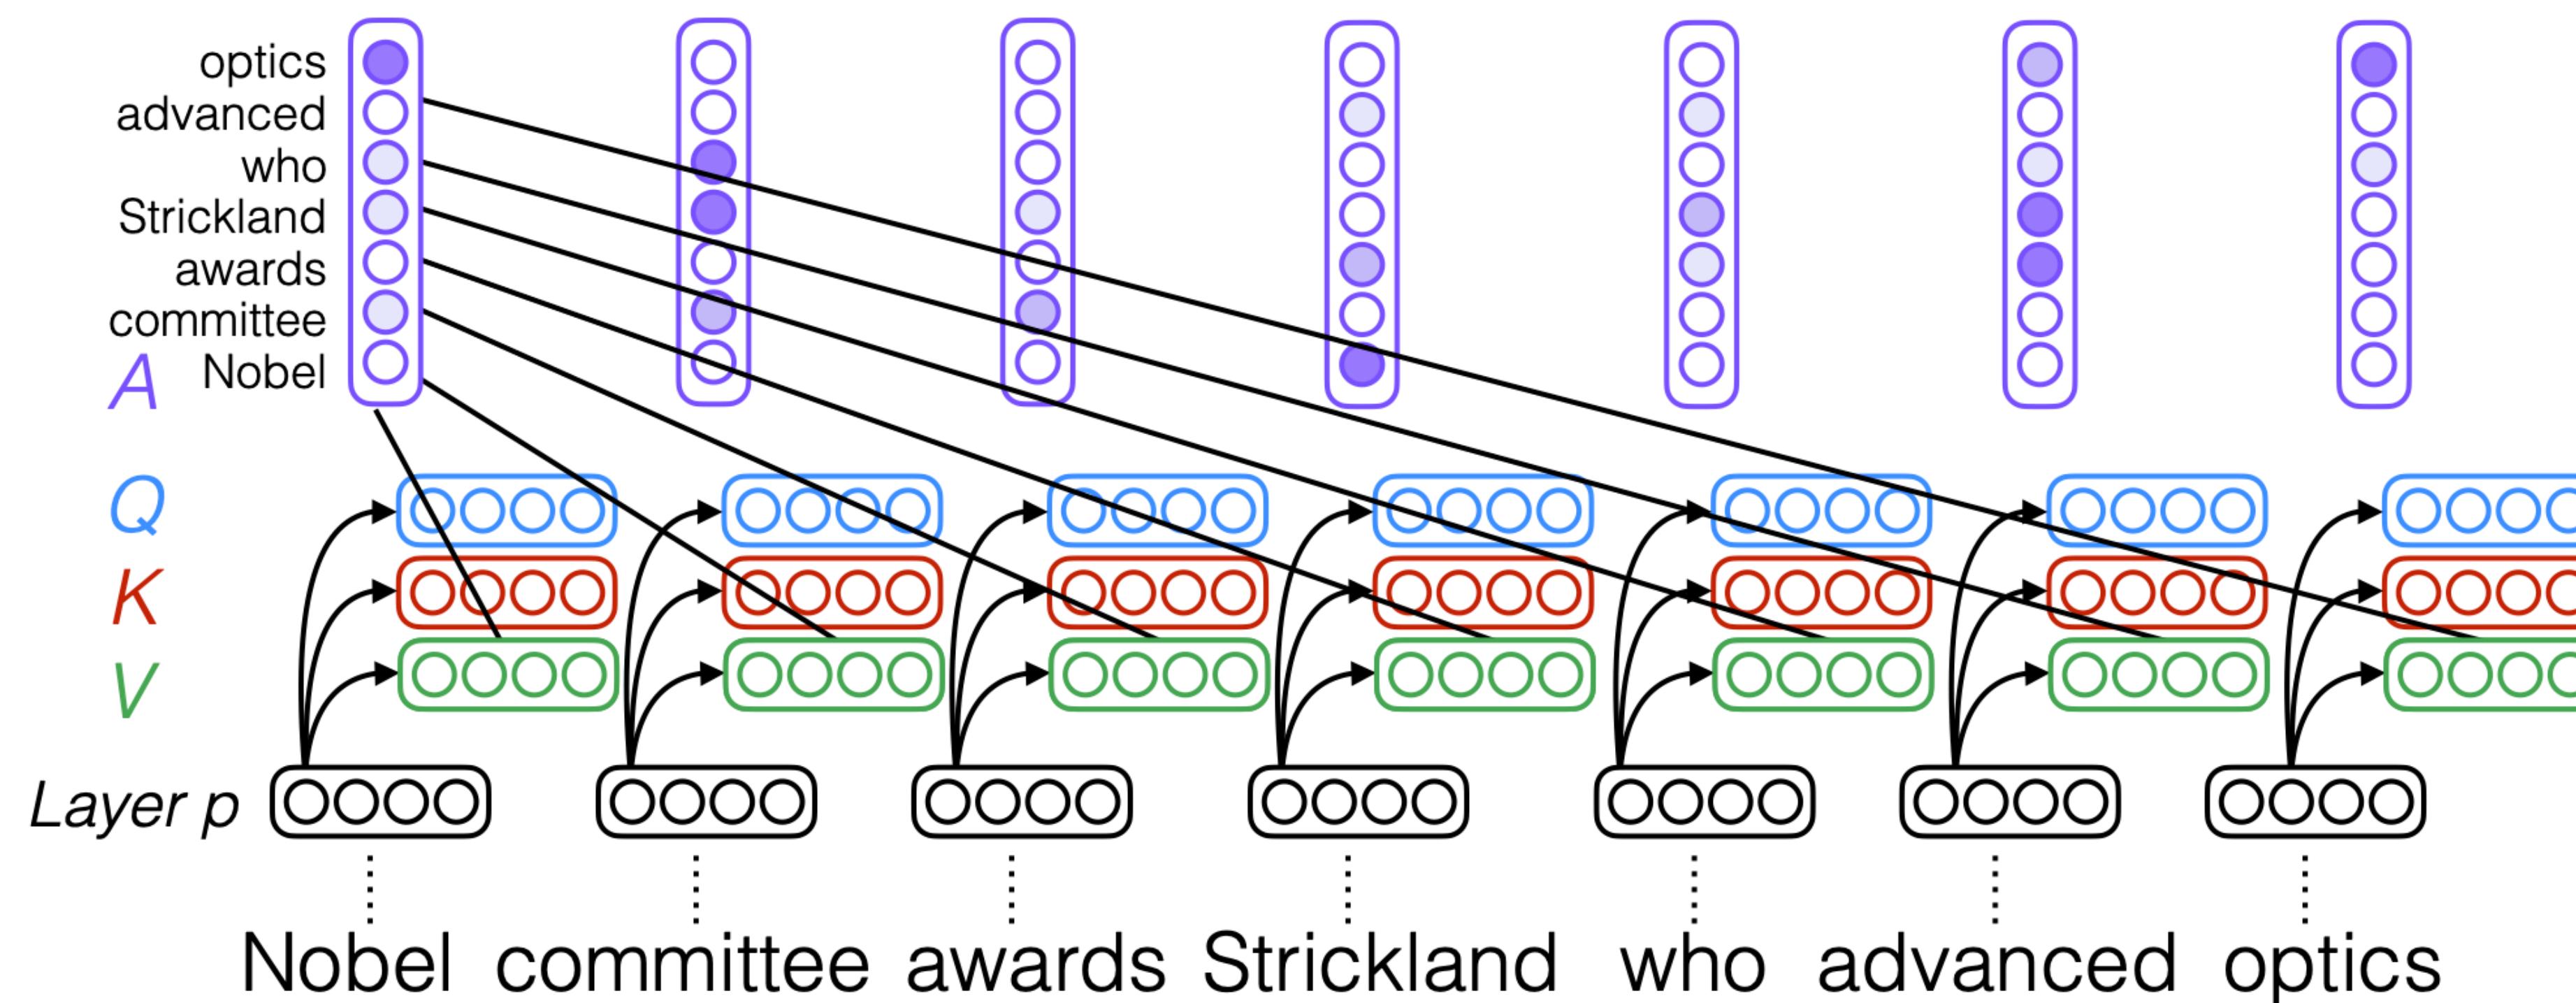


# Self-Attention

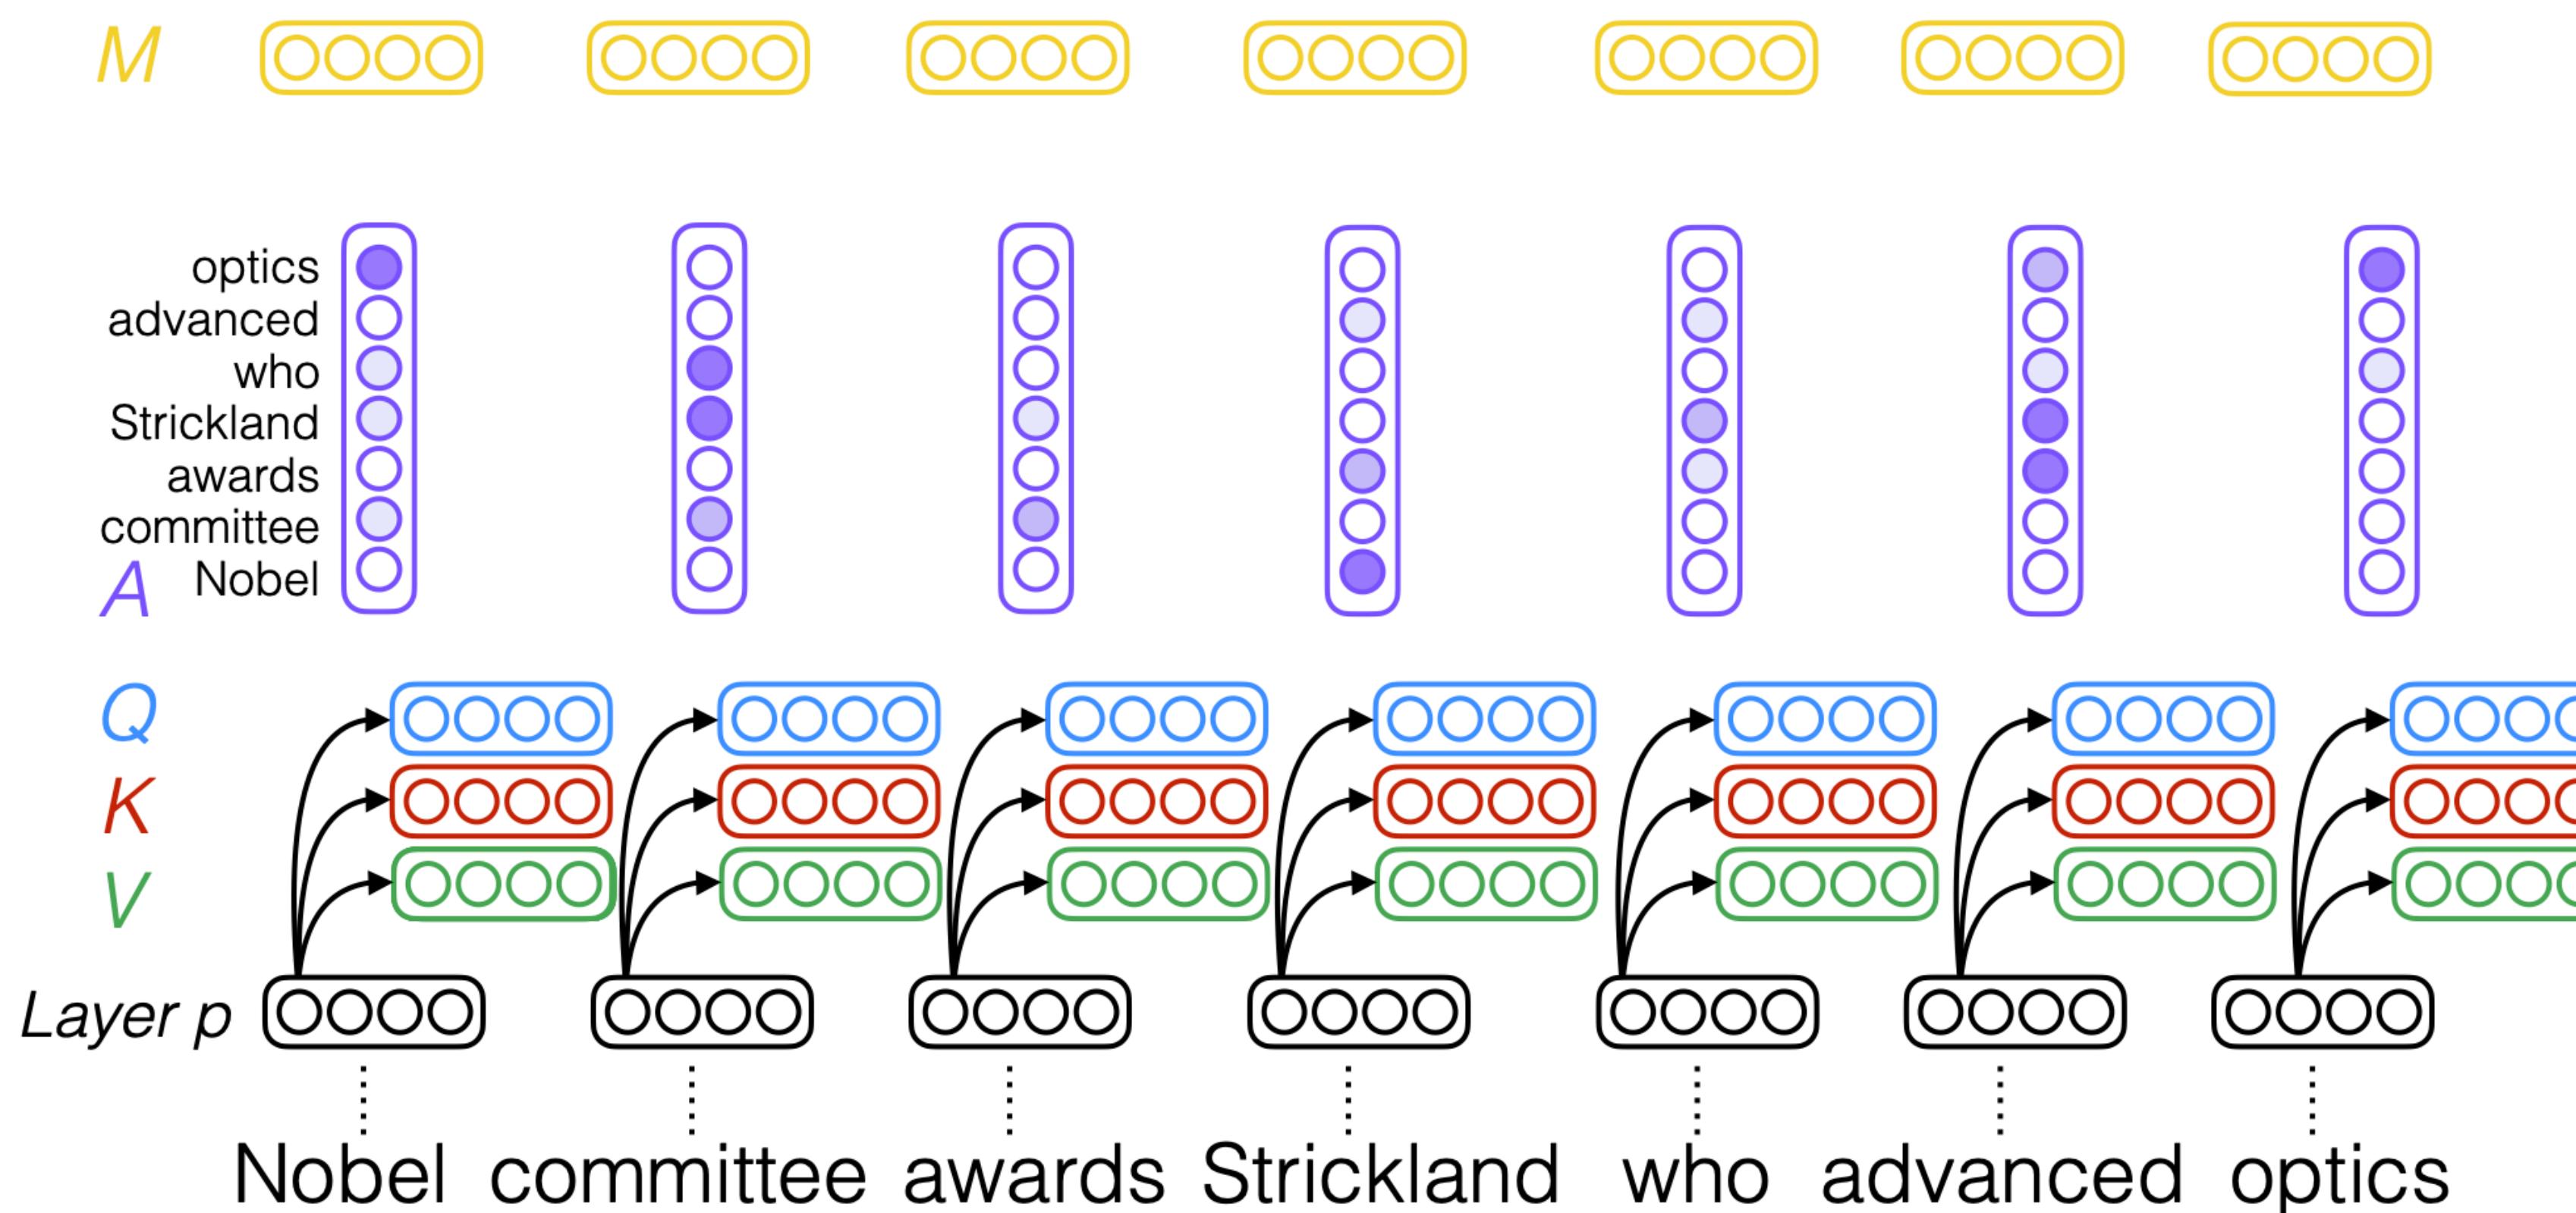
Attention weight on every word in the sequence



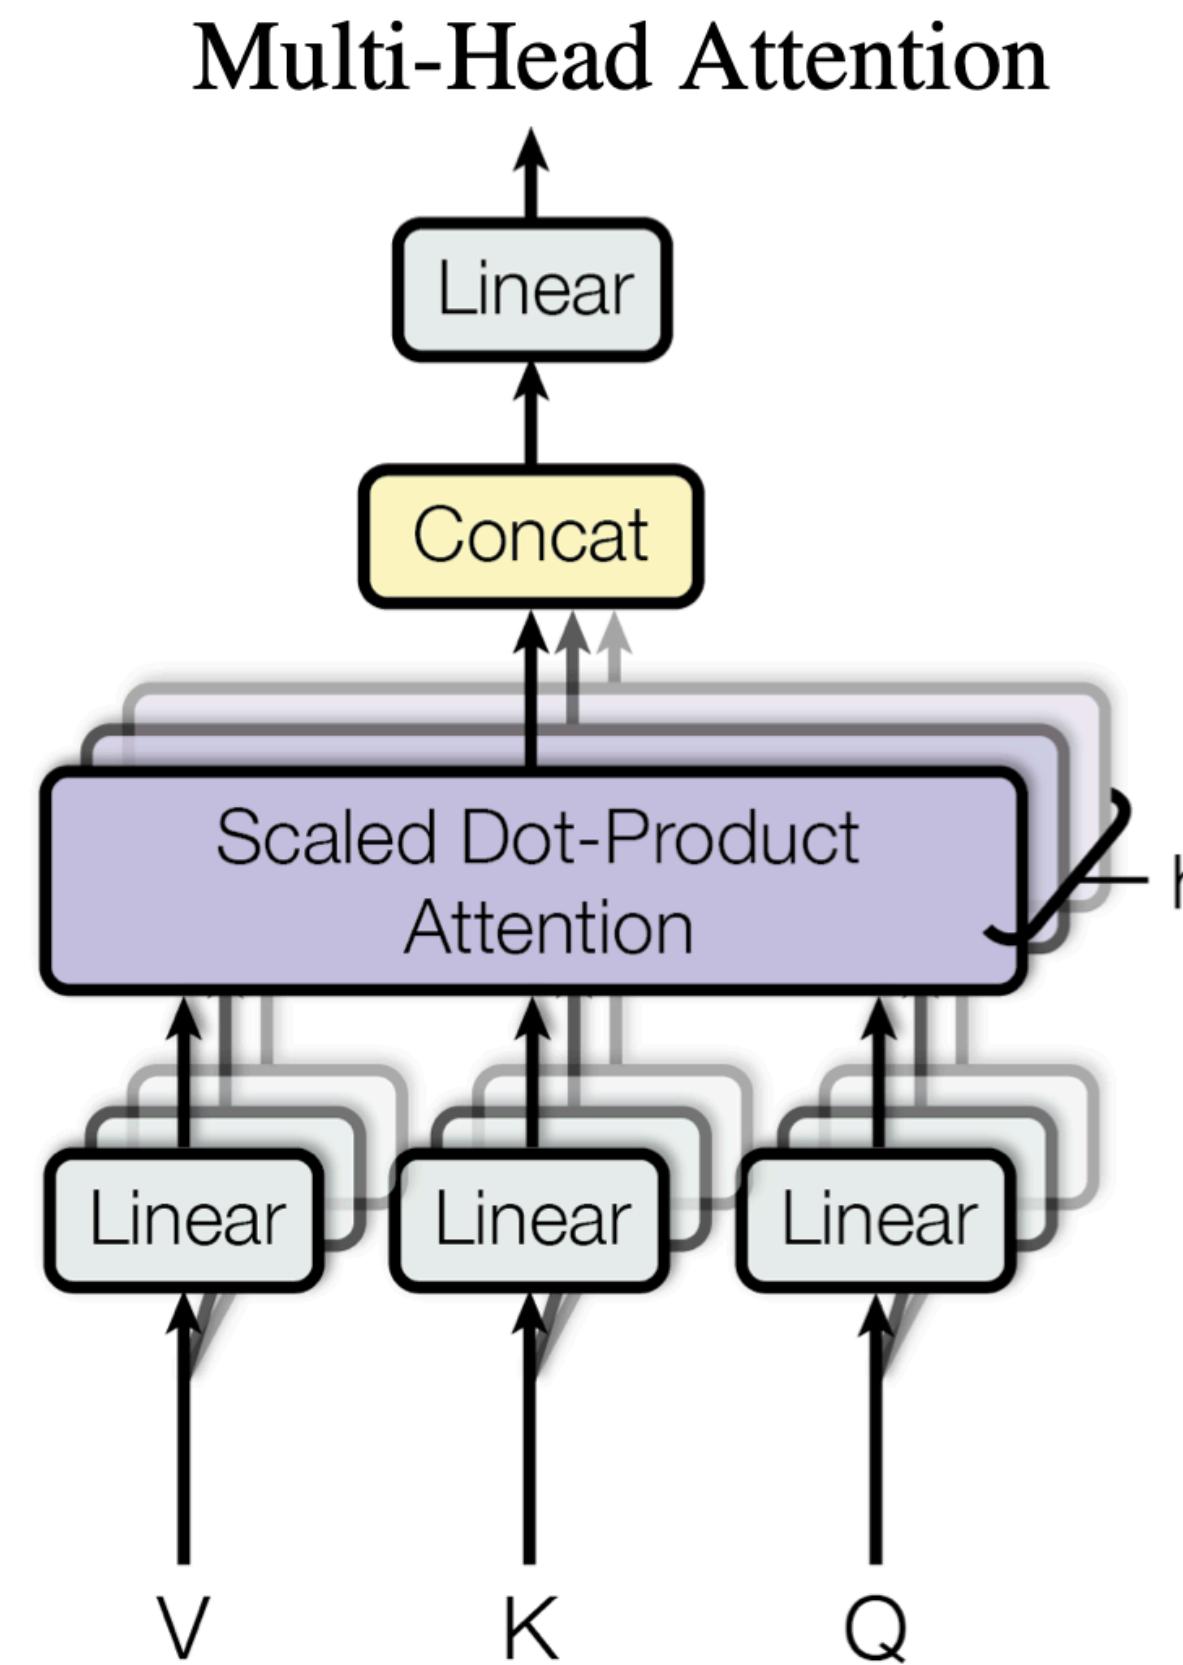
# Self-Attention



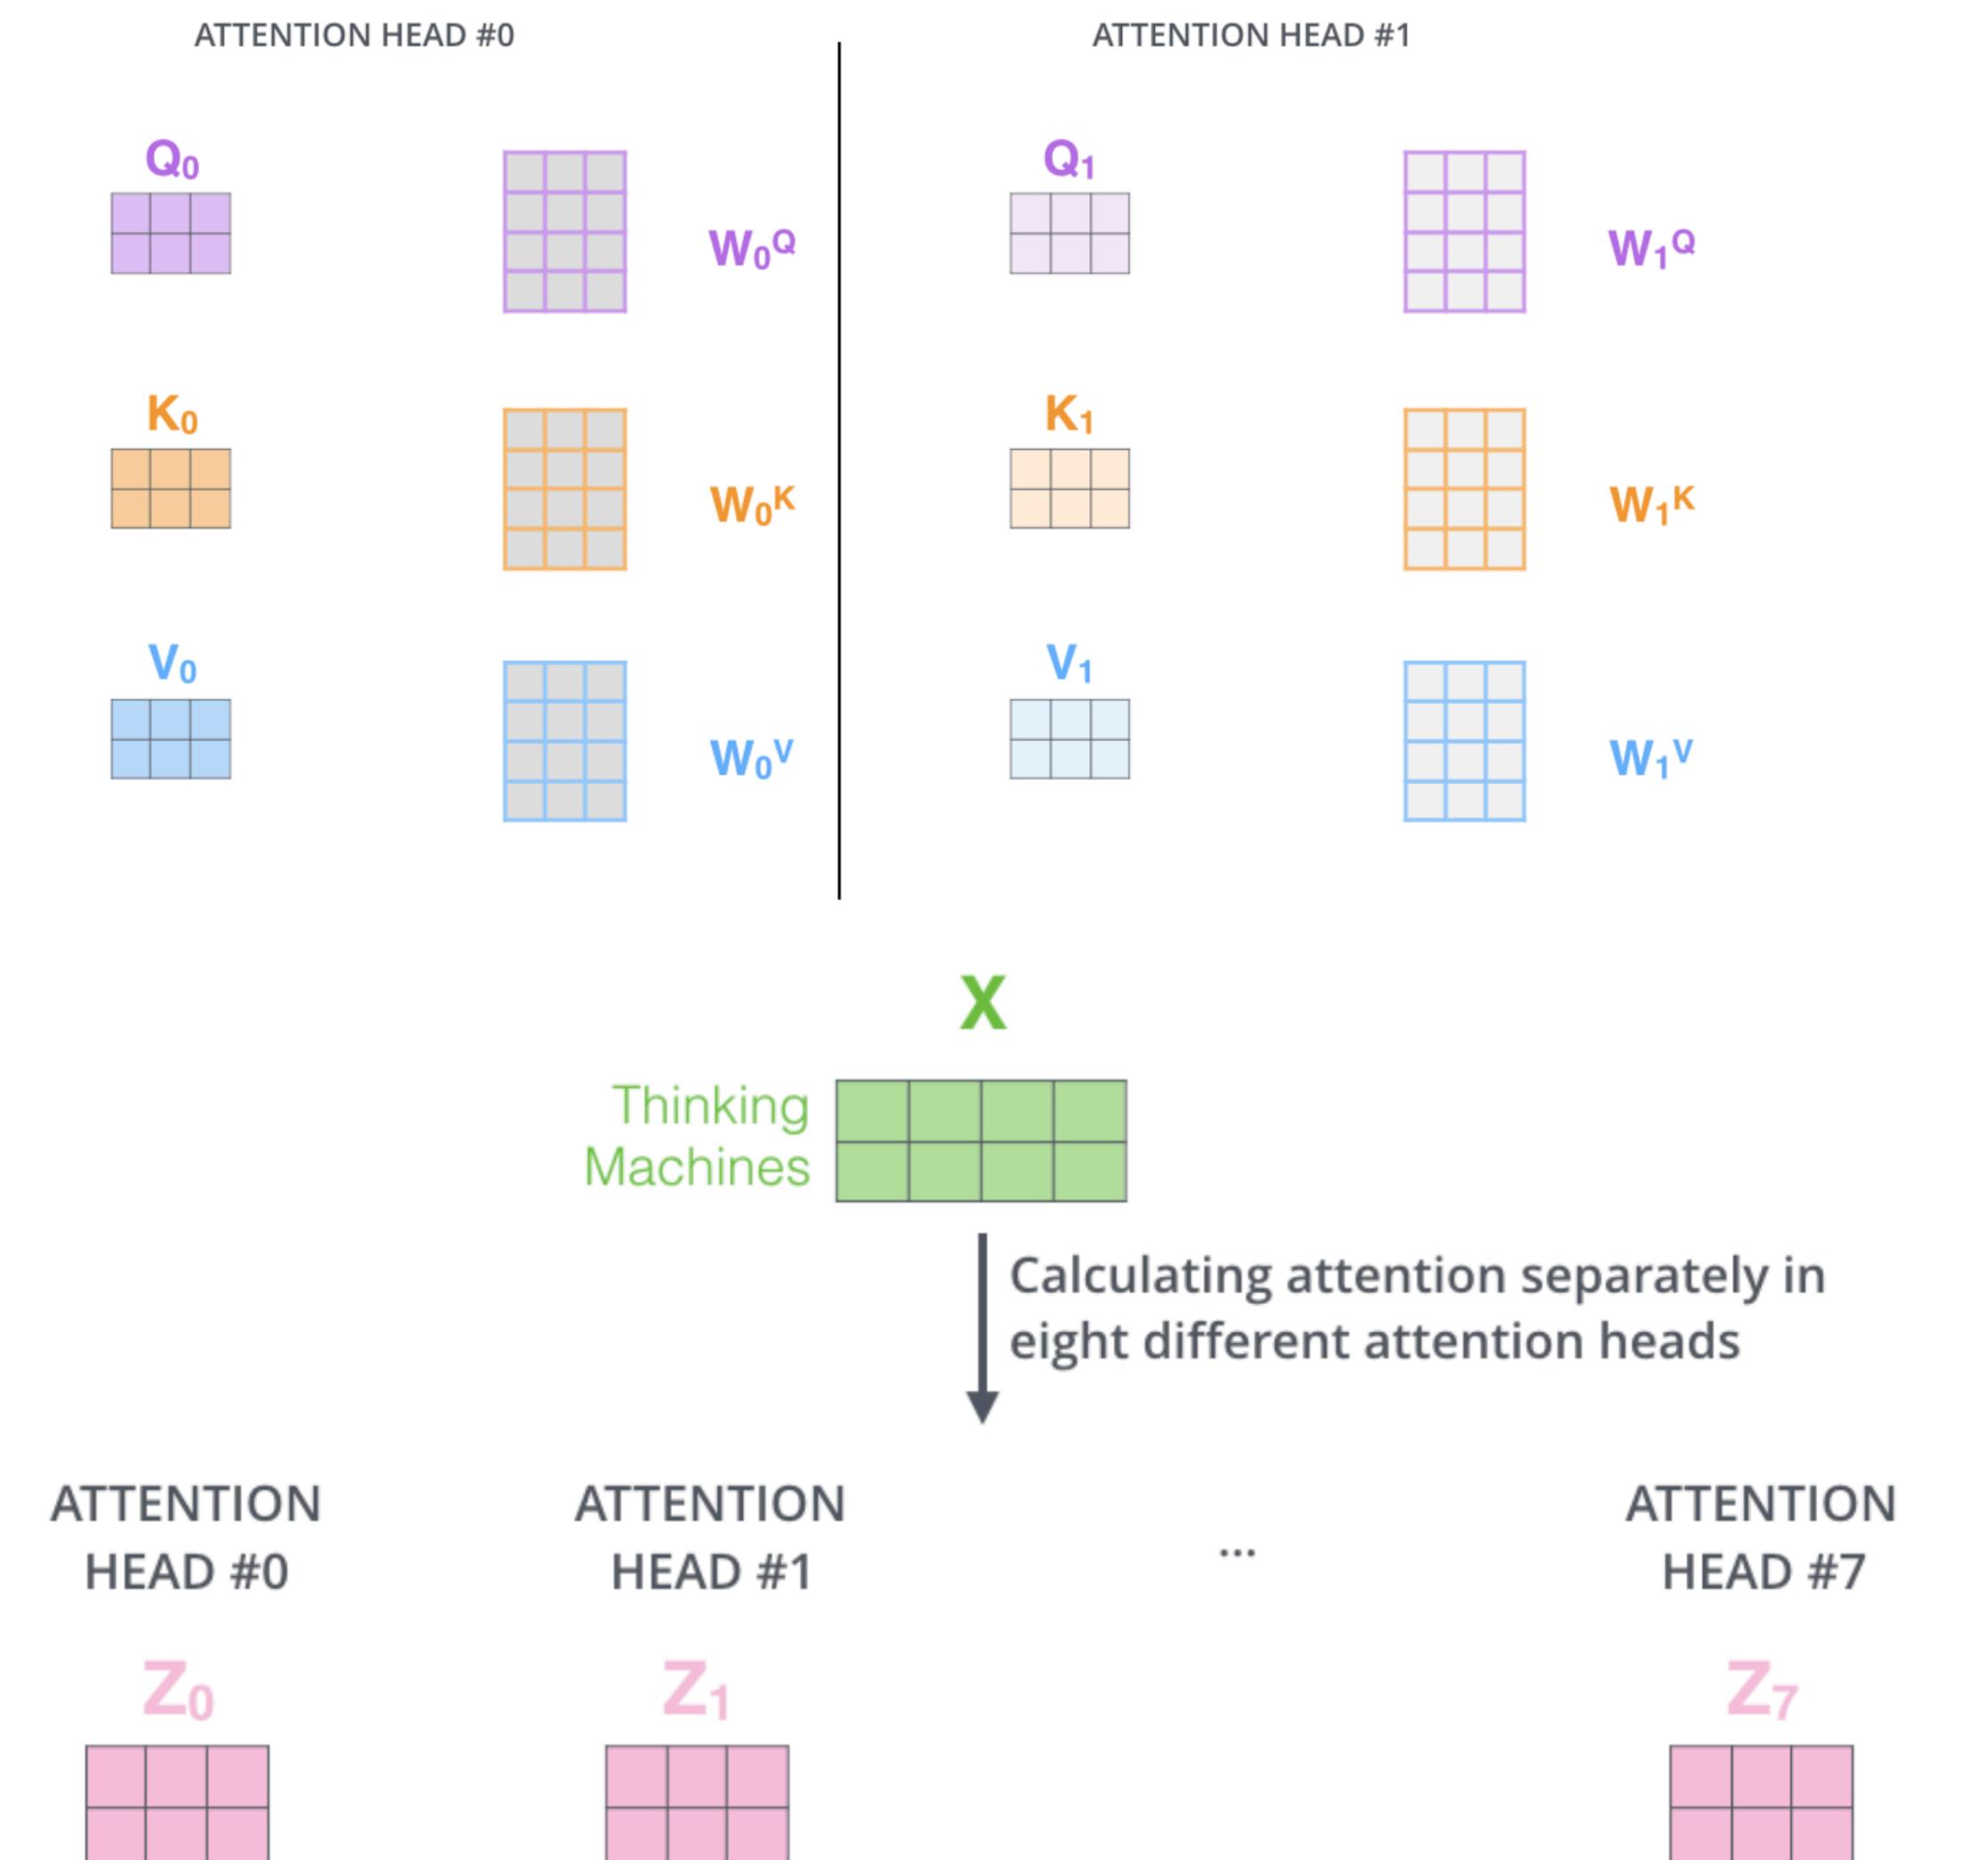
# Self-Attention



# Multi-Head Attention



# Multi-Head Self-Attention



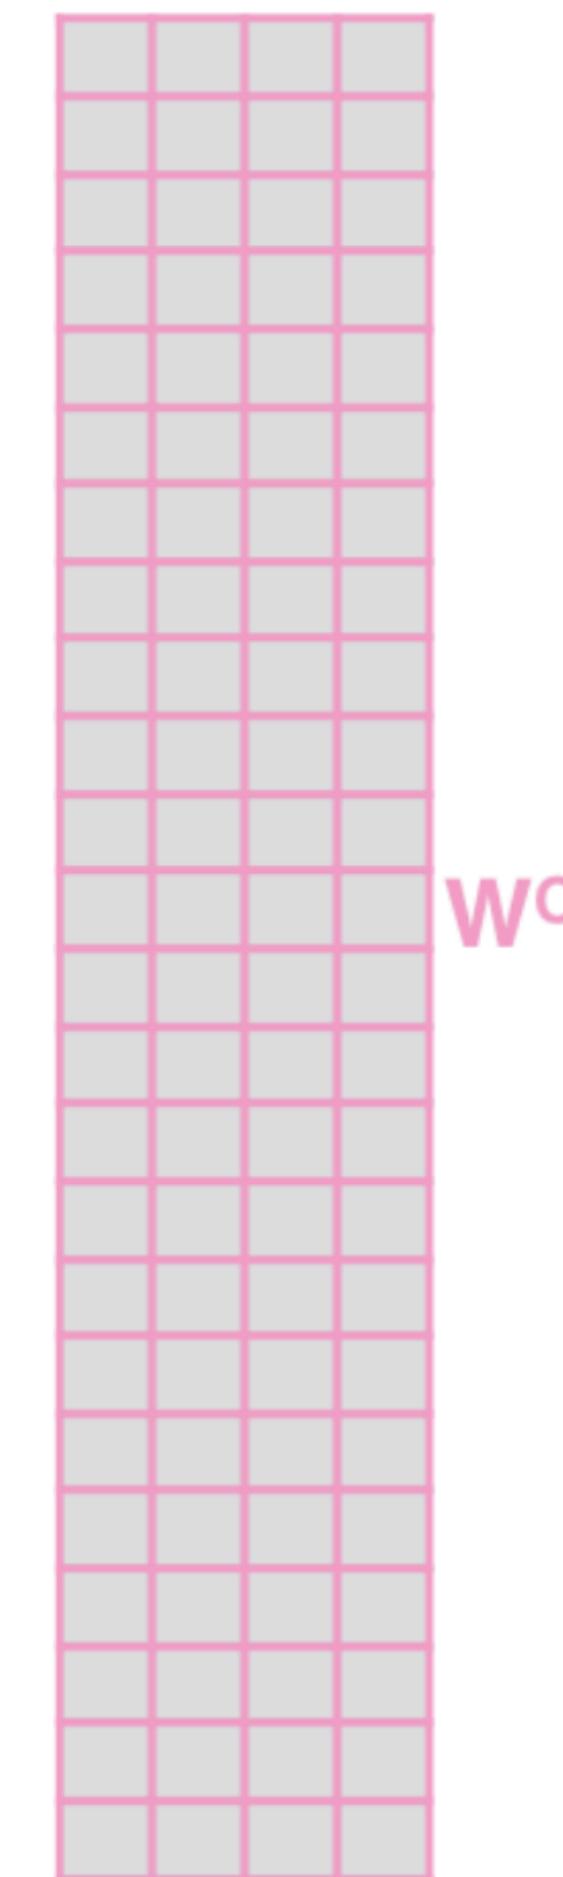
# Multi-Head Self-Attention

1) Concatenate all the attention heads



2) Multiply with a weight matrix  $W^o$  that was trained jointly with the model

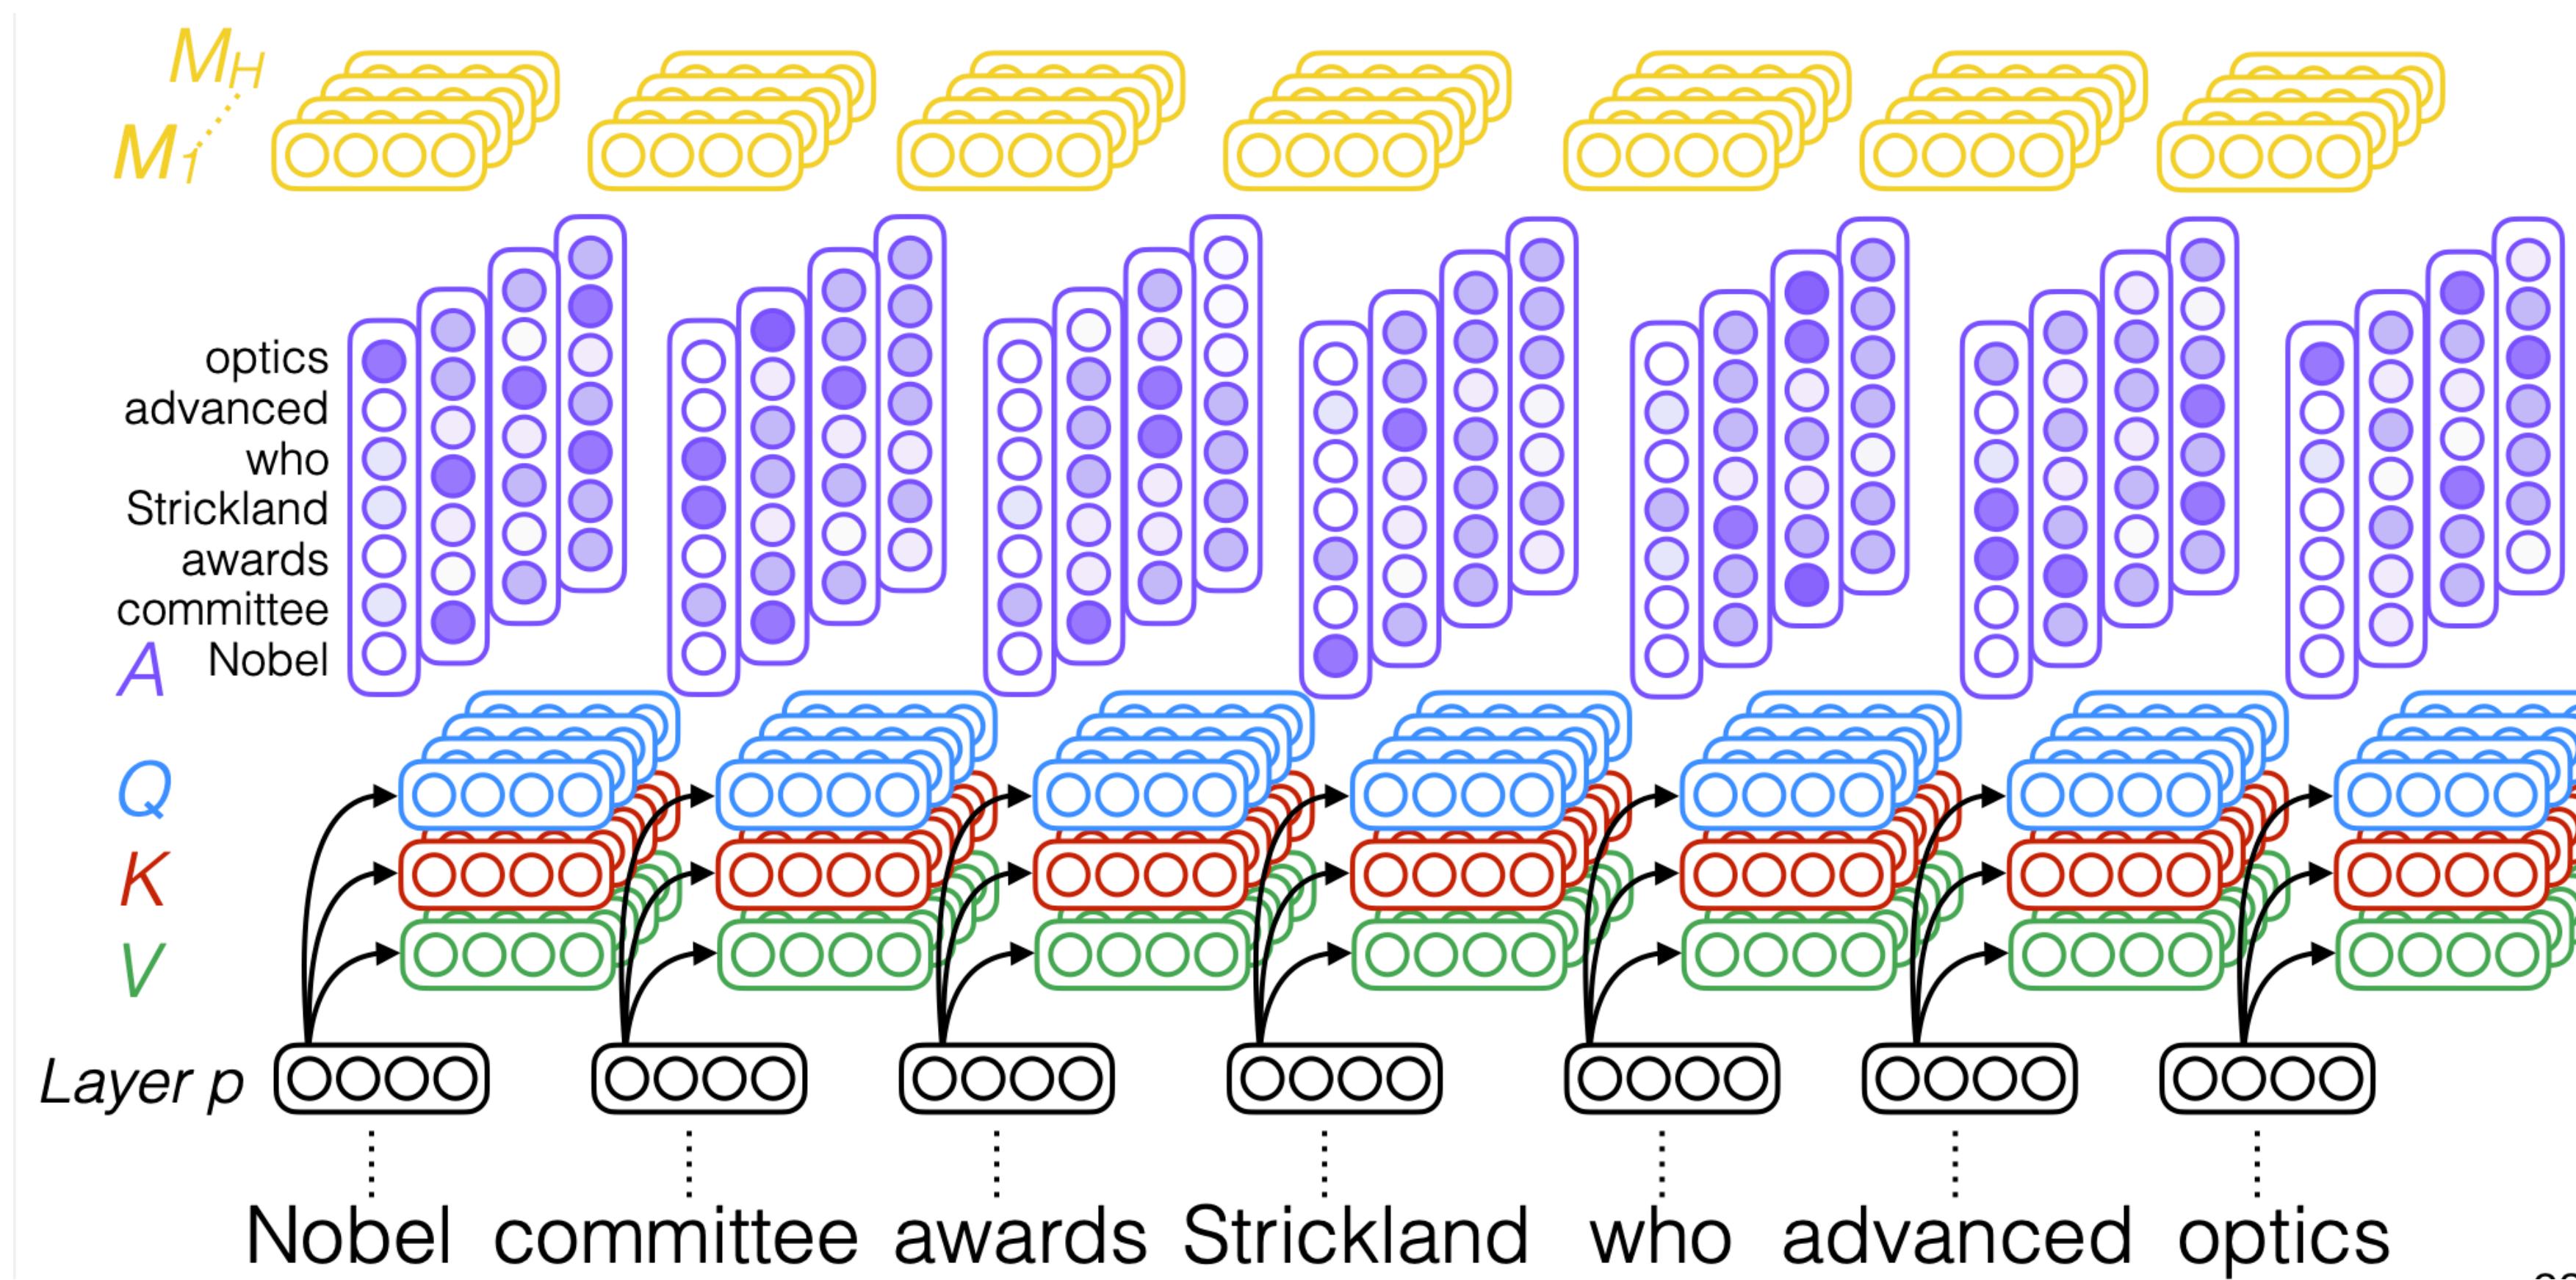
$X$



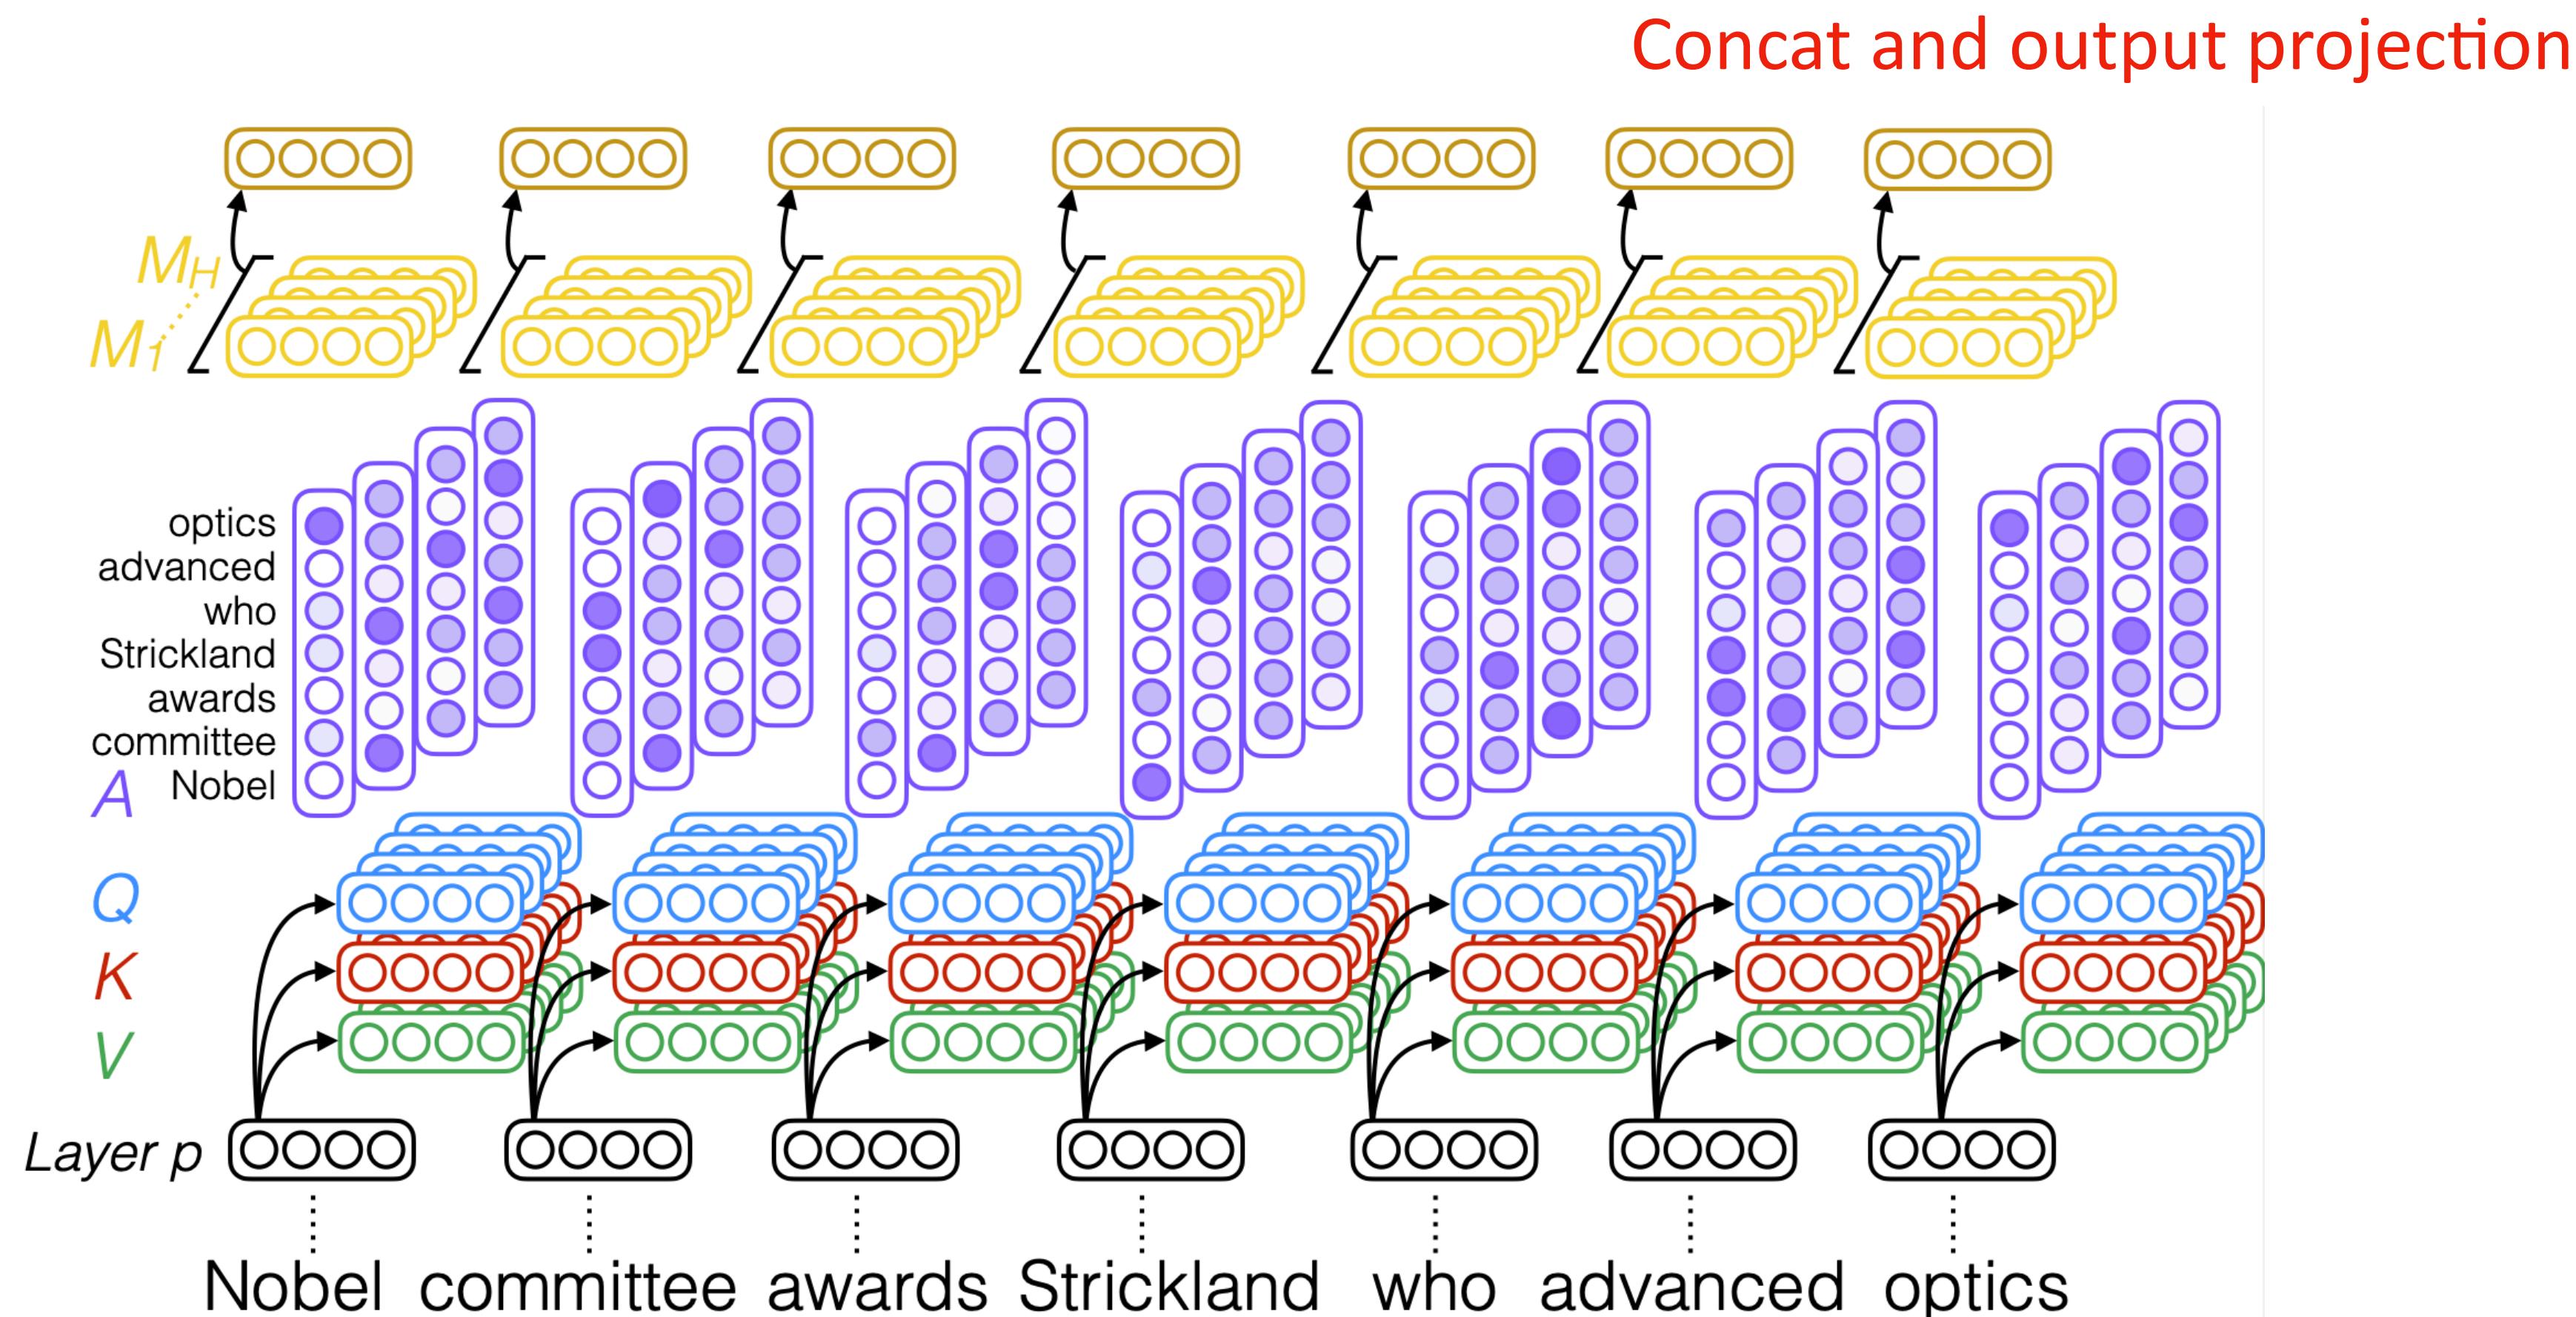
3) The result would be the  $Z$  matrix that captures information from all the attention heads. We can send this forward to the FFNN

$$= \begin{matrix} Z \\ \begin{matrix} \square & \square & \square & \square \end{matrix} \end{matrix}$$

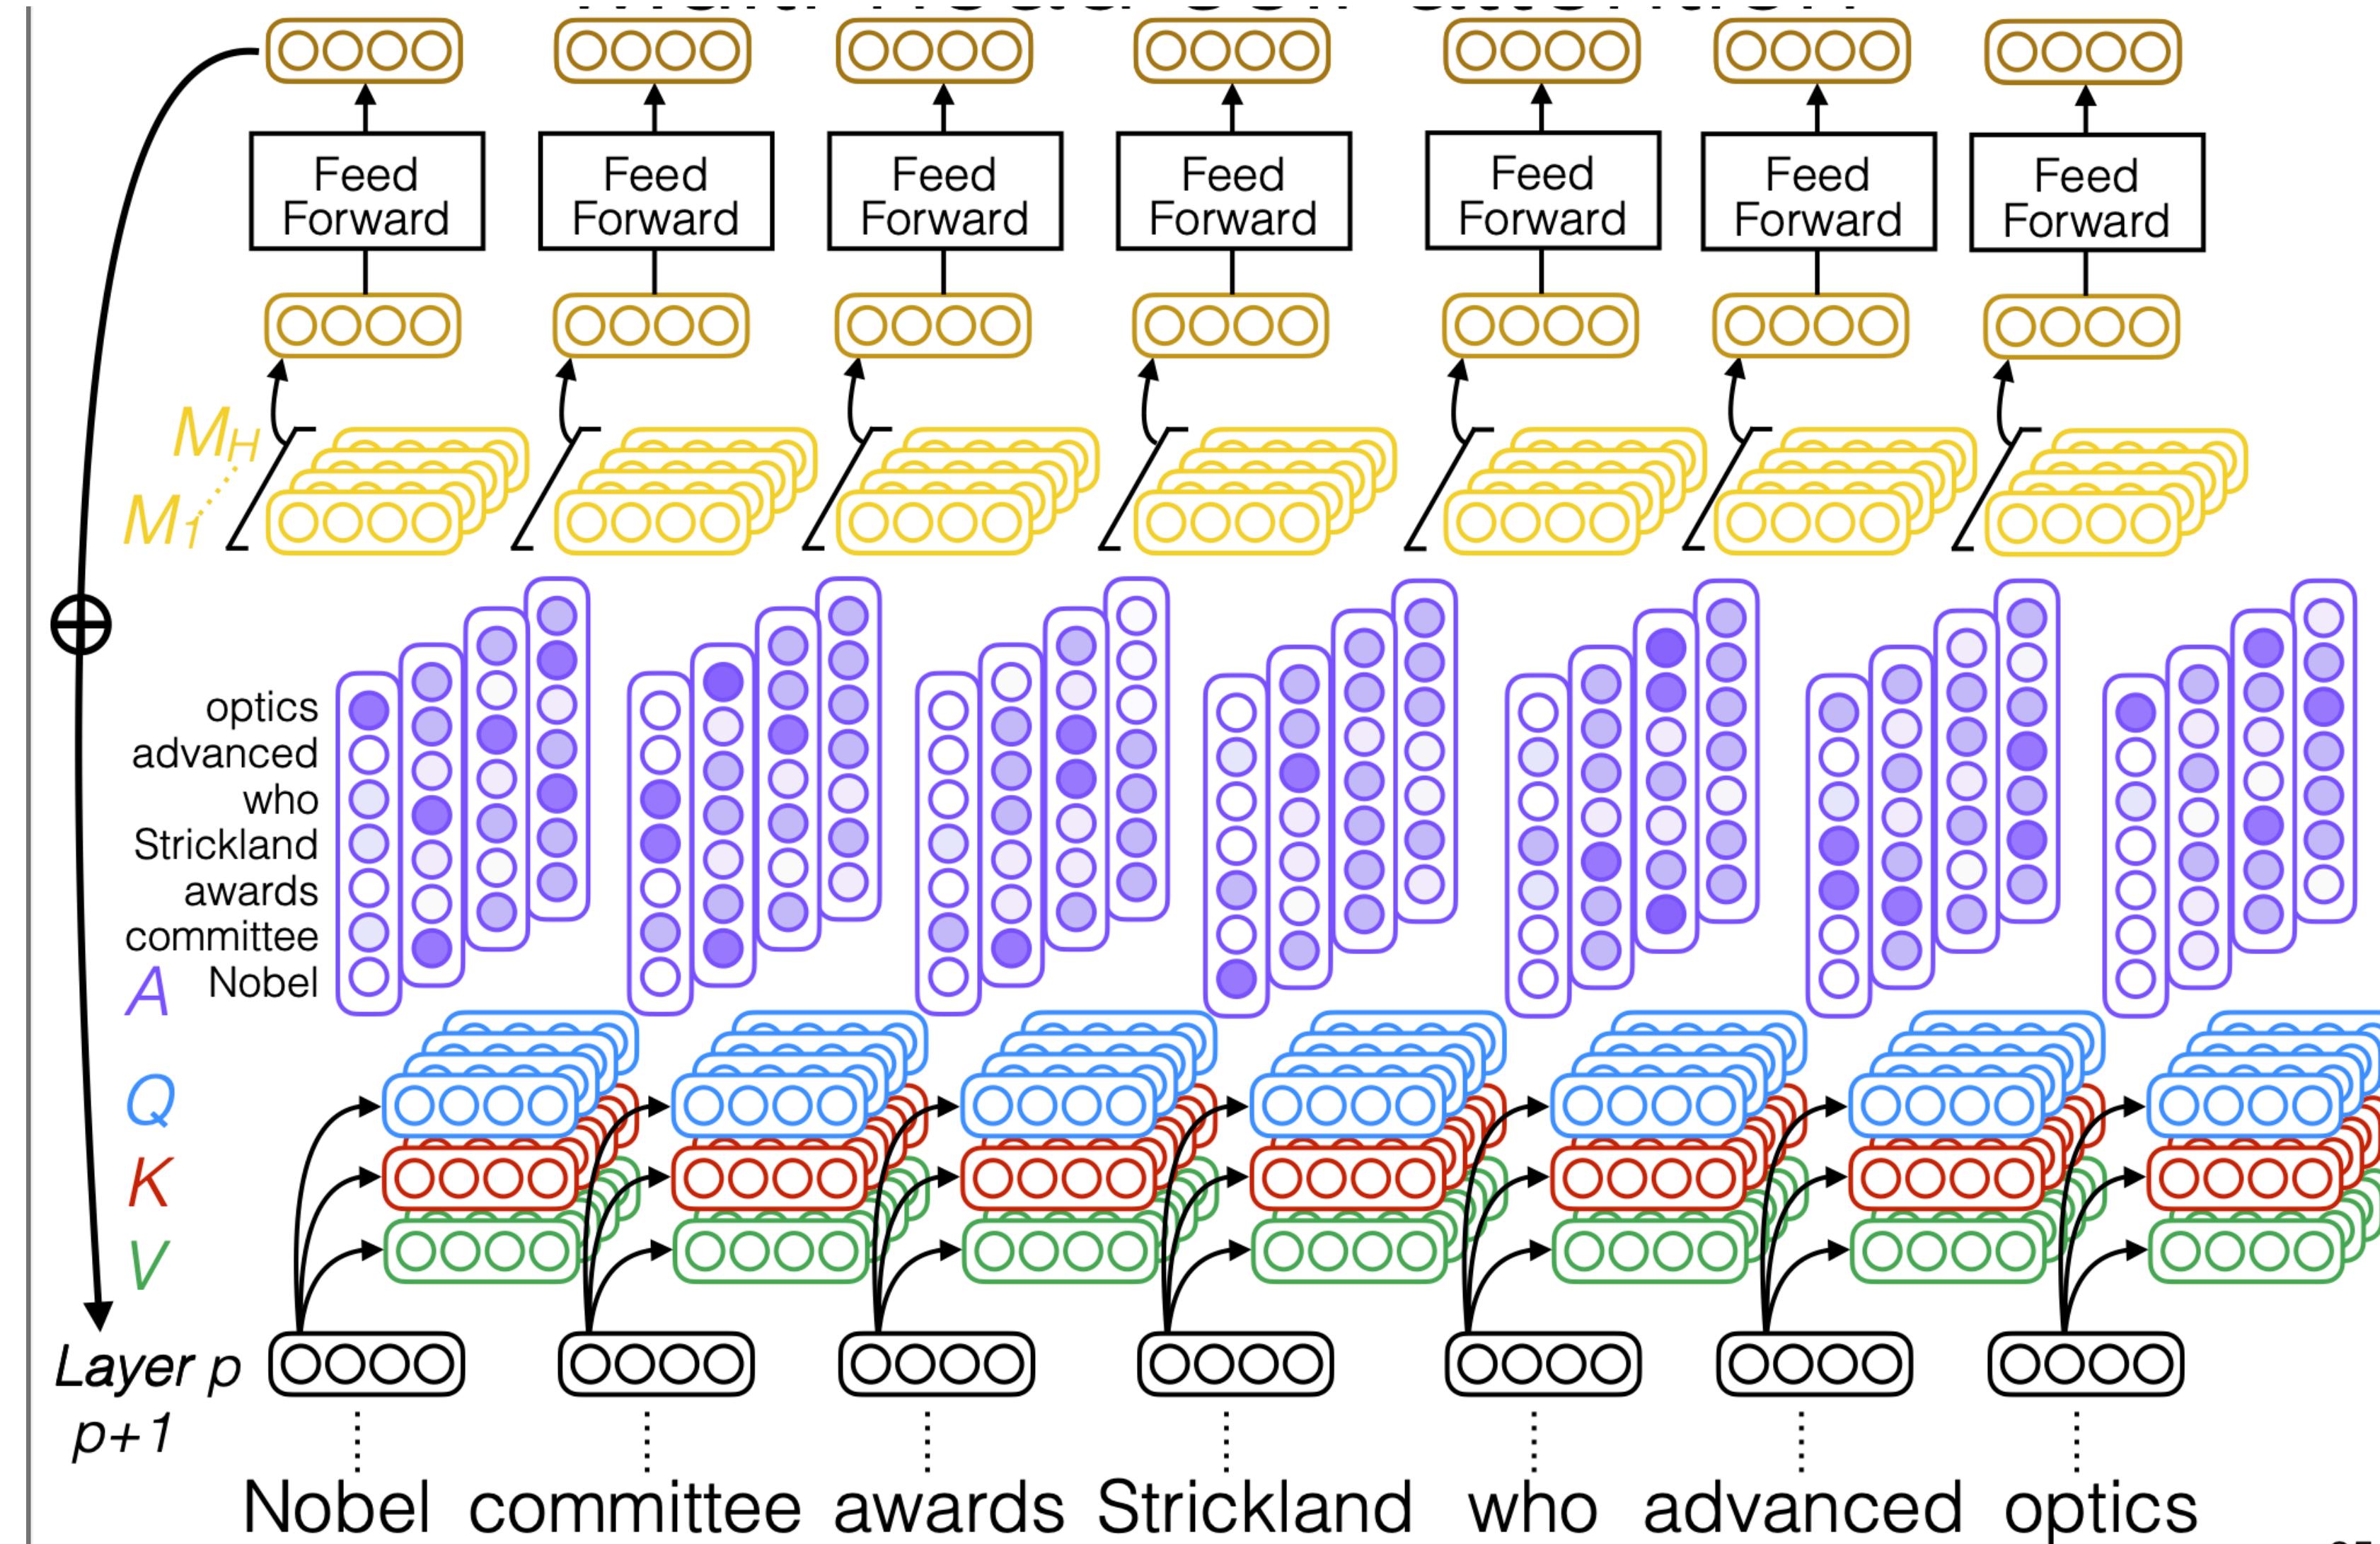
# Multi-head Self-Attention



# Multi-head Self-Attention

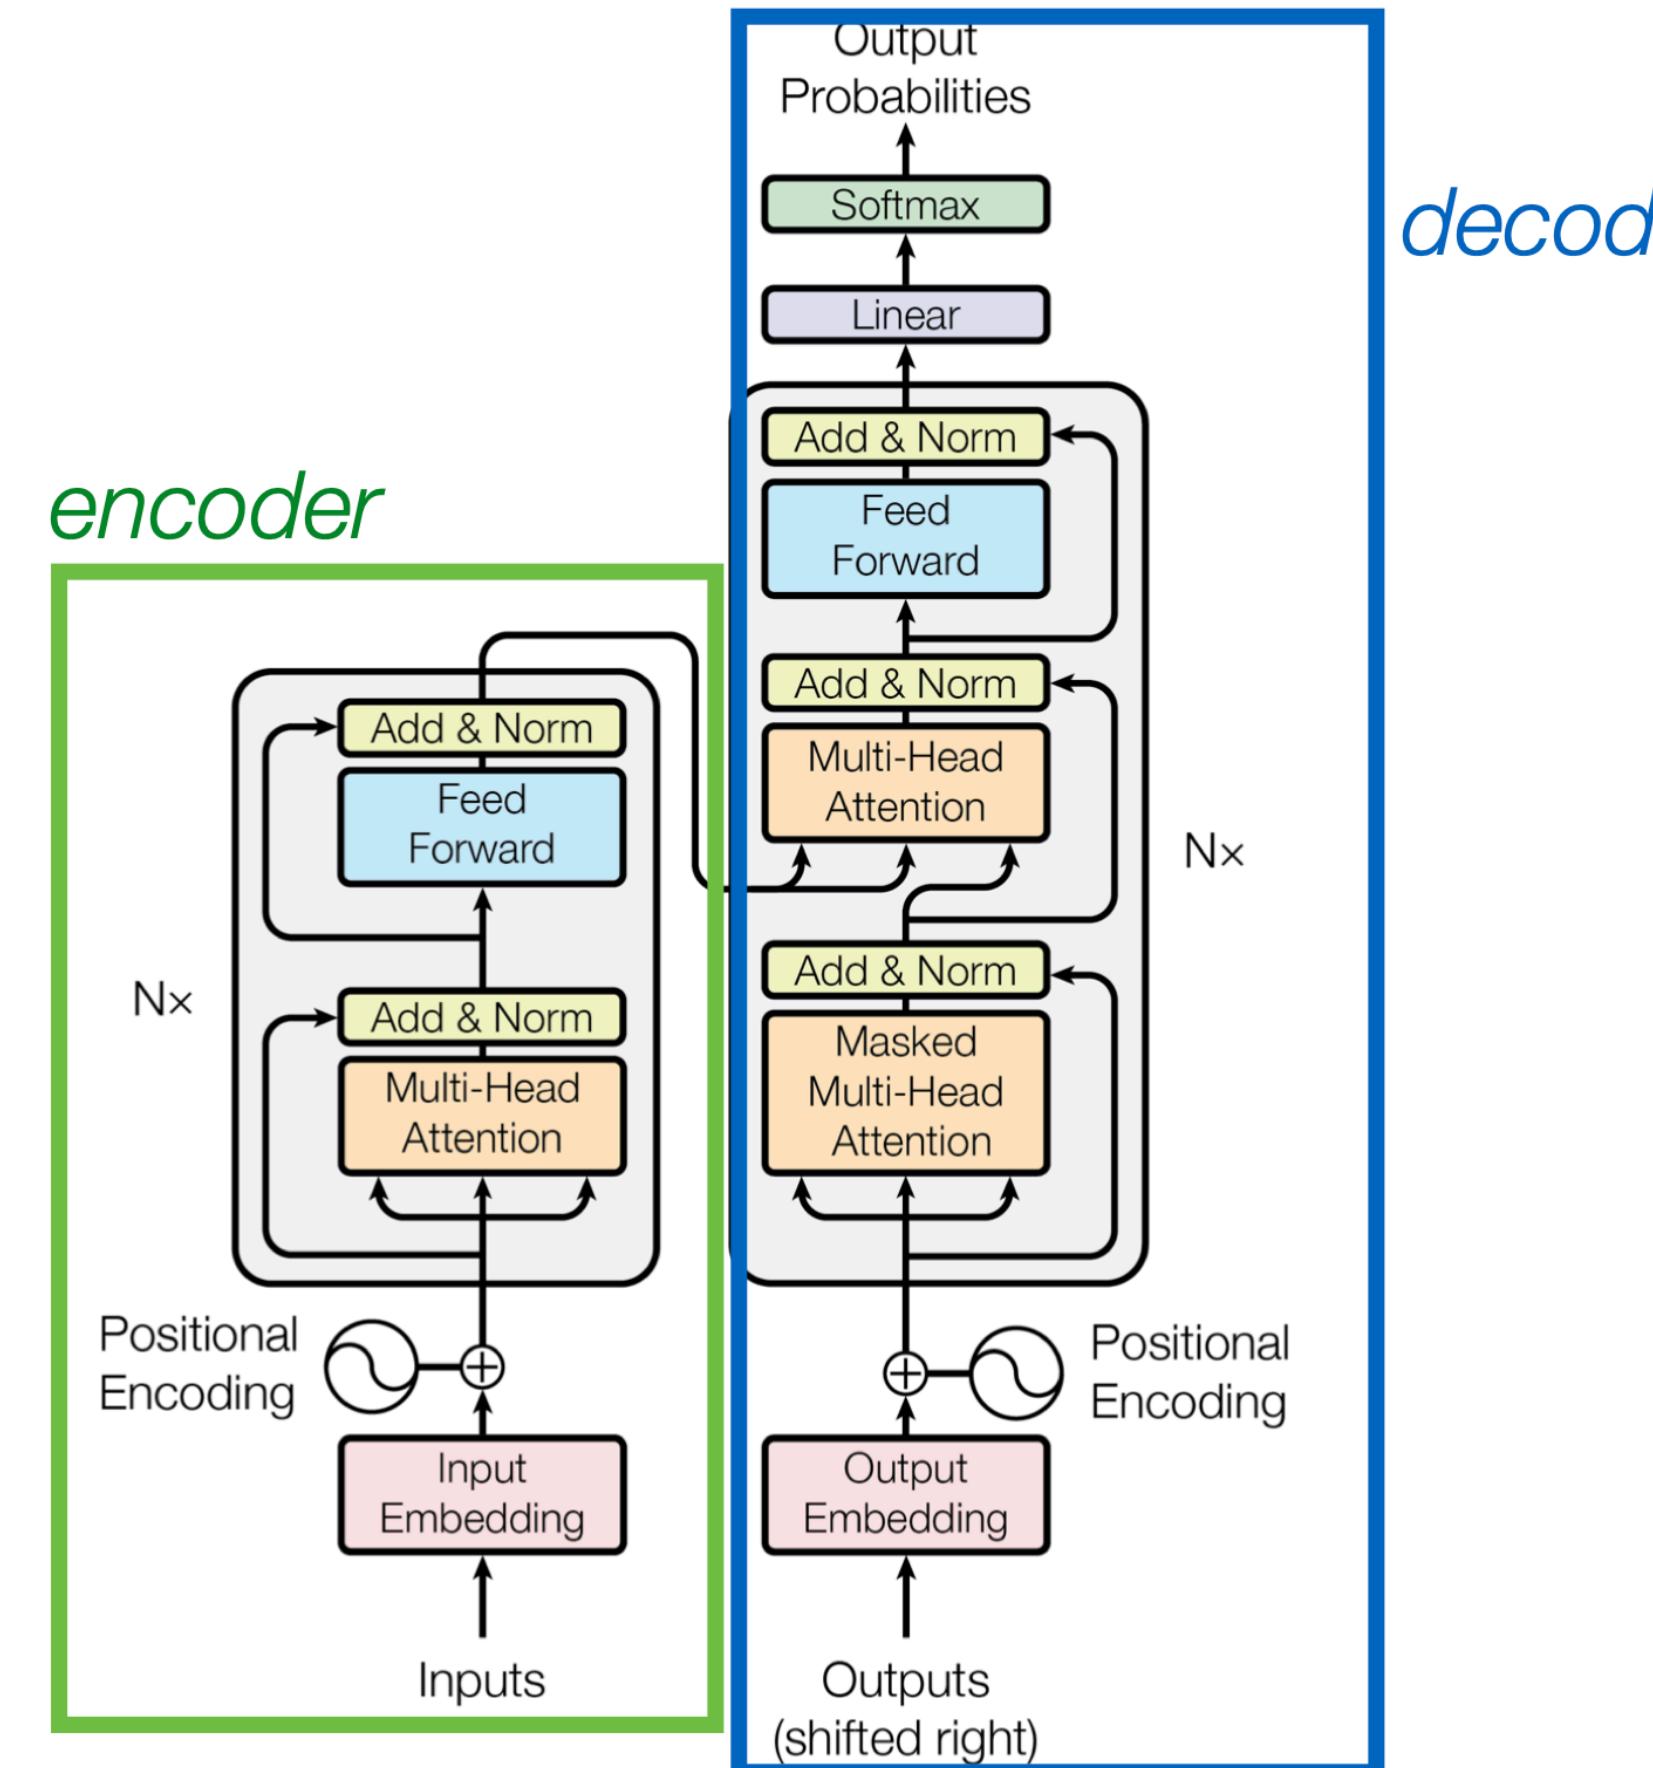
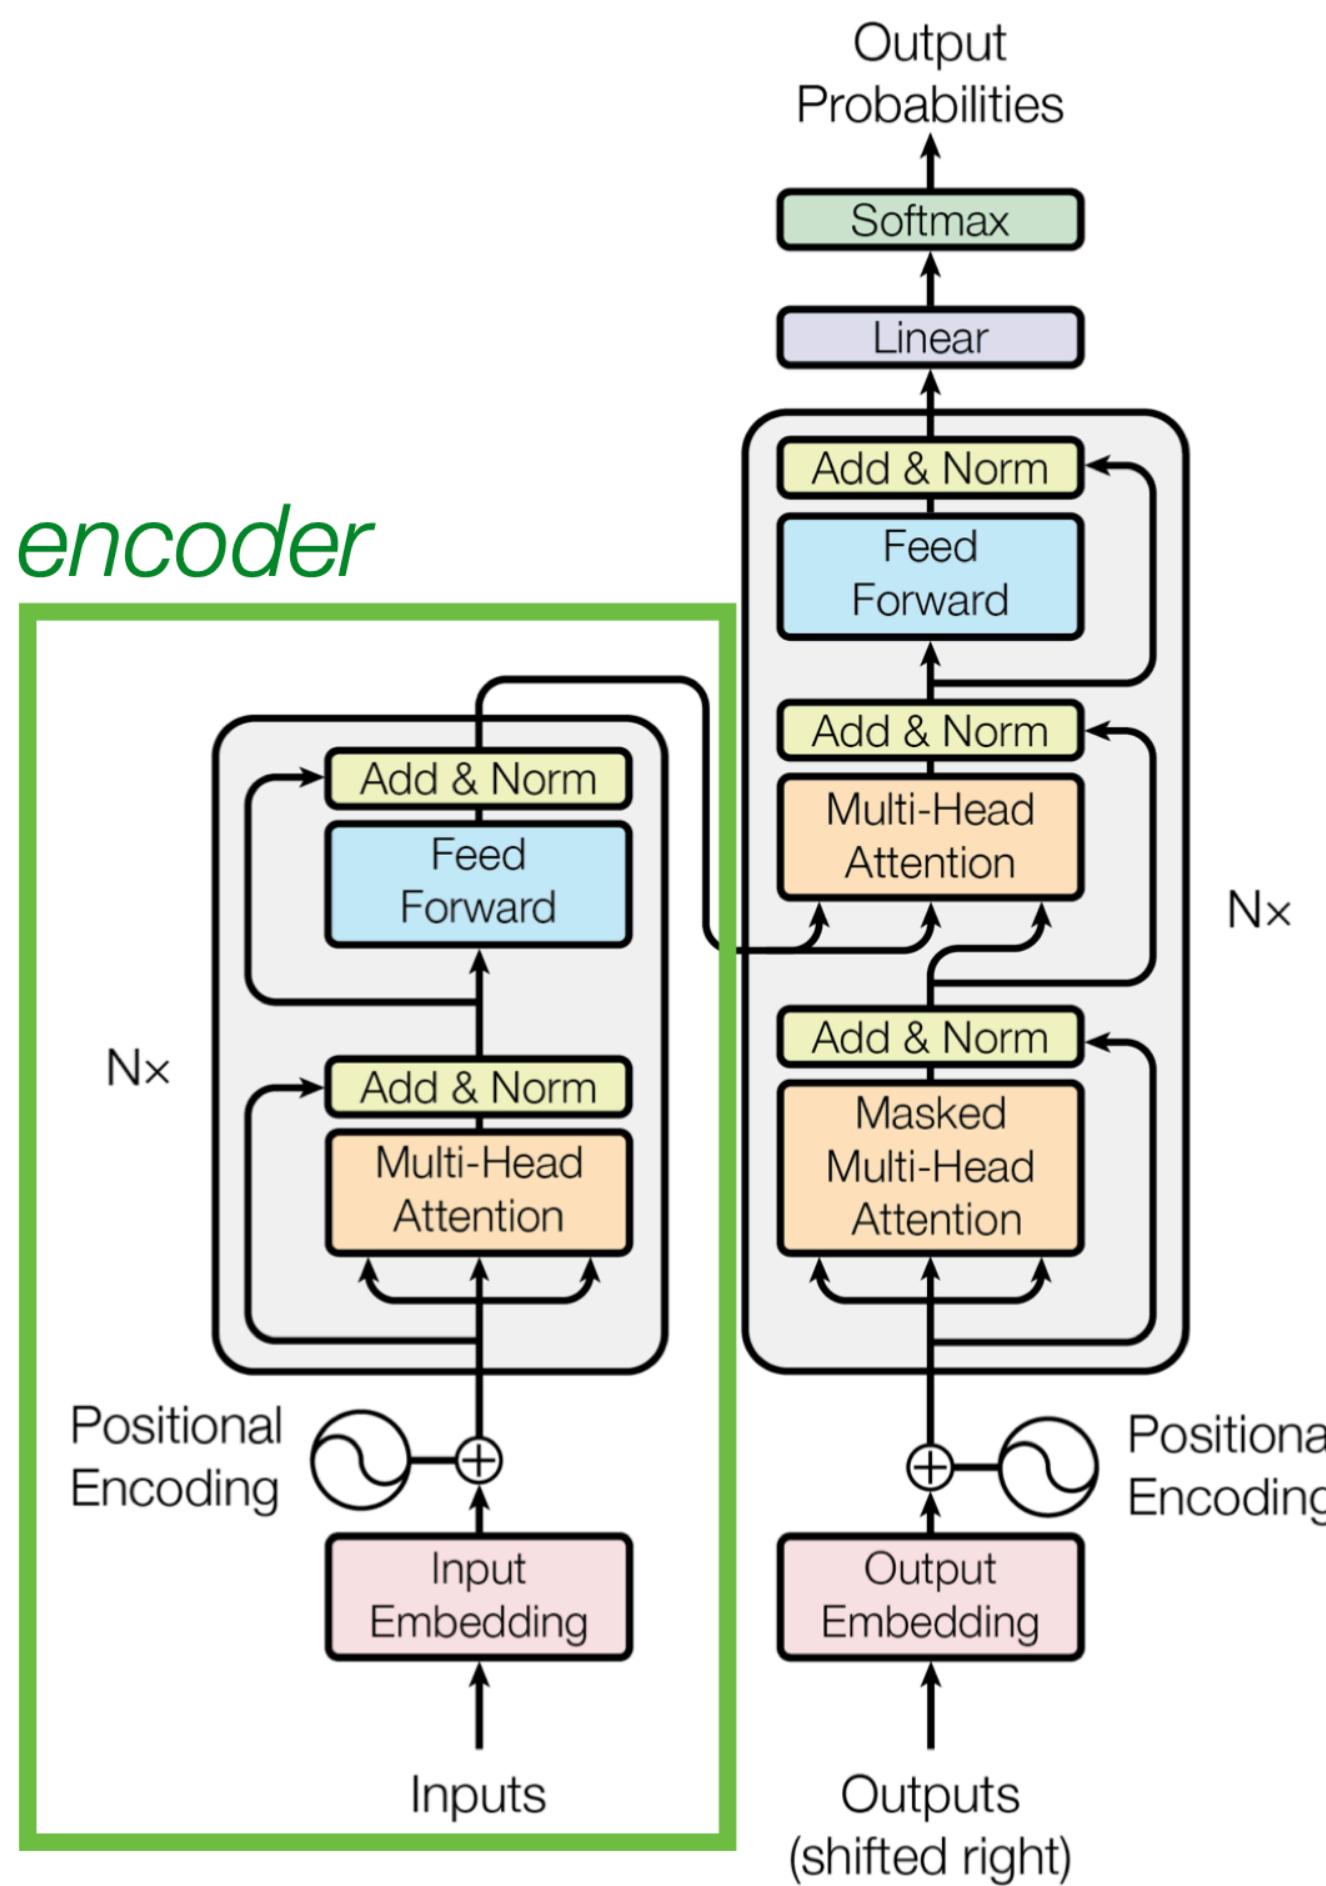


# Multi-head Self-Attention + FFN



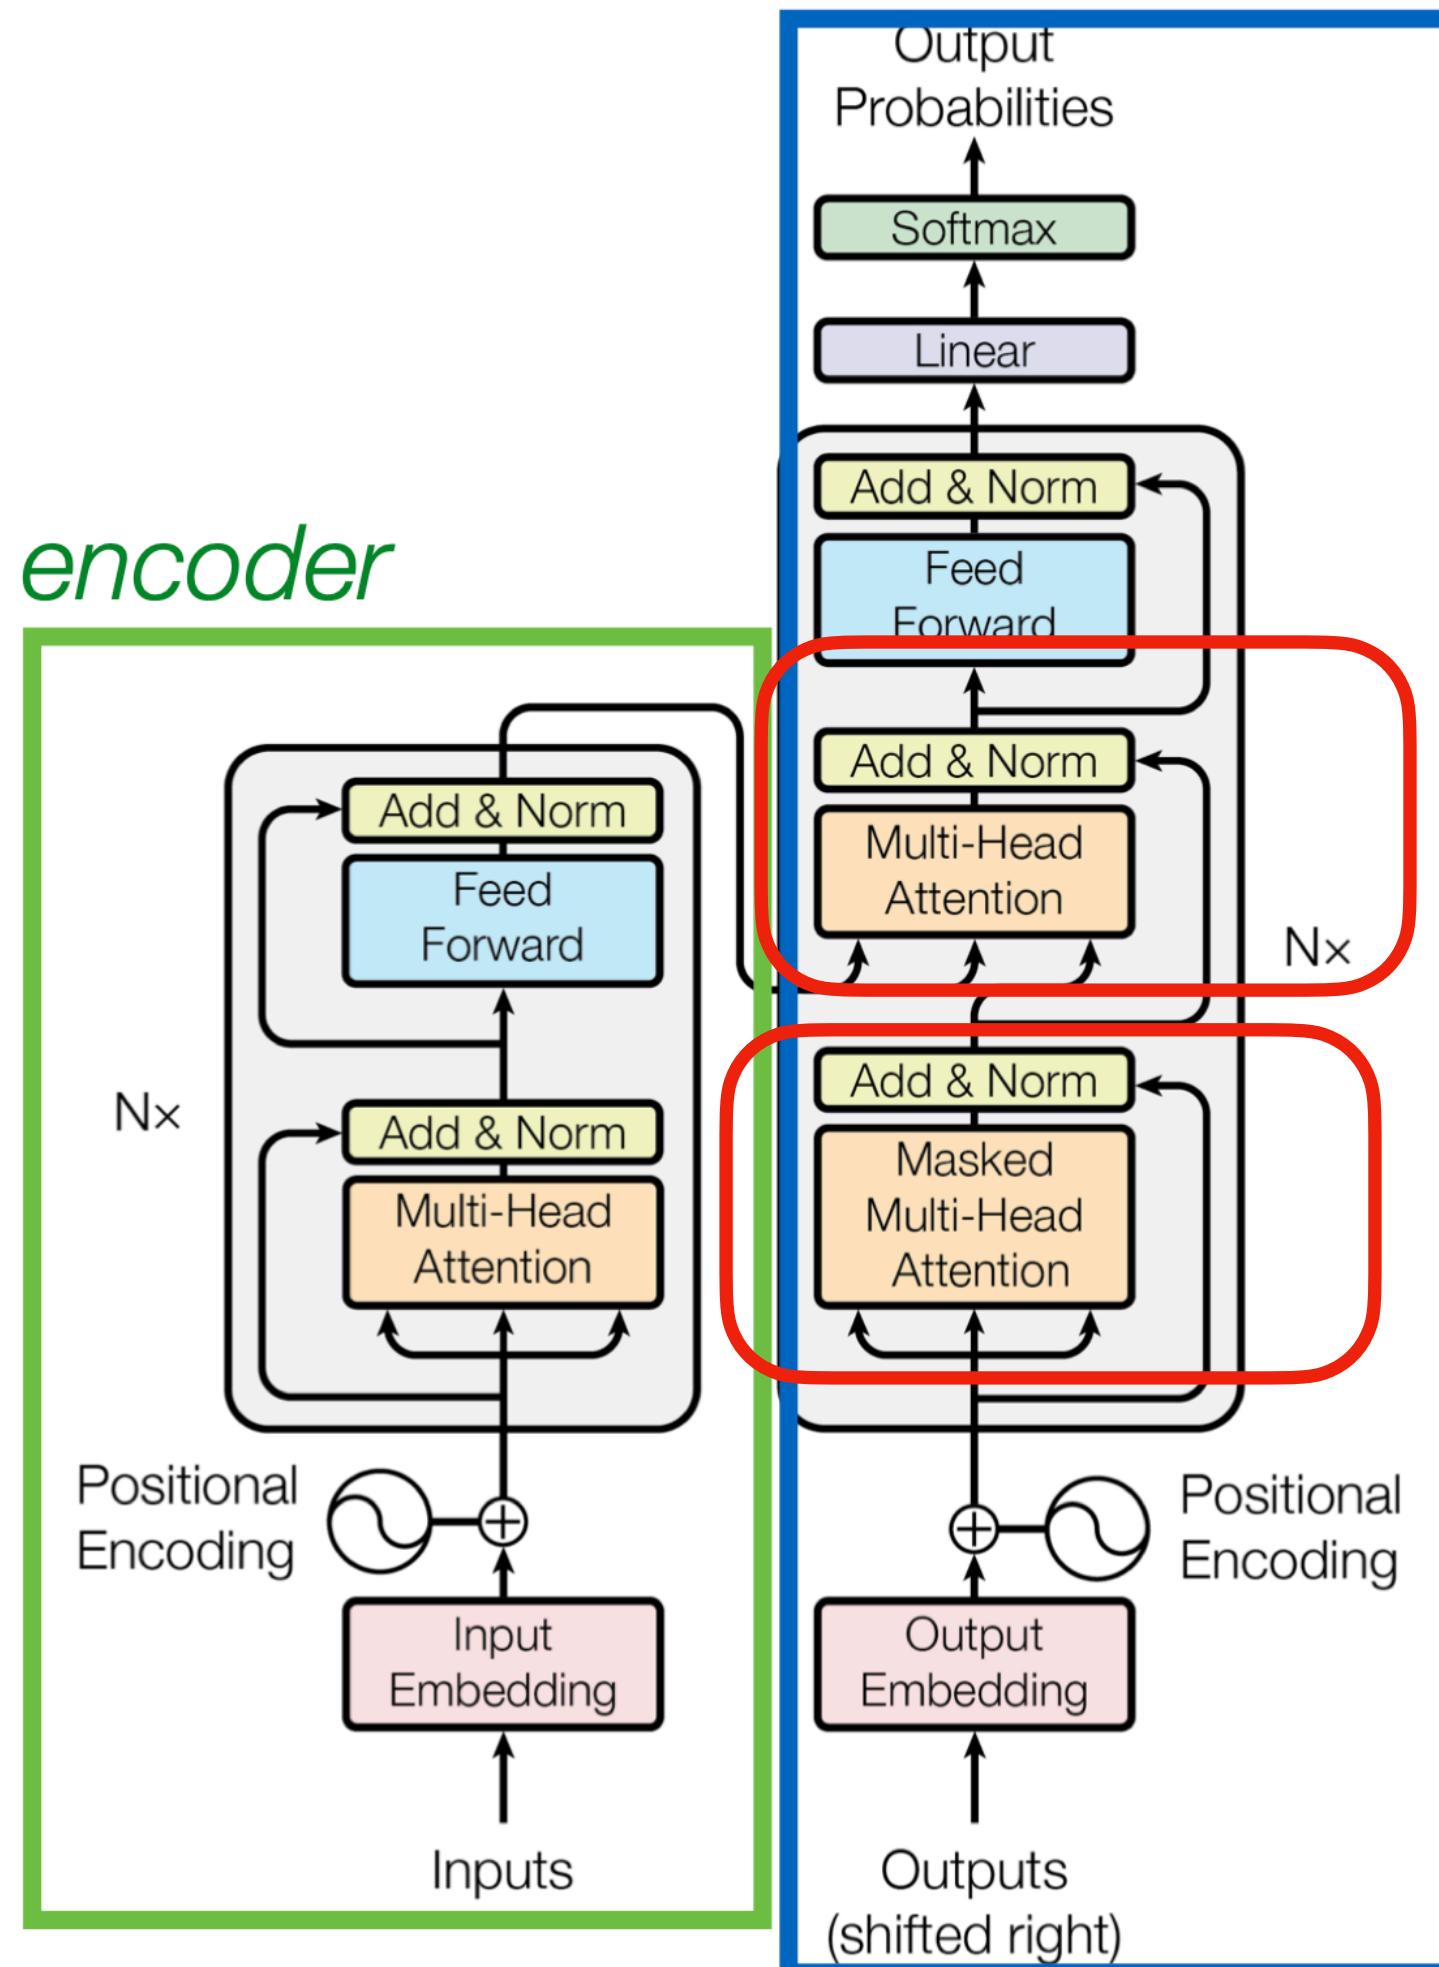
# Transformer Encoder

Currently we only cover the encoder side



This encoder-decoder arch is originally proposed as a seq2seq arch, for classification tasks, often only encoder is used. And language models often only have a decoder

# Transformer Decoder in Seq2Seq

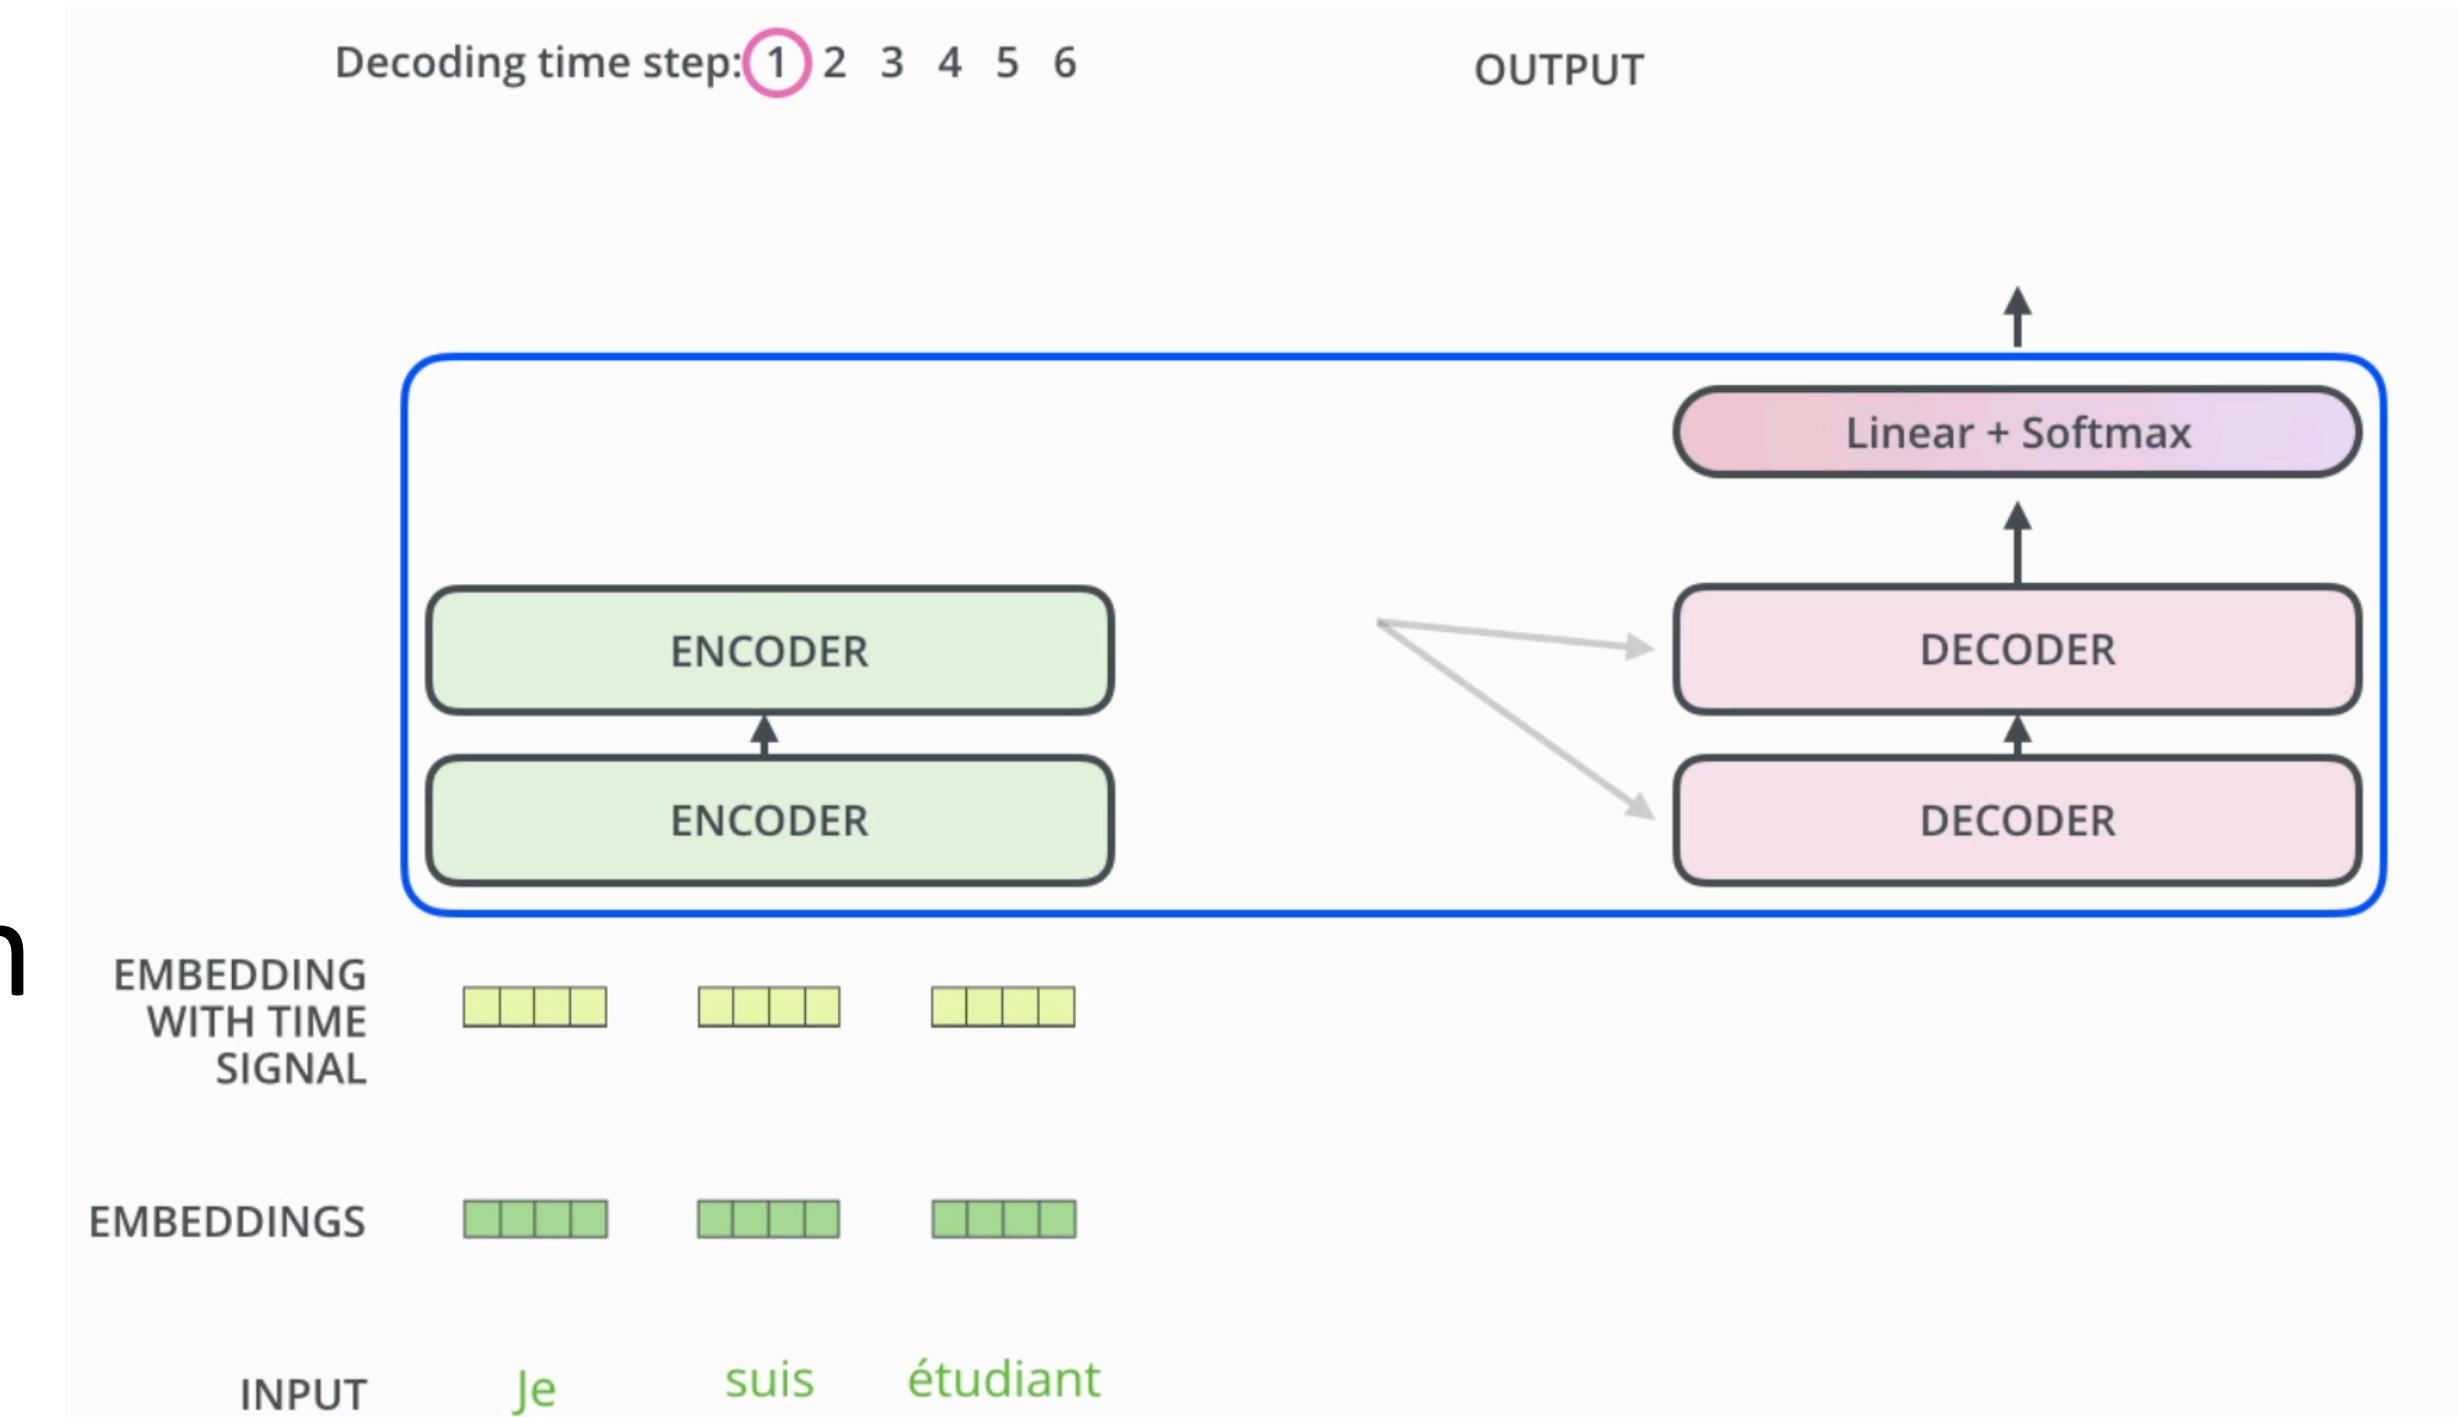


*decoder*

Cross-attention

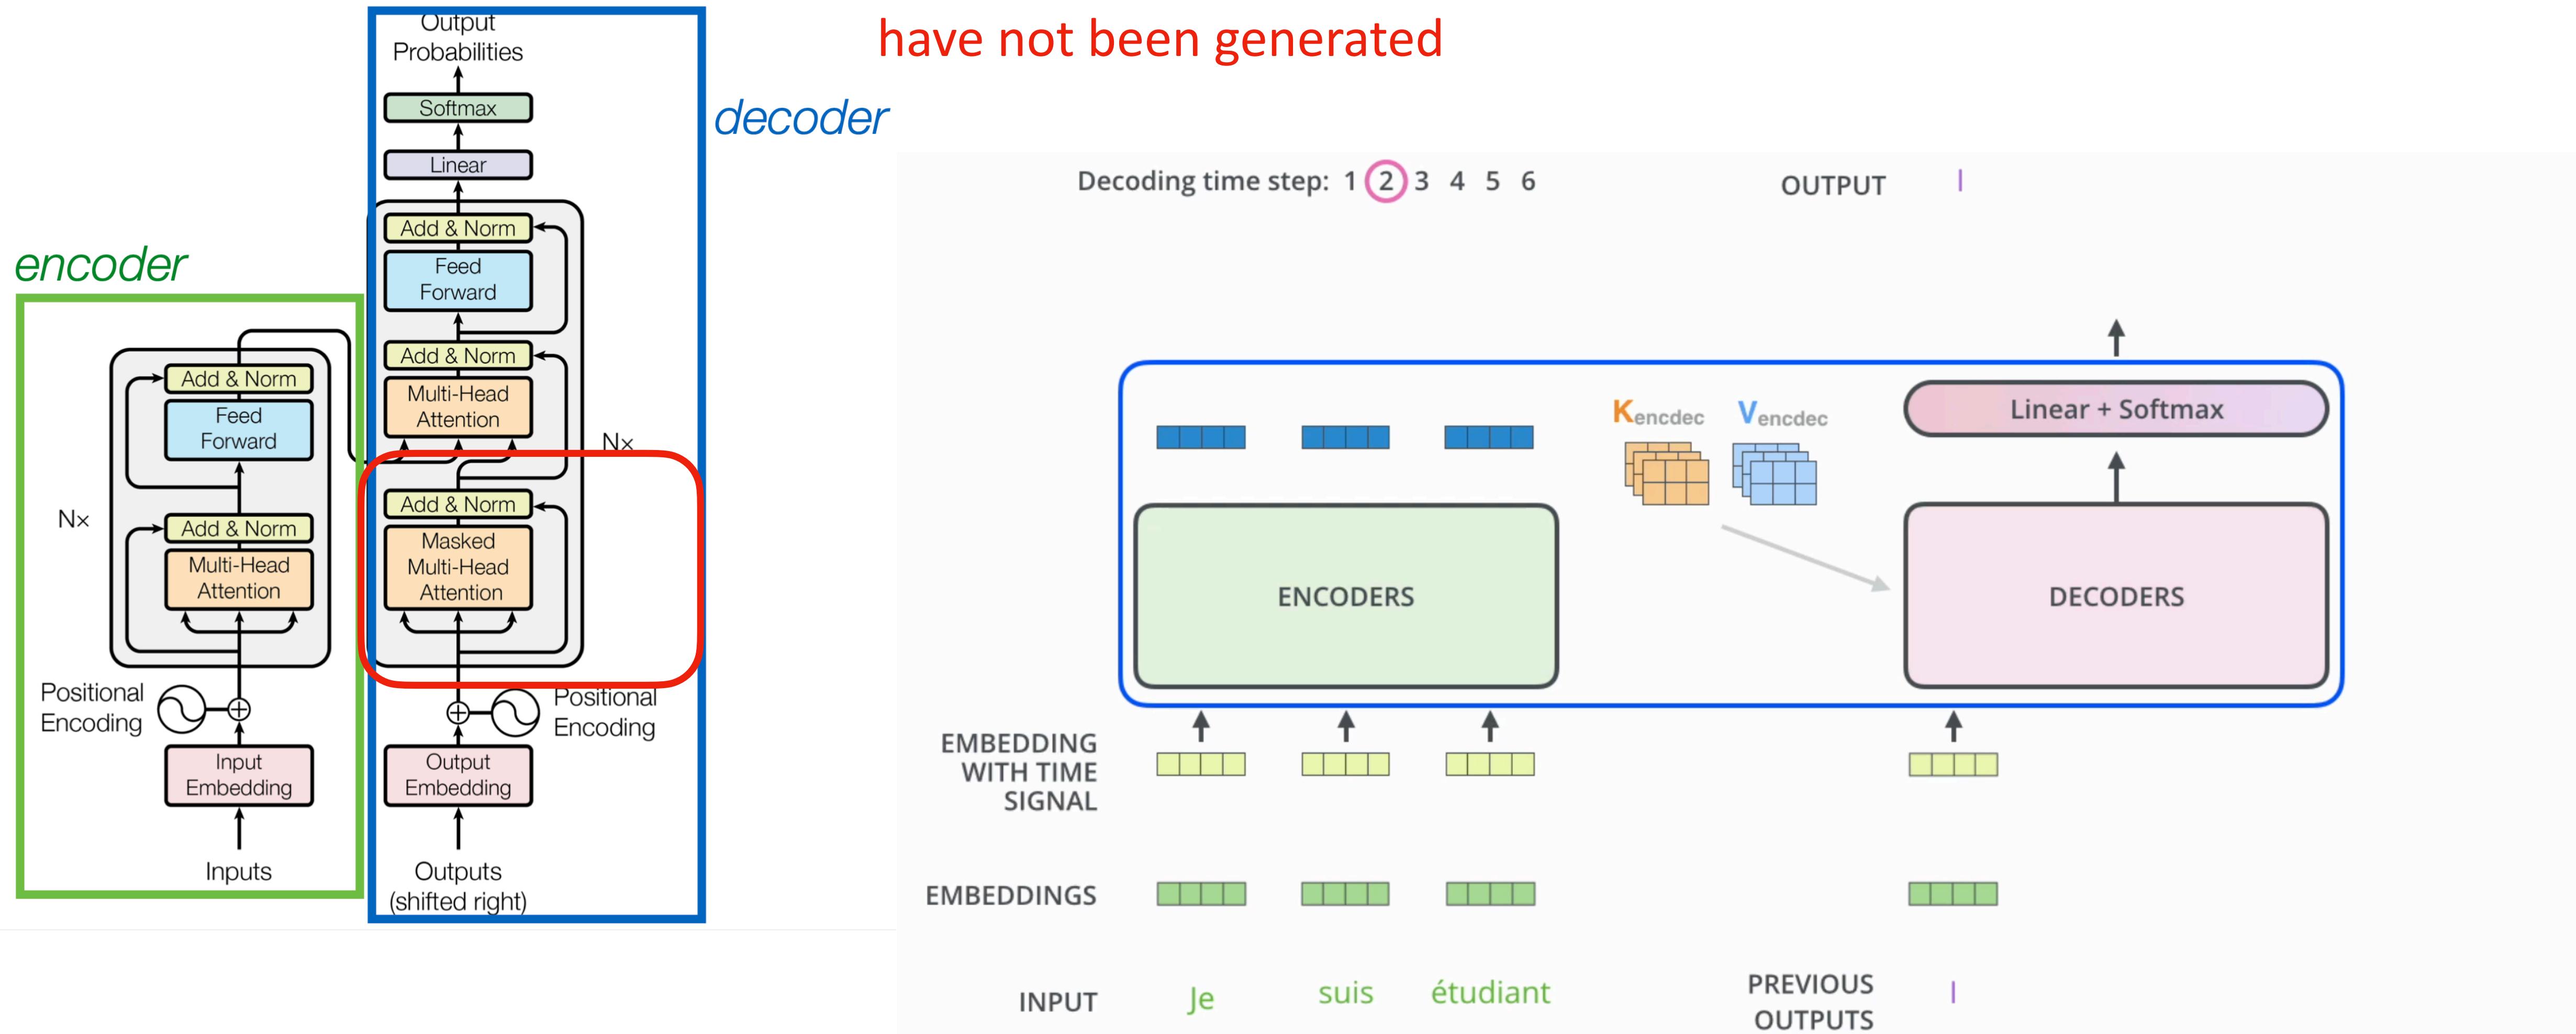
Self-attention

Cross-attention uses the output of encoder as input

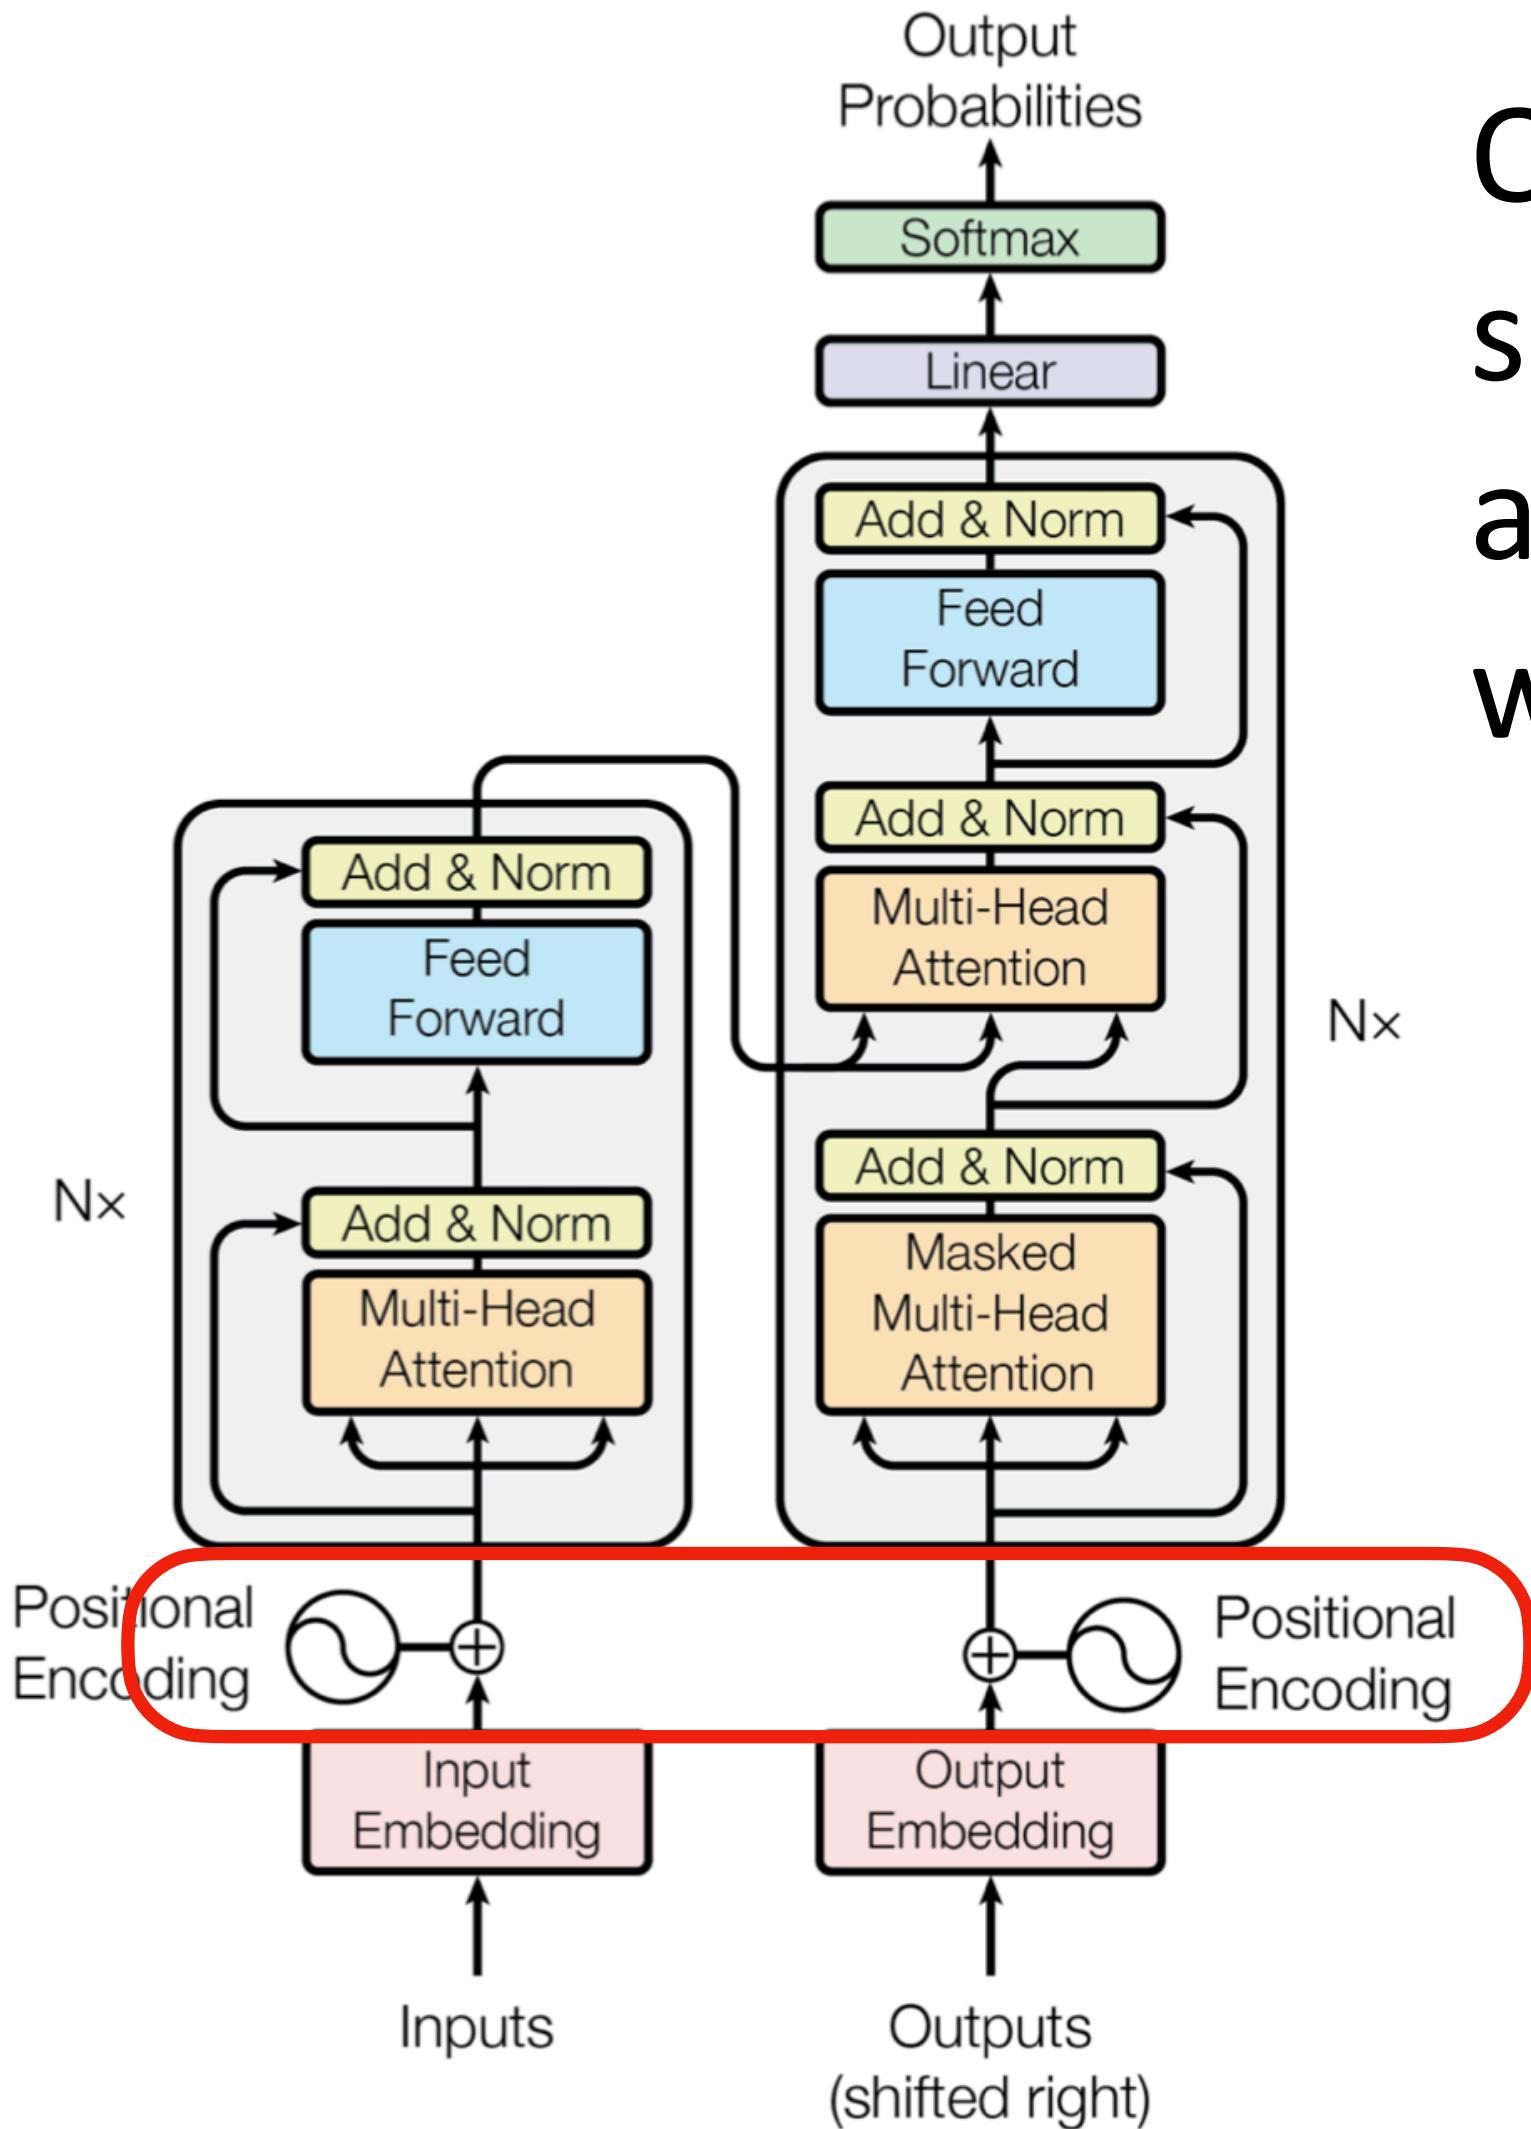


# Masked Attention

Typical attention attends to the entire sequence, while masked attention only attends to the ones on the left because future words have not been generated



# Position Embeddings



Question: If we shuffle the order of words in the sequence, will that change the attention output and feed forward output of the corresponding word?

Position embeddings are added to each word embedding, otherwise our model is unaware of the position of a word

# Positional Encoding

EMBEDDING  
WITH TIME  
SIGNAL

$$\mathbf{x}_1 \quad \begin{array}{|c|c|c|c|} \hline \text{light green} & \text{light green} & \text{light green} & \text{light green} \\ \hline \end{array}$$

=

POSITIONAL  
ENCODING

$$\mathbf{t}_1 \quad \begin{array}{|c|c|c|c|} \hline \text{yellow} & \text{yellow} & \text{yellow} & \text{yellow} \\ \hline \end{array}$$

+

EMBEDDINGS

$$\mathbf{x}_1 \quad \begin{array}{|c|c|c|c|} \hline \text{green} & \text{green} & \text{green} & \text{green} \\ \hline \end{array}$$

$$\mathbf{x}_2 \quad \begin{array}{|c|c|c|c|} \hline \text{light green} & \text{light green} & \text{light green} & \text{light green} \\ \hline \end{array}$$

=

$$\mathbf{t}_2 \quad \begin{array}{|c|c|c|c|} \hline \text{yellow} & \text{yellow} & \text{yellow} & \text{yellow} \\ \hline \end{array}$$

+

$$\mathbf{x}_3 \quad \begin{array}{|c|c|c|c|} \hline \text{light green} & \text{light green} & \text{light green} & \text{light green} \\ \hline \end{array}$$

=

$$\mathbf{t}_3 \quad \begin{array}{|c|c|c|c|} \hline \text{yellow} & \text{yellow} & \text{yellow} & \text{yellow} \\ \hline \end{array}$$

+

$$\mathbf{x}_3 \quad \begin{array}{|c|c|c|c|} \hline \text{green} & \text{green} & \text{green} & \text{green} \\ \hline \end{array}$$

INPUT

Je

suis

étudiant

# Transformer Positional Encoding

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

Positional encoding is a 512d vector  
 $i$  = a particular dimension of this vector  
 $pos$  = dimension of the word  
 $d_{model}$  = 512

# Complexity

| Layer Type                  | Complexity per Layer     | Sequential Operations |
|-----------------------------|--------------------------|-----------------------|
| Self-Attention              | $O(n^2 \cdot d)$         | $O(1)$                |
| Recurrent                   | $O(n \cdot d^2)$         | $O(n)$                |
| Convolutional               | $O(k \cdot n \cdot d^2)$ | $O(1)$                |
| Self-Attention (restricted) | $O(r \cdot n \cdot d)$   | $O(1)$                |

n is sequence length, d is embedding dimension.

Restricted self-attention means not attending all words in the sequence, but only a restricted field

Square complexity of sequence length is a major issue for transformers to deal with long sequence

# Thank You!