



香港科技大學  
THE HONG KONG  
UNIVERSITY OF SCIENCE  
AND TECHNOLOGY

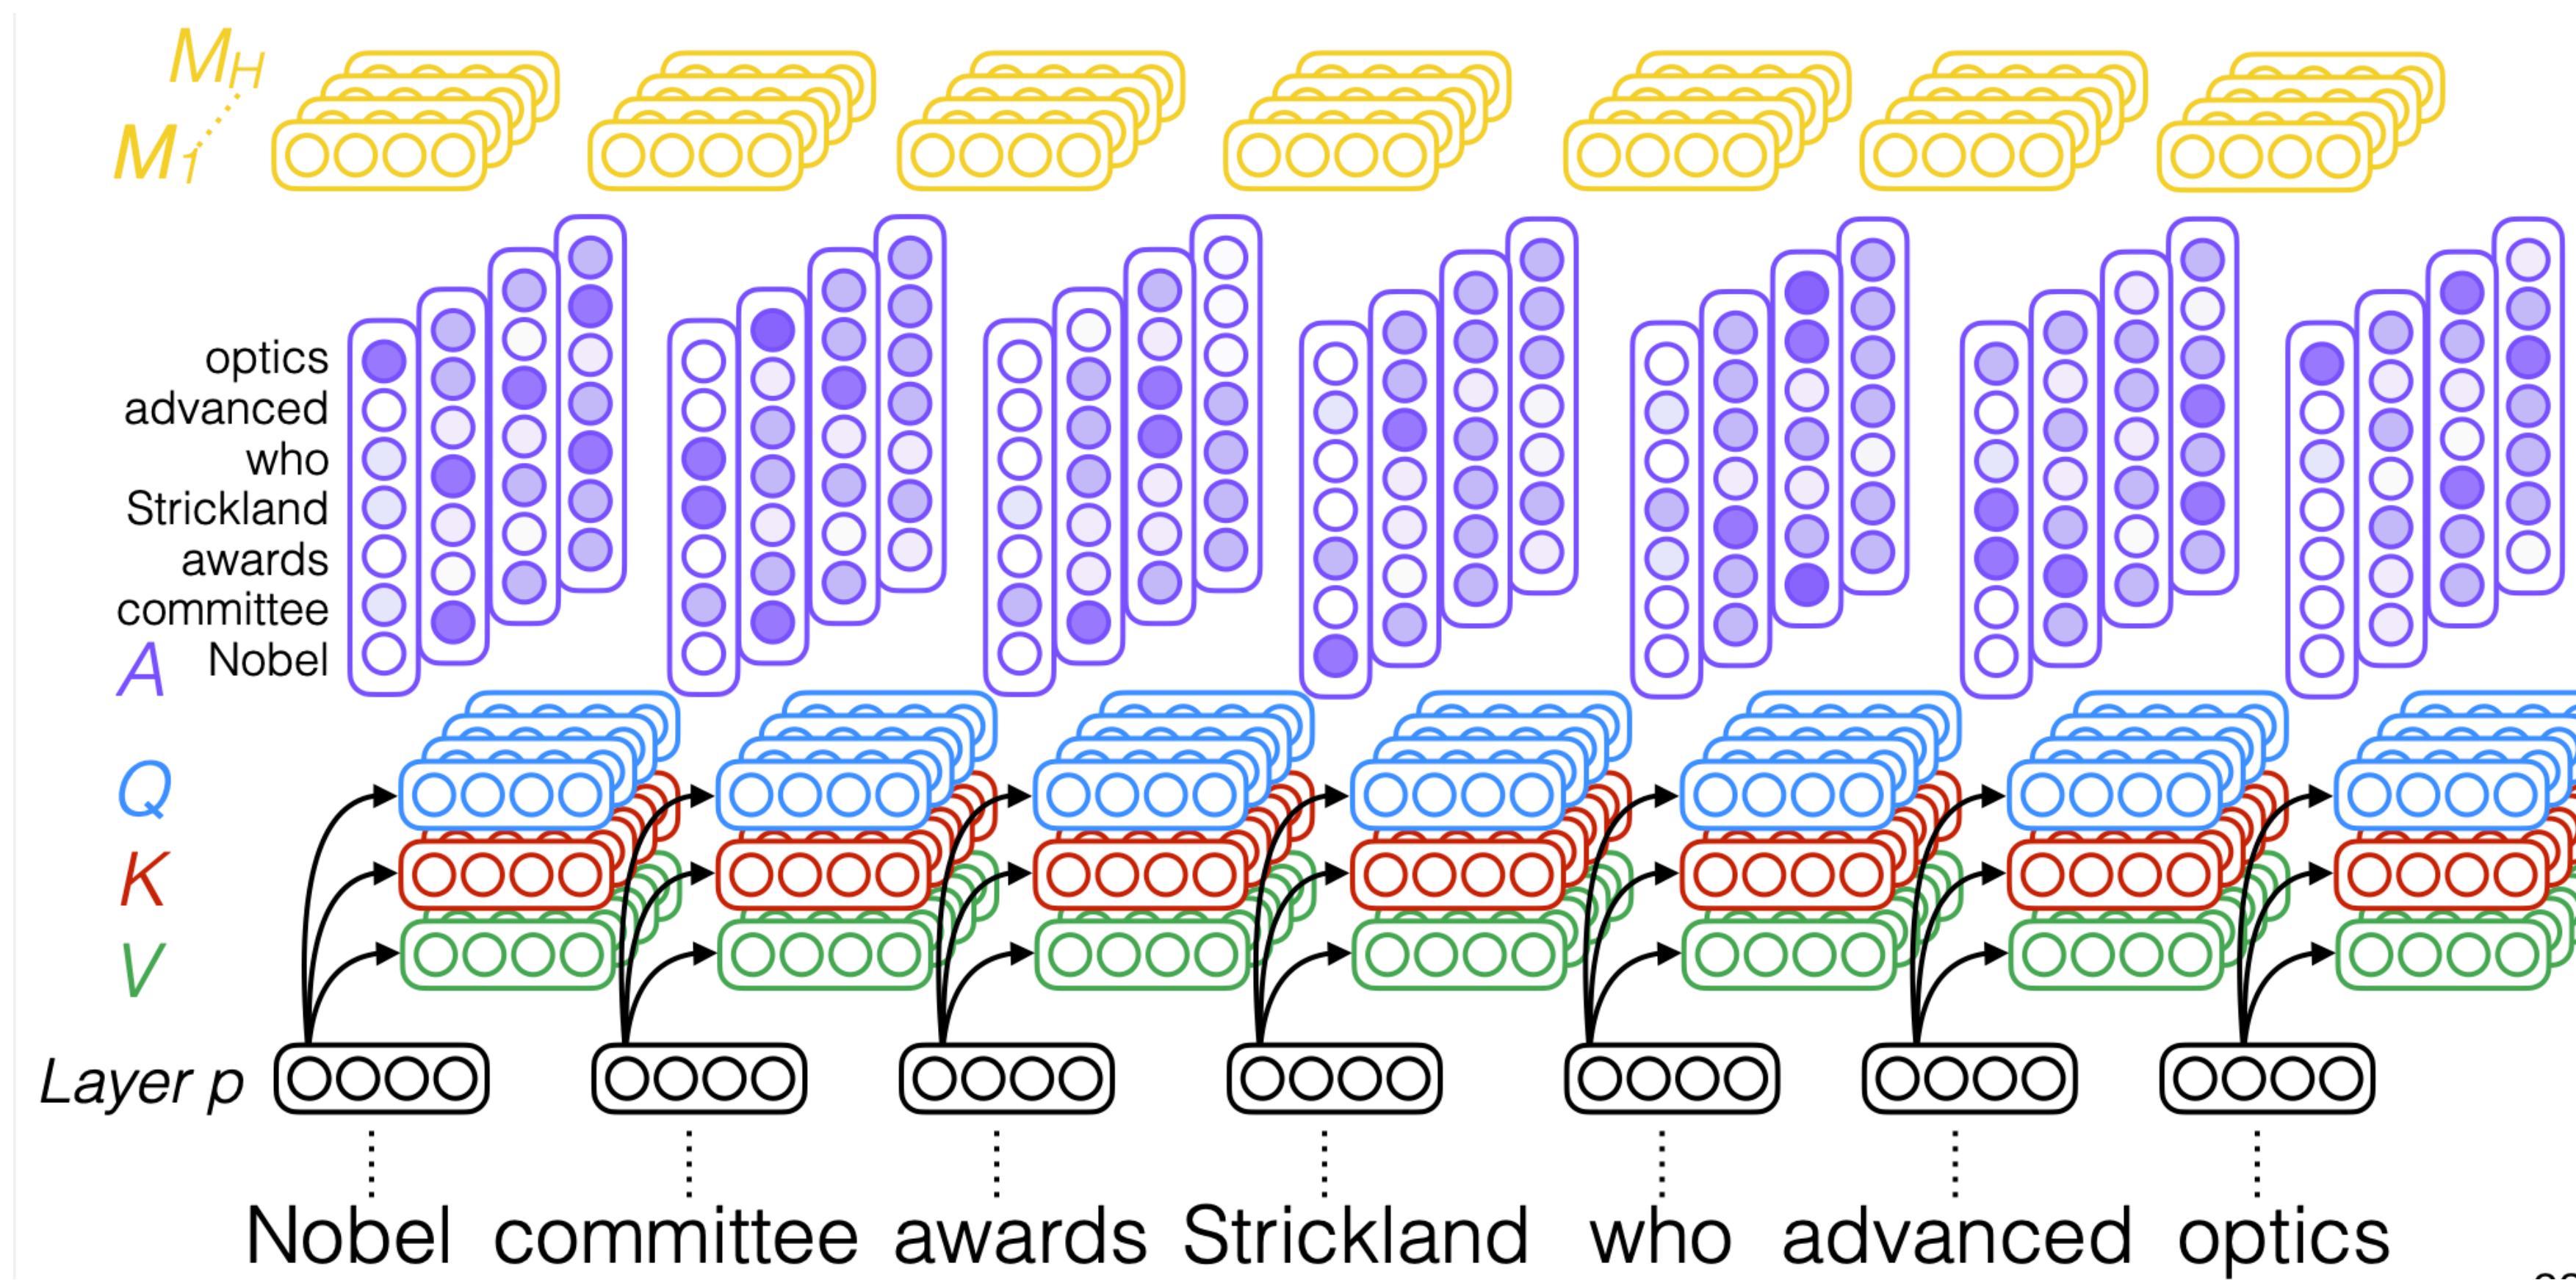
COMP 4901B  
Large Language Models

# Language Model Pretraining

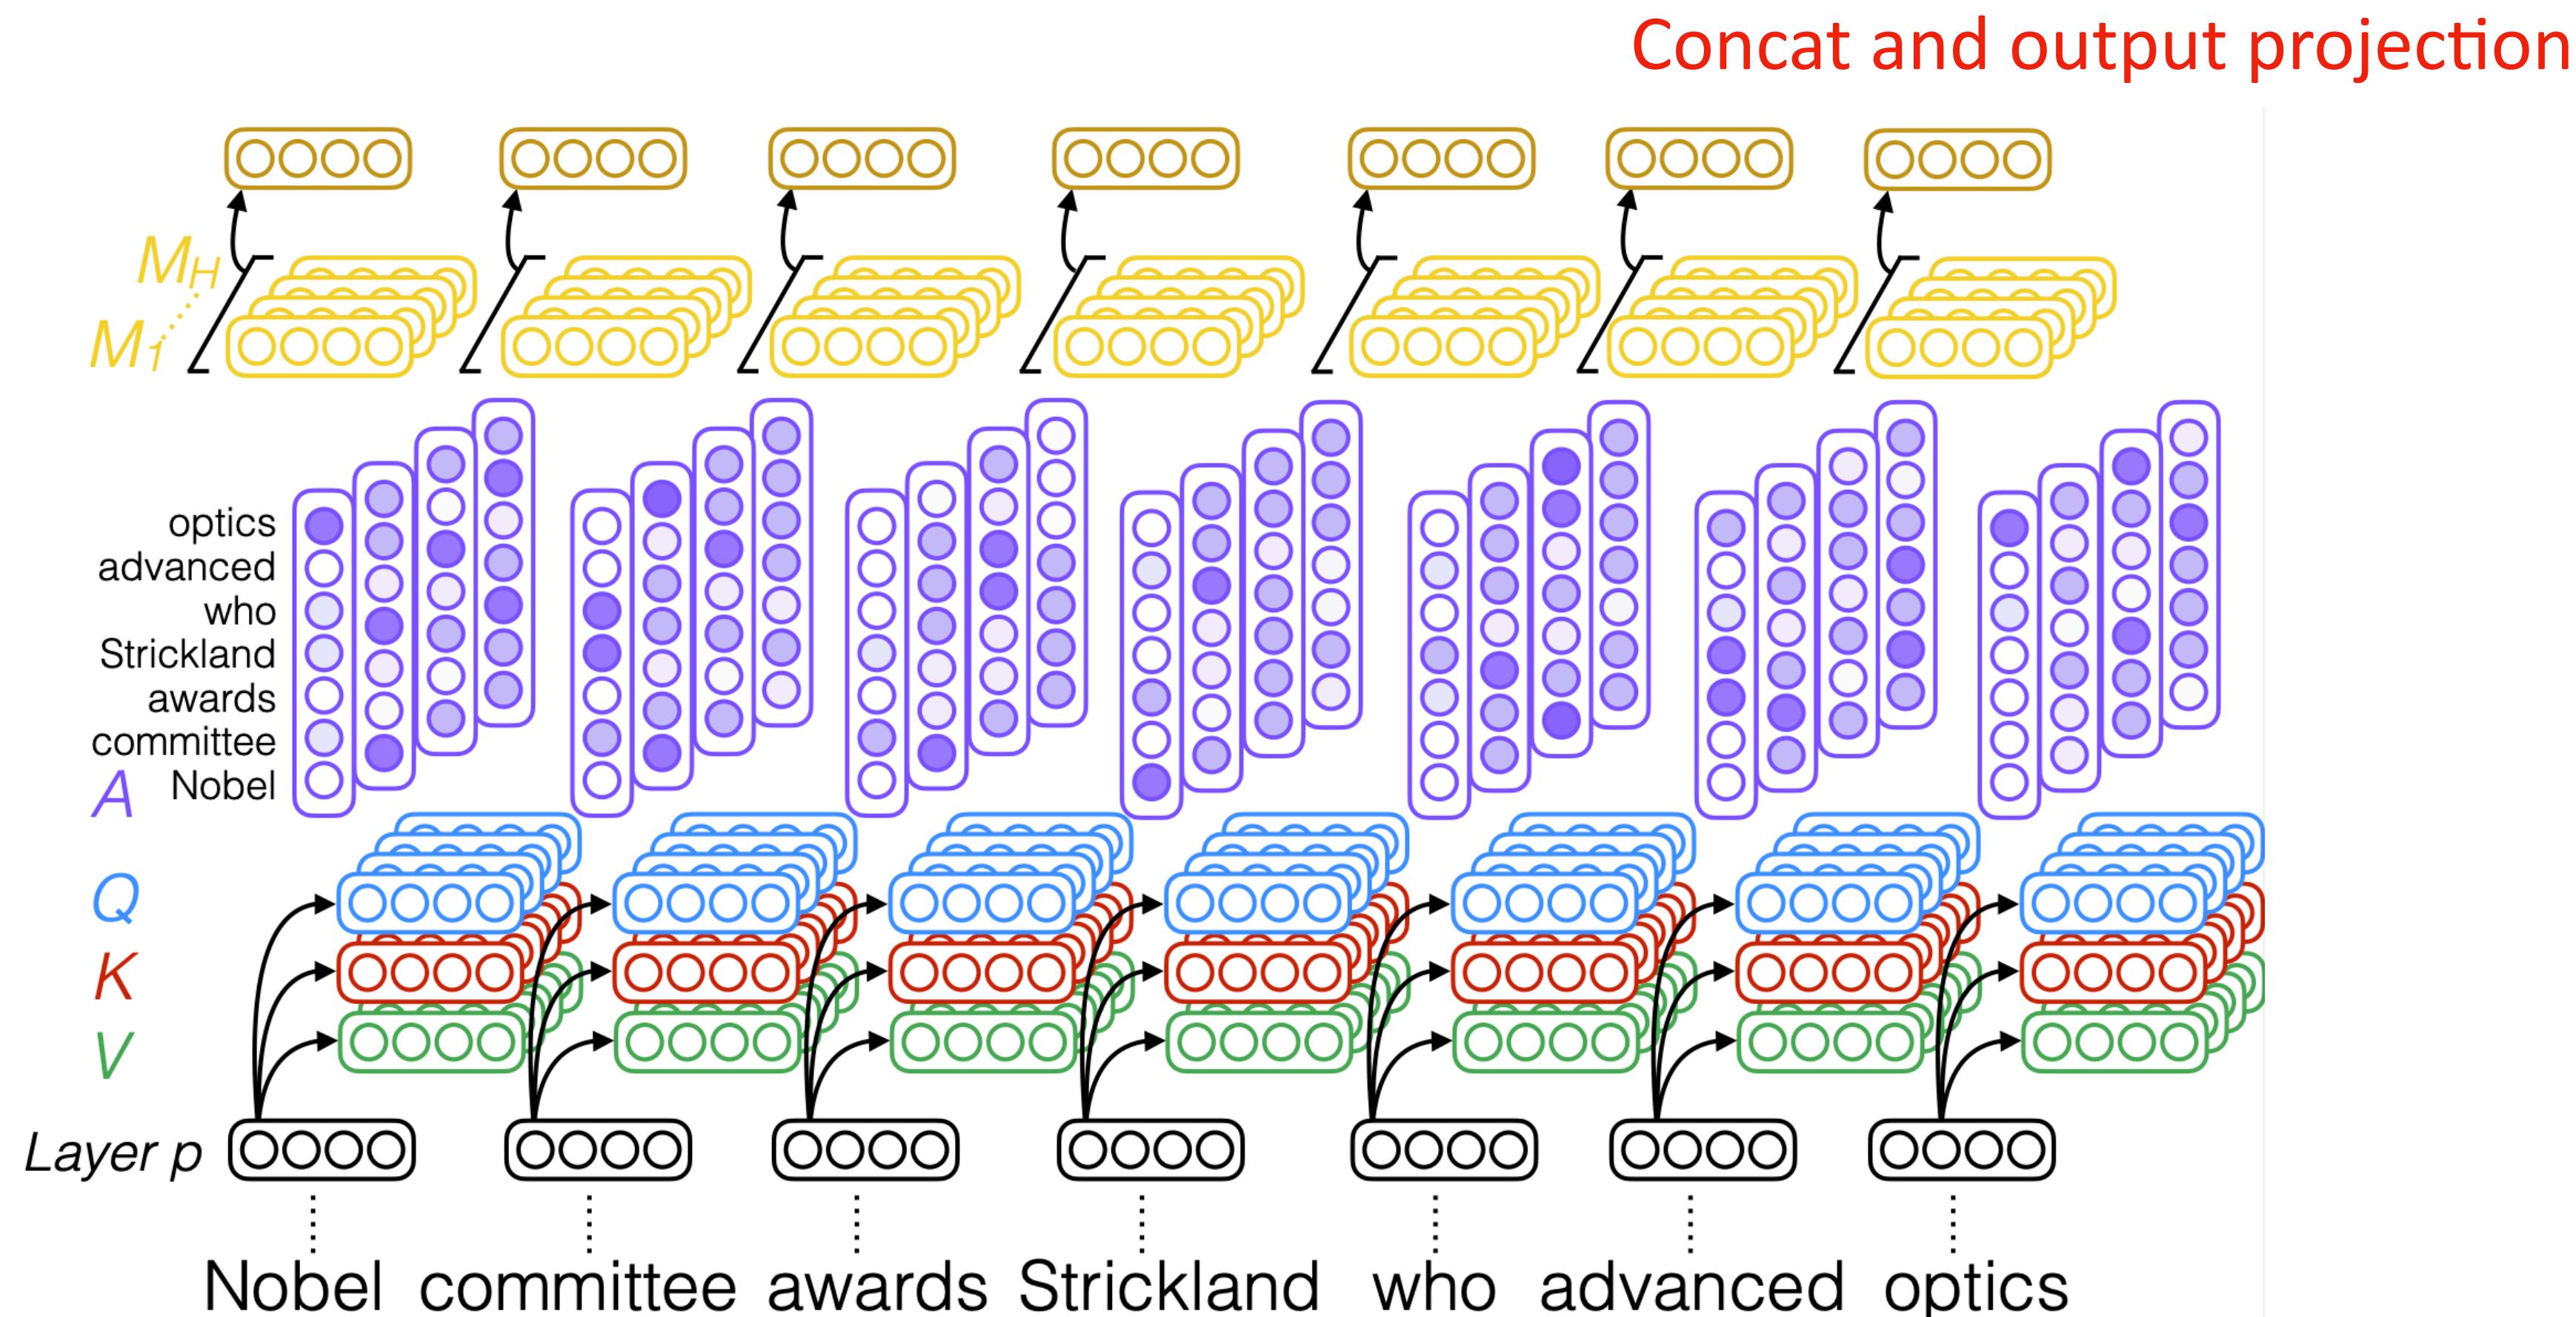
Junxian He

Sep 17, 2025

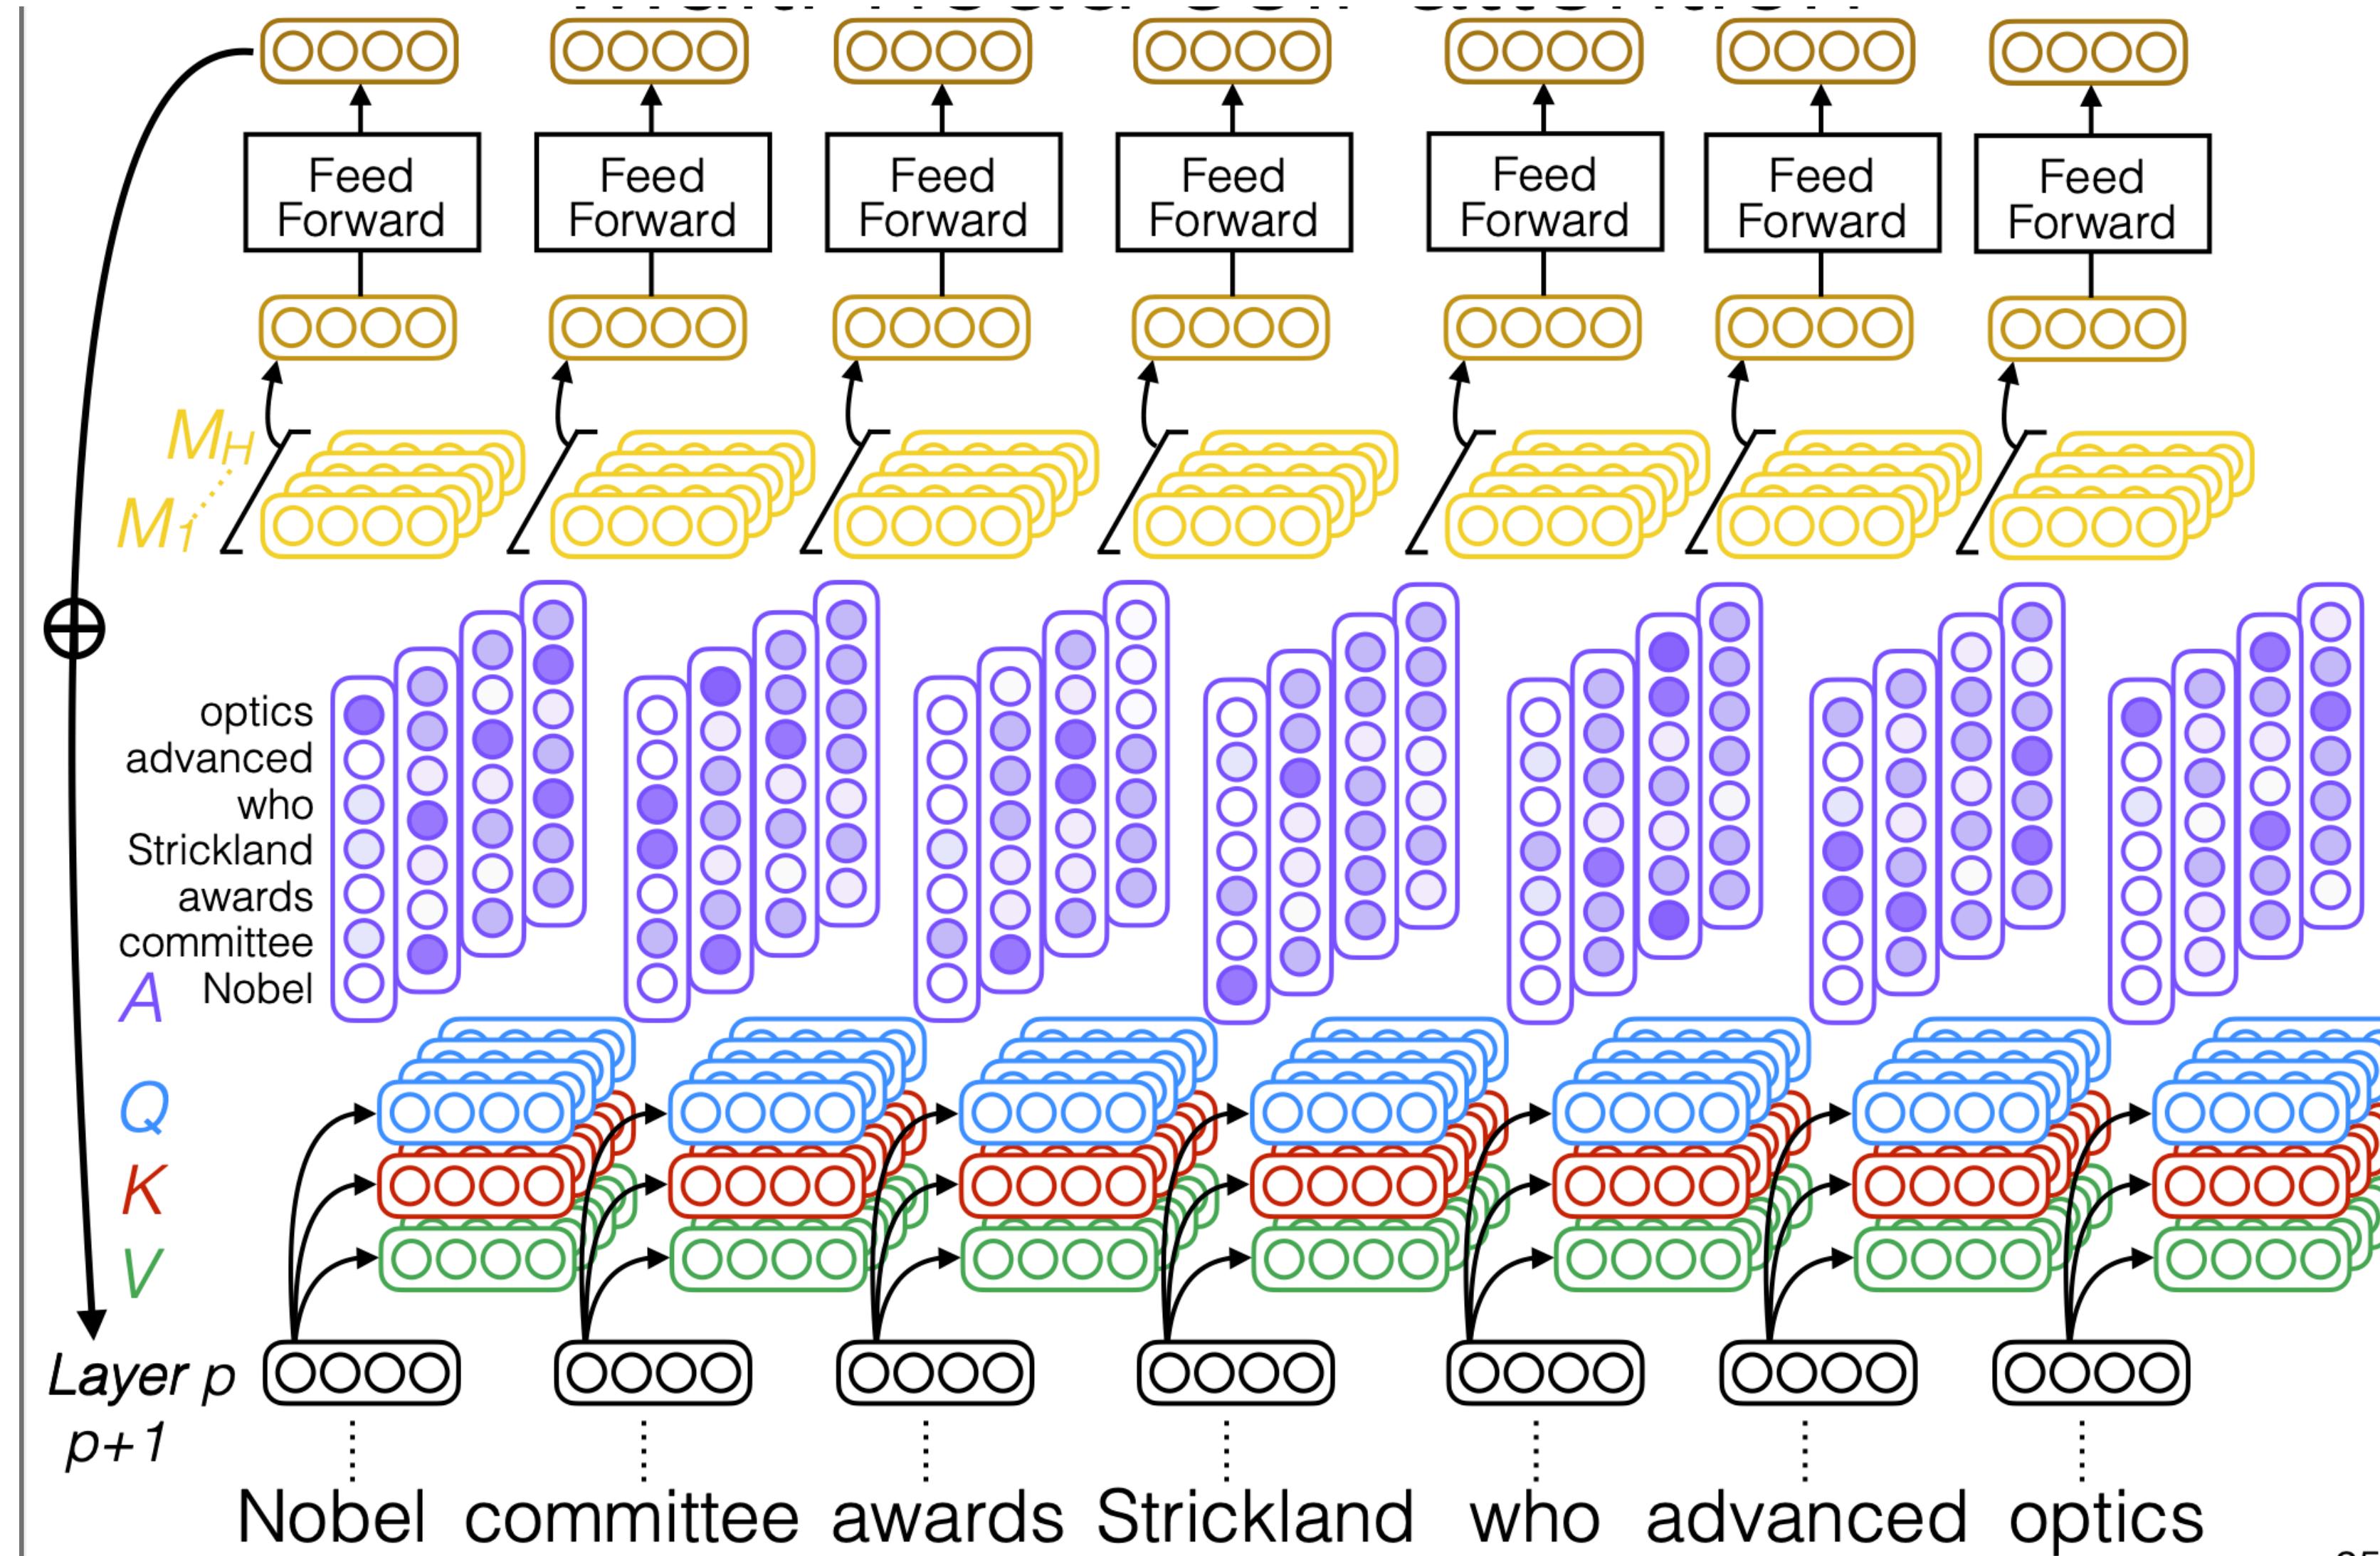
# Recap: Multi-head Self-Attention



# Recap: Multi-head Self-Attention

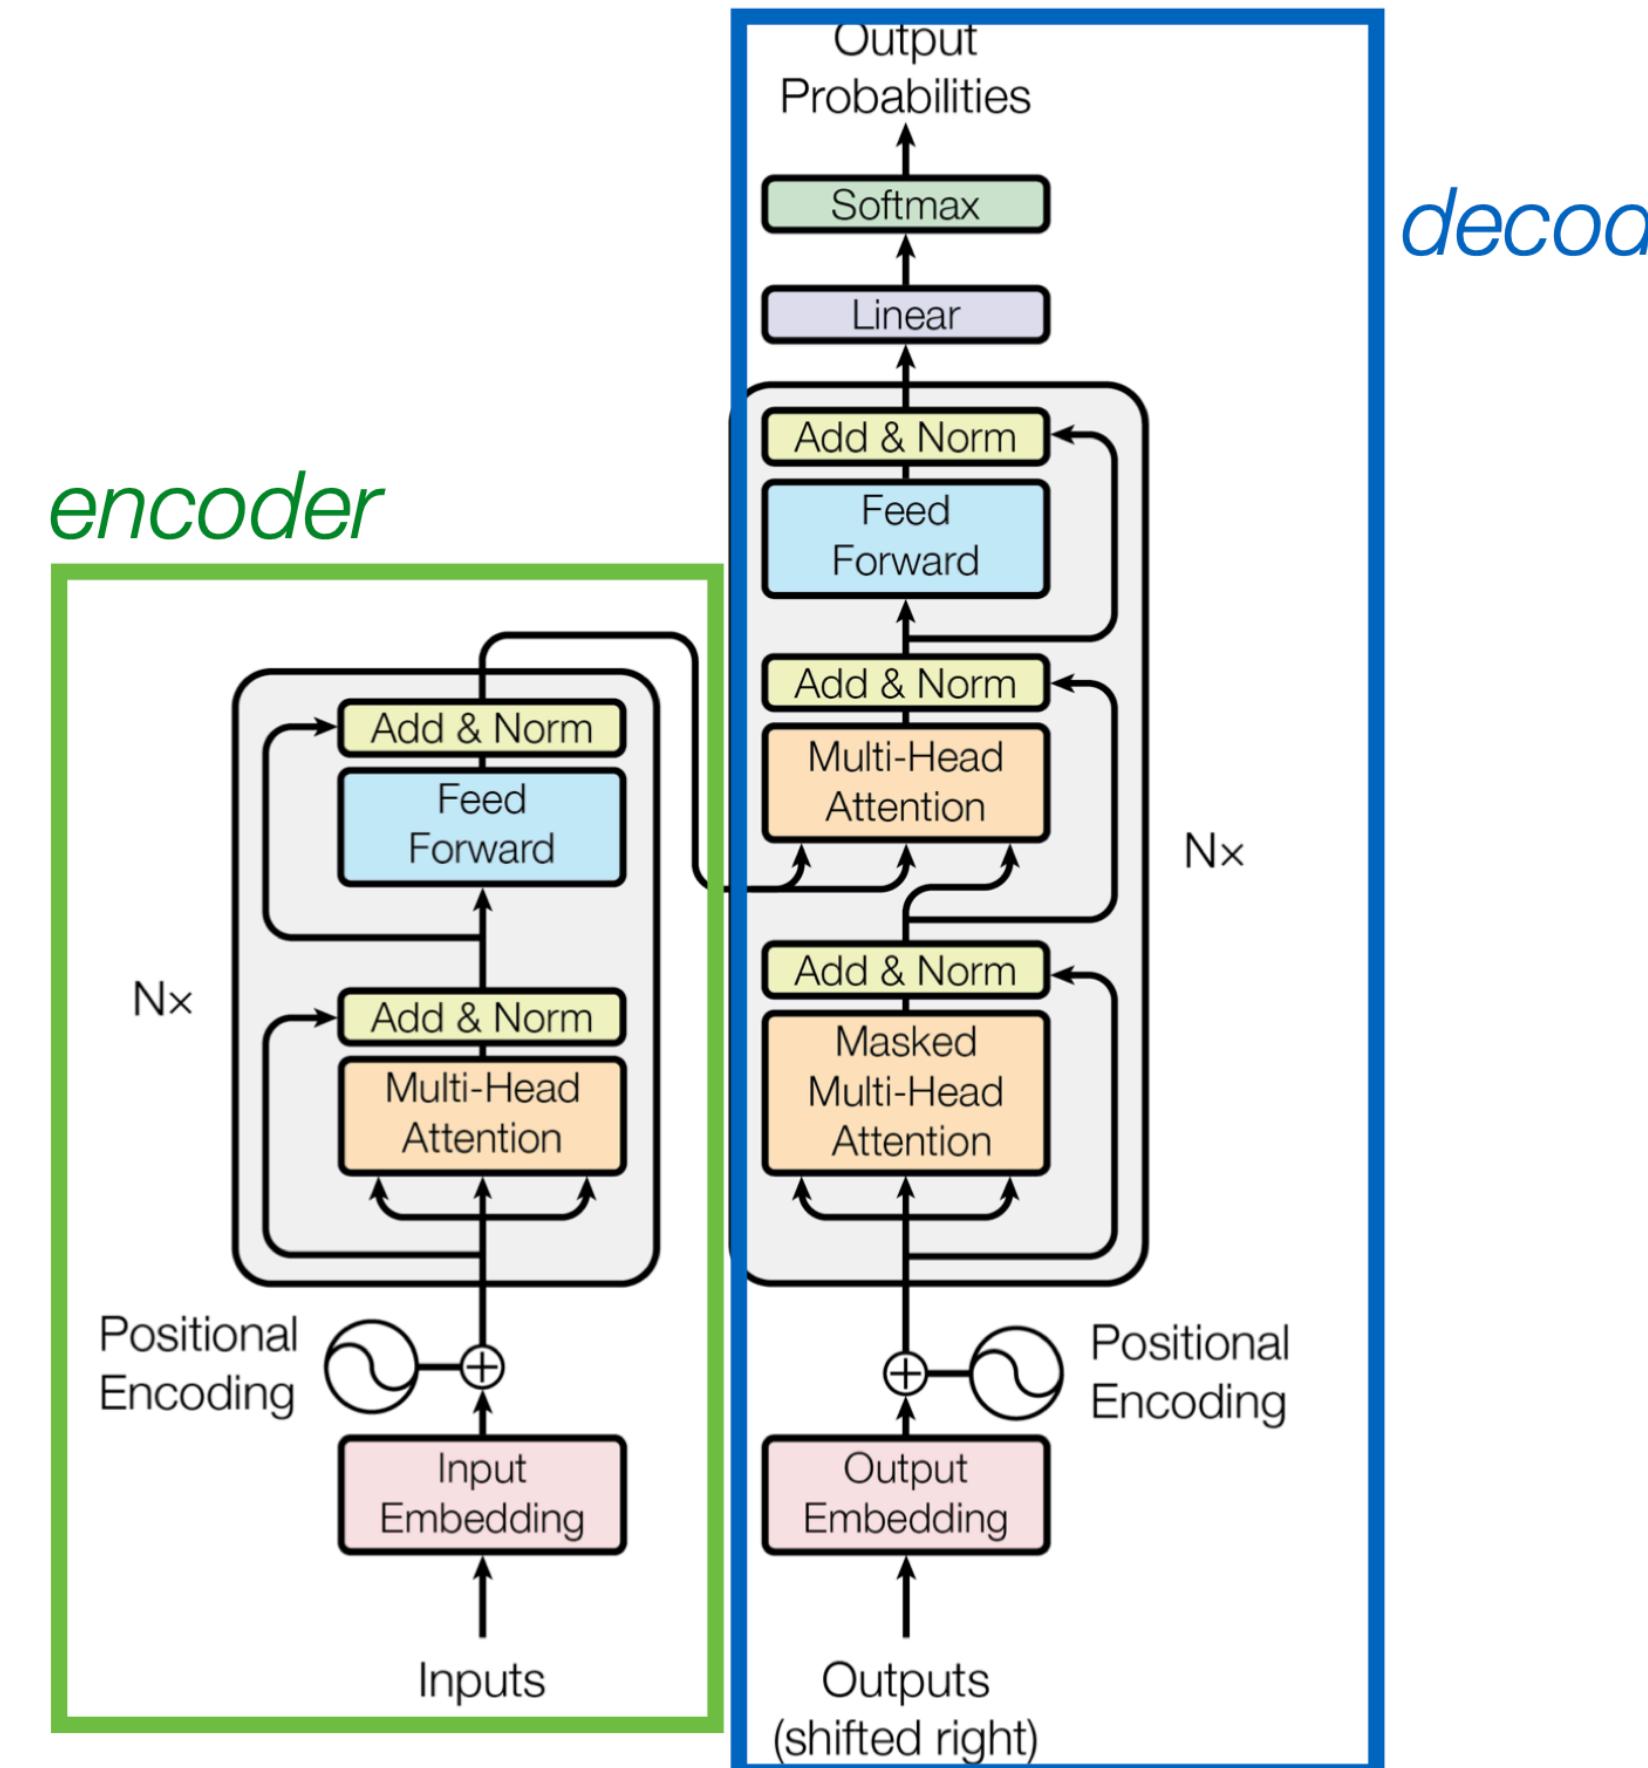
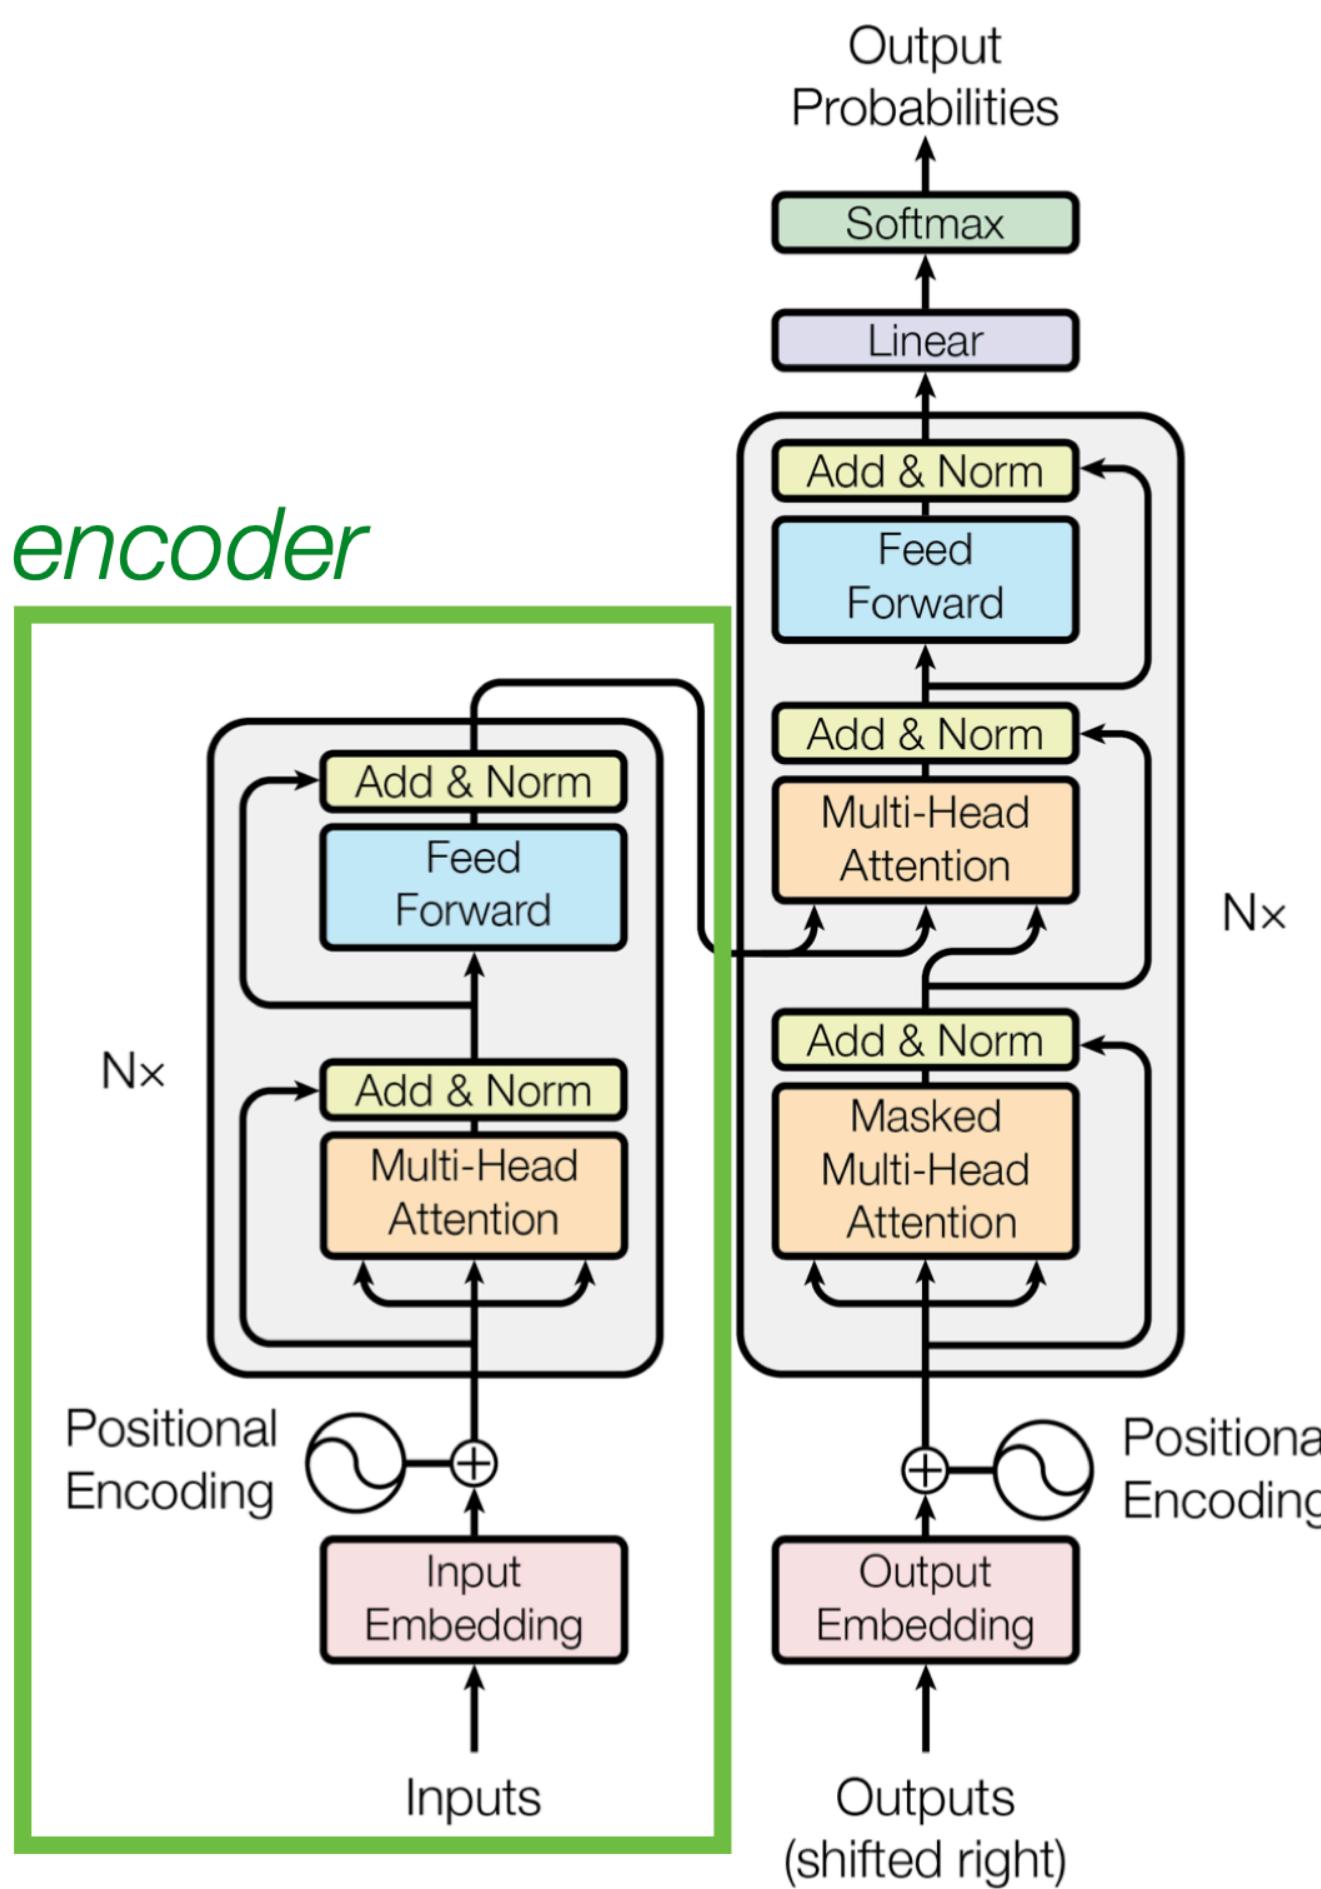


# Recap: Multi-head Self-Attention + FFN



# Recap: Transformer Encoder

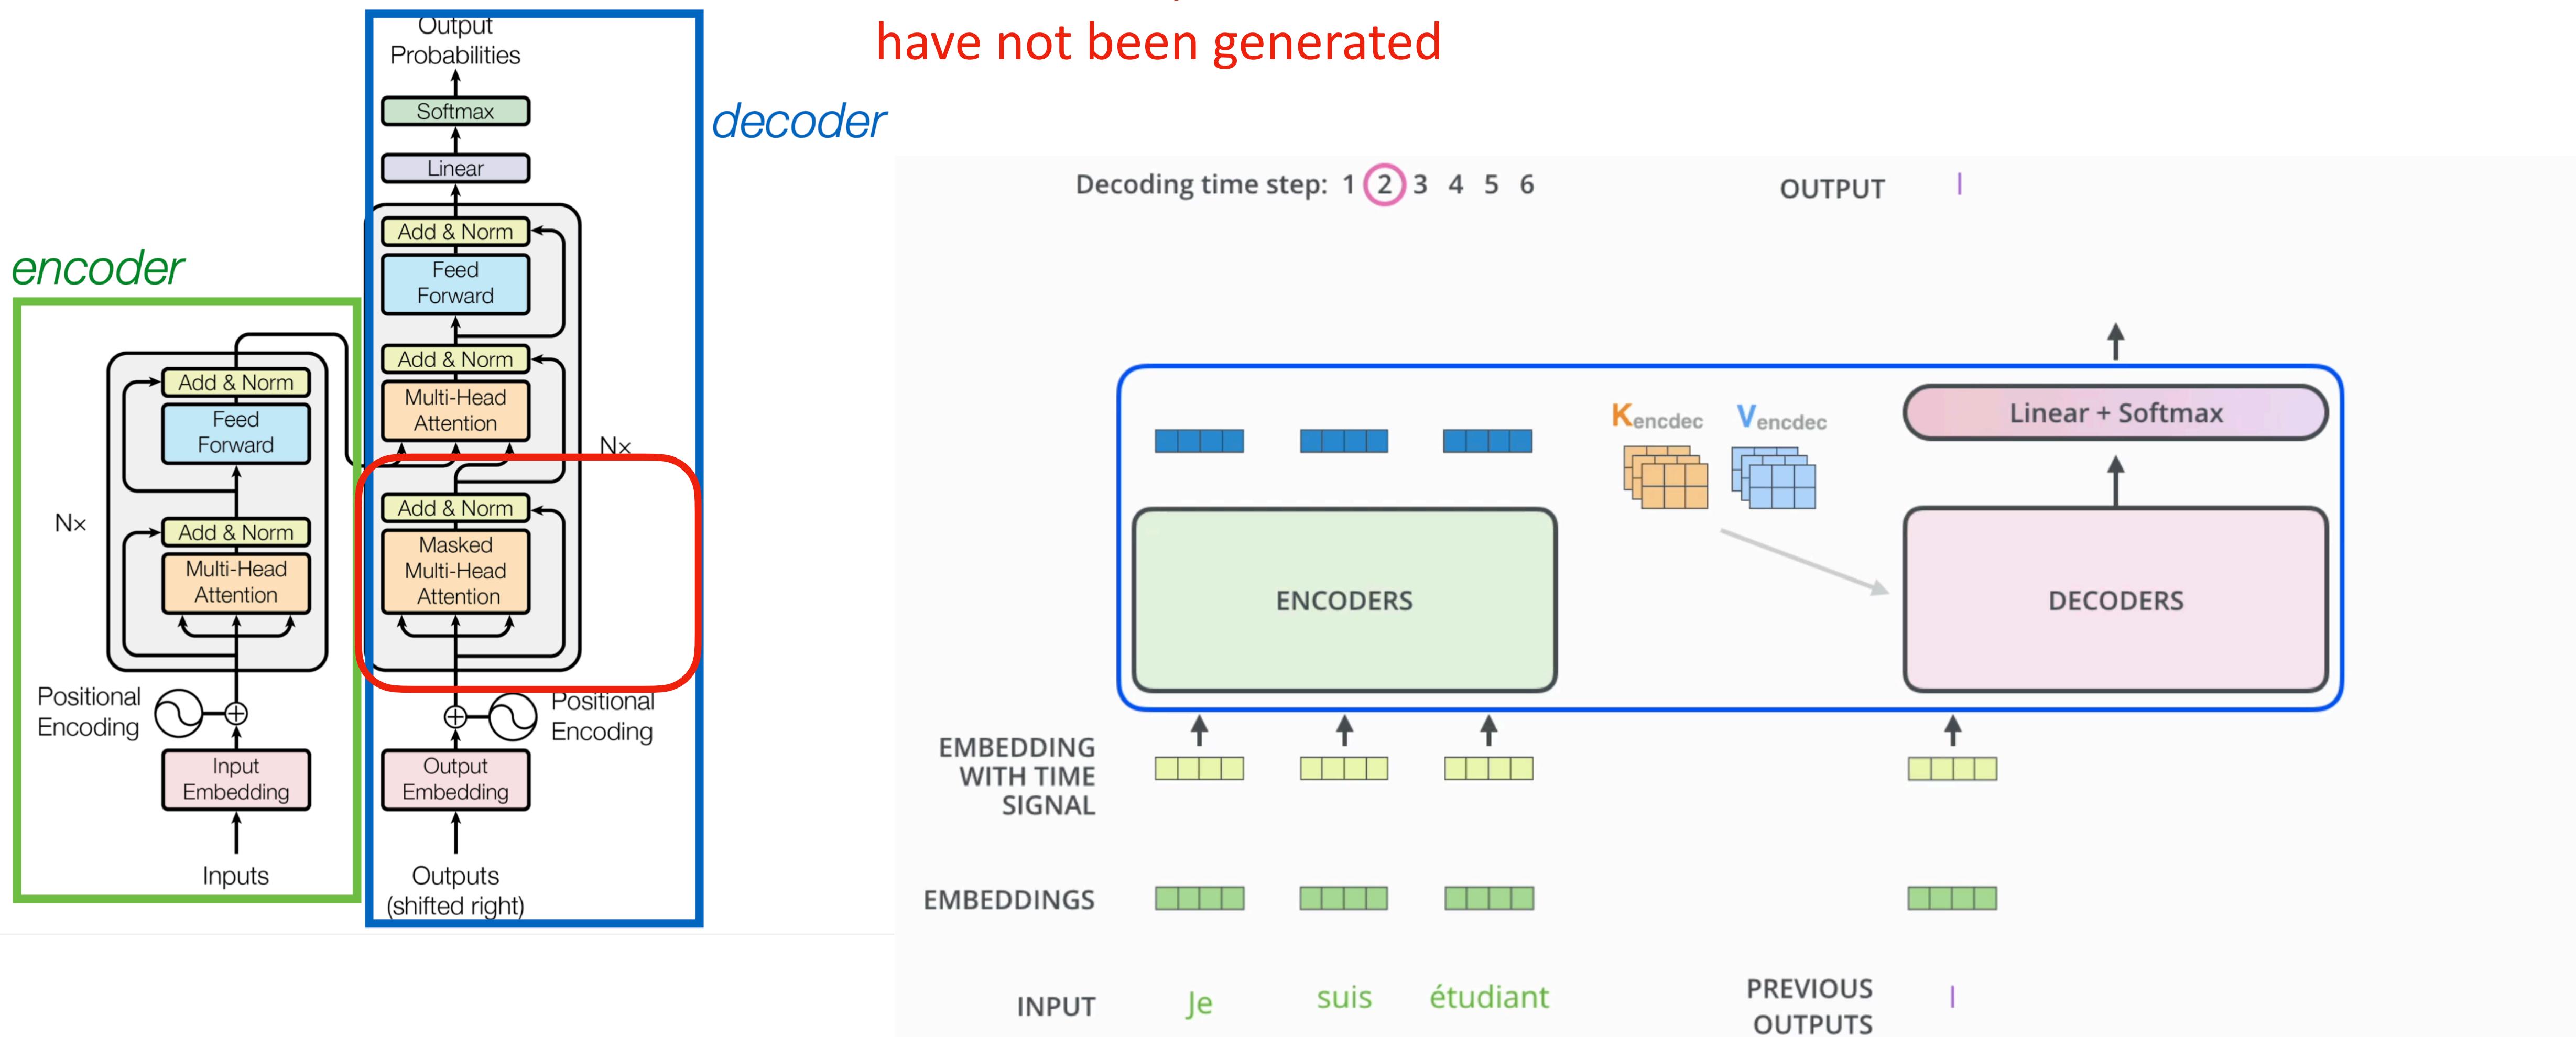
Currently we only cover the encoder side



This encoder-decoder arch is originally proposed as a seq2seq arch, for classification tasks, often only encoder is used. And language models often only have a decoder

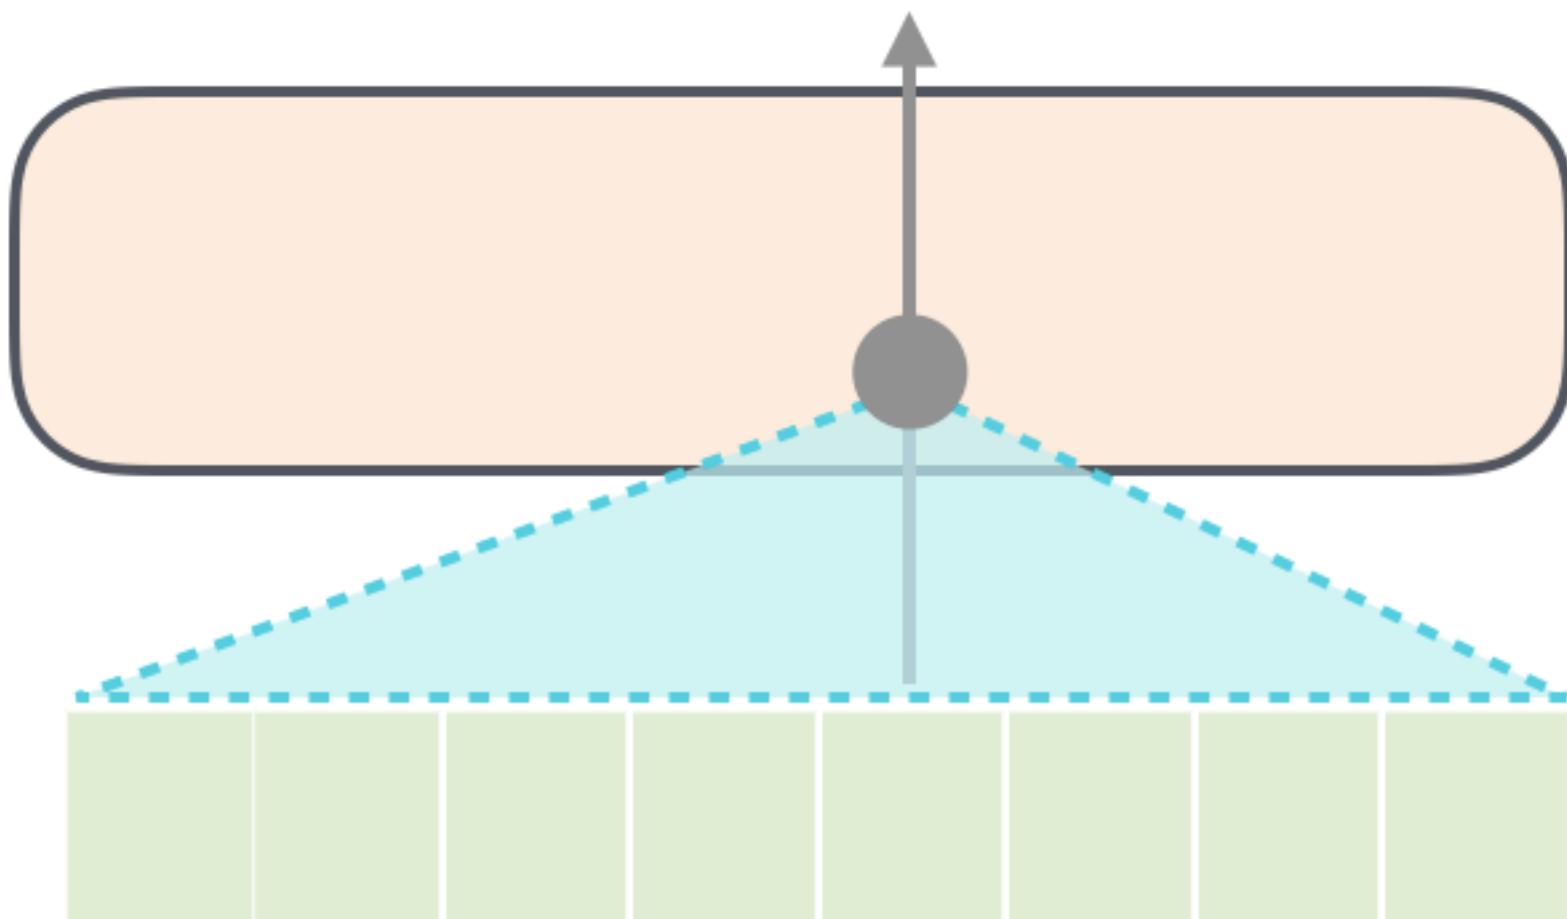
# Recap: Masked Attention

Typical attention attends to the entire sequence, while masked attention only attends to the ones on the left because future words have not been generated

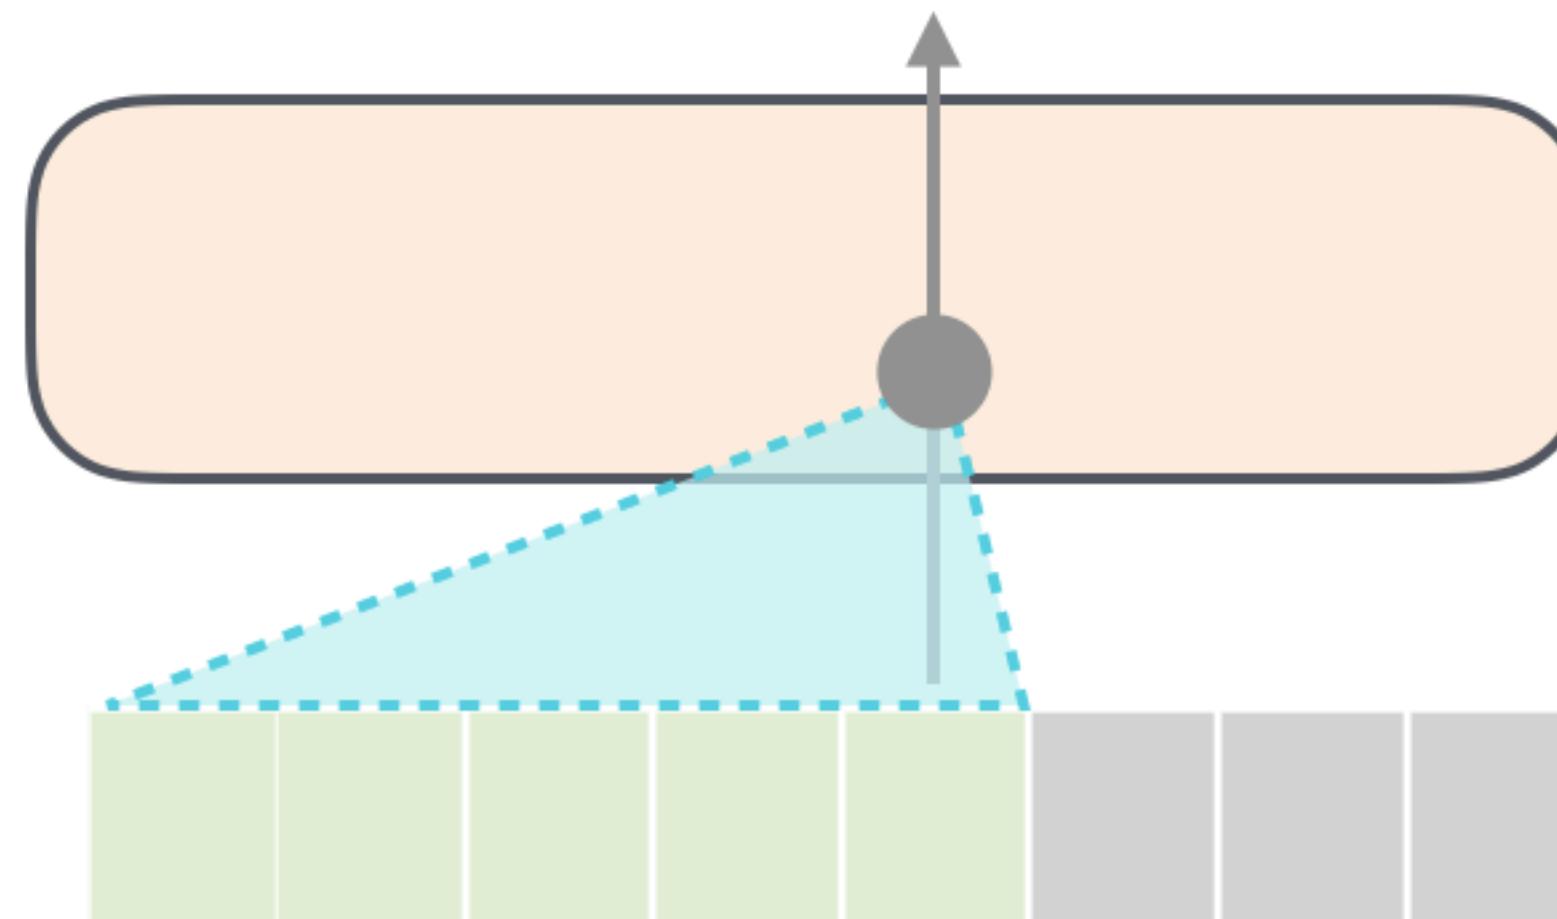


# Masked Attention

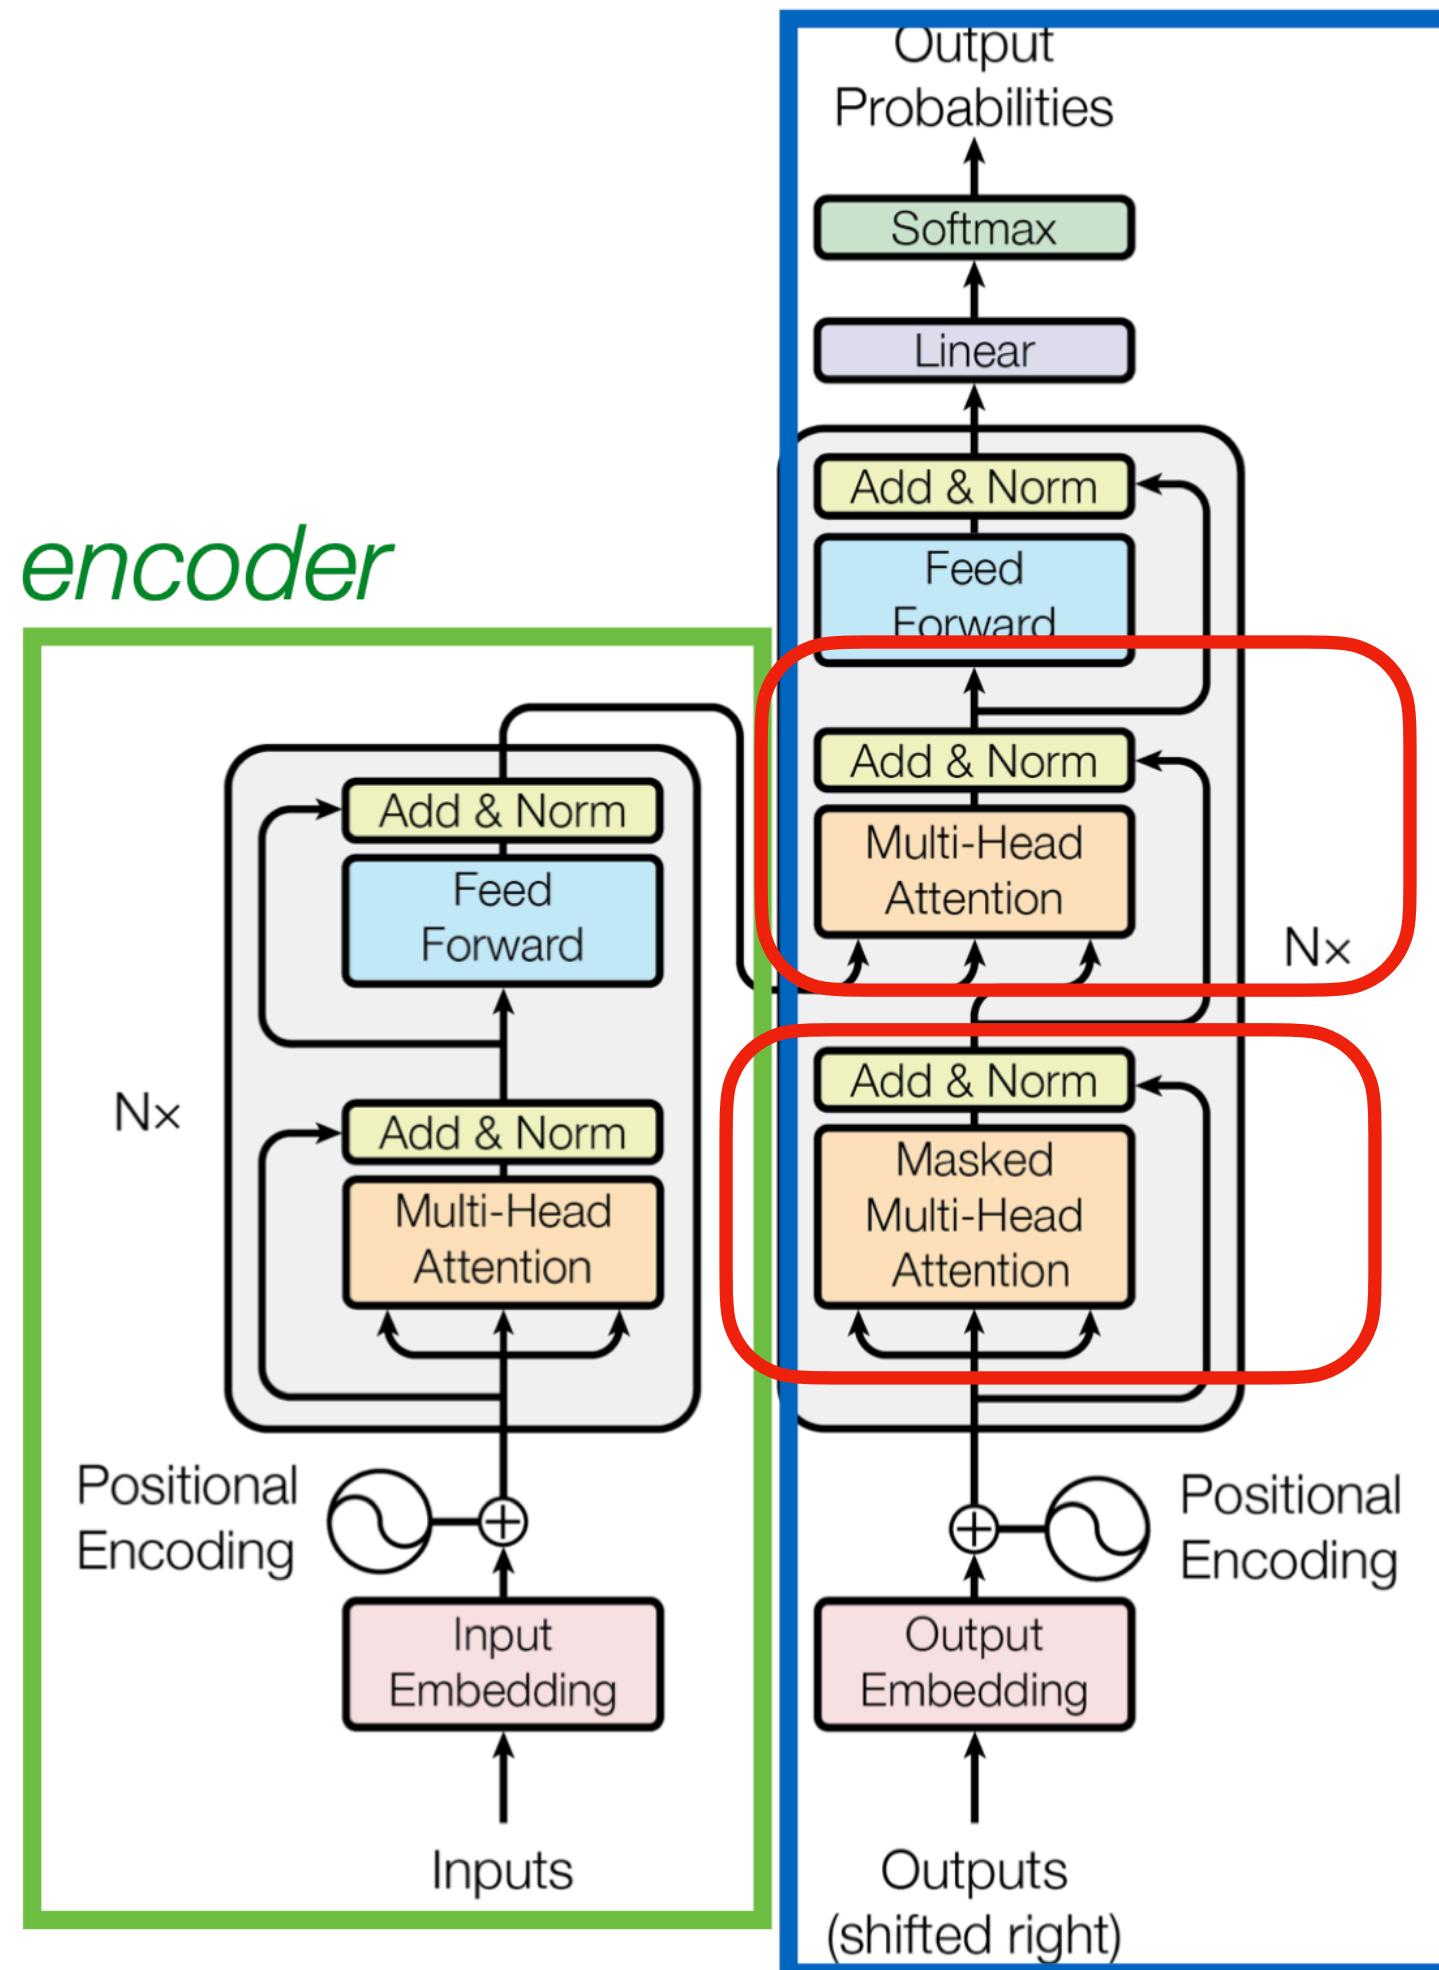
**Self-Attention**



**Masked Self-Attention**



# Transformer Decoder in Seq2Seq

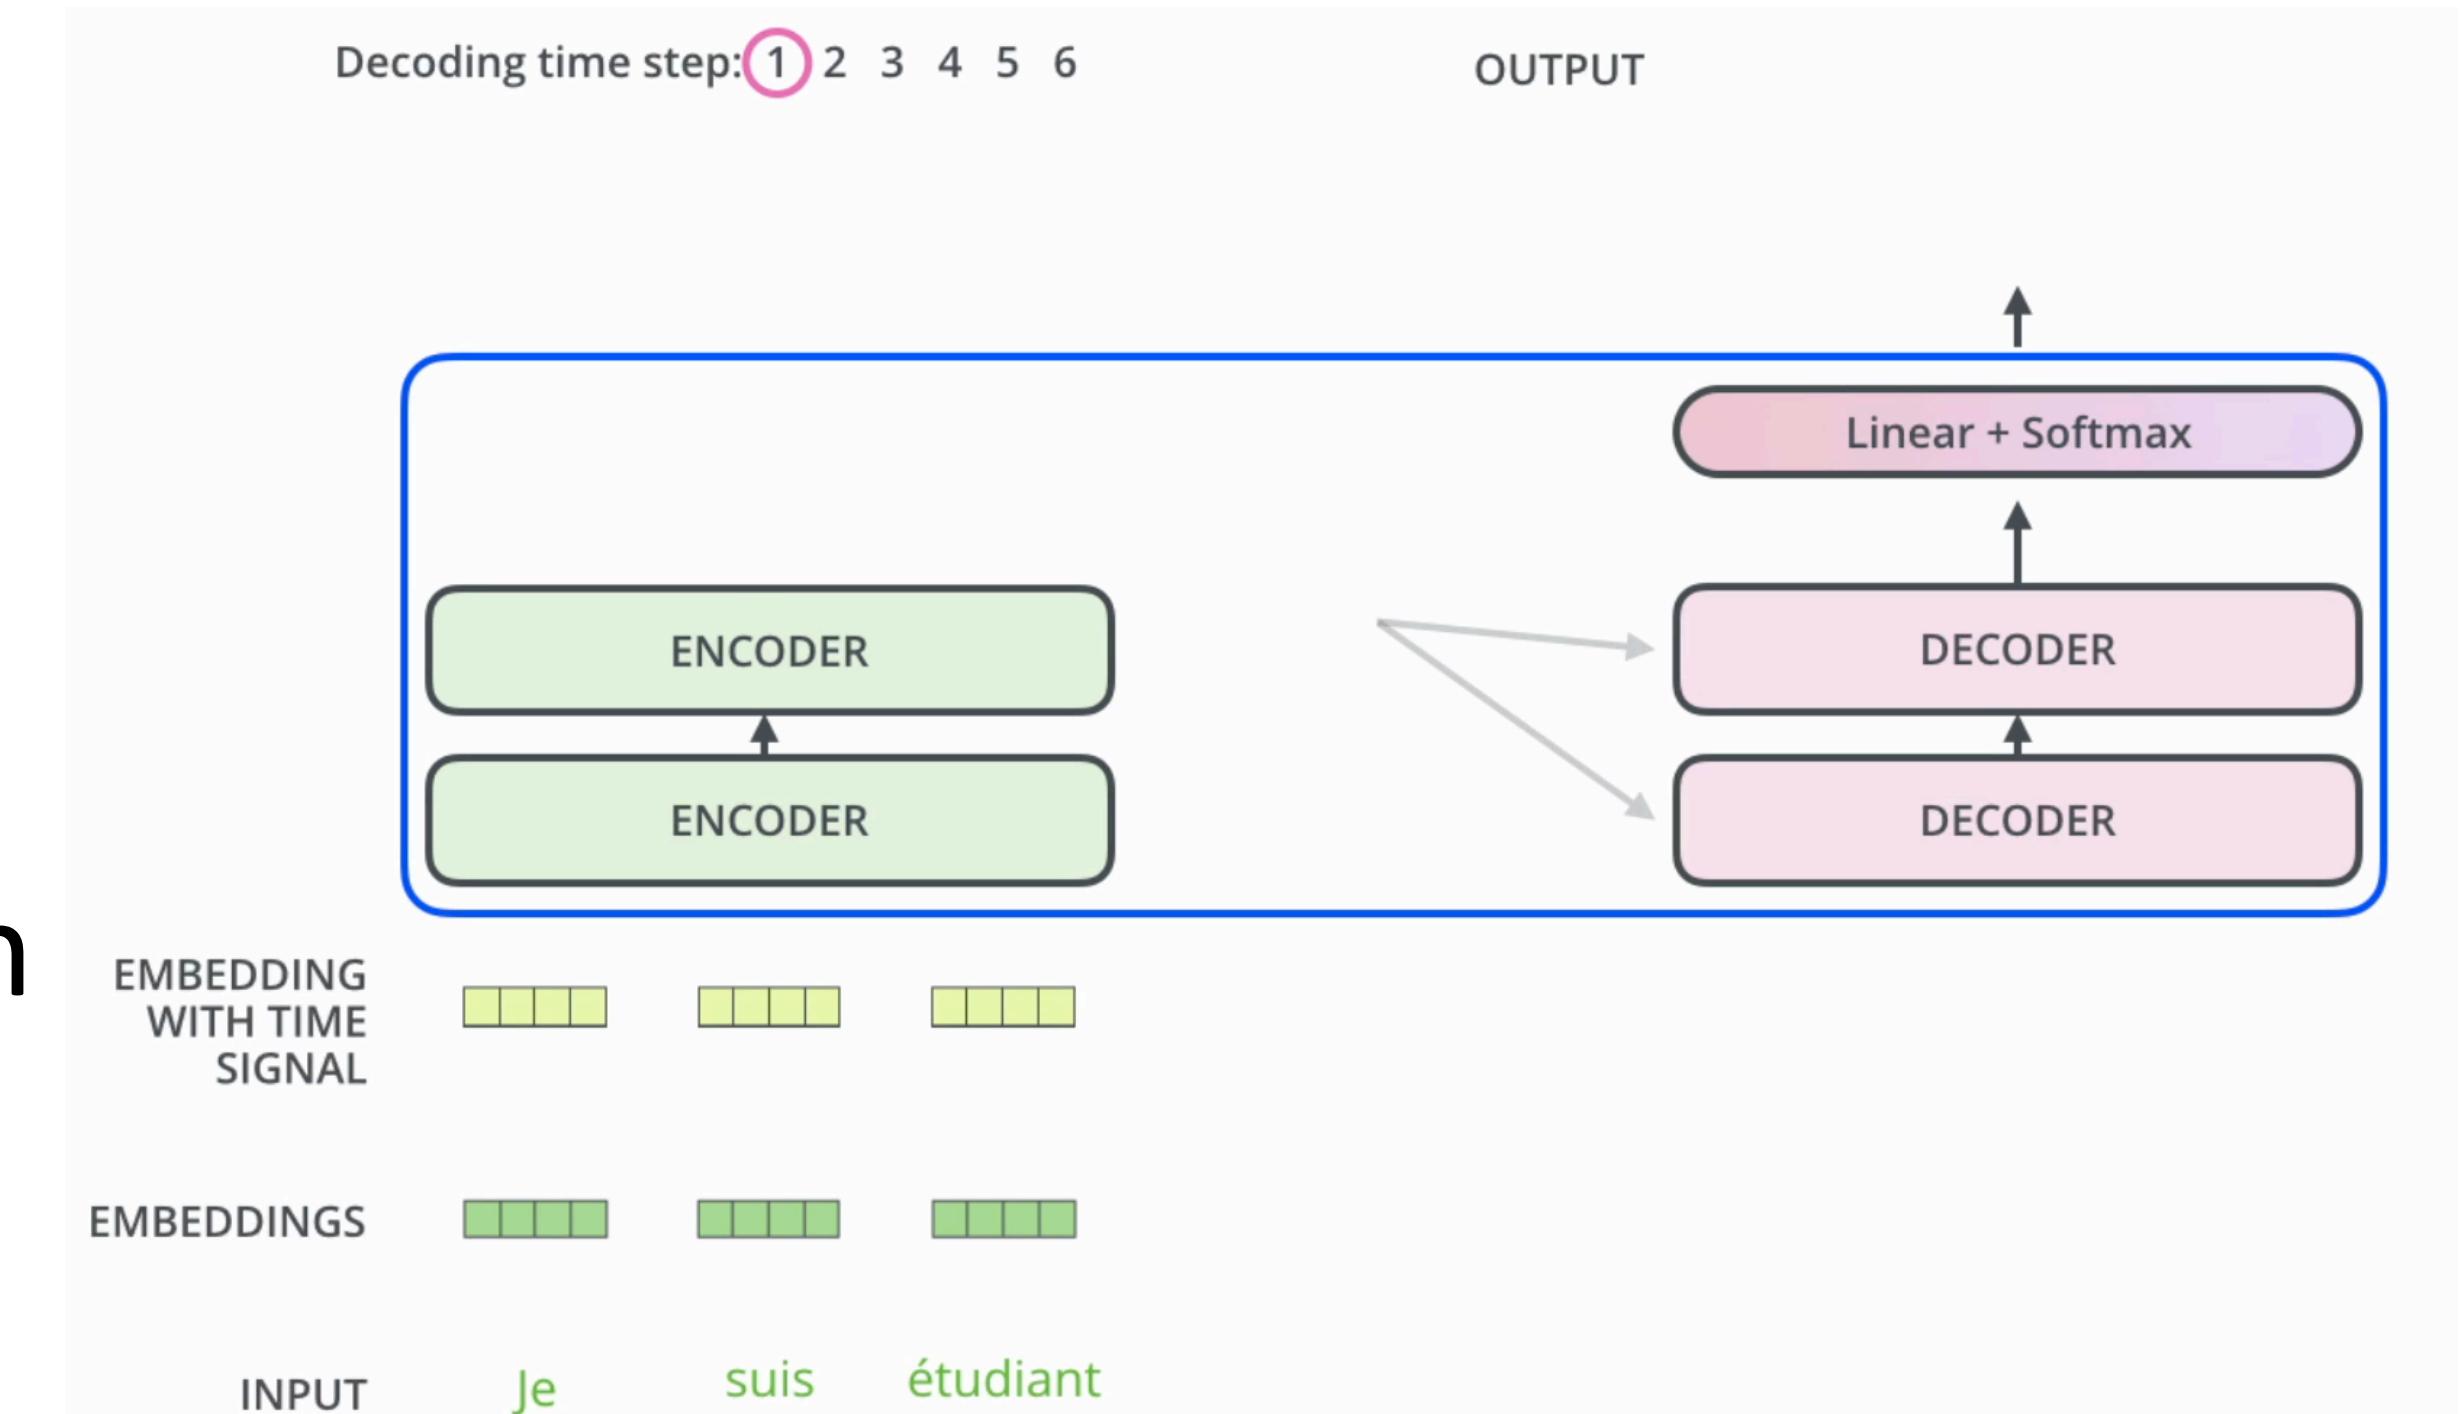


*decoder*

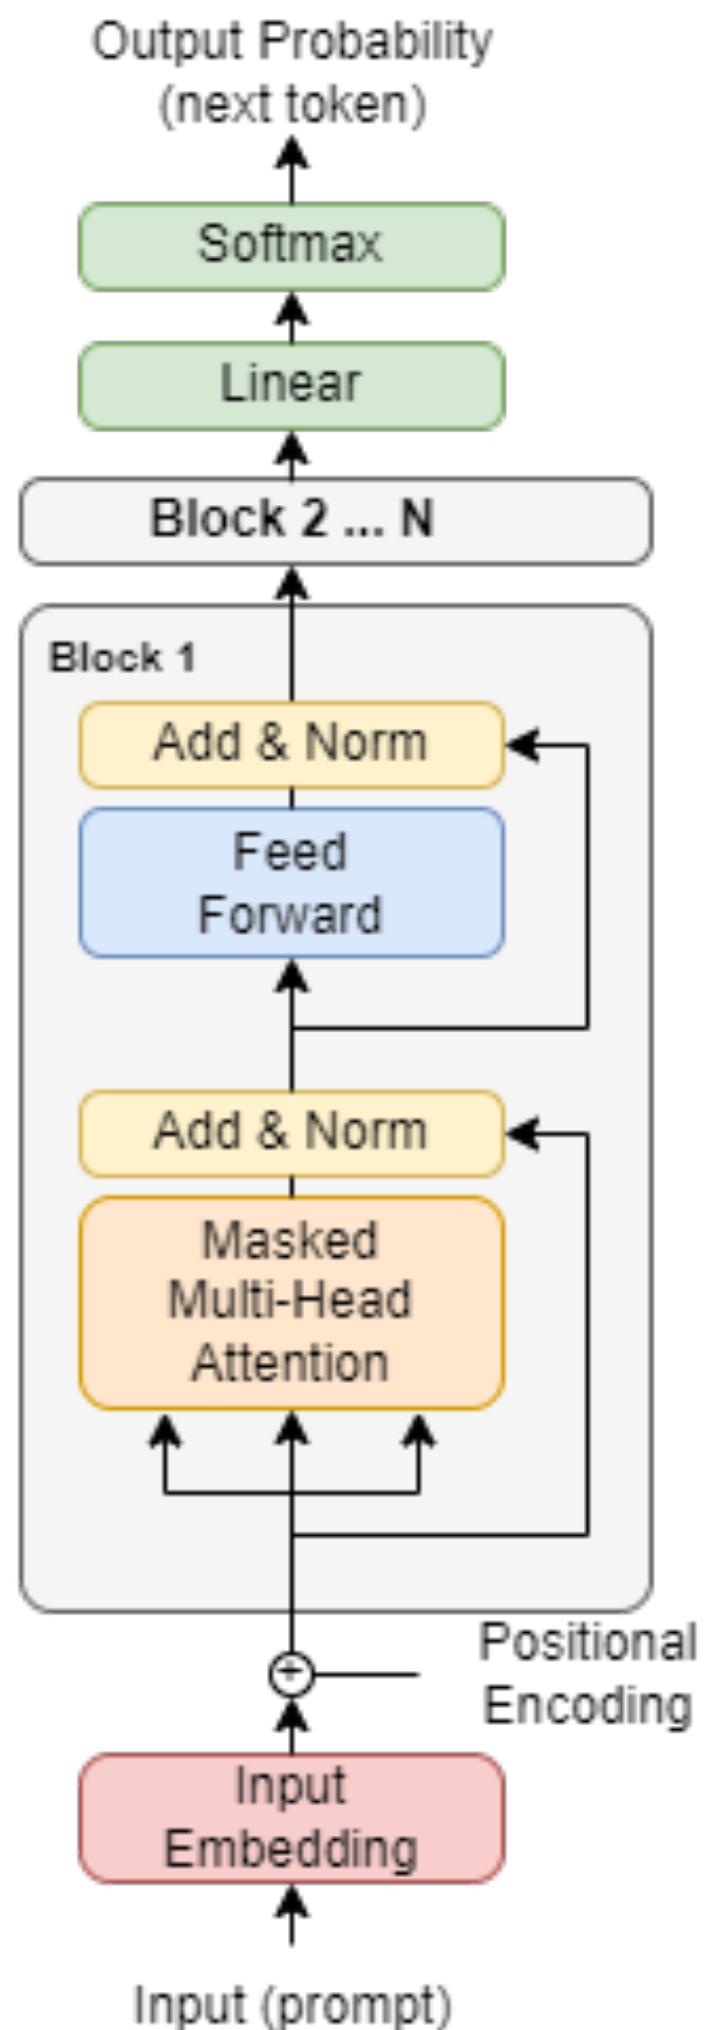
Cross-attention

Self-attention

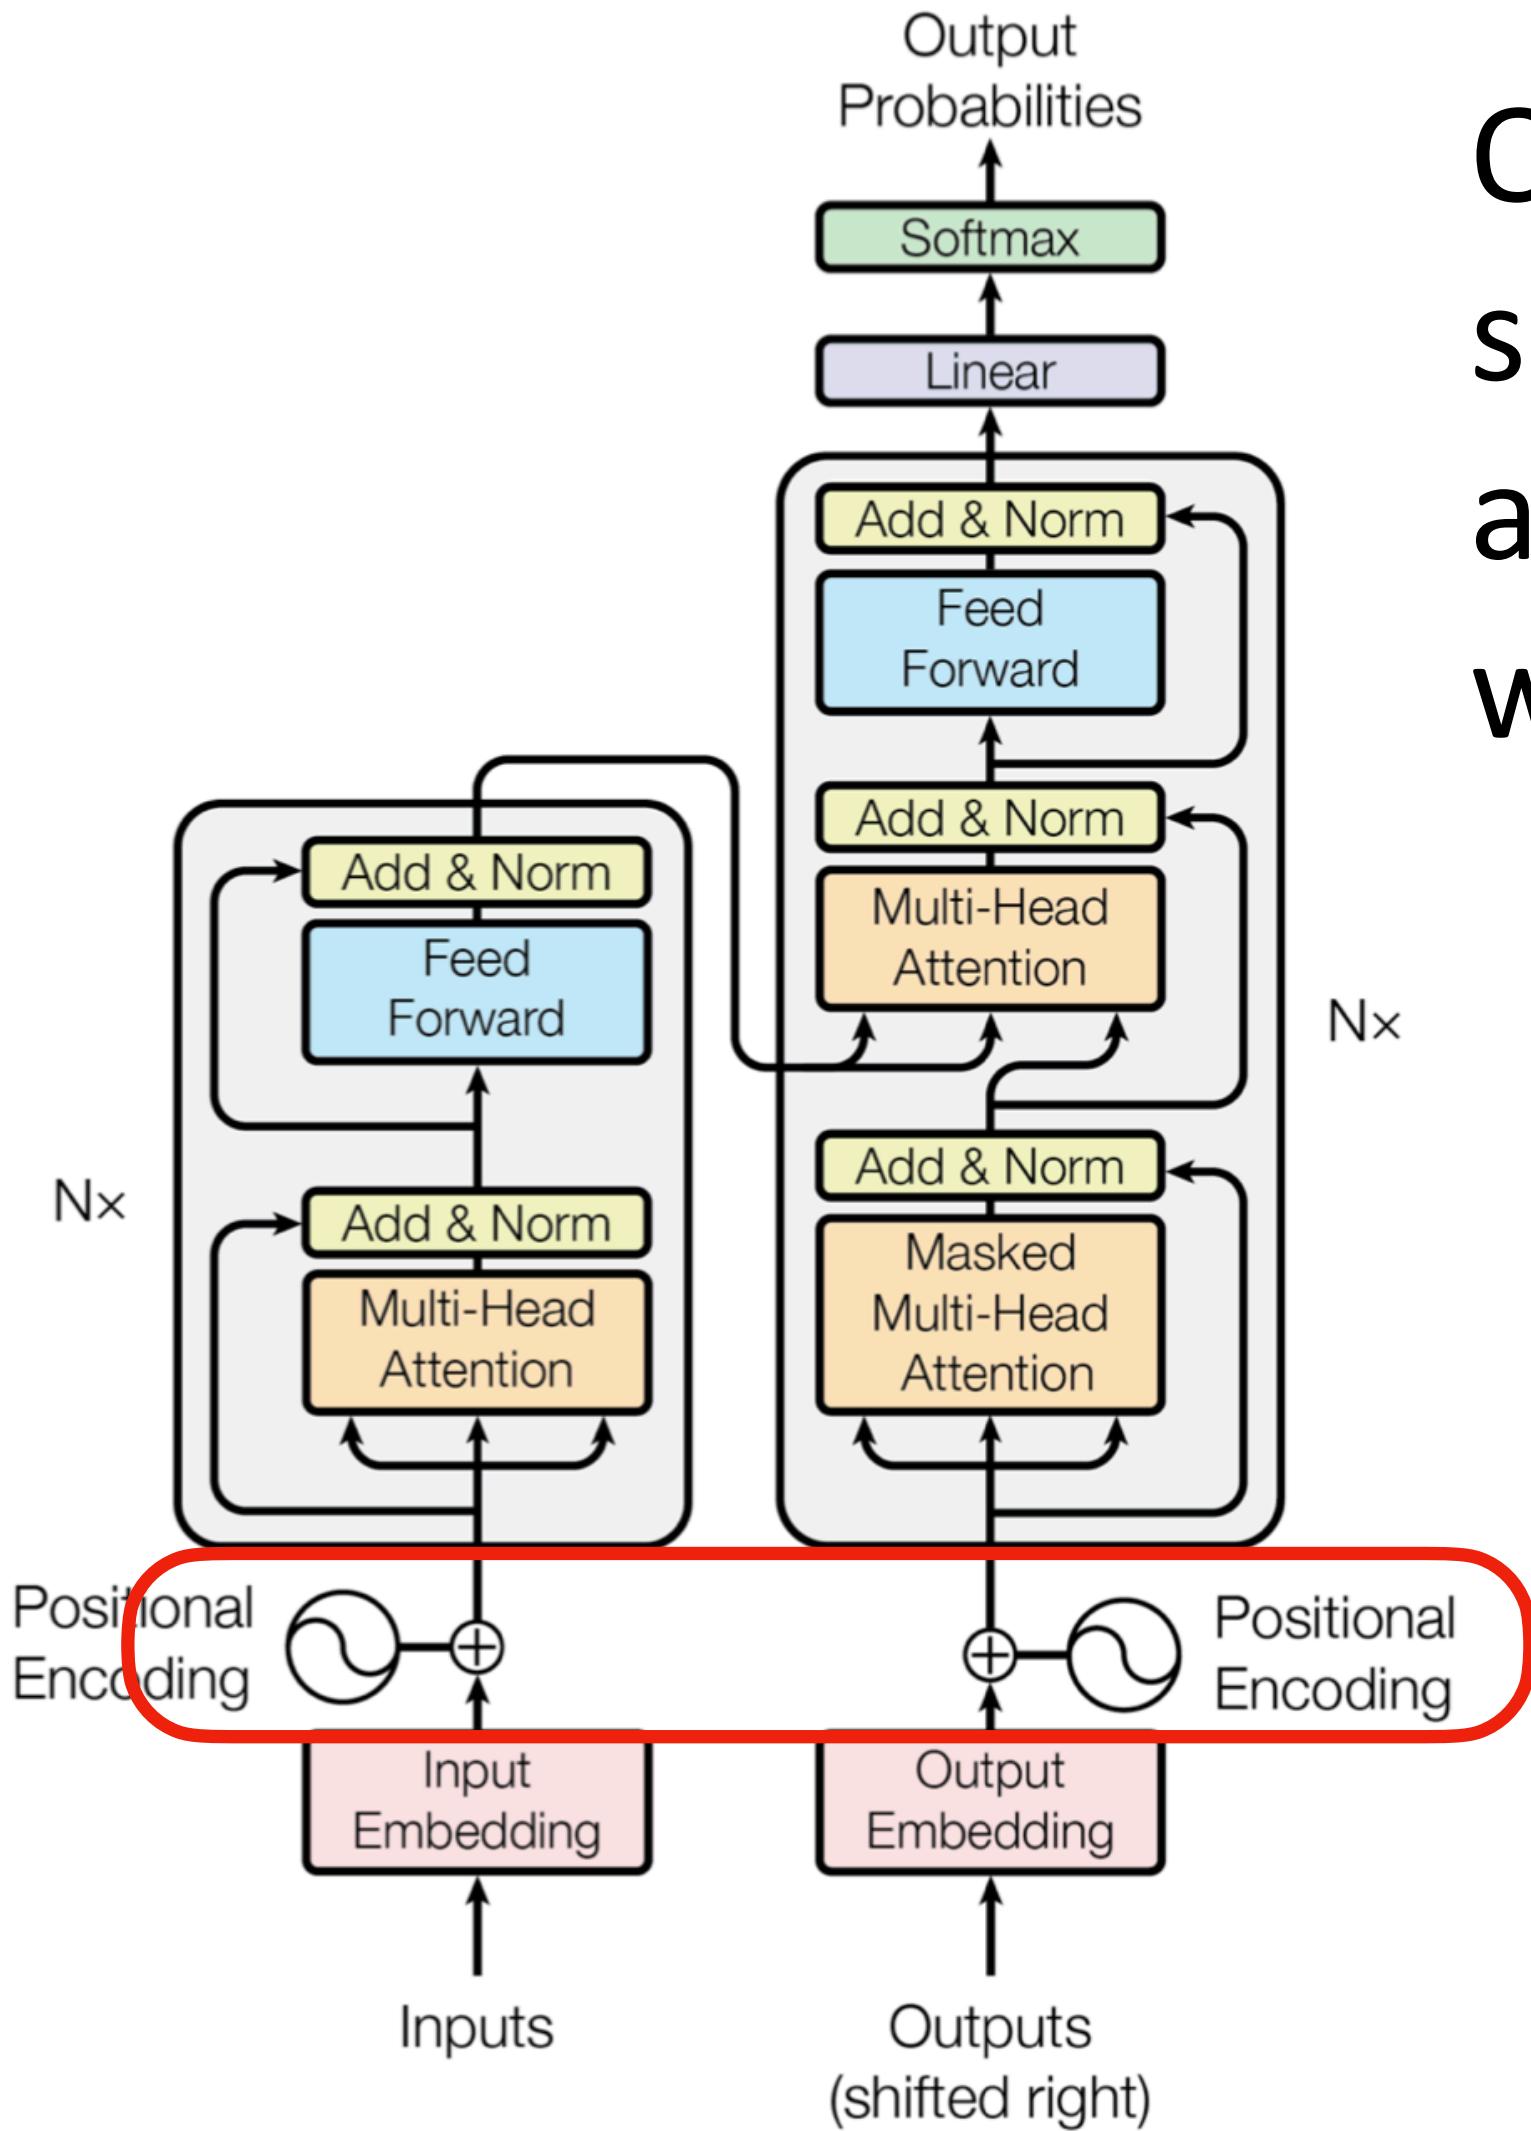
Cross-attention uses the output of encoder as input



# Transformer Language Model (e.g., ChatGPT)



# Position Embeddings



Question: If we shuffle the order of words in the sequence, will that change the attention output and feed forward output of the corresponding word?

Position embeddings are added to each word embedding, otherwise our model is unaware of the position of a word

# Positional Encoding

EMBEDDING  
WITH TIME  
SIGNAL

$$\mathbf{x}_1 \quad \begin{array}{|c|c|c|c|} \hline \text{light green} & \text{light green} & \text{light green} & \text{light green} \\ \hline \end{array}$$

=

POSITIONAL  
ENCODING

$$\mathbf{t}_1 \quad \begin{array}{|c|c|c|c|} \hline \text{yellow} & \text{yellow} & \text{yellow} & \text{yellow} \\ \hline \end{array}$$

+

EMBEDDINGS

$$\mathbf{x}_1 \quad \begin{array}{|c|c|c|c|} \hline \text{green} & \text{green} & \text{green} & \text{green} \\ \hline \end{array}$$

$$\mathbf{x}_2 \quad \begin{array}{|c|c|c|c|} \hline \text{light green} & \text{light green} & \text{light green} & \text{light green} \\ \hline \end{array}$$

=

$$\mathbf{t}_2 \quad \begin{array}{|c|c|c|c|} \hline \text{yellow} & \text{yellow} & \text{yellow} & \text{yellow} \\ \hline \end{array}$$

+

$$\mathbf{x}_3 \quad \begin{array}{|c|c|c|c|} \hline \text{light green} & \text{light green} & \text{light green} & \text{light green} \\ \hline \end{array}$$

=

$$\mathbf{t}_3 \quad \begin{array}{|c|c|c|c|} \hline \text{yellow} & \text{yellow} & \text{yellow} & \text{yellow} \\ \hline \end{array}$$

+

$$\mathbf{x}_3 \quad \begin{array}{|c|c|c|c|} \hline \text{green} & \text{green} & \text{green} & \text{green} \\ \hline \end{array}$$

INPUT

Je

suis

étudiant

# Transformer Positional Encoding

$$PE_{(pos,2i)} = \sin\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

$$PE_{(pos,2i+1)} = \cos\left(\frac{pos}{10000^{2i/d_{model}}}\right)$$

Positional encoding is a 512d vector  
 $i$  = a particular dimension of this vector  
 $pos$  = dimension of the word  
 $d_{model}$  = 512

# Complexity

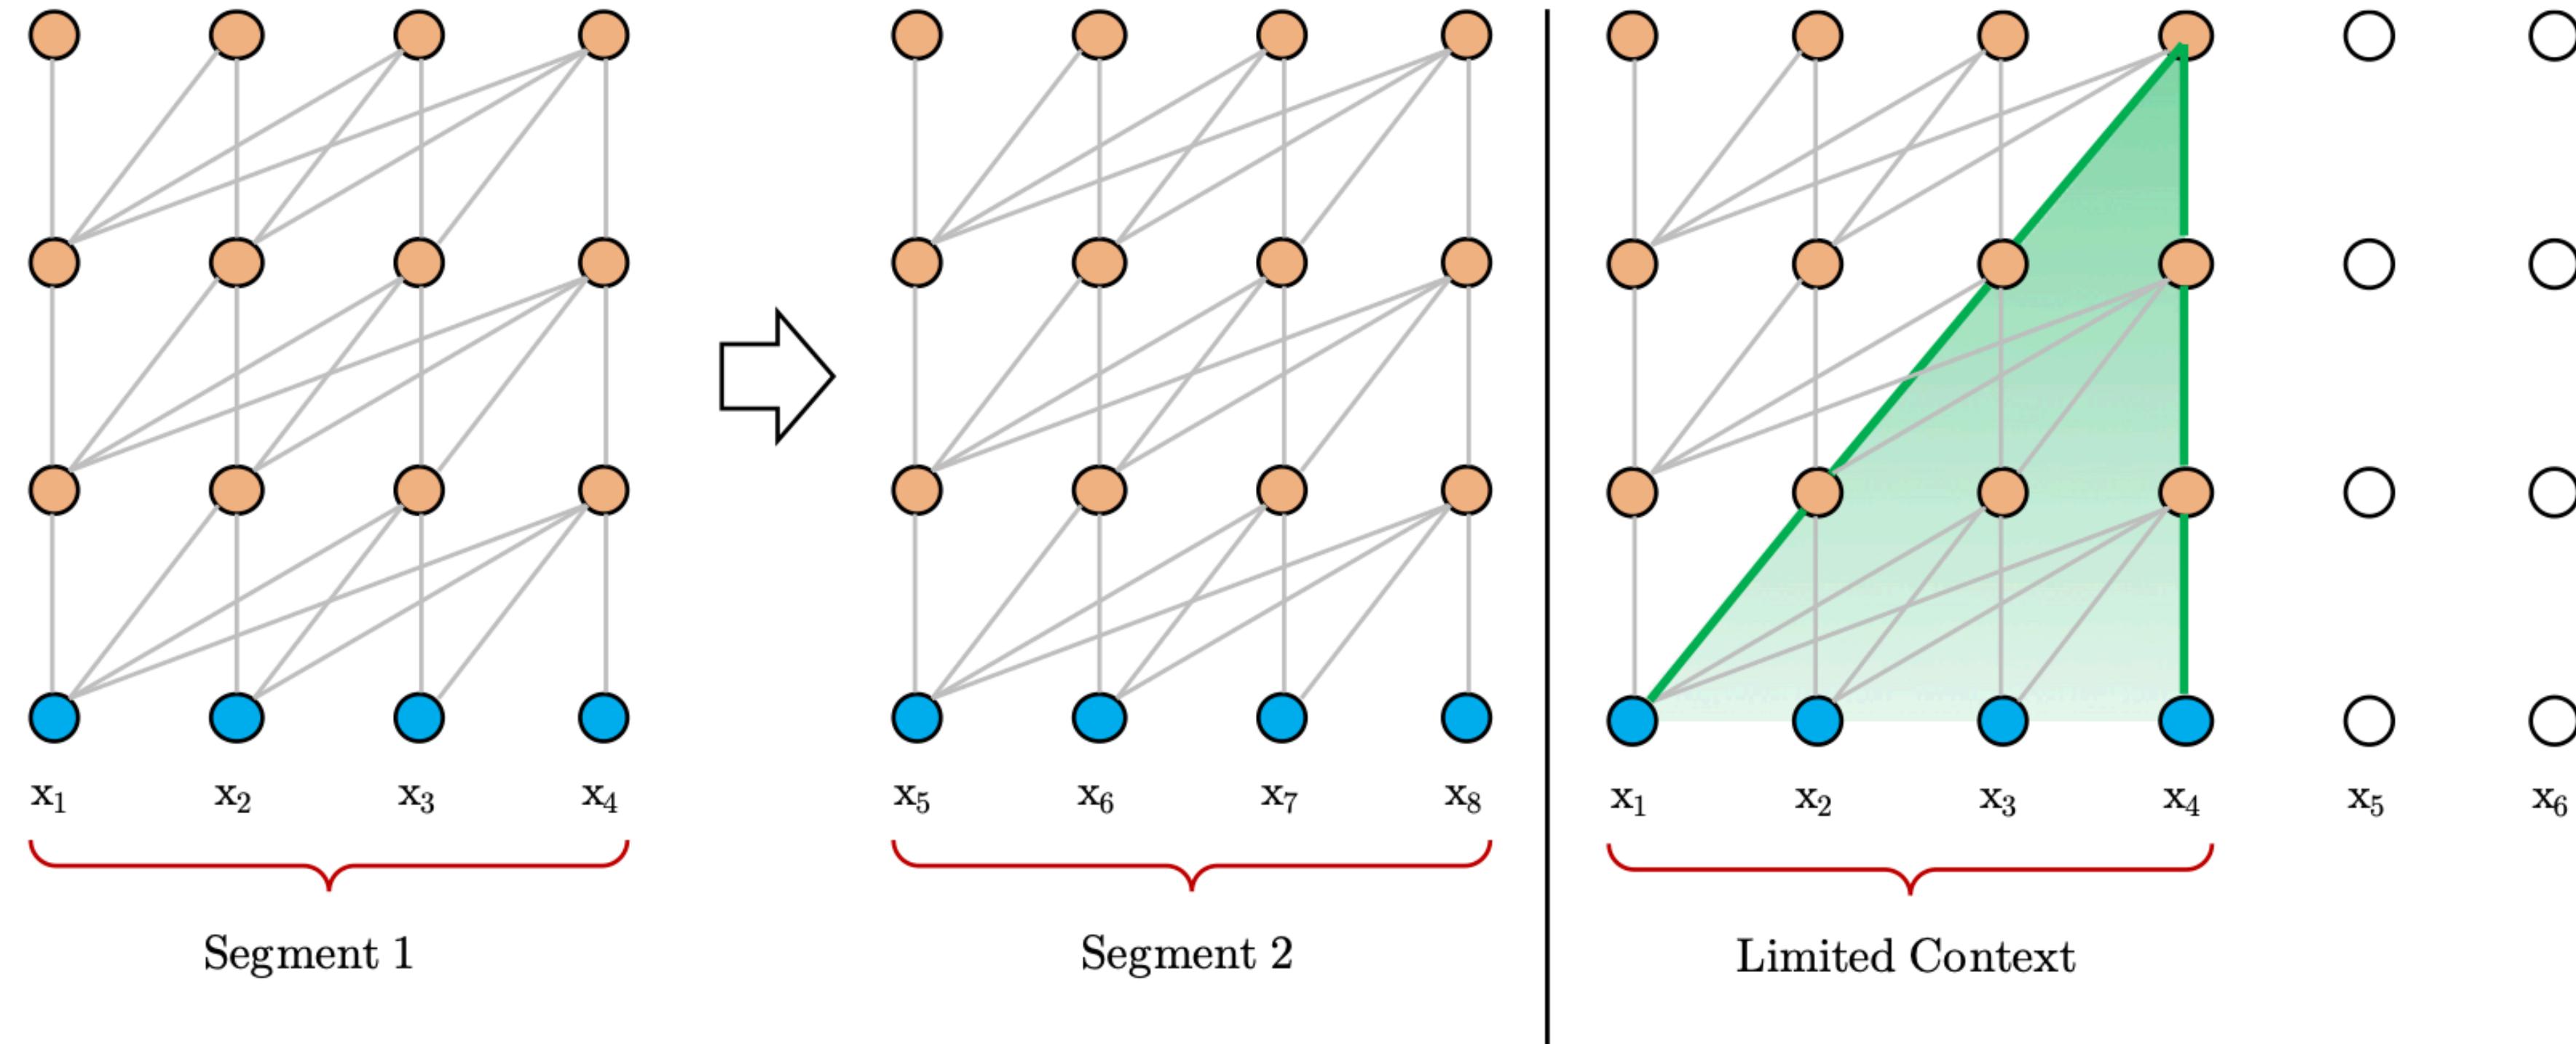
Layer Type	Complexity per Layer	Sequential Operations
Self-Attention	$O(n^2 \cdot d)$	$O(1)$
Recurrent	$O(n \cdot d^2)$	$O(n)$
Convolutional	$O(k \cdot n \cdot d^2)$	$O(1)$
Self-Attention (restricted)	$O(r \cdot n \cdot d)$	$O(1)$

n is sequence length, d is embedding dimension.

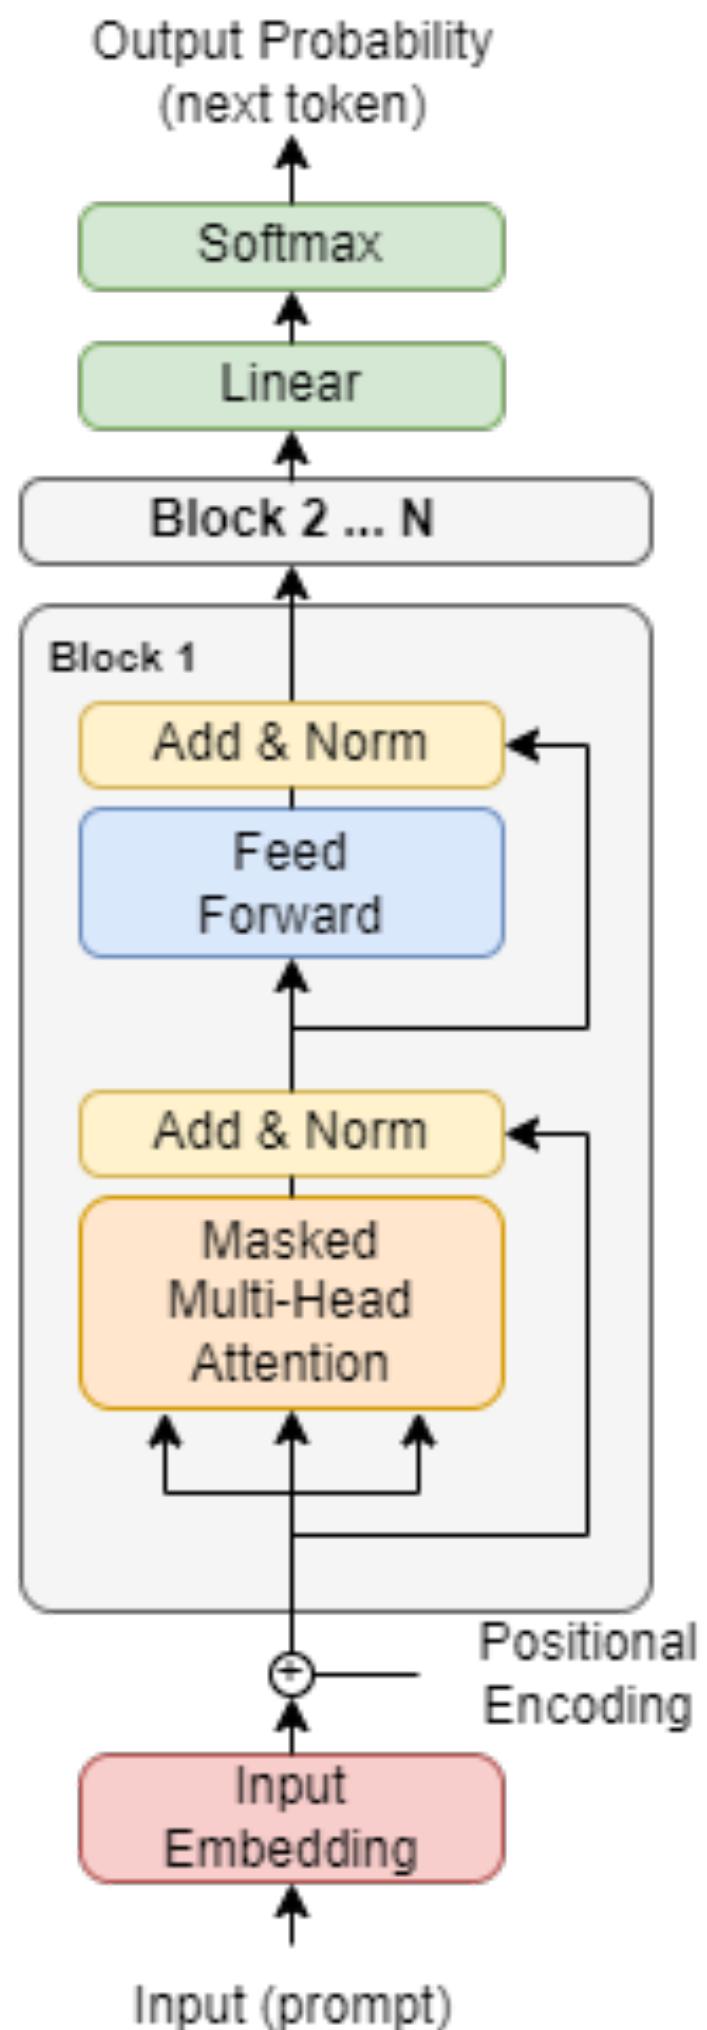
Restricted self-attention means not attending all words in the sequence, but only a restricted field

Square complexity of sequence length is a major issue for transformers to deal with long sequence

# Language Model Training with Limited Context



# Transformer Language Model (e.g., ChatGPT)





香港科技大學  
THE HONG KONG  
UNIVERSITY OF SCIENCE  
AND TECHNOLOGY

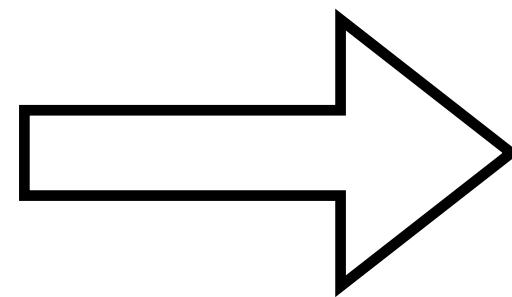
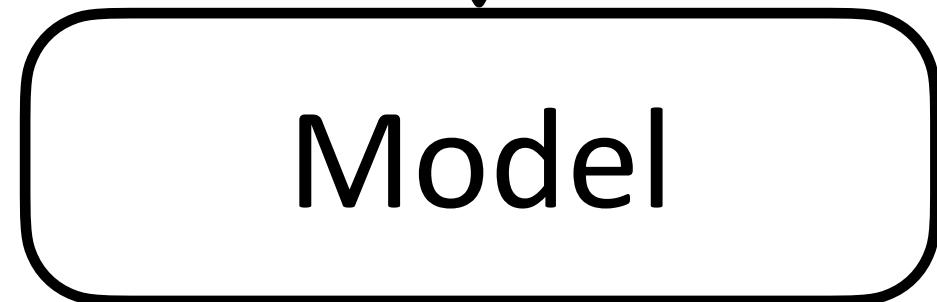
COMP 4901B  
Large Language Models

# Language Model Pretraining

# Pretraining

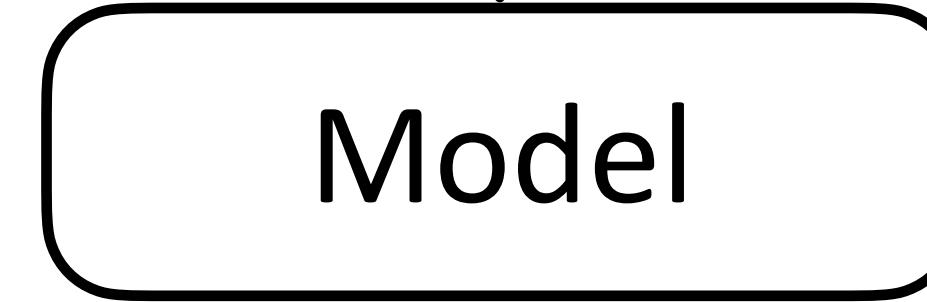
Source Data A (maybe a different task)

Train on data A first



Target Data B

Then train on data B



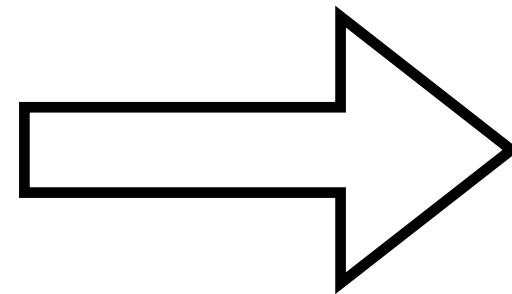
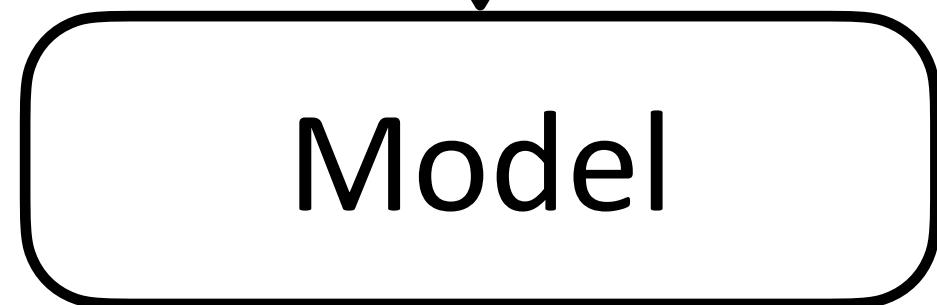
Classically, this is transfer Learning

It is now called pretraining because of the scale of A

# Pretraining

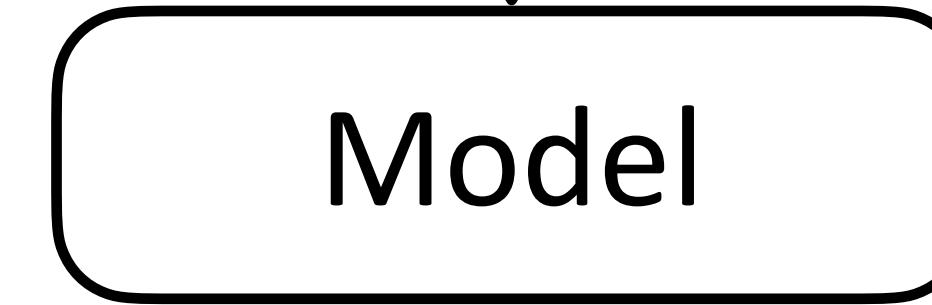
Source Data A (maybe a different task)

Train on data A first



Target Data B

Then train on data B

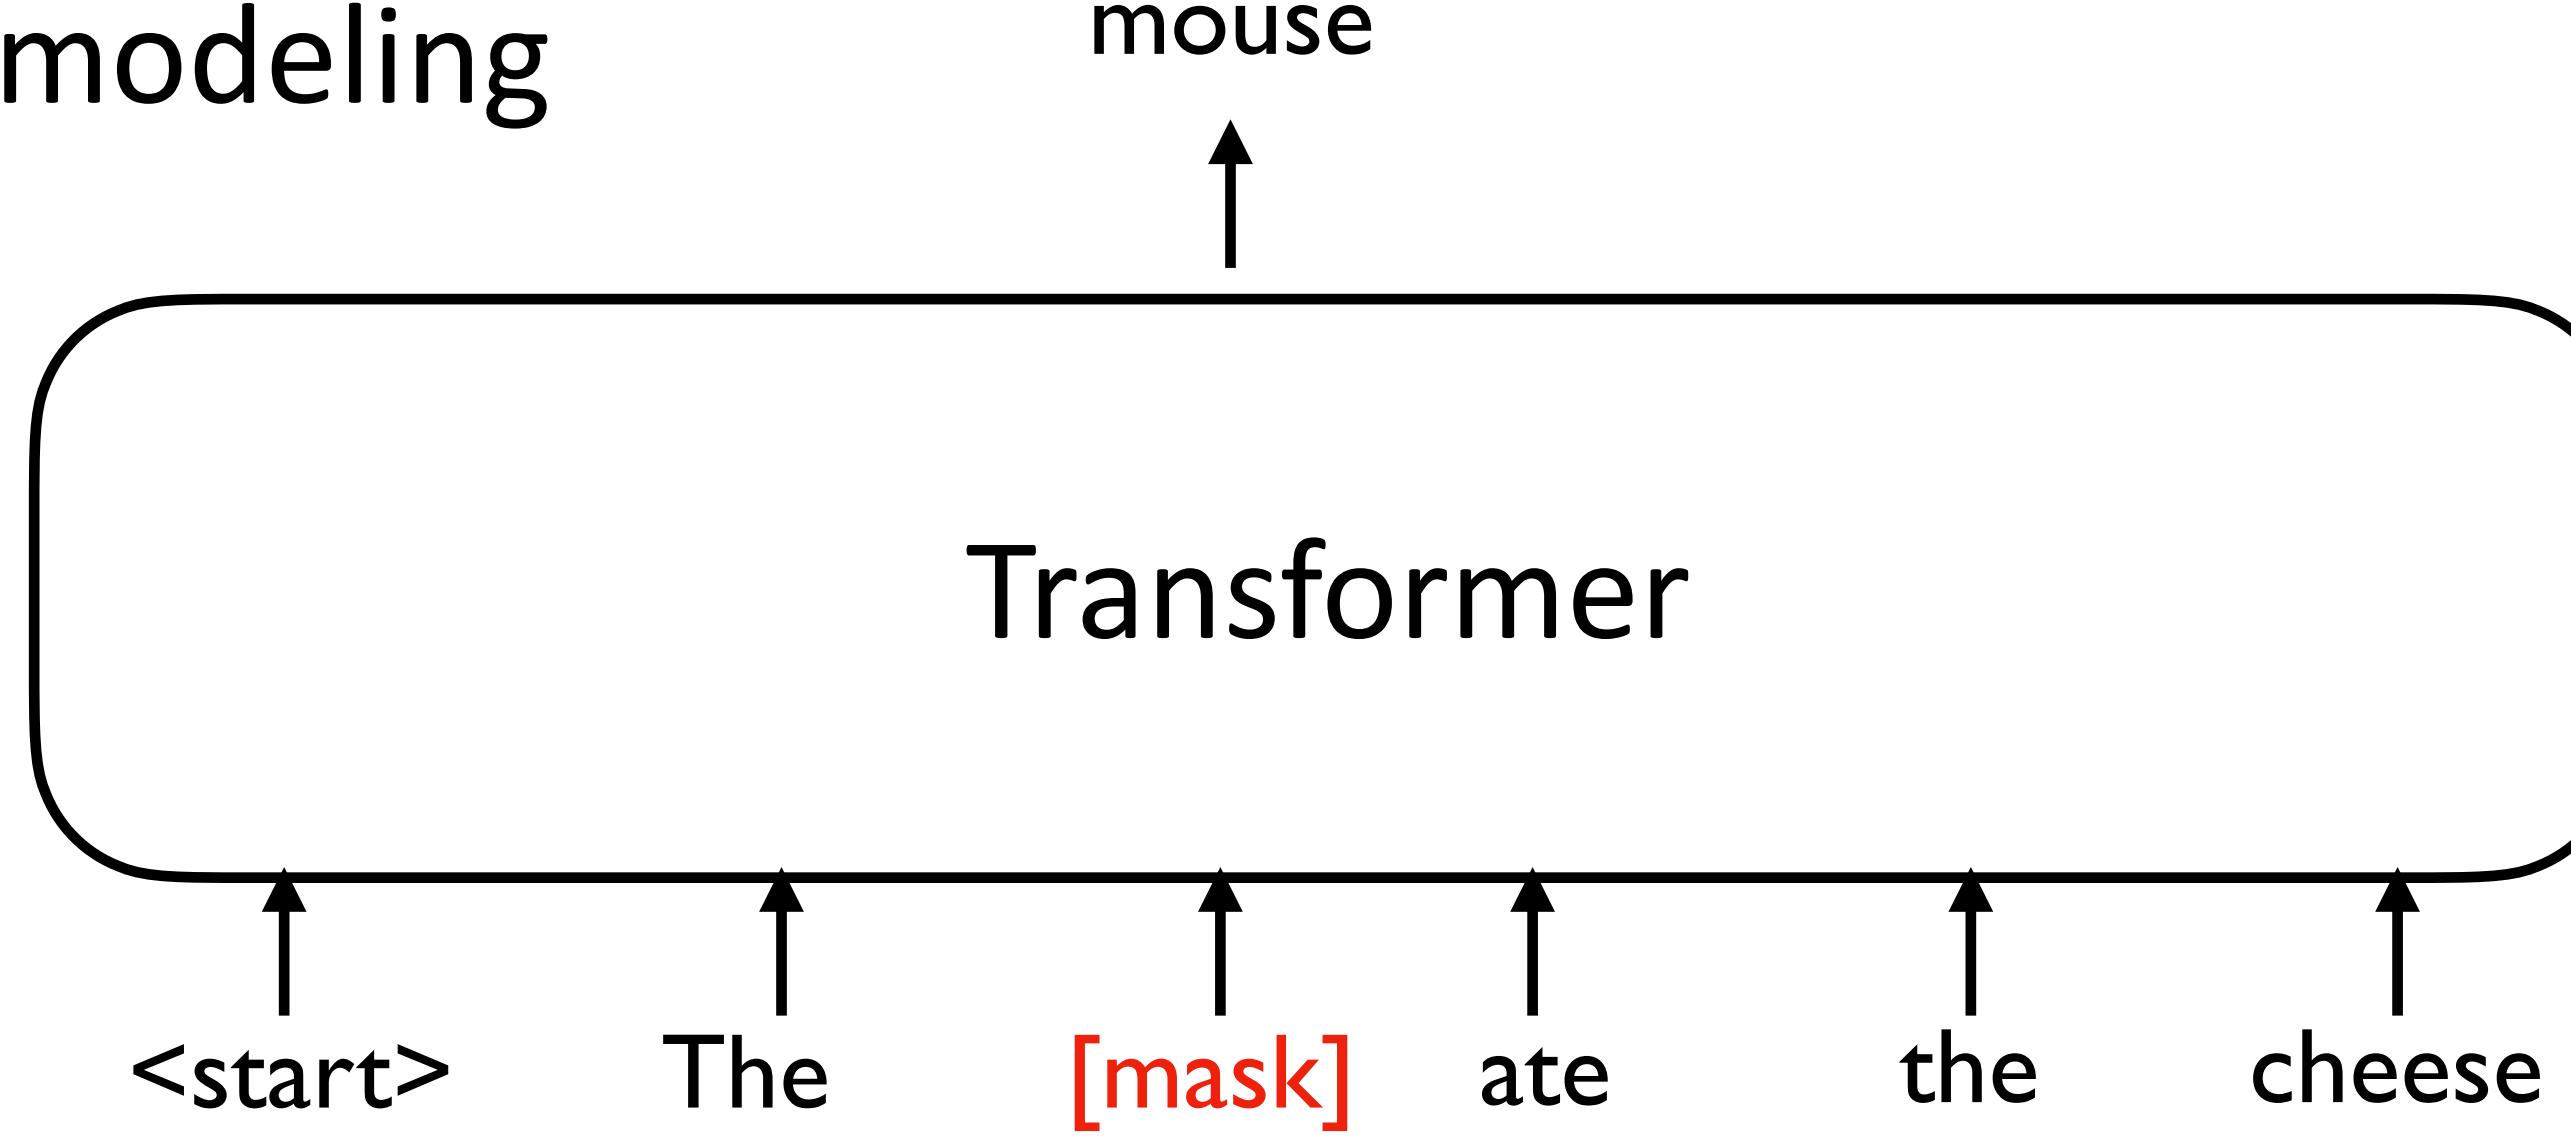


For supervised training, data A is often limited

How can we find large-scale data A to train?

# BERT

Mask language modeling



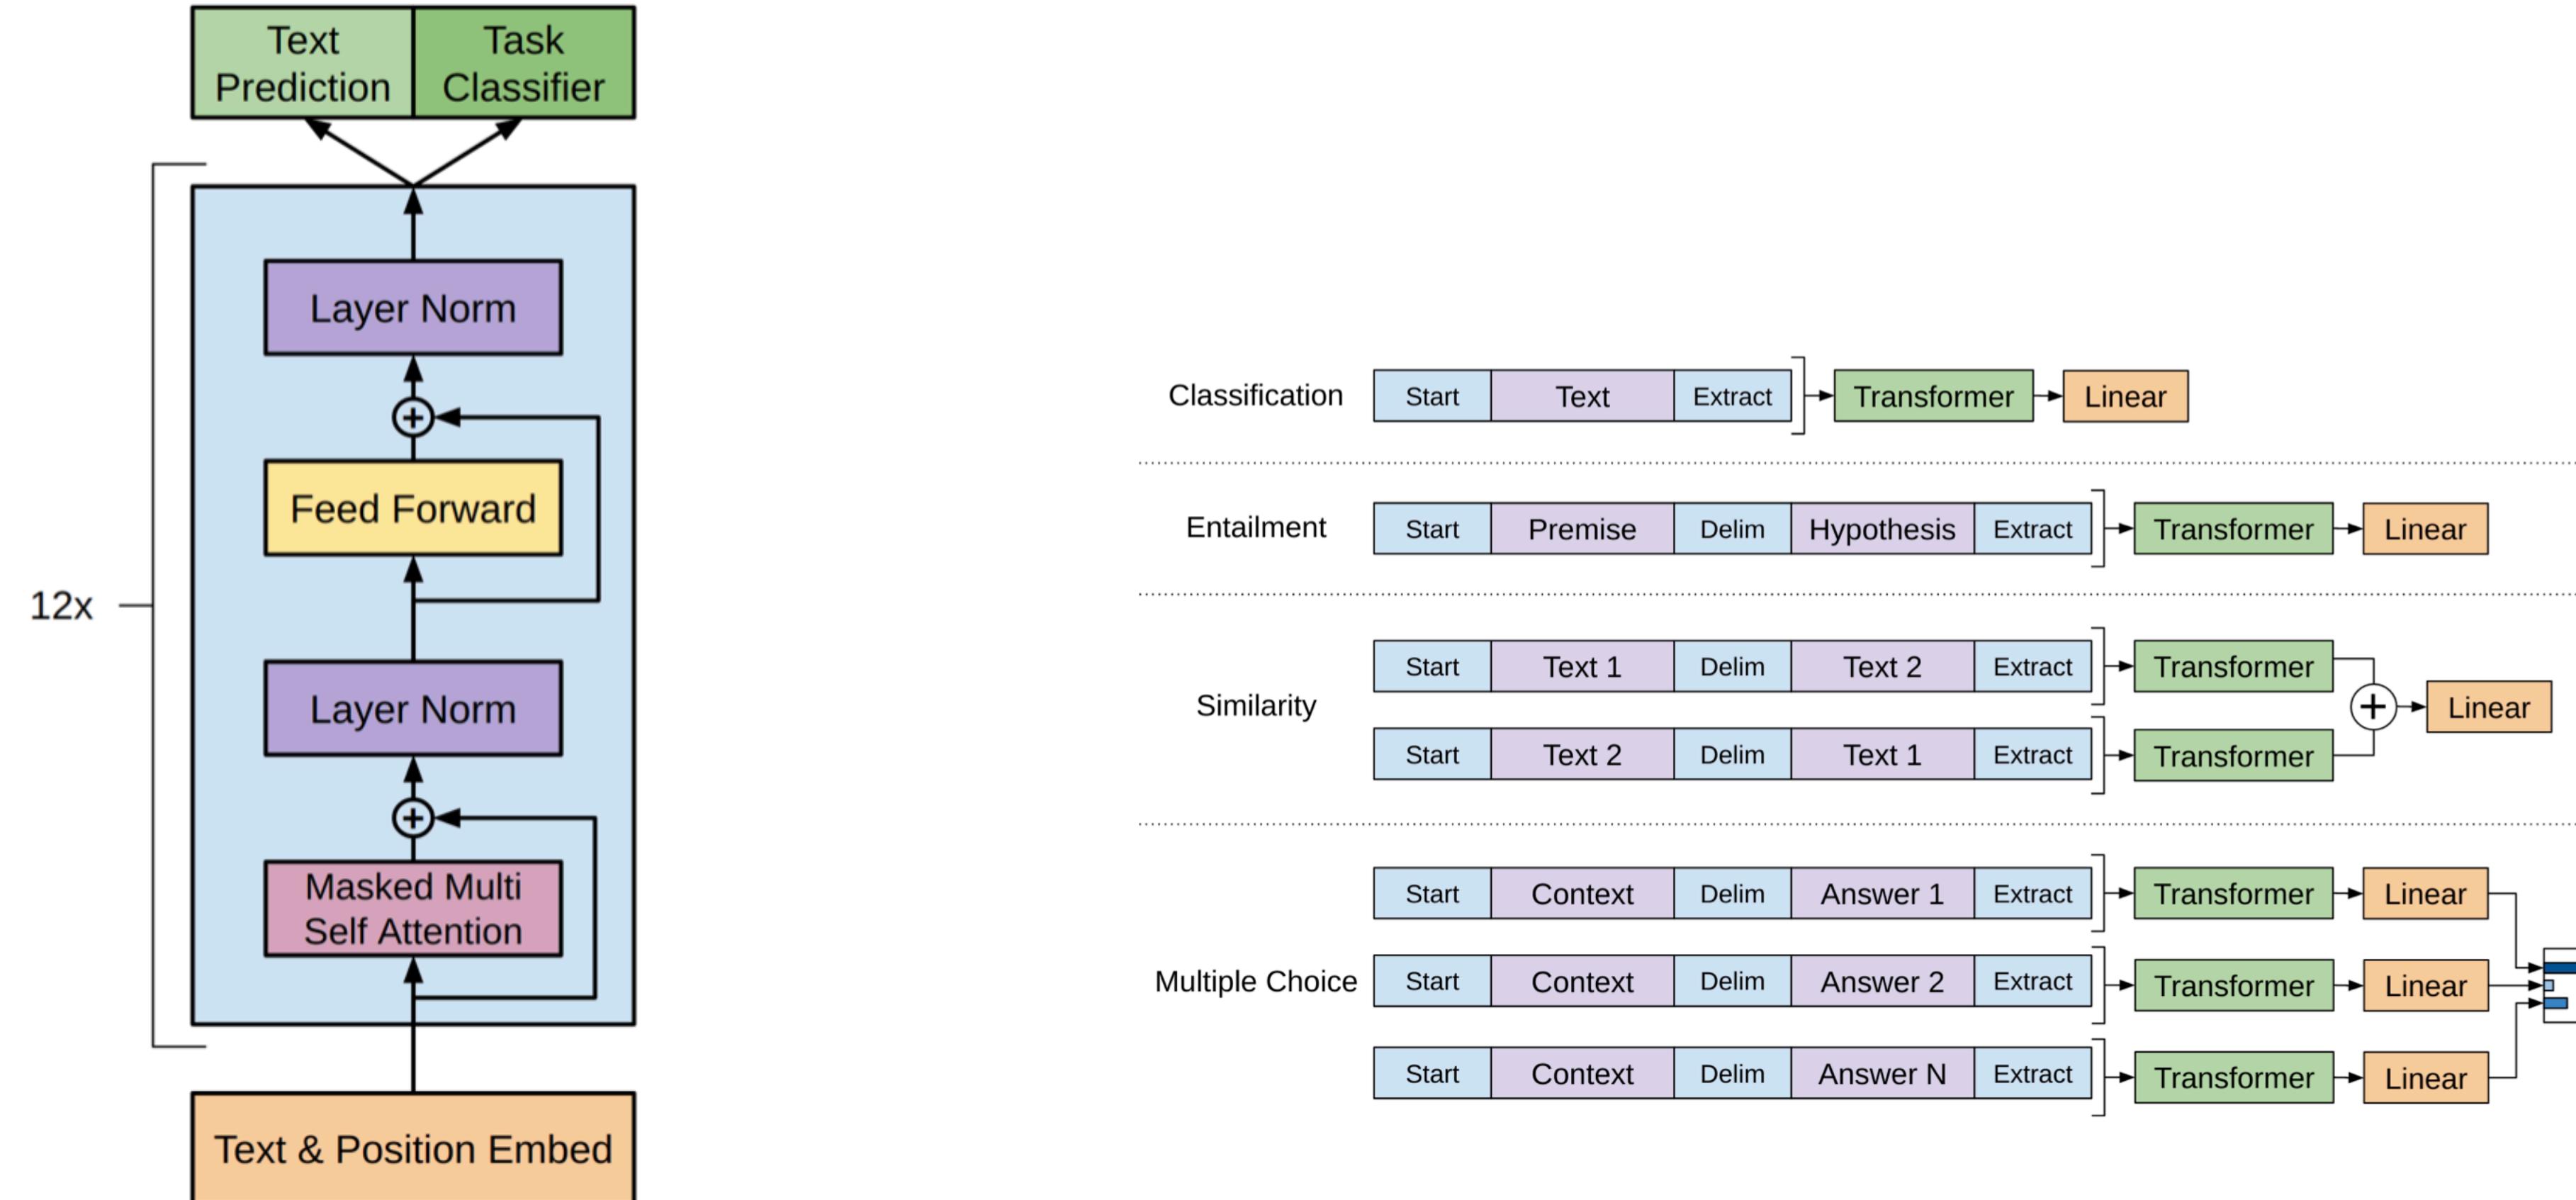
Self-supervised Learning

Construct a synthetic task from raw text only  
Can be made very large-scale

Is Bert a language model? Is it a generative model?

Devlin et al. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding. NAACL 2019.

# Generative Pre-Training (GPT)



Suppose I just want to do a sentence classification task, bidirectional or masked attention is better?

Suppose I just want to do a sentence classification task, bidirectional or masked attention is better for pretraining?

# Pretraining Data

We want to start with clean text

- Wikipedia
- Books

## History [edit]

In the late 1980s, the [Hong Kong Government](#) anticipated a strong demand for university graduates to fuel an economy increasingly based on services. [Sir Sze-Yuen Chung](#) and the territory's governor, [Sir Edward Youde](#), conceived the idea of establishing a third university, in addition to the pre-existing [University of Hong Kong](#) and [Chinese University of Hong Kong](#).<sup>[7]</sup>

Planning for the "Third University", as the university was known provisionally, began in 1986. On 8 November 1989, [Charles, Prince of Wales](#) (now King Charles III) laid the foundation stone of the campus,<sup>[8]</sup> which was constructed at the [Kohima Barracks](#) site in [Tai Po Tsai](#) on the [Clear Water Bay Peninsula](#). The site was earmarked for the construction of a new British Army garrison to house the [2nd King Edward VII's Own](#) and [7th Duke of Edinburgh's Own Gurkha Rifles](#),<sup>[9]</sup> but plans for its construction were shelved after the 1984 signing of the [Sino-British Joint Declaration](#) resulted in the downsizing of army presence in Hong Kong.<sup>[10]</sup>

Originally scheduled to finish in 1994, the planning committee for the university decided in 1987 that the new institution should open its doors three years early, in keeping with the community's need and in fulfilment of the wishes of Youde, who died in 1986.<sup>[11][12]</sup> The university was officially opened by Youde's successor as governor, [Sir David Wilson](#), on 10 October 1991.<sup>[13]</sup> Several leading scientists and researchers took up positions at the university in its early years, including physicist [Leroy Chang](#) who arrived in 1993 as Dean of Science and went on to become vice-president for academic affairs.<sup>[14]</sup> Thomas E. Stelson was also a founding member of the administration.<sup>[15]</sup>

# Pretraining Data Reality

In practice, the **web** is the most viable option for data collection.

In the digital era, this is the go-to place for general domain human knowledge.

But web data can be challenging to work with

- Copyright and usage constraints, privacy
- Data is noisy, dirty, and biased

# Example Noisy Web Data

```
<html ⚡ lang="en">
<head>
<meta http-equiv="Content-Type" content="text/html; charset=windows-1252">
<title>Best Coffee Beans 2025 | Best Coffee Beans 2025 | Buy Coffee Now!</title>
<meta name="description" content="best coffee beans best coffee beans best coffee be">
<link rel="canonical" href="http://example.com/best-coffee?utm_source=spam&utm_campa">
<meta property="og:title" content="Best Coffee Beans 2025">
<script type="application/ld+json">
{@context:"http://schema.org","@type":"Article","headline":"Best Coffee Beans 2025"
</script>
<script>
// tracking & A/B test noise
(function(){try{var u='https://trk.example.net/p.js?id=UA-XXX';var s=document.create
window.__AB__={"exp":"homepage-v17","bucket":Math.random()<0.5?'A':'B'};
</script>
<style>
/* inline CSS with dead classes */
.hero {background:url(data:image/png;base64,iVBORw0KGgoAAAANSUhEUgAA...);height:48
.hidden{display:none} .cookie{position:fixed;bottom:0;background:#000;color:#fff;p
@media (max-width: 600px){ .table{display:block;overflow:auto} }
</style>
</head>
<body id="top" class="post post-1234" oncopy="return false">
<!-- BEGIN Cookie banner, duplicated --&gt;
&lt;div class="cookie" role="dialog" aria-live="polite"&gt;
We use cookies to improve your experience. &lt;button id="ok"&gt;OK&lt;/button&gt;
&lt;/div&gt;
&lt;div class="cookie" style="display:none"&gt;We use cookies &lt;a href="/privacy?ref=popup"
&lt!-- END Cookie banner --&gt;

&lt;noscript&gt;&lt;img src="https://track.example.org/pixel.gif?id=abc" alt=""&gt;&lt;/noscript&gt;

&lt;header&gt;
&lt;h1&gt;Best Coffee Beans 2025&lt;/h1&gt;
&lt;div class="rating"&gt;★★★★★ 4.9/5 (3,214)&lt;/div&gt;
&lt;div class="breadcrumbs"&gt;
&lt;a href="/"&gt;Home&lt;/a&gt; &gt; &lt;a href="/category?c=cof%66ee"&gt;Coffee&lt;/a&gt; &gt; Best Coffee B
&lt;/div&gt;</pre>
```

# Web Data Pipeline

- Content is posted to the web
- Web crawlers identify and download a portion of the content
- The data is filtered and cleaned

# Web Data Pipeline

## General Idea

1. Start with a set of seed websites
2. Explore outward by following all hyperlinks on the webpage.
3. Systematically download each webpage and extract the raw text.



# How to Clean Text Data

1. Remove noisy, spammy, templated, and fragmented texts
2. Select higher quality texts from a massive candidate pool
3. Avoid toxic and biased content

# What Defines Good Pretraining Data?

1. Clean (fluent)
2. Diverse (covers many domains)
3. Non-trivial (a trivial case is to learn from massive documents and each has no more than 20 words)
4. “high-quality”

“intelligence” of the data is high, generally requiring a lot of knowledge and reasoning to predict the next word

# How to Identify High-Quality Content

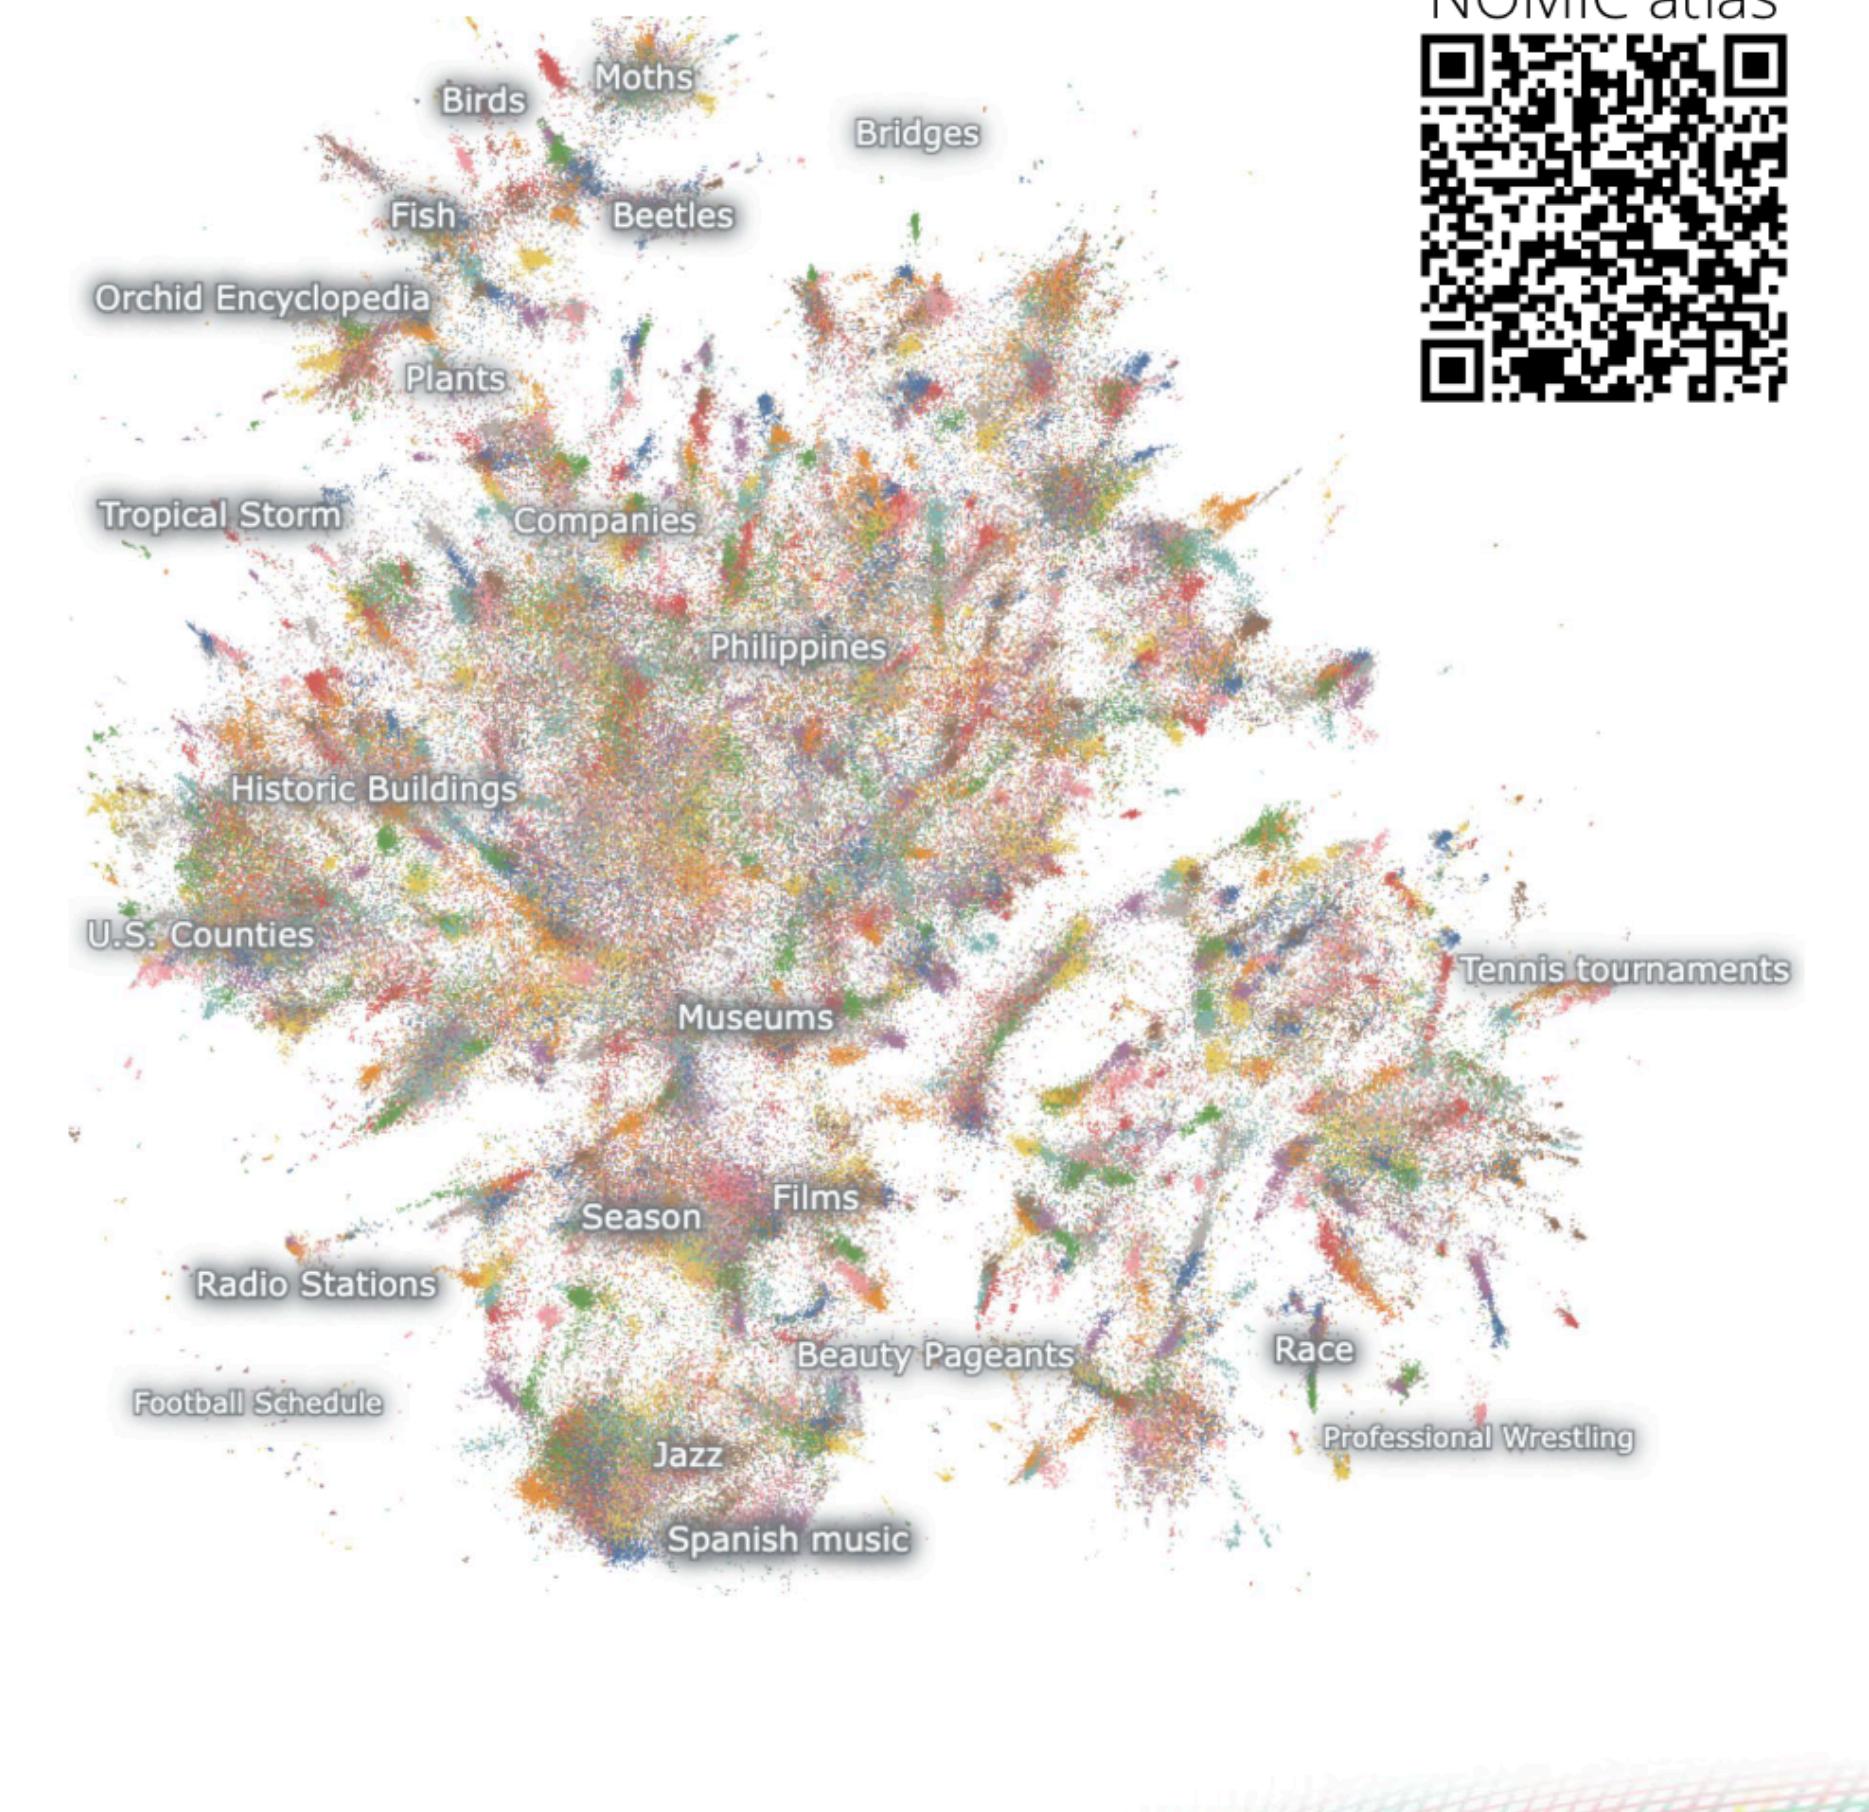
- Rule-based Heuristics
- Classifiers (how to use GPT4 to help train GPT5?)

# Notable Datasets

- Wikipedia dataset
- CommonCrawl
- Colossal Clean Crawled Corpus (C4)
- FineWeb
- Dolma

# Wikipedia Dataset

- Contains cleaned articles (65M) written in many languages (~350).
- The dataset is built from the Wikipedia dumps and split per language.
- Each example contains a cleaned article with stripped markdown and unwanted sections.
- The data fields are **id**, **url**, **title**, and **text**.
- Conveniently available on HuggingFace.

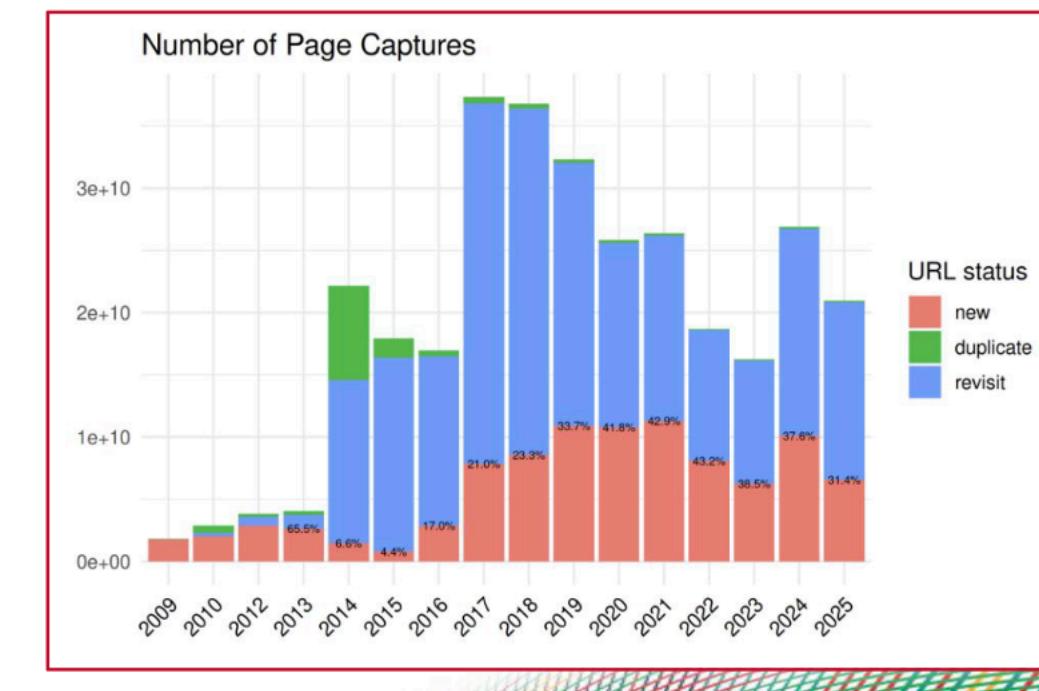
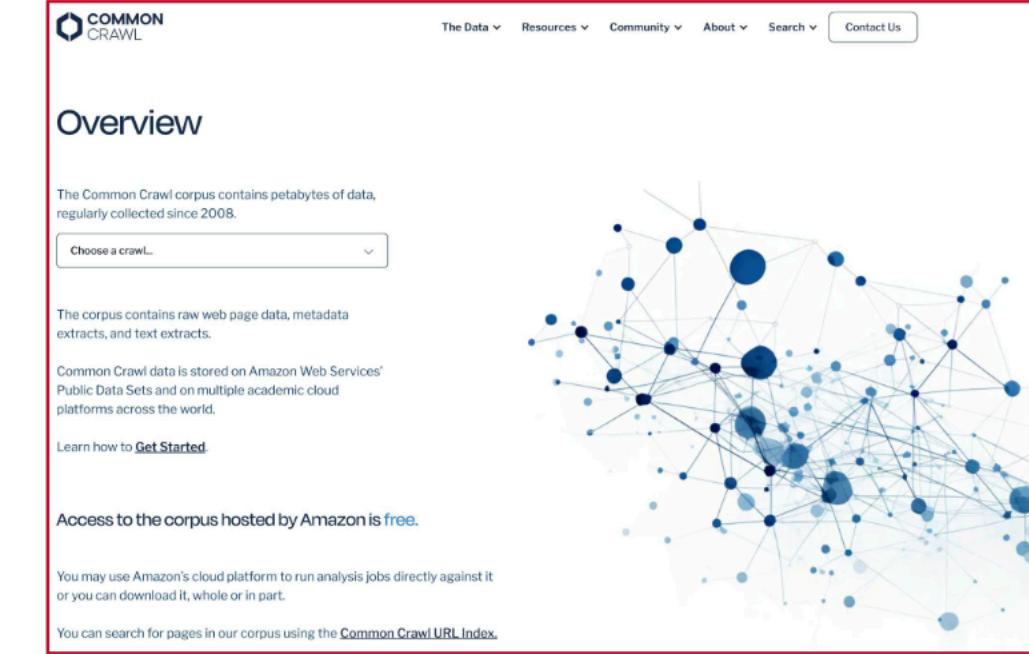


NOMIC atlas



# CommonCrawl

- Non-profit organization that provides open access to large scale web crawls
- Petabytes of web pages are available
- Monthly crawls and dumps
  - Re-crawled web pages and fresh dumps (bi)monthly
  - The dumps are ~k billion pages
- Dates back to 2008



# Preprocessing Clean Text

After the text is cleaned, now we need to convert it into a batch of training data

The Steelers enjoy a large,  
widespread fanbase  
nicknamed Steeler Nation.  
They currently play their  
home games at Acrisure  
Stadium.

Raw Clean Text

Tokenization

'\_The', '\_Steel', 'ers', '\_enjoy', '\_a',  
'\_large', ',', '\_wide', 'spre', 'ad', '\_fan',  
'base', '\_nick', 'na', 'med', '\_Steel', 'er',  
'\_Nation', '.', '\_They', '\_currently',  
'\_play', '\_their', '\_home', '\_games',  
'\_at', '\_A', 'cris', 'ure', '\_Stadium', ''

Tokenized

Batching

[580, 109027, 1313, 25224, 9,  
21333, 3, 38133, 21328, 711, 1206,  
37381, 128910, 75, 4805, 109027,  
55, 82580, 4, 0]  
[10659, 82423, 11300, 2362,  
5367, 27527, 98, 61, 58531, 3407,  
88259, 4, 0, 0, 0, 0, 0, 0, 0]

Tensor

# Tokenizing Text

A **tokenizer** takes text and turns it into a sequence of discrete **tokens**

A **vocabulary** is a list of all available tokens

Example: “A hippopotamus ate my homework”

Vocab Type	Example	Length
character-level	<code>['A', ' ', 'h', 'i', 'p', 'p', 'o', 'p', 'o', 't', 'a', 'm', 'u', 's', ' ', 'a', 't', 'e', ' ', 'm', 'y', ' ', 'h', 'o', 'm', 'e', 'w', 'o', 'r', 'k', '.']</code>	31
subword-level	<code>['A', 'hip', '#pop', '#ota', '#mus', 'ate', 'my', 'homework', '.']</code>	9
word-level	<code>['A', 'hippopotamus', 'ate', 'my', 'homework', '.']</code>	6

# Word-Level Tokenization

rule-based (split text by spaces, punctuation, and other similar heuristics)

## Challenges

- Open vocabulary problem
  - Many words may never appear in training data (becomes [UNK])
  - This is more severe in other low-resource languages
- Words with typos also get tokenized as [UNK]

# Character-Level Tokenization

Vocab Type	Example	Length
character-level	<code>['A', ' ', 'h', 'i', 'p', 'p', 'o', 'p', 'o', 't', 'a', 'm', 'u', 's', ' ', 'a', 't', 'e', ' ', 'm', 'y', ' ', 'h', 'o', 'm', 'e', 'w', 'o', 'r', 'k', '.']</code>	31
subword-level	<code>['A', 'hip', '#pop', '#ota', '#mus', 'ate', 'my', 'homework', '.']</code>	9
word-level	<code>['A', 'hippopotamus', 'ate', 'my', 'homework', '.']</code>	6

Pro: No unseen tokens anymore

Con: Sequence is unnecessarily long, expensive to work with

# Sub-word Tokenization

- Words get split into multiple tokens
- Vocabulary is build dynamically
  - Frequent words get assigned their own tokens
  - Rare words are split into subwords

Vocab Type	Example	Length
character-level	<code>['A', ' ', 'h', 'i', 'p', 'p', 'o', 'p', 'o', 't', 'a', 'm', 'u', 's', ' ', 'a', 't', 'e', ' ', 'm', 'y', ' ', 'h', 'o', 'm', 'e', 'w', 'o', 'r', 'k', '.']</code>	31
subword-level	<code>['A', 'hip', '#pop', '#ota', '#mus', 'ate', 'my', 'homework', '.']</code>	9
word-level	<code>['A', 'hippopotamus', 'ate', 'my', 'homework', '.']</code>	6

# Byte Pair Encoding (BPE)

## Main Idea

- Construct subword vocabulary by learning to merge characters
- Inspiration comes from compression algorithms

## Training Steps

1. Initialize the vocabulary with characters as tokens (e.g., in English: alphabet, numbers, punctuation)
2. Merge the most frequent token pair in the corpus (vocabulary size +1)
3. Re-tokenize the corpus with the merged subword pair
4. Repeat steps 2 and 3 until the target vocabulary size is reached

# Advantages of Subword Tokenization

- Controlled vocabulary size
- Strike a good balance between word-level and character-level
  - Frequent words kept whole
  - Tail words split to sub-words
  - More observations on sub-words
  - Utilization of morphology information

# Batching Data

The Steelers enjoy a large, widespread fanbase nicknamed Steeler Nation. They currently play their home games at Acrisure Stadium.

**Raw Clean Text**

**Tokenization**

'\_The', '\_Steel', 'ers', '\_enjoy', '\_a',  
'\_large', '!', '\_wide', 'spre', 'ad', '\_fan',  
'base', '\_nick', 'na', 'med', '\_Steel', 'er',  
'\_Nation', '!', '\_They', '\_currently',  
'\_play', '\_their', '\_home', '\_games',  
'\_at', '\_A', 'cris', 'ure', '\_Stadium', ''

**Tokenized**

**Batching**

[580, 109027, 1313, 25224, 9,  
21333, 3, 38133, 21328, 711, 1206,  
3 381, 128910, 75, 4805, 109027,  
55, 82580, 4, 0]  
[10659, 82423, 11300, 2362,  
5367, 27527, 98, 61, 58531, 3407,  
88259, 4, 0, 0, 0, 0, 0, 0, 0, 0]

**Tensor**

# Thank You!