

Language Model Bias and Safety

Junxian He

Nov 14, 2025

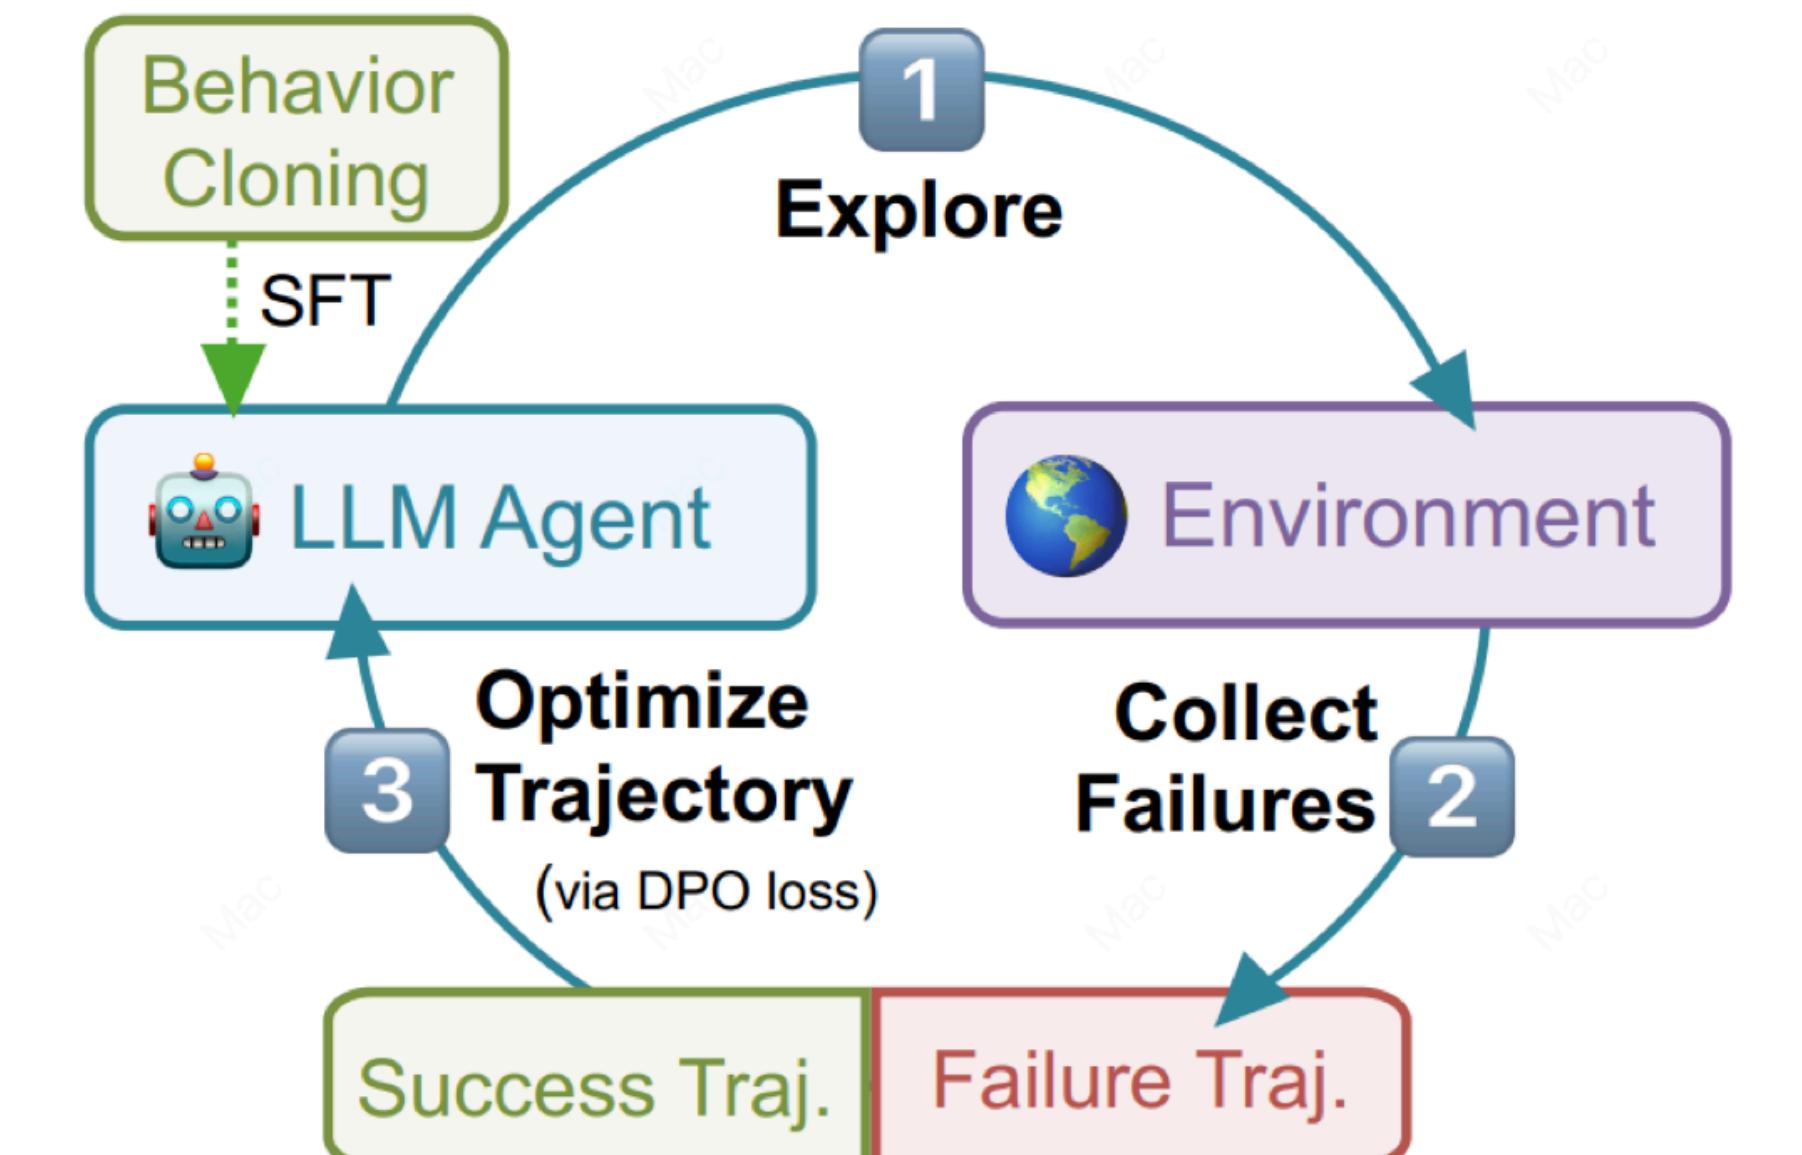
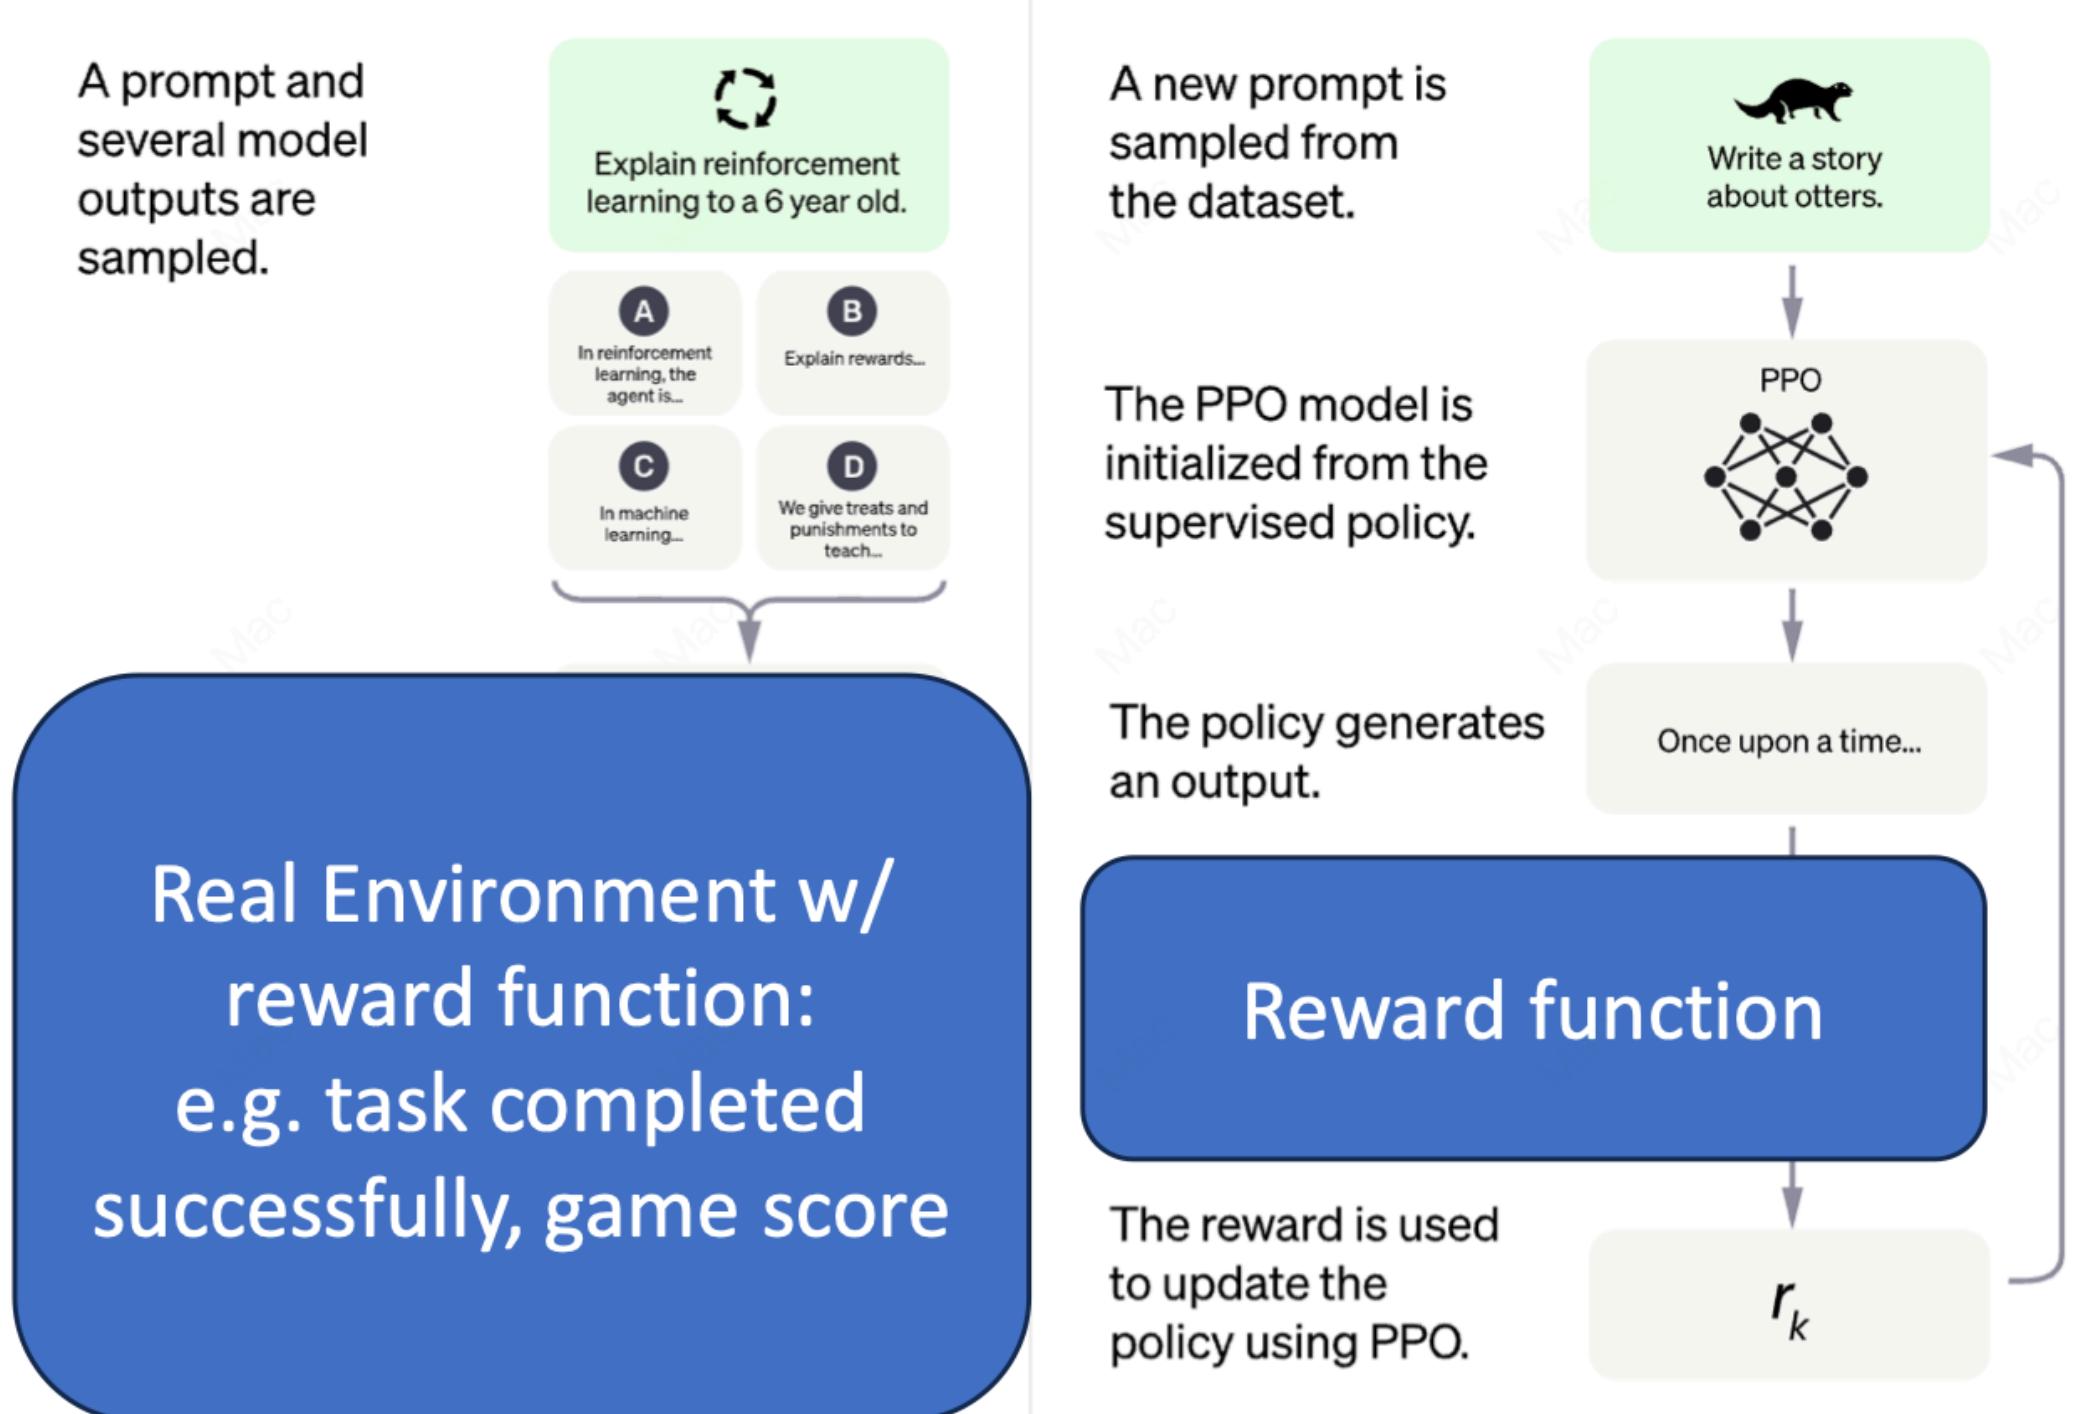
Recap: Learning of LLM Agents

- In-Context Learning – Learning from few-shot exemplars
- Supervised Finetuning – Learning From *Experts*
- Reinforcement Learning – Learning from *Environment*

Recap: Reinforcement Learning

Compared to RLHF:

Given environment, *reward function*
(trajectory, reward) pairs without human



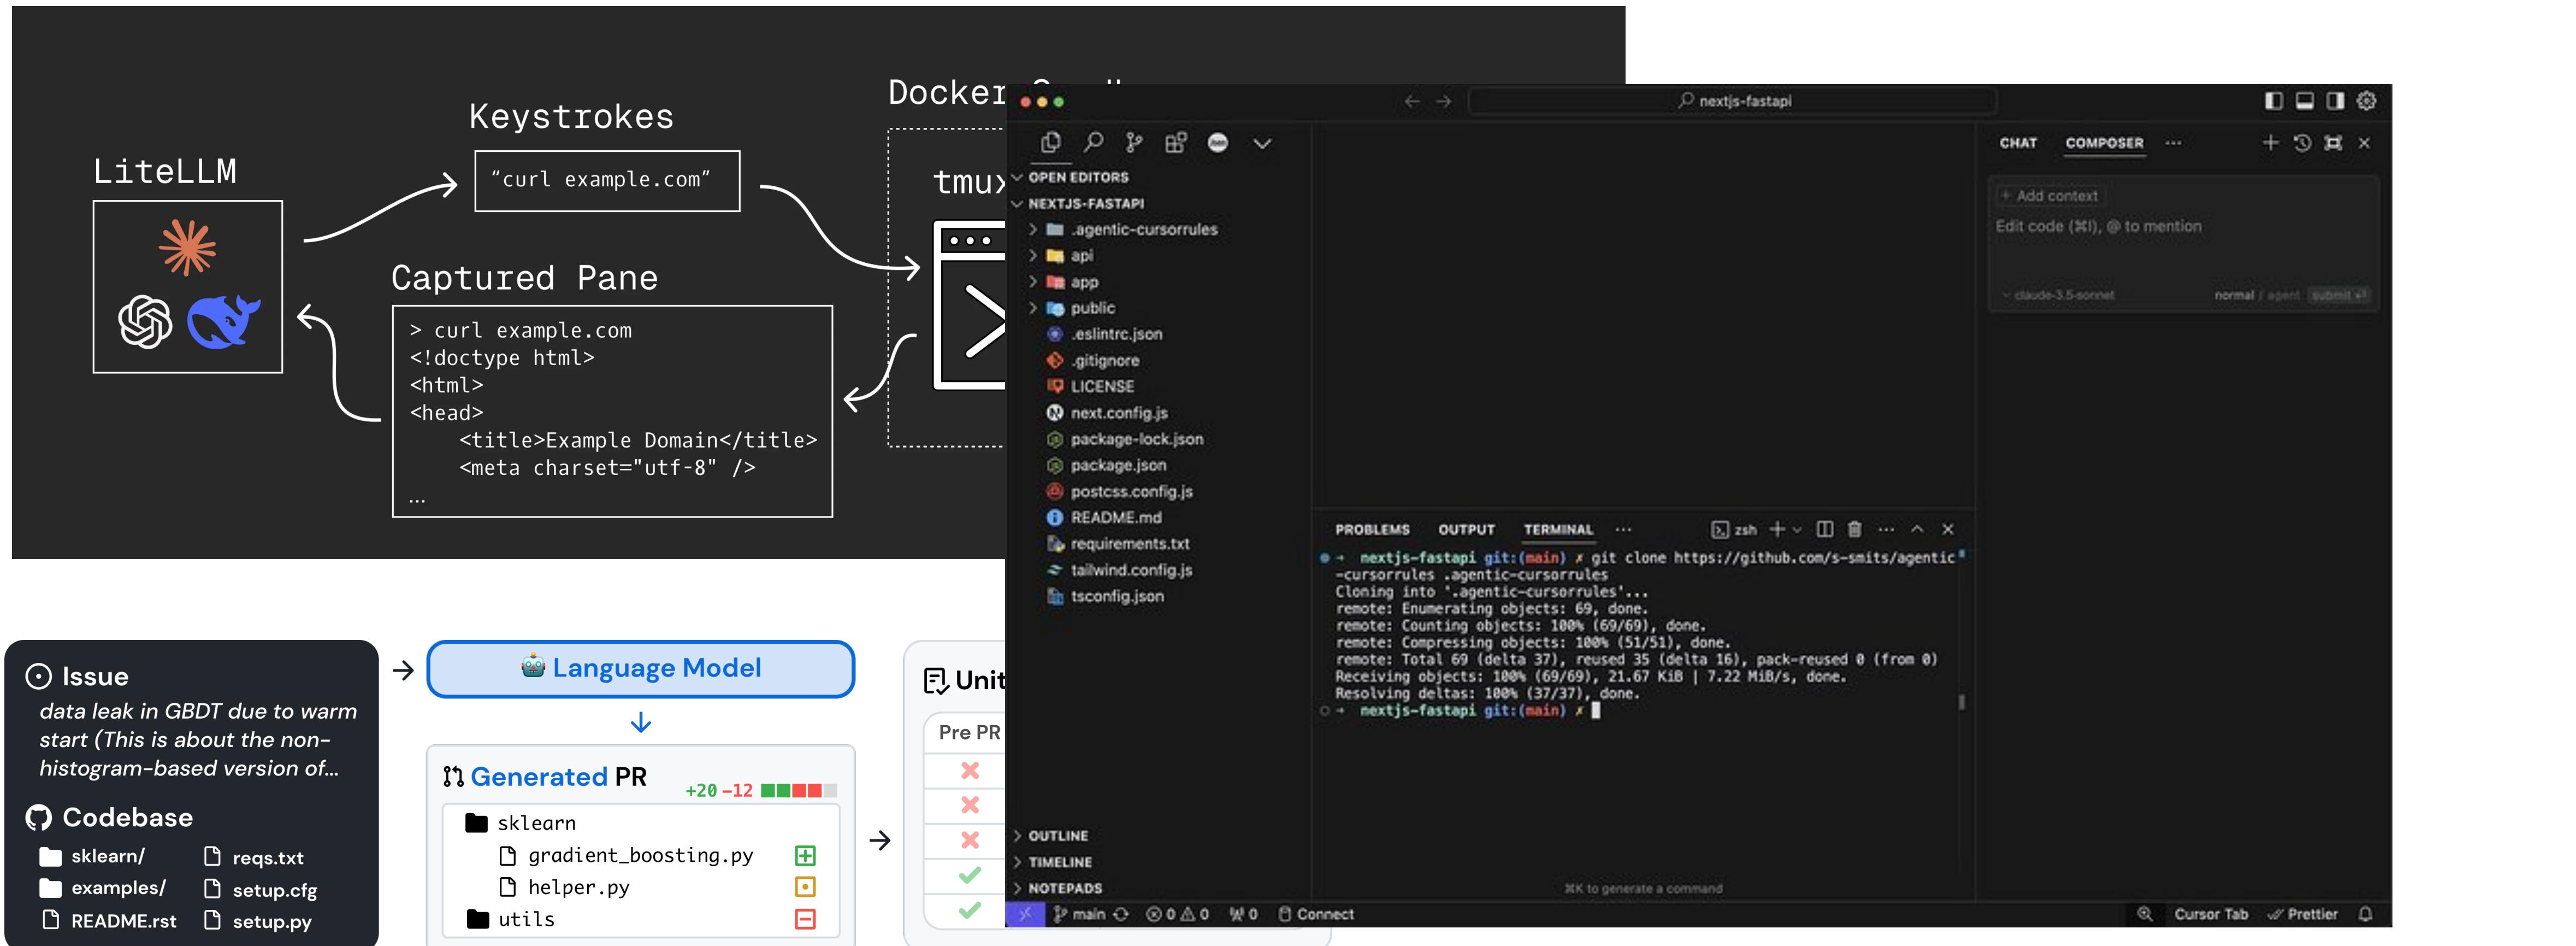
Trial and Error (Song et al. 24')

Reinforcement Learning

- Closed loop, interactive environment
- Need good reward functions
 - What if the task success/fail is not easy to automatically assess?
- Need good initial models
 - Has decent basic knowledge ability, sparse rewards
- Scalability
 - The environment takes 10 seconds to env.step()
 - The reward function takes 100 seconds to get a scalar reward

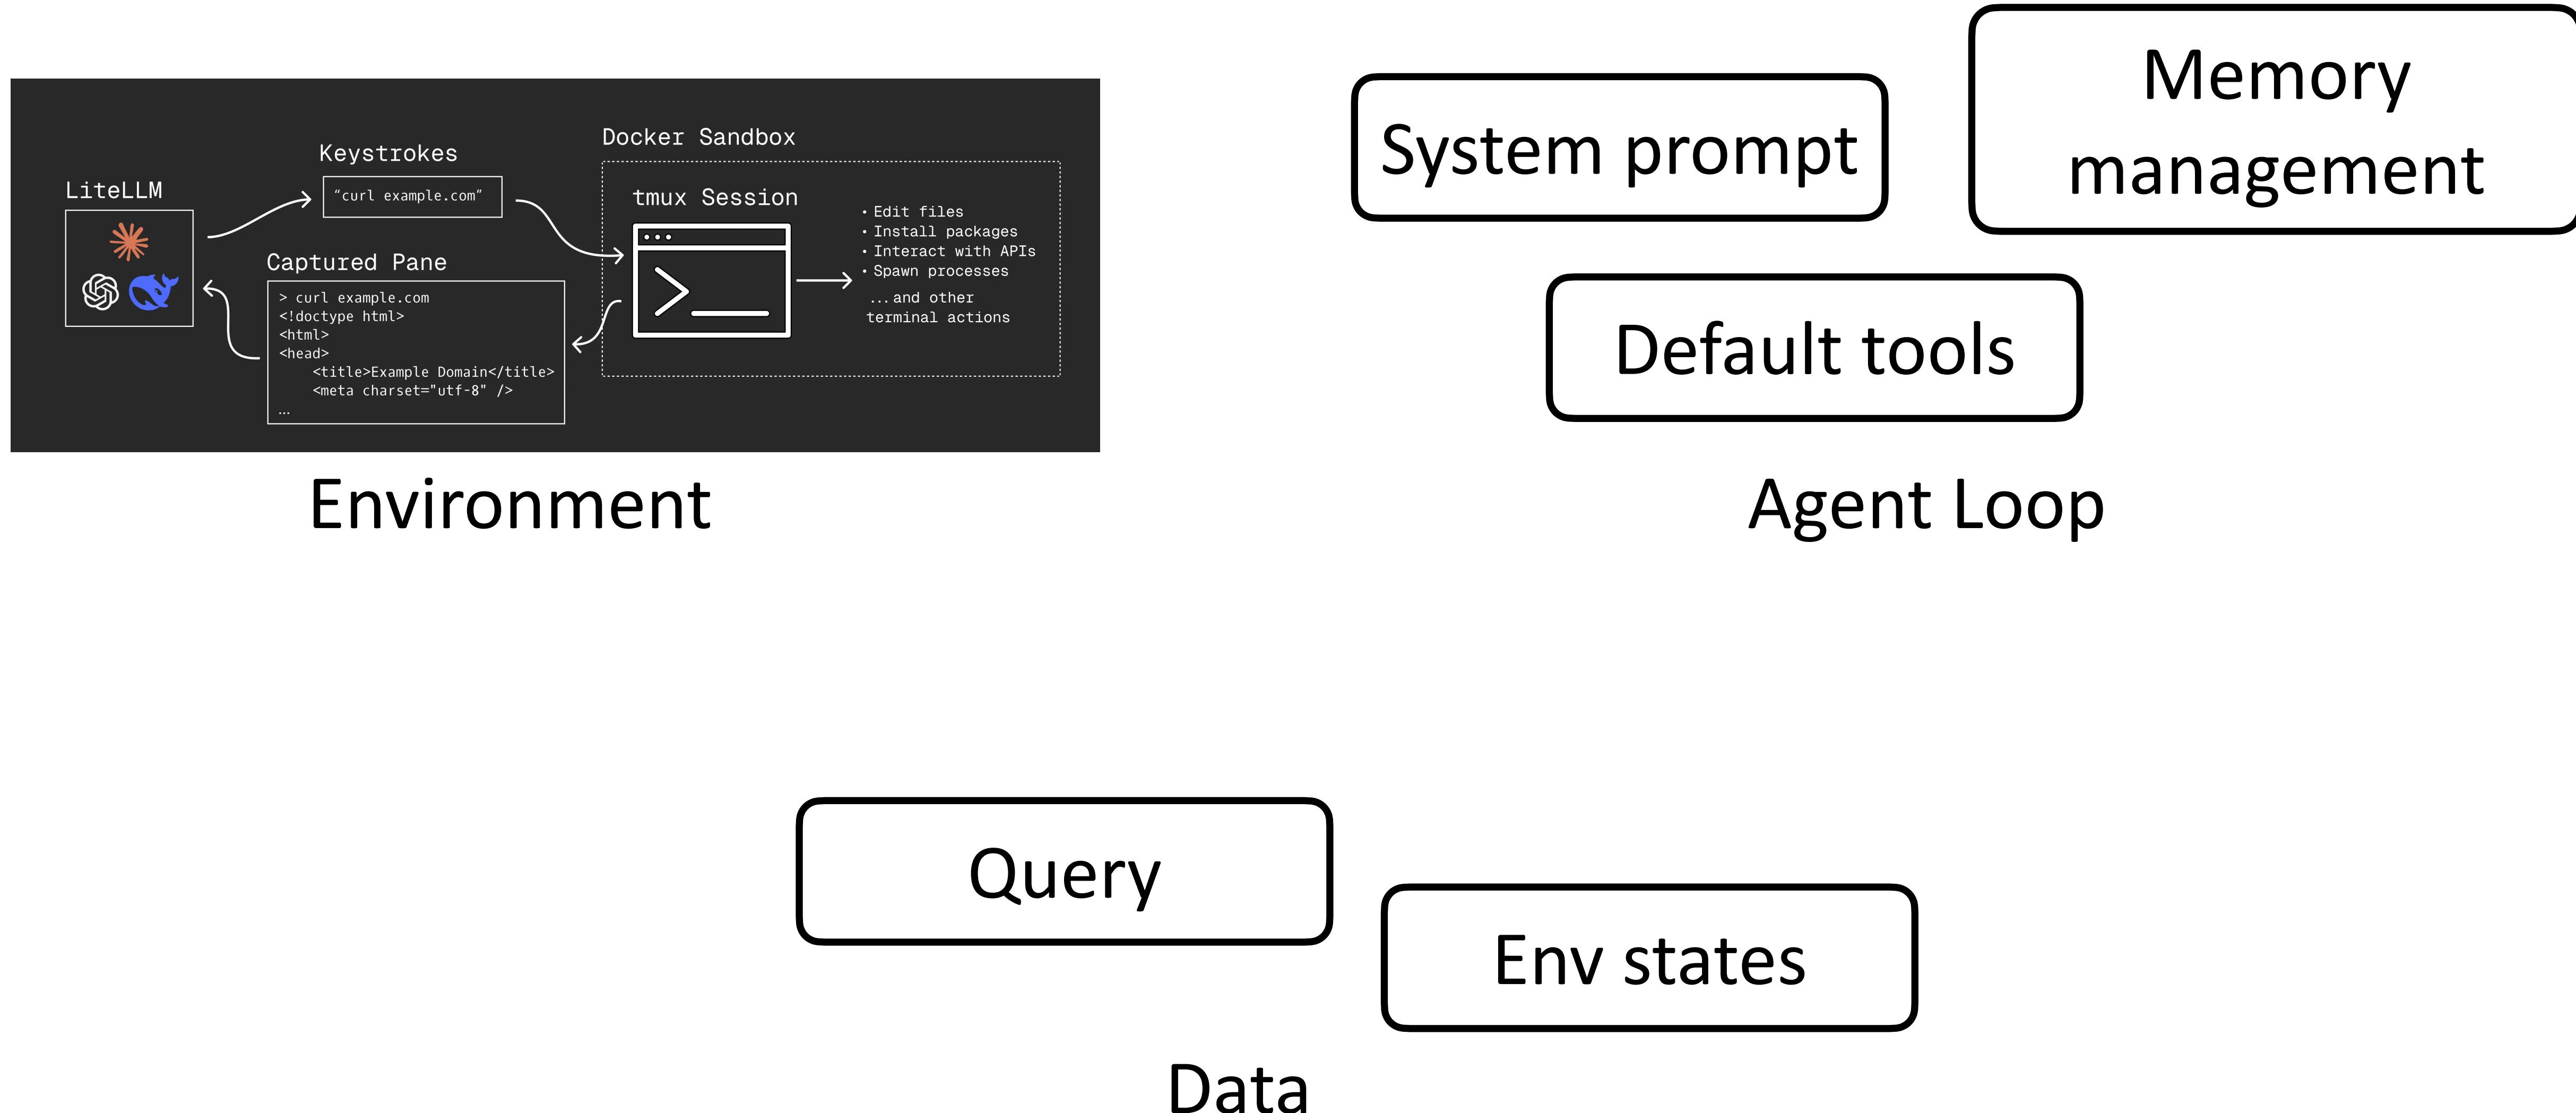
RL Environments

Environments and benchmarks typically come together



Research and Products are really close nowadays, and we can directly RL in real, product-level environments

A Holistic Framework of Agents



Implement a Live Example to Understand Actual Model Input/Output

```
tokenizer = AutoTokenizer.from_pretrained("Qwen/Qwen3-Coder-30B-A3B-Instruct")

# Define the available tools
tools = [
    {
        "type": "function",
        "function": {
            "name": "get_weather",
            "description": "Get the current weather information for a specified location",
            "parameters": {
                "type": "object",
                "properties": {
                    "location": {
                        "type": "string",
                        "description": "The city name, e.g. San Francisco, Tokyo"
                    },
                    "unit": {
                        "type": "string",
                        "enum": ["celsius", "fahrenheit"],
                        "description": "The temperature unit to use"
                    }
                },
                "required": ["location"]
            }
        }
    },
    {
        "type": "function",
        "function": {
            "name": "get_current_time",
            "description": "Get the current time for a specified timezone",
            "parameters": {
                "type": "object",
                "properties": {
                    "timezone": {
                        "type": "string",
                        "description": "The timezone identifier, e.g. Asia/Tokyo, America/New_York"
                    }
                },
                "required": ["timezone"]
            }
        }
    }
]
```

Tool Definition

Implement a Live Example to Understand Actual Model Input/Output

```
conversation = [
    {
        "role": "system",
        "content": "You are a helpful AI assistant with access to tools."
    },
    {
        "role": "user",
        "content": "What's the weather like in San Francisco?"
    },
    {
        "role": "assistant",
        "content": "I'll check the weather in San Francisco for you.",
        "tool_calls": [
            {
                "id": "call_123",
                "type": "function",
                "function": {
                    "name": "get_weather",
                    "arguments": json.dumps({"location": "San Francisco", "unit": "celsius"})
                }
            }
        ]
    },
    {
        "role": "tool",
        "tool_call_id": "call_123",
        "name": "get_weather",
        "content": json.dumps({"temperature": 18, "condition": "partly cloudy", "humidity": 65})
    },
    {
        "role": "assistant",
        "content": "The weather in San Francisco is currently 18°C and partly cloudy with 65% humidity."
    },
    {
        "role": "user",
        "content": "How about in Tokyo? Also, what time is it there?"
    },
    {
        "role": "assistant",
        "content": "Let me check both the weather and time in Tokyo.",
        "tool_calls": [
            {
                "id": "call_456",
                "type": "function",
                "function": {
                    "name": "get_weather",
                    "arguments": json.dumps({"location": "Tokyo", "unit": "celsius"})
                }
            }
        ]
    }
]
```

Example Conversation Data

Implement a Live Example to Understand Actual Model Input/Output

```
formatted_text = tokenizer.apply_chat_template(  
    conversation,  
    tools=tools,  
    tokenize=False,  
    add_generation_prompt=False  
)
```

Live Example after Applying Template

```
[  
  {  
    "role": "system",  
    "content": "You are a helpful AI assistant with access to tools."  
  },  
  {  
    "role": "user",  
    "content": "What's the weather like in San Francisco?"  
  },  
  {  
    "role": "assistant",  
    "content": "I'll check the weather in San Francisco for you.",  
    "tool_calls": [  
      {  
        "id": "call_123",  
        "type": "function",  
        "function": {  
          "name": "get_weather",  
          "arguments": "{\"location\": \"San Francisco\", \"unit\": \"celsius\"}"  
        }  
      }  
    ]  
  },  
  {  
    "role": "tool",  
    "tool_call_id": "call_123",  
    "name": "get_weather",  
    "content": "{\"temperature\": 18, \"condition\": \"partly cloudy\", \"humidity\": 65}"  
  },  
  {  
    "role": "assistant",  
    "content": "The weather in San Francisco is currently 18\u00b0C and partly cloudy with 65% humidity."  
  },  
  {  
    "role": "user",  
    "content": "How about in Tokyo? Also, what time is it there?"  
  },  
  {  
    "role": "assistant",  
    "content": "Let me check both the weather and time in Tokyo.",  
    "tool_calls": [  
      {  
        "id": "call_456",  
        "type": "function",  
        "function": {  
          "name": "get_weather",  
          "arguments": "{\"location\": \"Tokyo\", \"unit\": \"celsius\"}"  
        }  
      },  
      {  
        "id": "call_789",  
        "type": "function",  
        "function": {  
          "name": "get_current_time",  
          "arguments": "{\"timezone\": \"Asia/Tokyo\"}"  
        }  
      }  
    ]  
  }]
```

Json formatted context

Live Example after Applying Template

```
<|im_start|>system
You are a helpful AI assistant with access to tools.

# Tools

You may call one or more functions to assist with the user query.

You are provided with function signatures within <tools></tools> XML tags:
<tools>
{
  "type": "function",
  "function": {
    "name": "get_weather",
    "description": "Get the current weather information for a specified location",
    "parameters": {
      "type": "object",
      "properties": {
        "location": {
          "type": "string",
          "description": "The city name, e.g. San Francisco, Tokyo"
        },
        "unit": {
          "type": "string",
          "enum": [
            "celsius",
            "fahrenheit"
          ],
          "description": "The temperature unit to use"
        }
      },
      "required": [
        "location"
      ]
    }
  }
}

  "type": "function",
  "function": {
    "name": "get_current_time",
    "description": "Get the current time for a specified timezone",
    "parameters": {
      "type": "object",
      "properties": {
        "timezone": {
          "type": "string",
          "description": "The timezone identifier, e.g. Asia/Tokyo, America/New_York"
        }
      },
      "required": [
        "timezone"
      ]
    }
  }
</tools>
```

Qwen2.5-coder tokenizer
After applying chat templates

System prompt

Live Example after Applying Template

```
For each function call, return a json object with function name and arguments within <tool_call></tool_call> XML tags:  
<tool_call>  
{"name": <function-name>, "arguments": <args-json-object>}  
</tool_call><|im_end|>  
<|im_start|>user  
What's the weather like in San Francisco?<|im_end|>  
<|im_start|>assistant  
I'll check the weather in San Francisco for you.  
<tool_call>  
{  
    "name": "get_weather",  
    "arguments": "{\"location\": \"San Francisco\", \"unit\": \"celsius\"}"  
}  
</tool_call><|im_end|>  
<|im_start|>user  
<tool_response>  
{  
    "temperature": 18,  
    "condition": "partly cloudy",  
    "humidity": 65  
}  
</tool_response><|im_end|>  
<|im_start|>assistant  
The weather in San Francisco is currently 18°C and partly cloudy with 65% humidity.<|im_end|>  
<|im_start|>user  
How about in Tokyo? Also, what time is it there?<|im_end|>  
<|im_start|>assistant  
Let me check both the weather and time in Tokyo.  
<tool_call>  
{  
    "name": "get_weather",  
    "arguments": "{\"location\": \"Tokyo\", \"unit\": \"celsius\"}"  
}  
</tool_call>  
<tool_call>  
{  
    "name": "get_current_time",  
    "arguments": "{\"timezone\": \"Asia/Tokyo\"}"  
}  
</tool_call><|im_end|>  
<|im_start|>user  
<tool_response>  
{  
    "temperature": 15,  
    "condition": "clear",  
    "humidity": 45  
}  
</tool_response>  
<tool_response>  
{  
    "time": "14:30",  
    "timezone": "Asia/Tokyo",  
    "date": "2025-01-15"  
}
```

After applying chat templates

Live Example after Applying Template

```
<|im_start|>system  
You are a helpful AI assistant.  
  
# Tools  
  
You have access to the following functions:  
  
<tools>  
<function>  
<name>get_weather</name>  
<description>Get the current weather information for a specified location</description>  
<parameters>  
<parameter>  
<name>location</name>  
<type>string</type>  
<description>The city name, e.g. San Francisco, Tokyo</description>  
</parameter>  
<parameter>  
<name>unit</name>  
<type>string</type>  
<description>The temperature unit to use</description>  
<enum>["celsius", "fahrenheit"]</enum>  
</parameter>  
<required>["location"]</required>  
</parameters>  
</function>  
<function>  
<name>get_current_time</name>  
<description>Get the current time for a specified timezone</description>  
<parameters>  
<parameter>  
<name>timezone</name>  
<type>string</type>  
<description>The timezone identifier, e.g. Asia/Tokyo, America/New_York</description>  
</parameter>  
<required>["timezone"]</required>  
</parameters>  
</function>  
</tools>
```

If you choose to call a function ONLY reply in the following format with NO suffix:

```
<tool_call>  
<function=example_function_name>  
<parameter=example_parameter_1>  
value_1  
</parameter>  
<parameter=example_parameter_2>  
This is the value for the second parameter  
that can span  
multiple lines  
</parameter>  
</function>  
</tool_call>
```

Qwen3-Coder tokenizer template

Switch from json tool call to xml

Live Example after Applying Template

```
<IMPORTANT>
Reminder:
- Function calls MUST follow the specified format: an inner <function=...></function> block must be nested within <tool_call></tool_call> XML tags
- Required parameters MUST be specified
- You may provide optional reasoning for your function call in natural language BEFORE the function call, but NOT after
- If there is no function call available, answer the question like normal with your current knowledge and do not tell the user about function calls
</IMPORTANT><|im_end|>
<|im_start|>user
What's the weather like in San Francisco?<|im_end|>
<|im_start|>assistant
I'll check the weather in San Francisco for you.

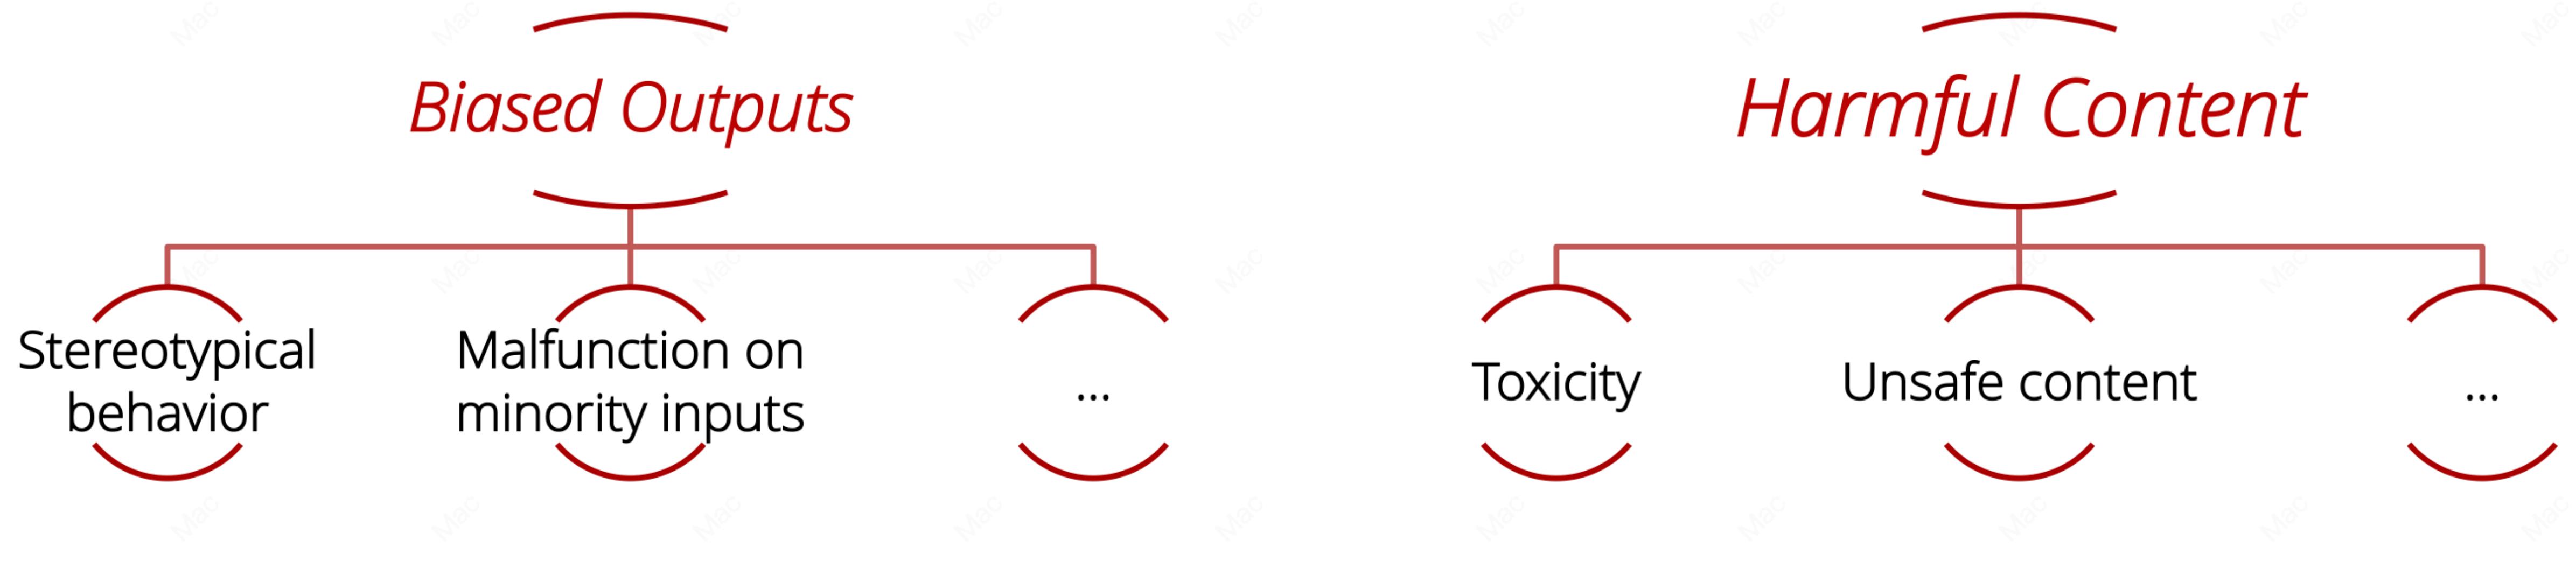
<tool_call>
<function=get_weather>
<parameter=location>
San Francisco
</parameter>
<parameter=unit>
celsius
</parameter>
</function>
</tool_call><|im_end|>
<|im_start|>user
<tool_response>
{
  "temperature": 18,
  "condition": "partly cloudy",
  "humidity": 65
}
</tool_response>
<|im_end|>
<|im_start|>assistant
The weather in San Francisco is currently 18°C and partly cloudy with 65% humidity.<|im_end|>
<|im_start|>user
How about in Tokyo? Also, what time is it there?<|im_end|>
<|im_start|>assistant
Let me check both the weather and time in Tokyo.

<tool_call>
<function=get_weather>
<parameter=location>
Tokyo
</parameter>
<parameter=unit>
celsius
</parameter>
</function>
</tool_call>
<tool_call>
<function=get_current_time>
<parameter=timezone>
Asia/Tokyo
</parameter>
</function>
</tool_call><|im_end|>
<|im_start|>user
<tool_response>
```

Qwen3-Coder tokenizer template

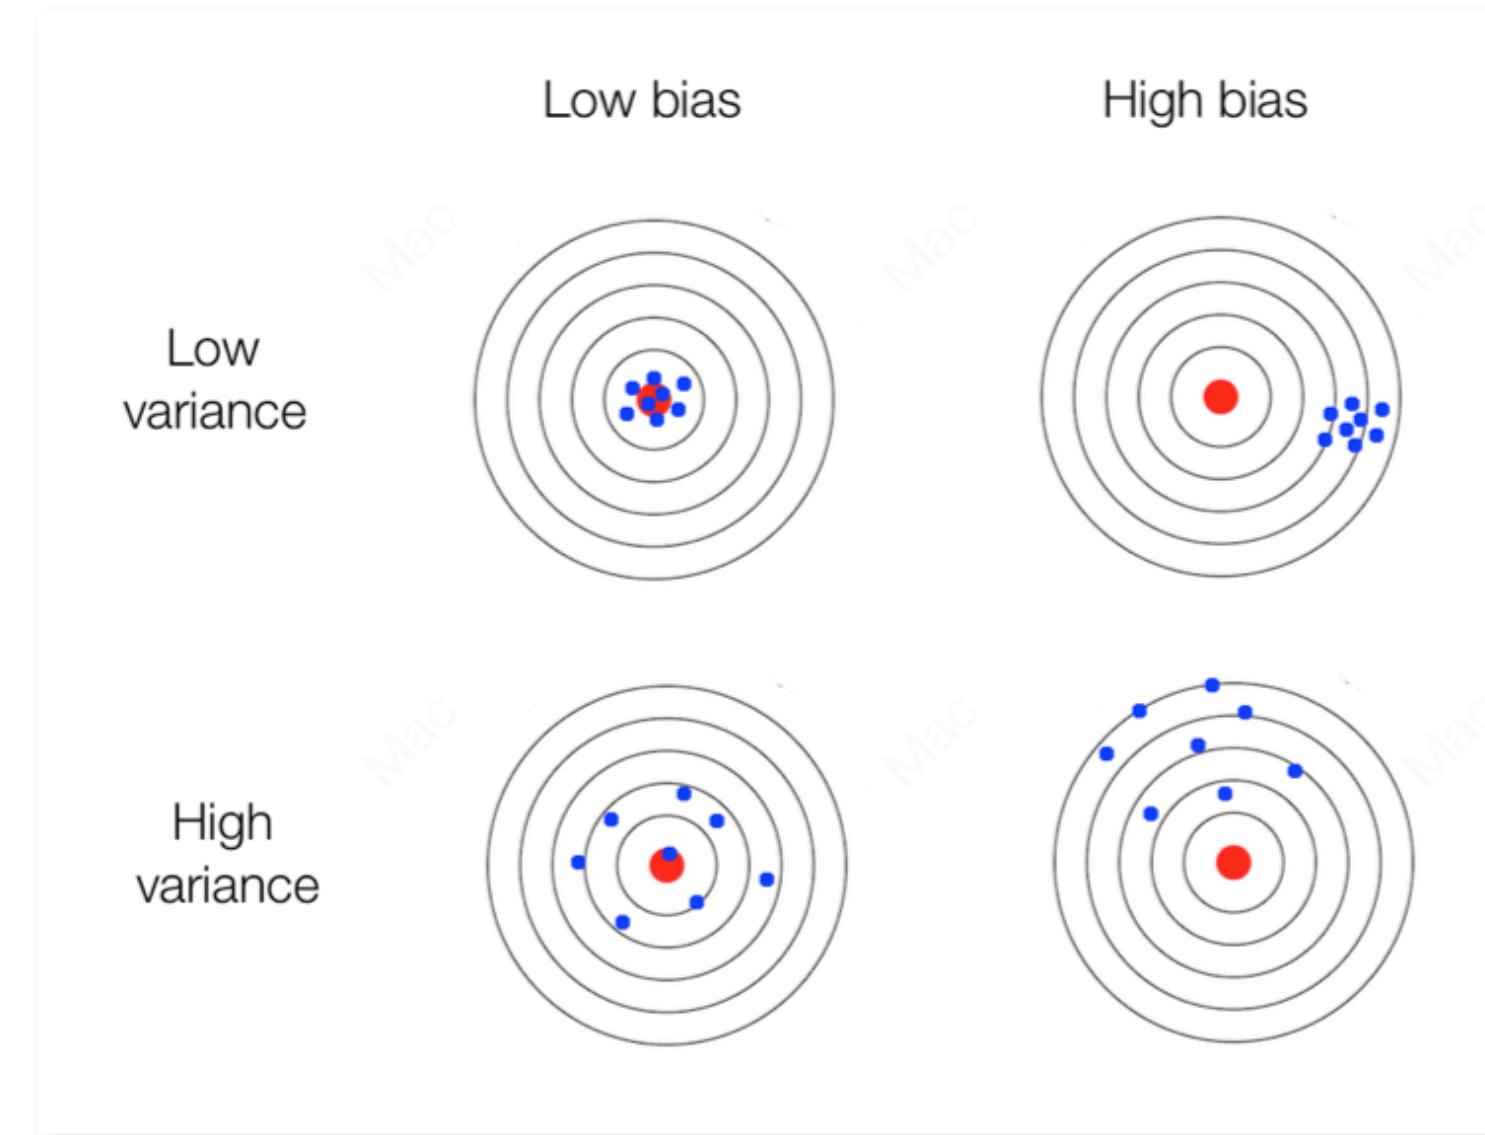
Switch from json tool call to xml

LLM Safety and Bias



Some Definition of Bias

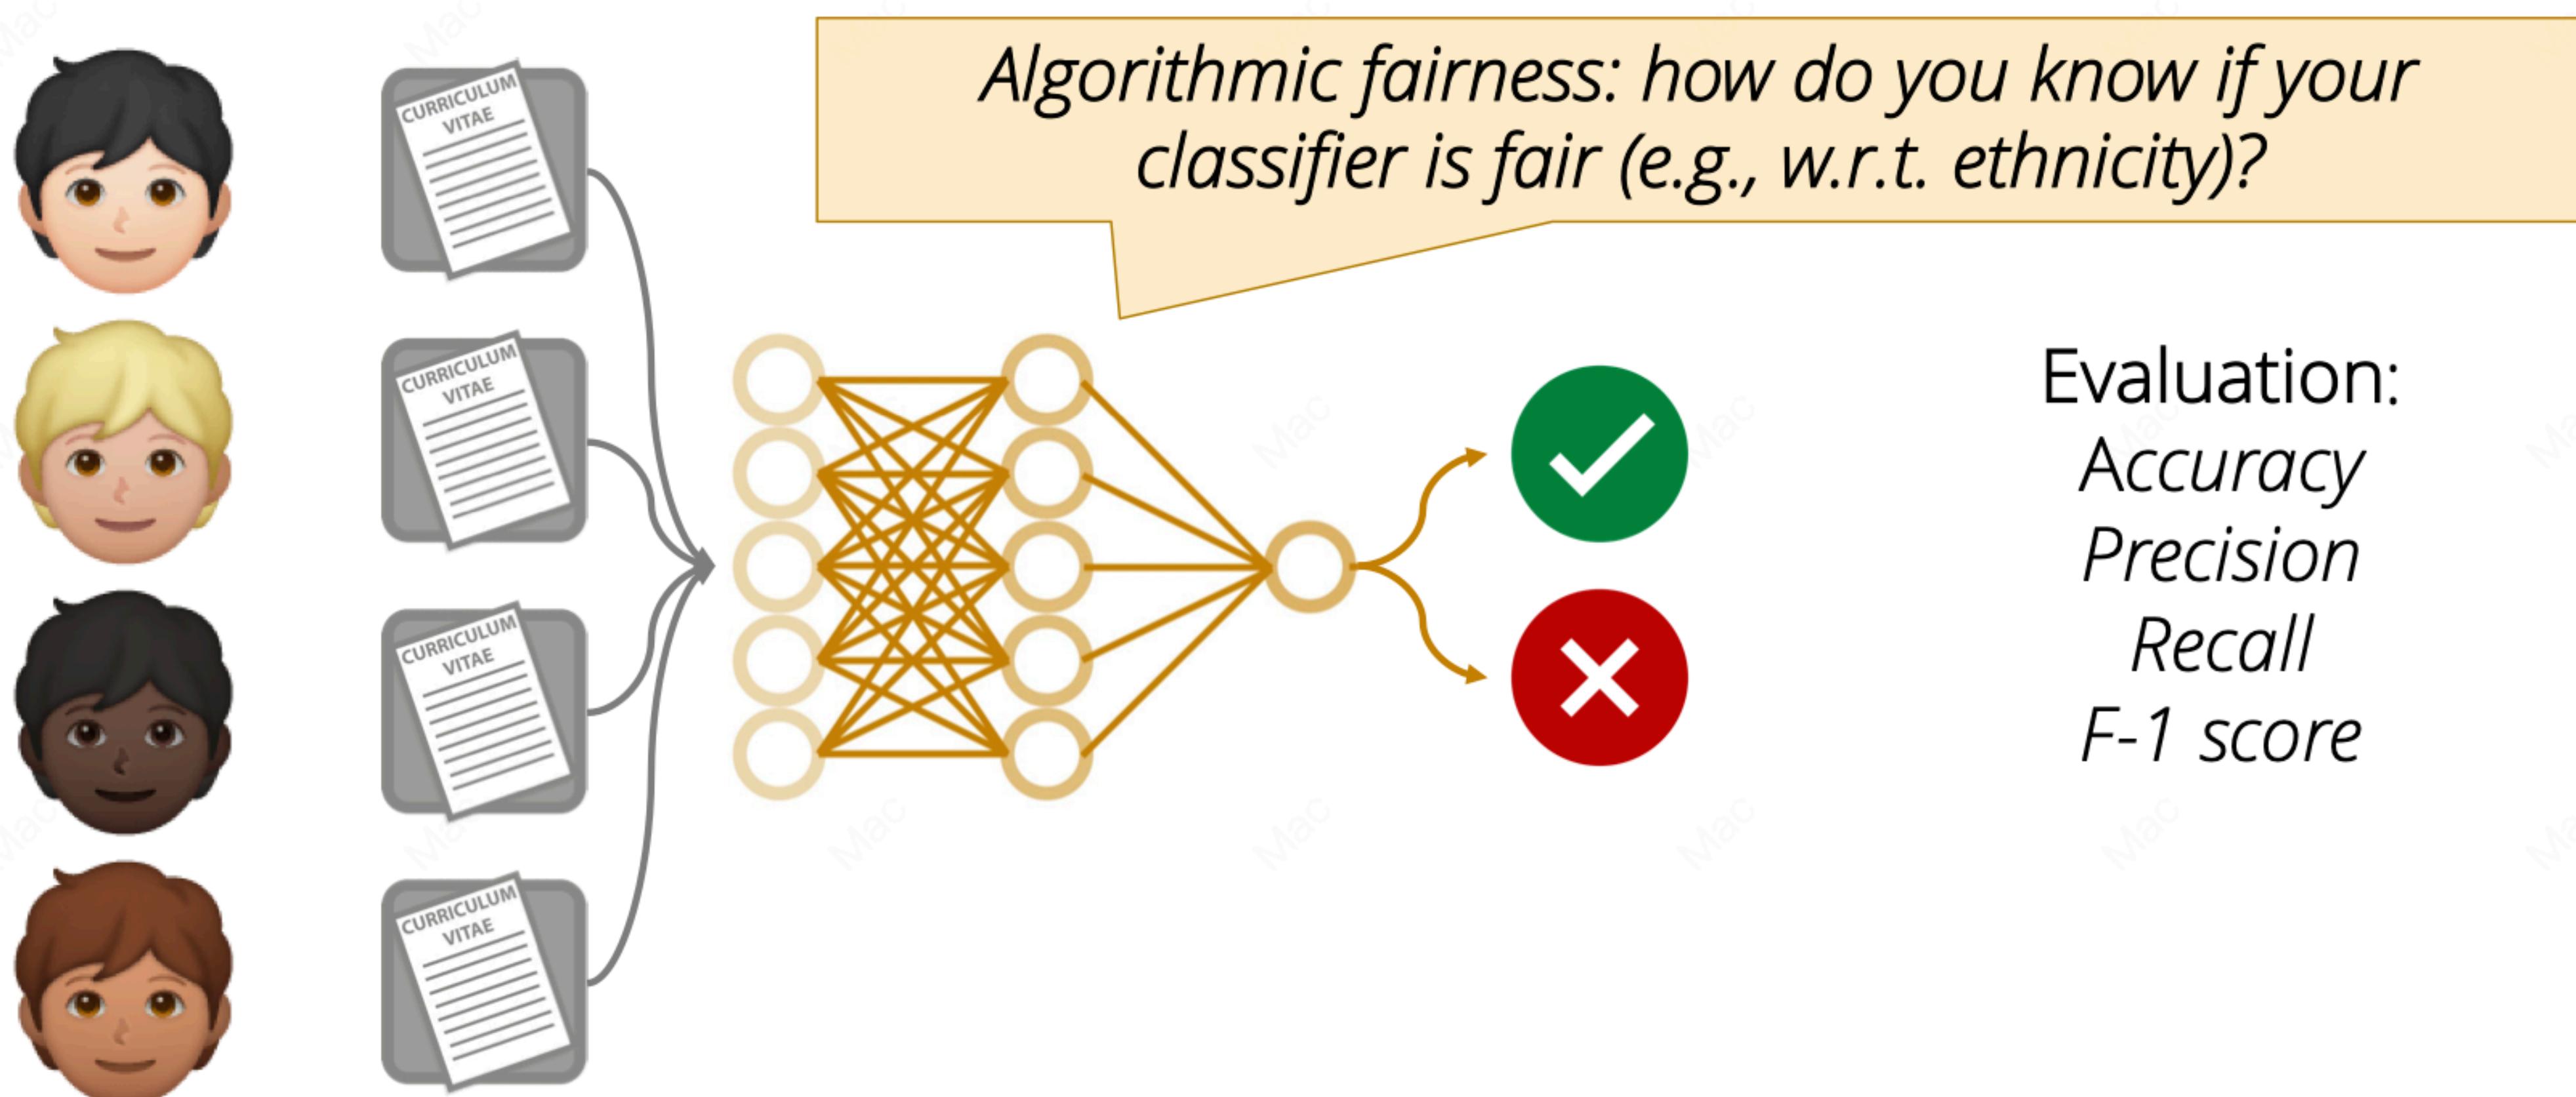
- Bias [statistics]: systematic tendency causing differences between model estimates / predictions
- Bias [general]: "*disproportionate weight in favor of or against an idea or thing, usually in a way that is closed-minded, prejudicial, or unfair*" - Wikipedia



Presence of bias \simeq absence of fairness
Algorithmic fairness: attempts to correct biases in ML systems
But... how is fairness defined?

Algorithmic Fairness

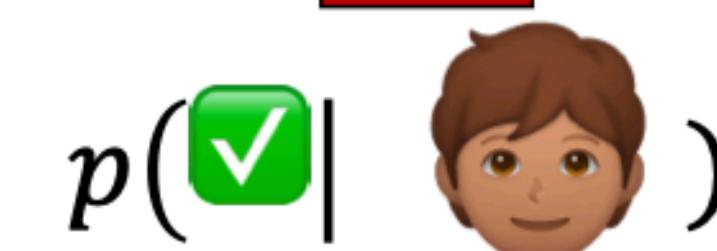
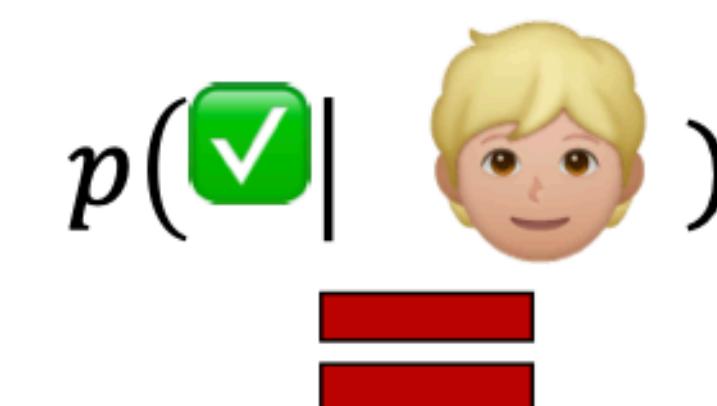
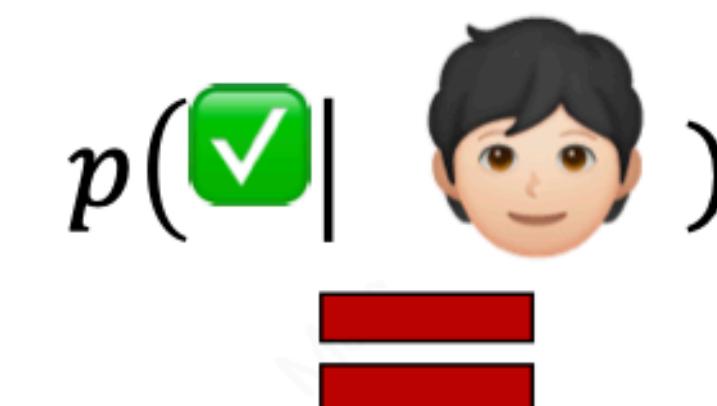
Let's assume a toy task: given a resumé, predict whether a candidate is qualified



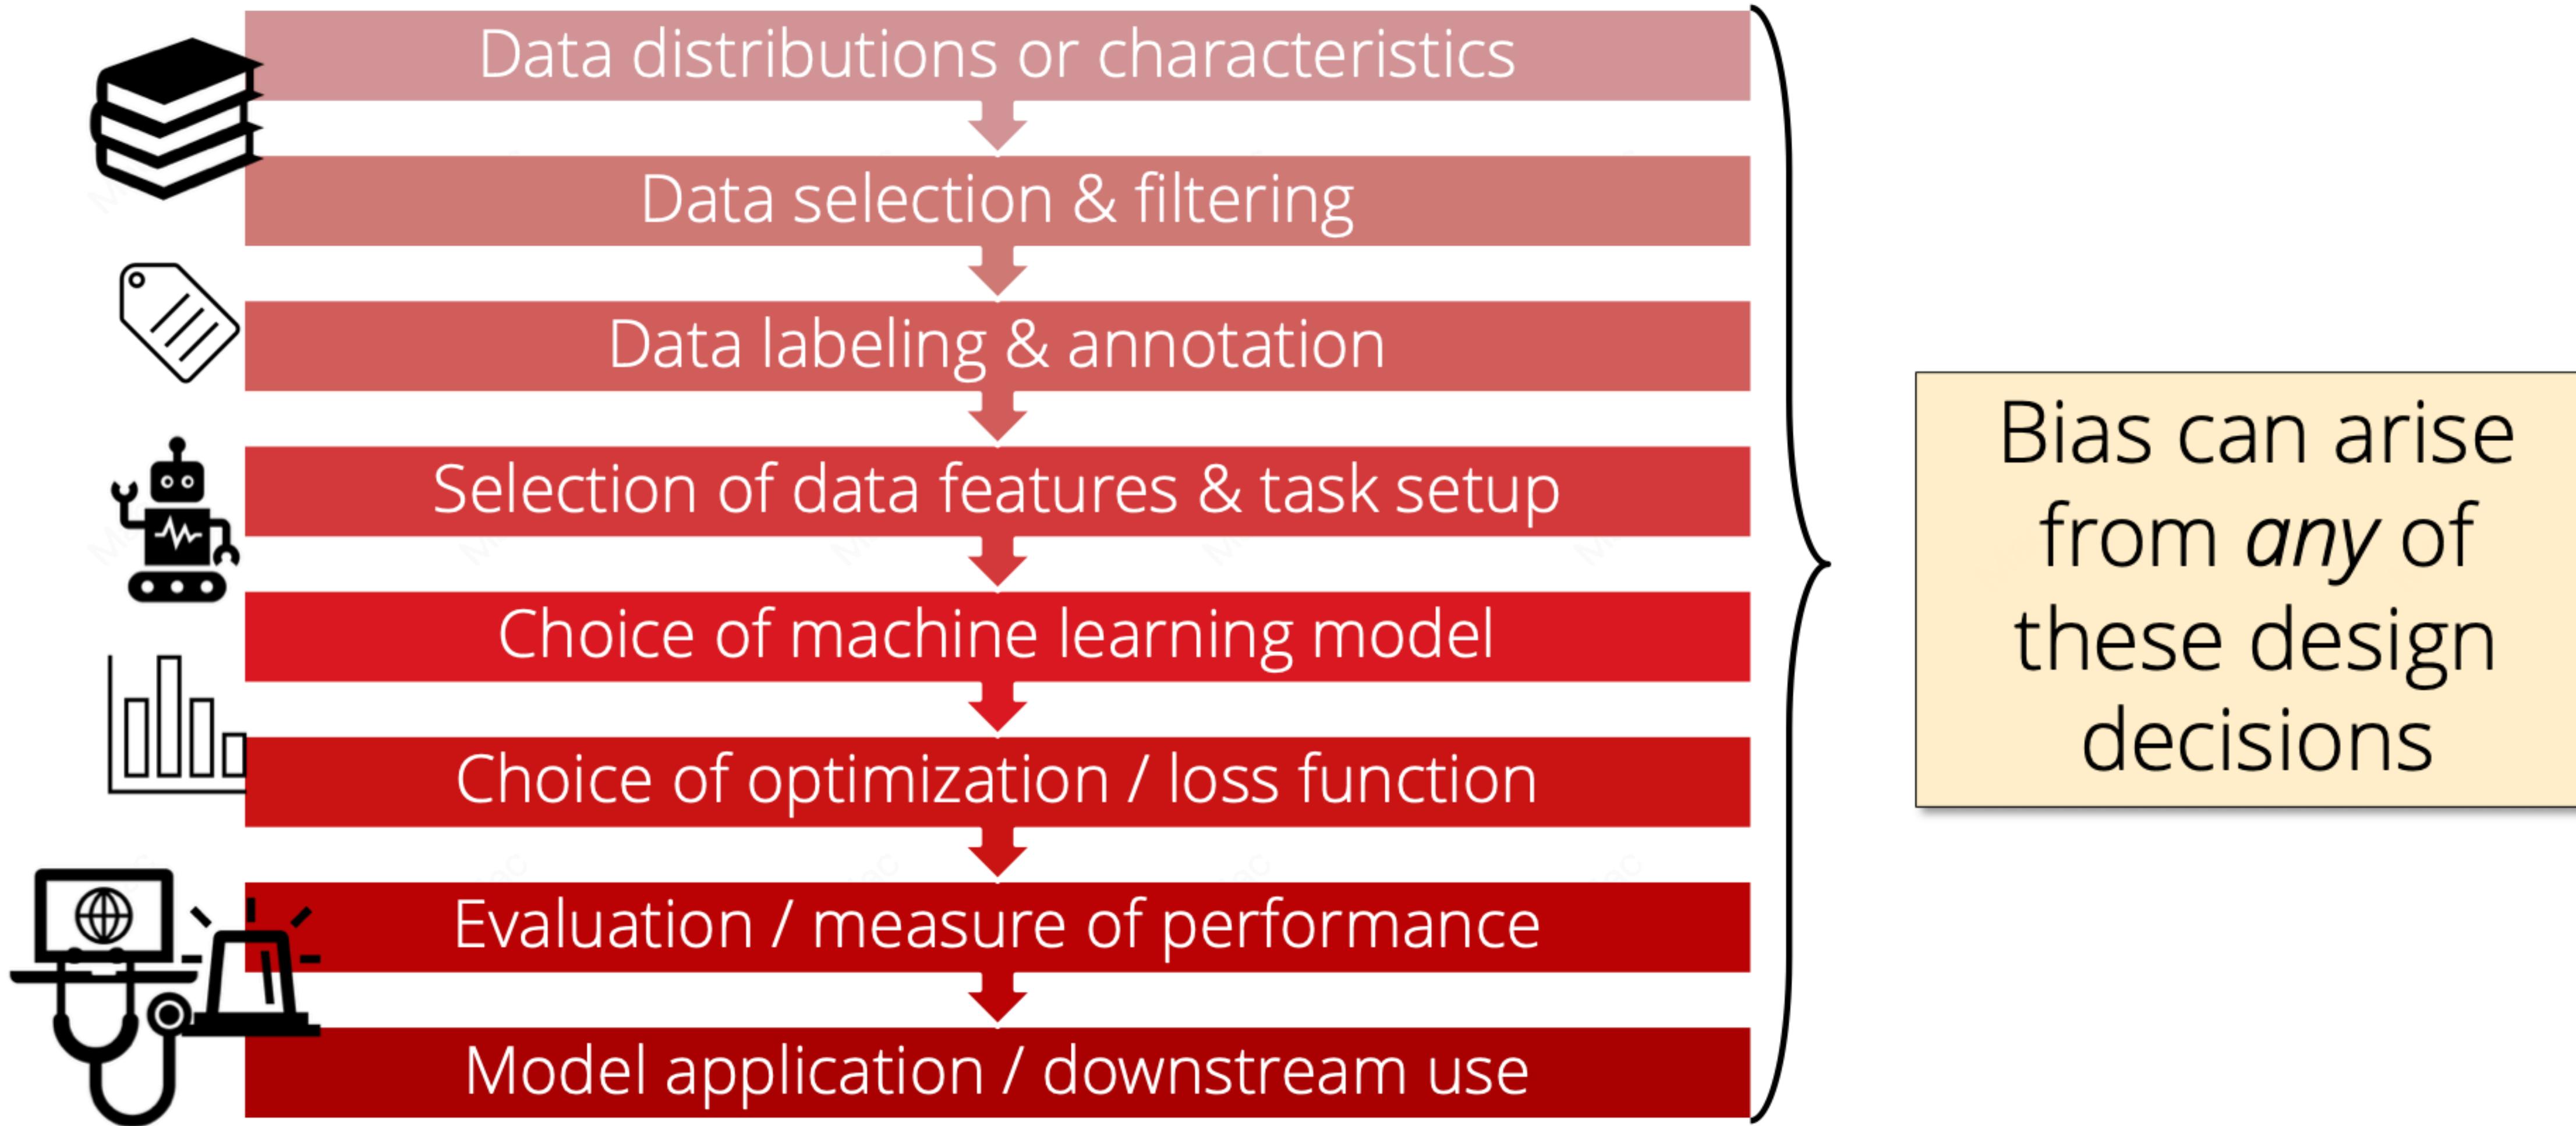
Fairness Metrics

To control variables, let's use the same set of resume, but only change the candidates' ethnic group information on the resume

- Accuracy quality: a classifier is fair if the people from different groups have the same accuracy
- Statistical parity: groups should have the same probability of being assigned positive class

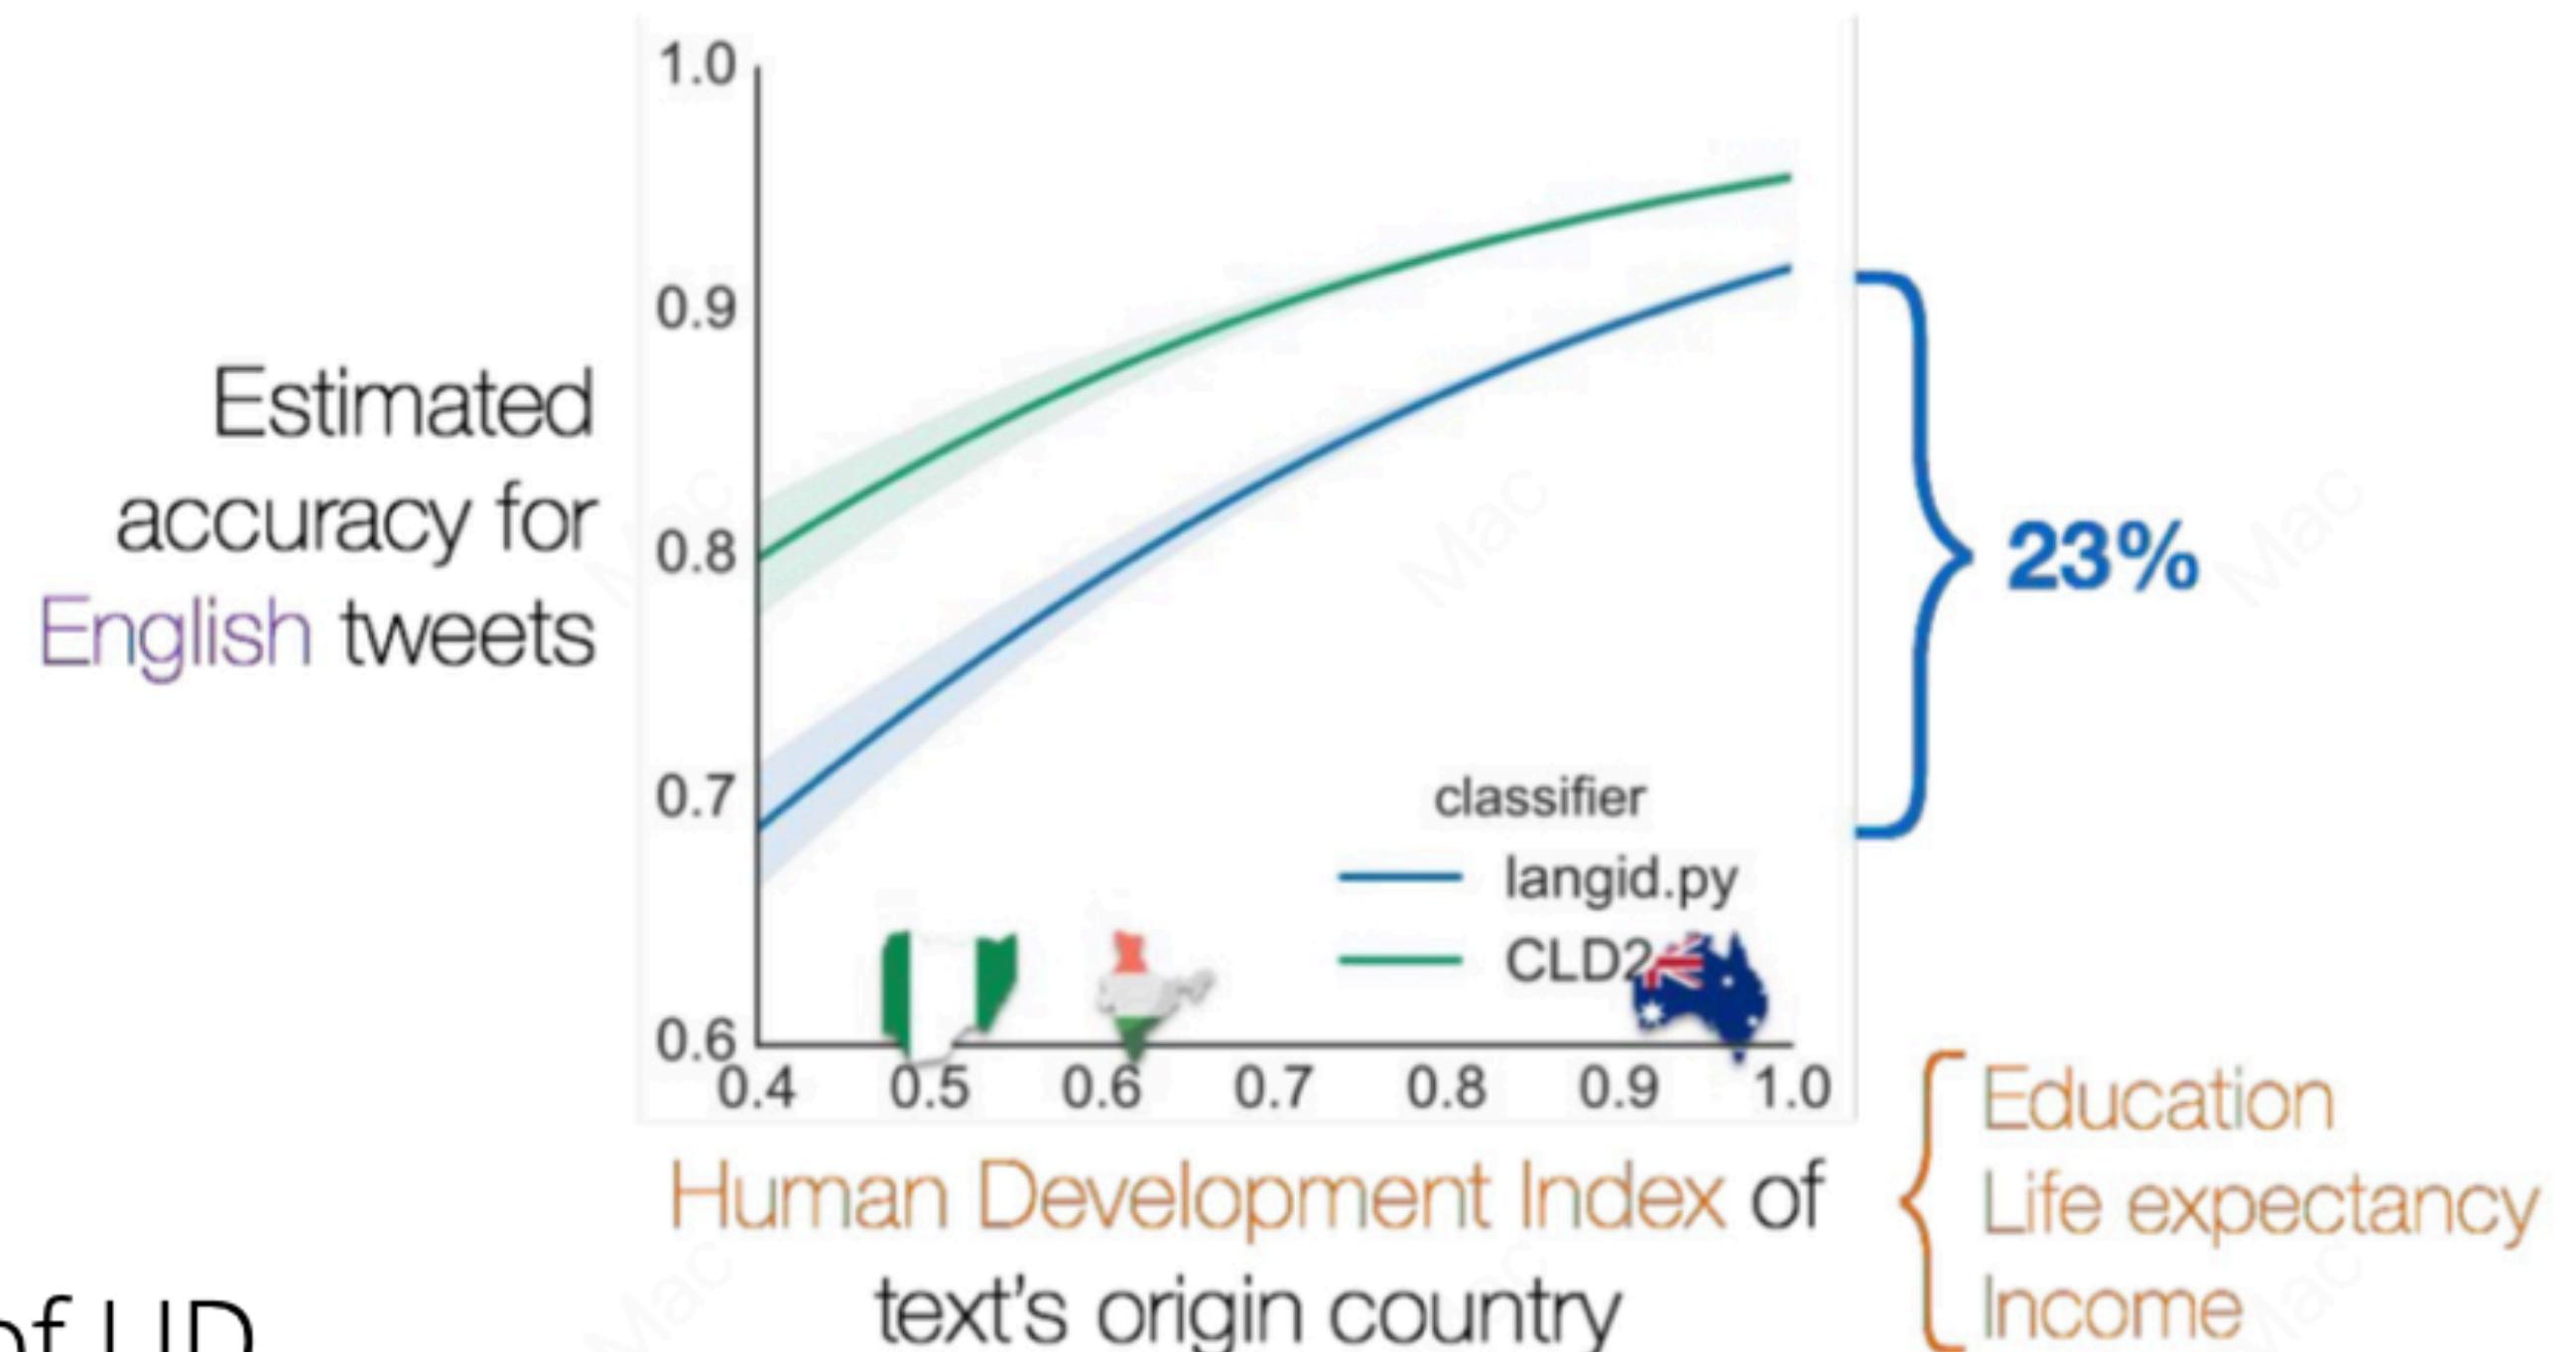


Where Does Bias Come From?



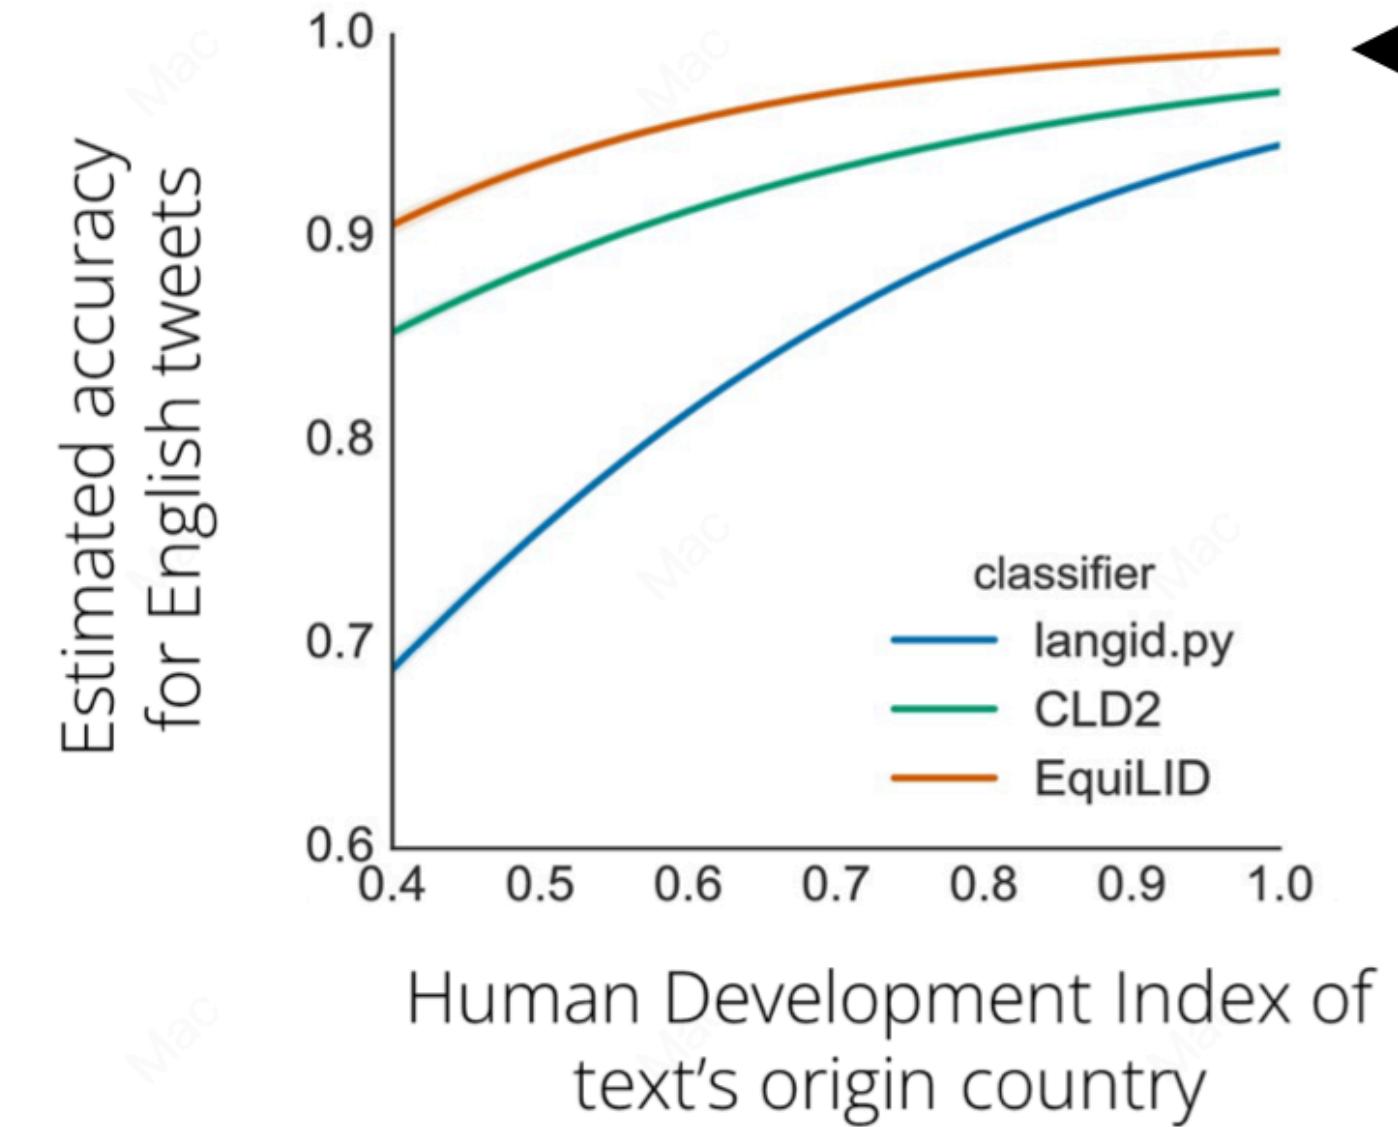
Example of Bias from Data: LangID Tool

- LangID task: determine which language an input text is in
 - Considered a “*a solved problem* suitable for undergraduate instruction” (McNamee, 2005)
- Often a first step in most NLP and CSS preprocessing pipelines
 - e.g., filtering LLM pretraining data
- *But*, many variations of English in the world
 - *Int'l*: Nigerian English, Indian English, etc.
 - *Within US*: African American English, etc.
- [Jurgens et al. \(2017\)](#) found that accuracy of LID tool correlated with wealth/development level of country; works worse for low HDI countries



Example of Bias from Data: LangID Tool

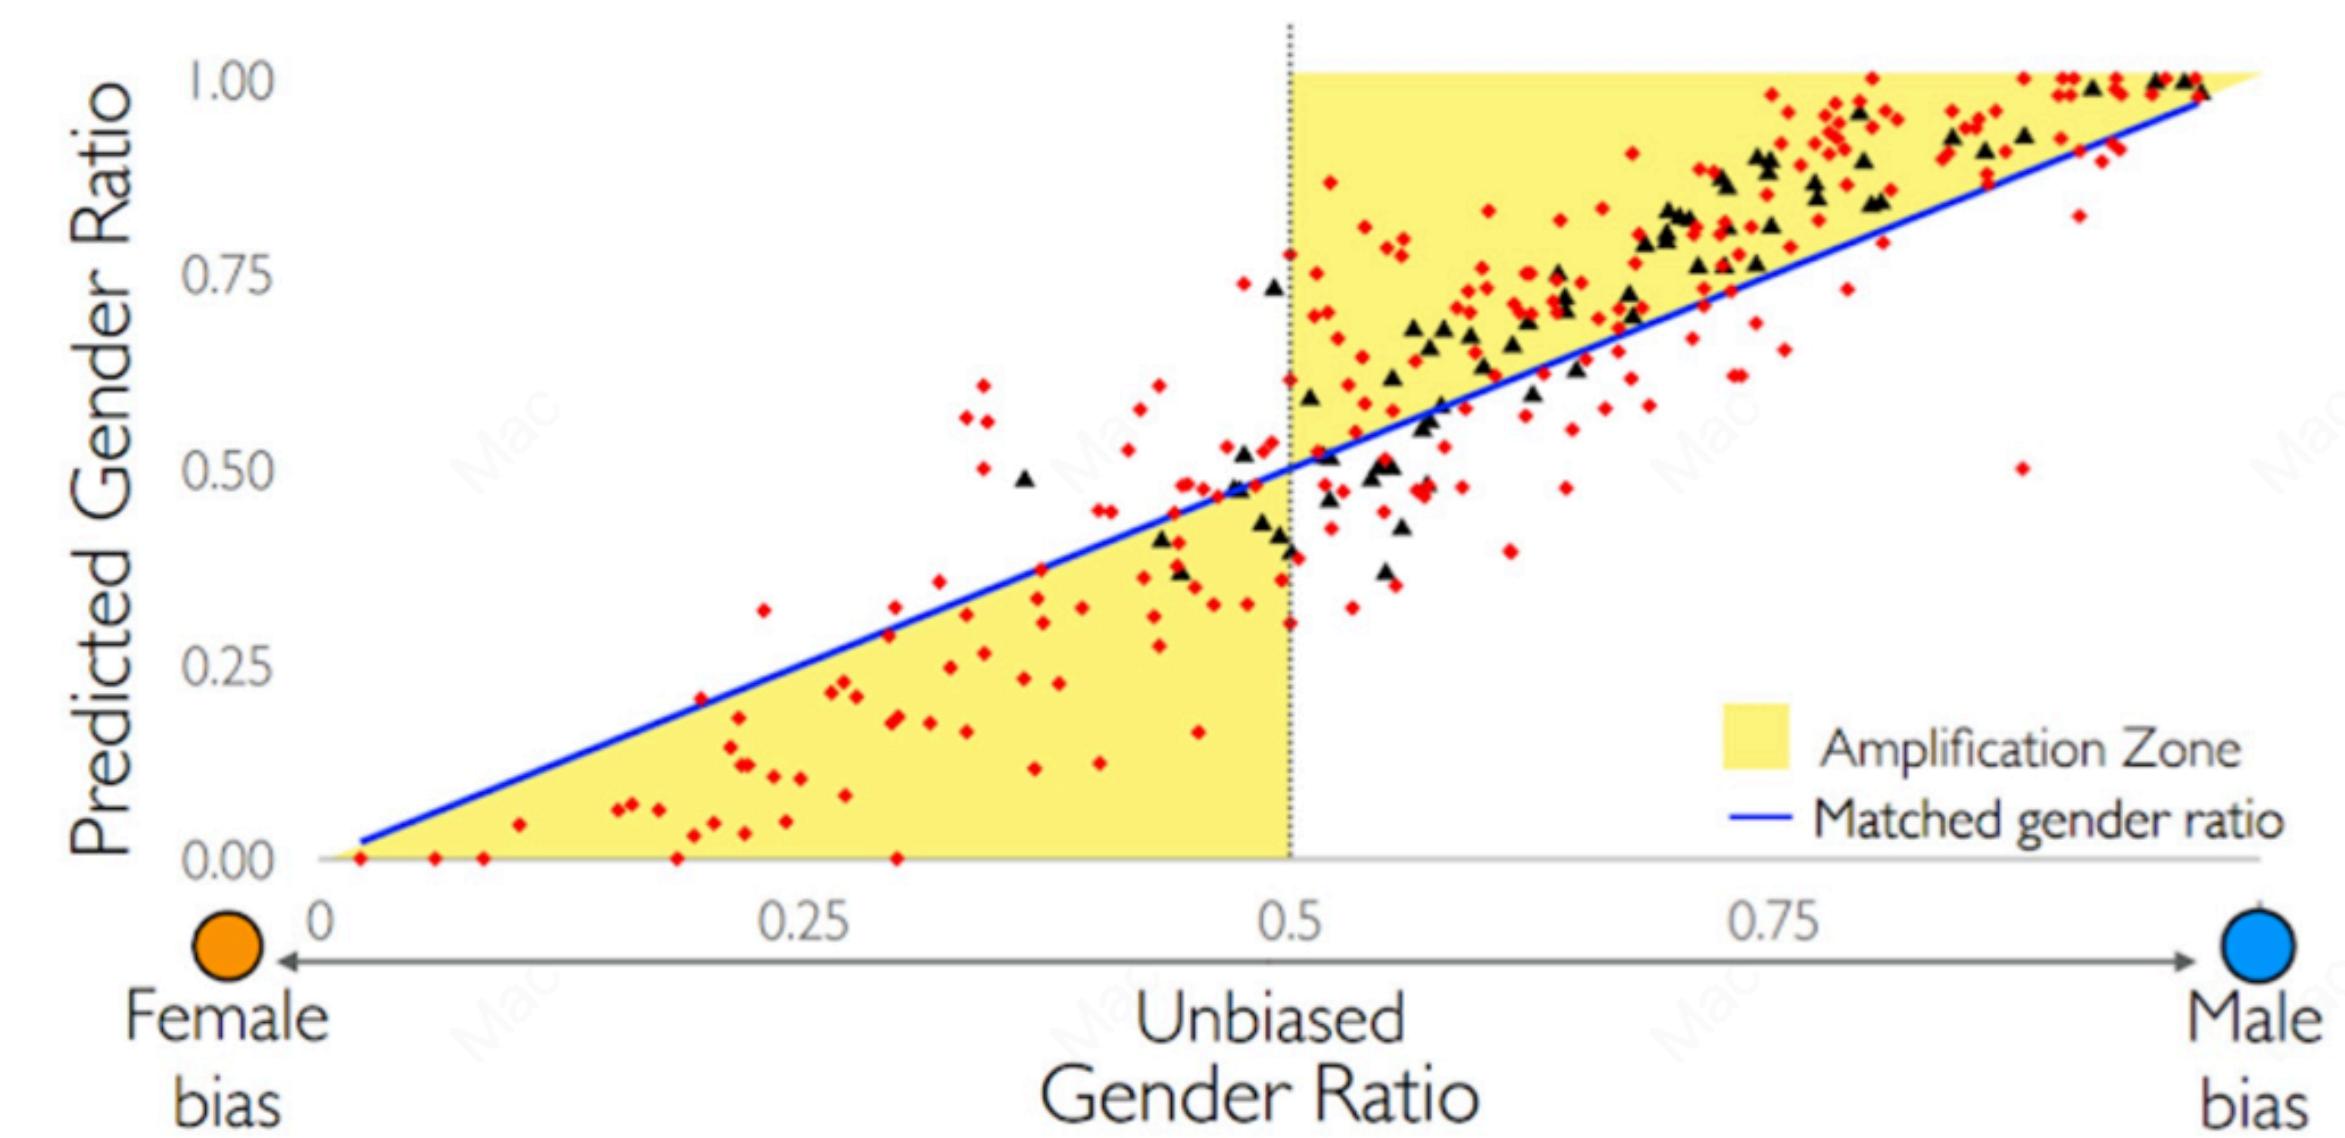
- Jurgens et al (2017) introduce EquiLID
 - Trained by sampling more variety of data, topically, socially, geographically diverse, and even multilingual data
- Find that tool works much better than original LID systems
 - Bonus: even improved accuracy on highly developed countries!



**Takeaway: Bias can be mitigated by making better data choices
But this is not the only source....**

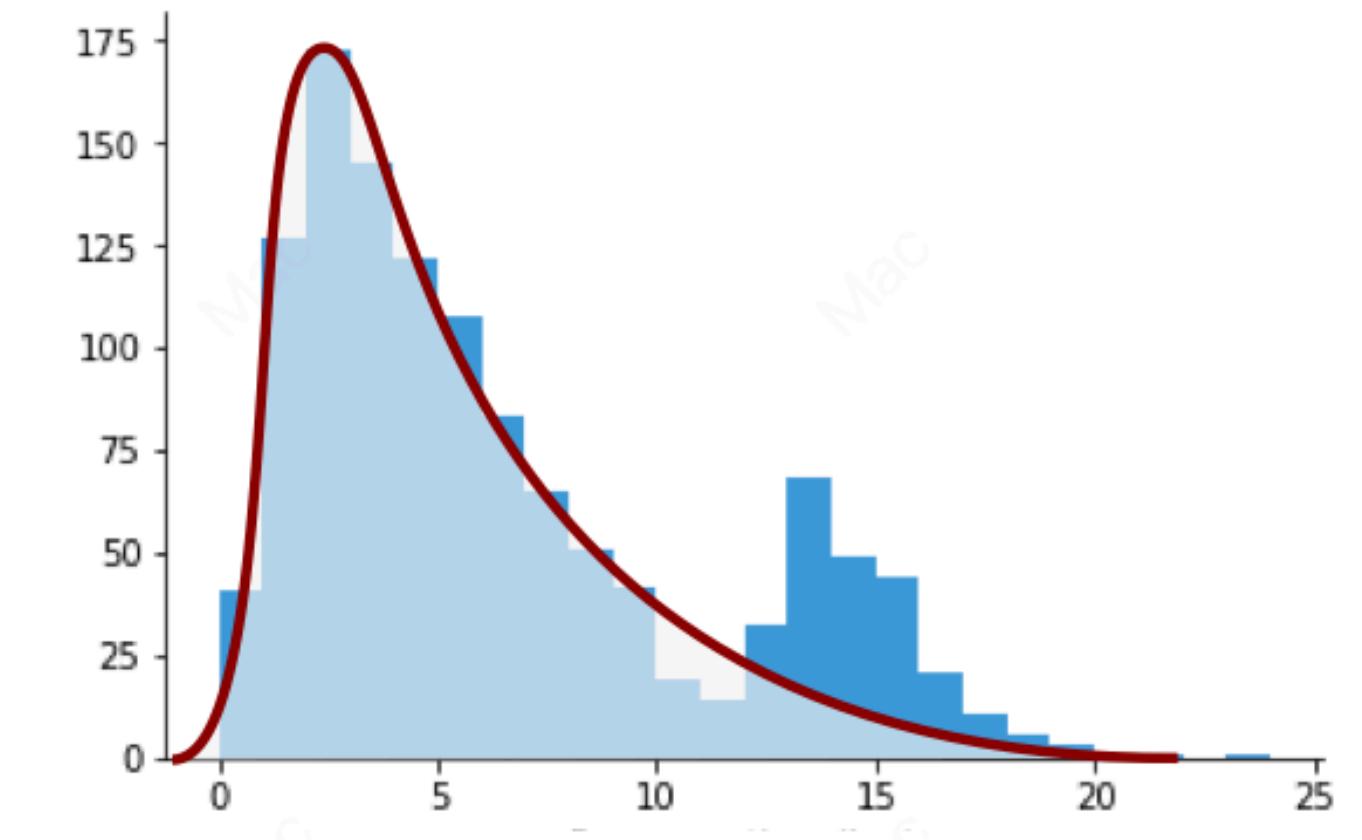
Bias amplification from models

- Zhao et al (2017) examined visual semantic role labeling task
 - Given an image, predict various semantic roles, including agent (person doing the action)
- Found skews in training dataset
 - E.g., 66% of training cooking images had agent=woman
- Found that models *amplified* biases
 - E.g., 84% of test cooking images predicted as agent=woman (~18% men mis-labeled)
- Showed that prediction / inference functions can mitigate this bias



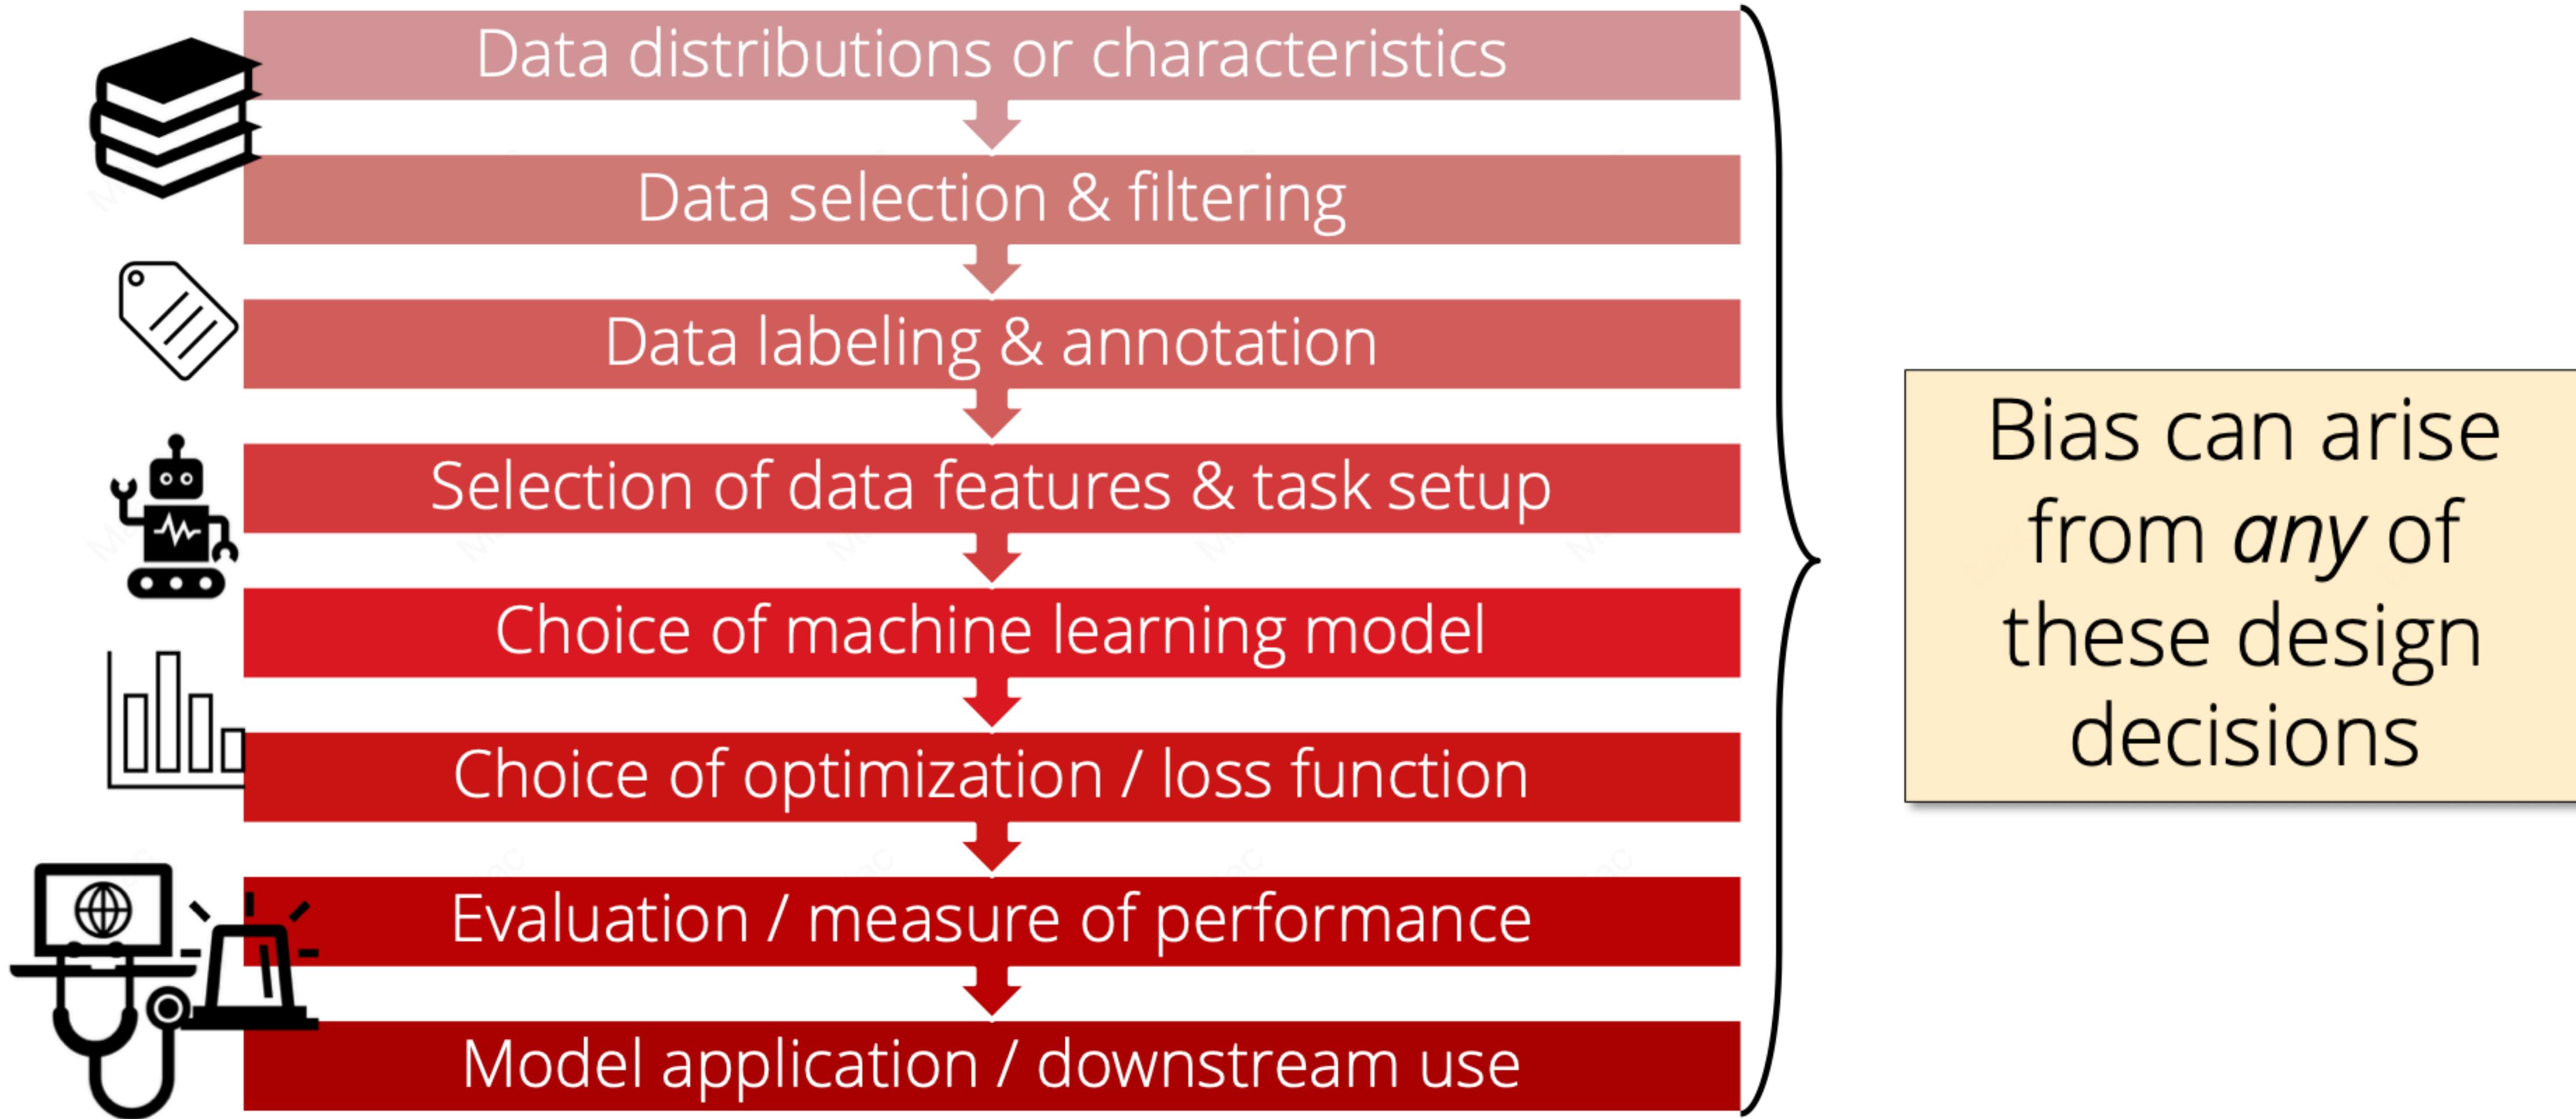
Model Biases: Mathematical Links

- *Competing losses*: objective functions aim to minimize loss globally → learns to predict most frequent class
 - Often at the expense of less frequent classes (e.g. minority groups)

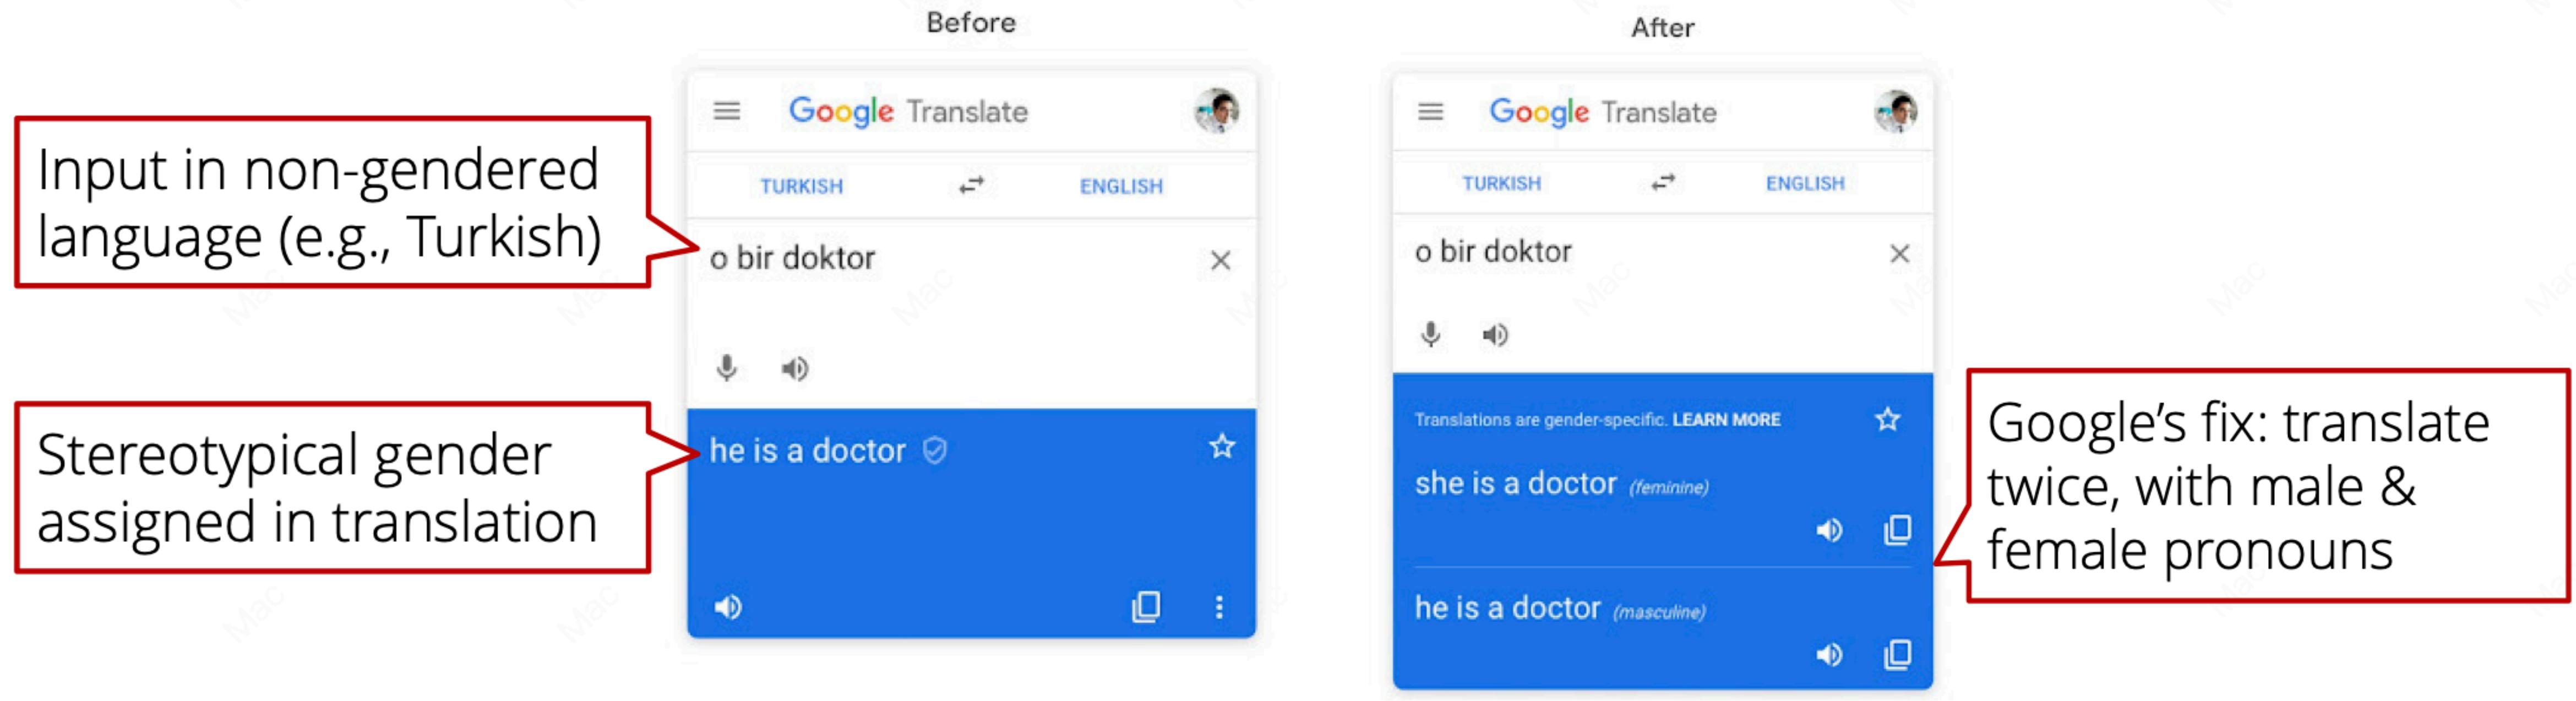


Suppose you have a machine translation dataset covering multiple languages, but most of them are about English-Spanish translation, and some of them are about other languages. Models will prioritize decrease losses of English-Spanish translation

Where Does Bias Come From?

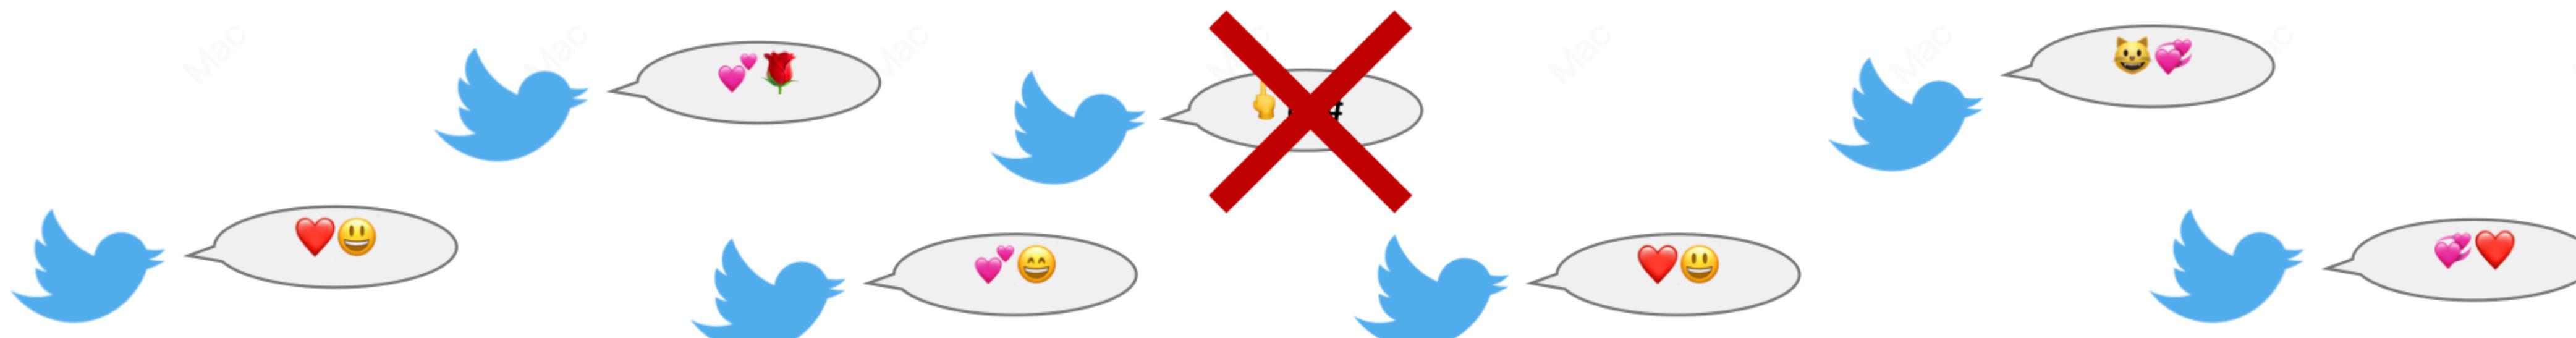


Google Translate Issue



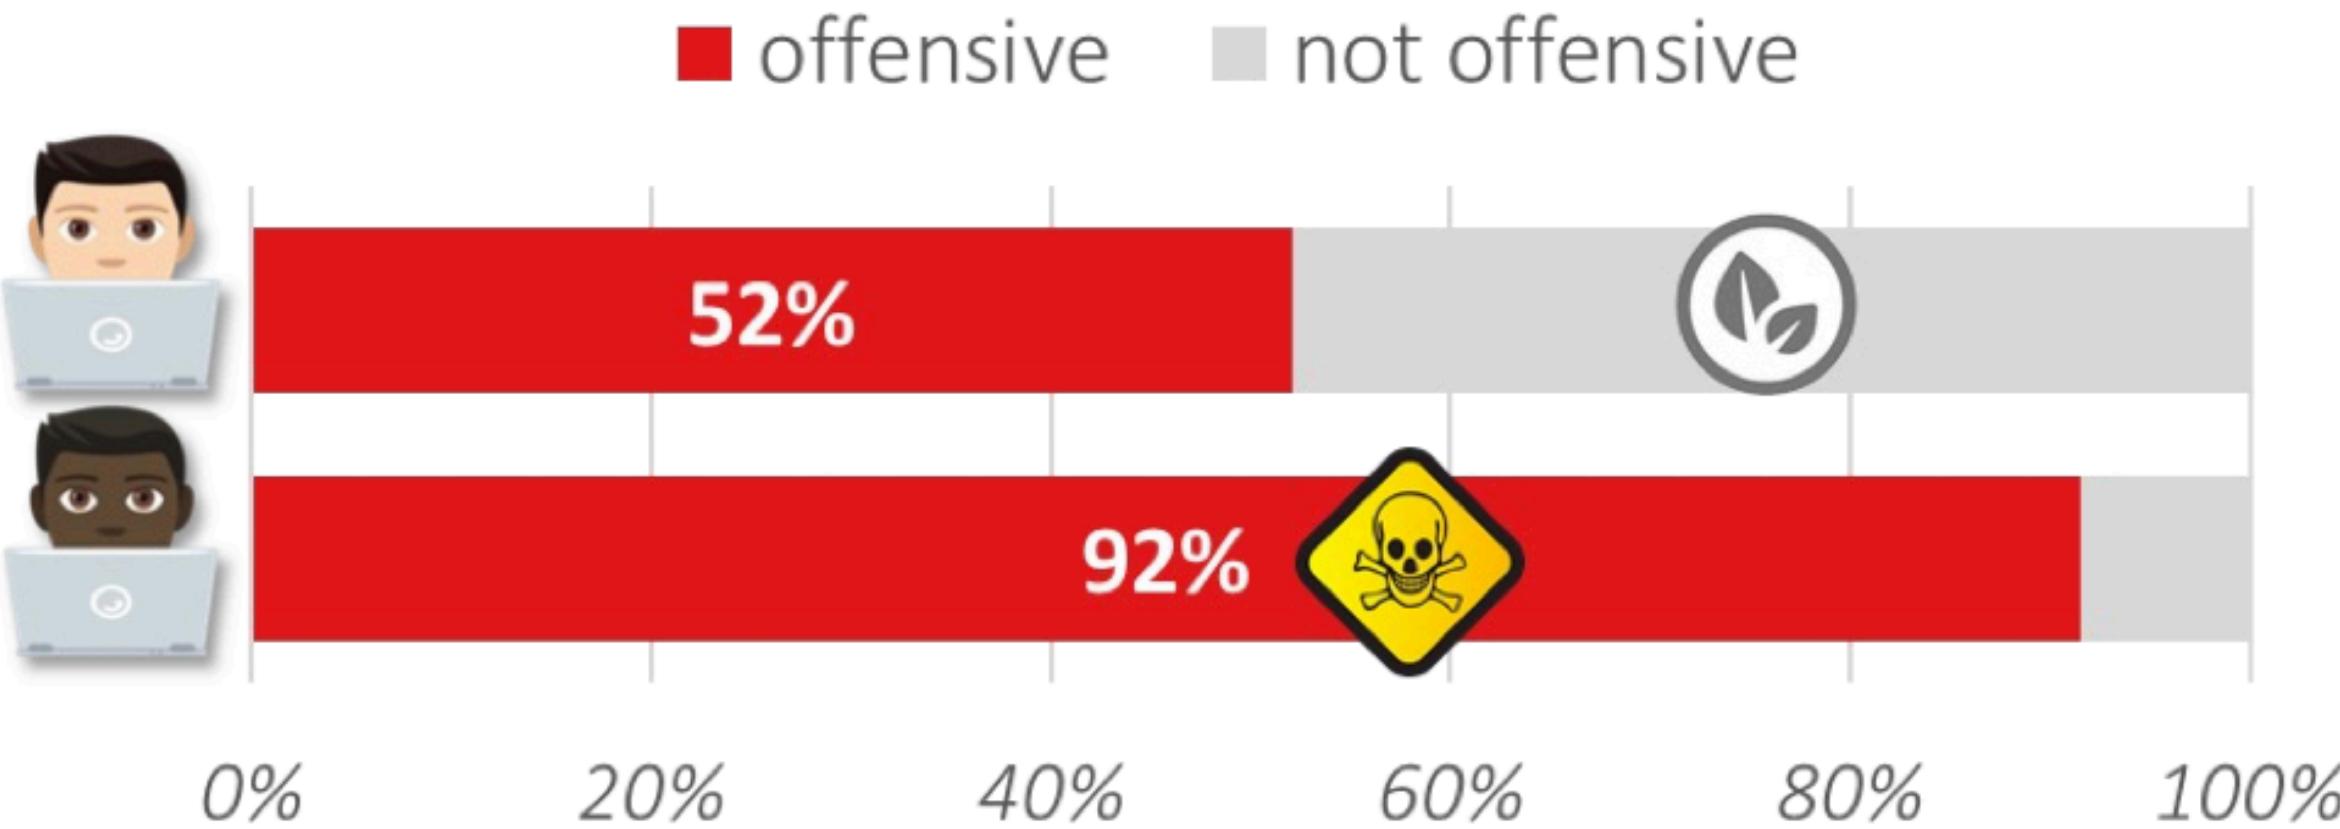
- Takeaways: mitigating bias may involve *system-level changes* to UI, input processing, output formatting, etc. while underlying AI model is similar

Hate Speech or Toxic Language Detection



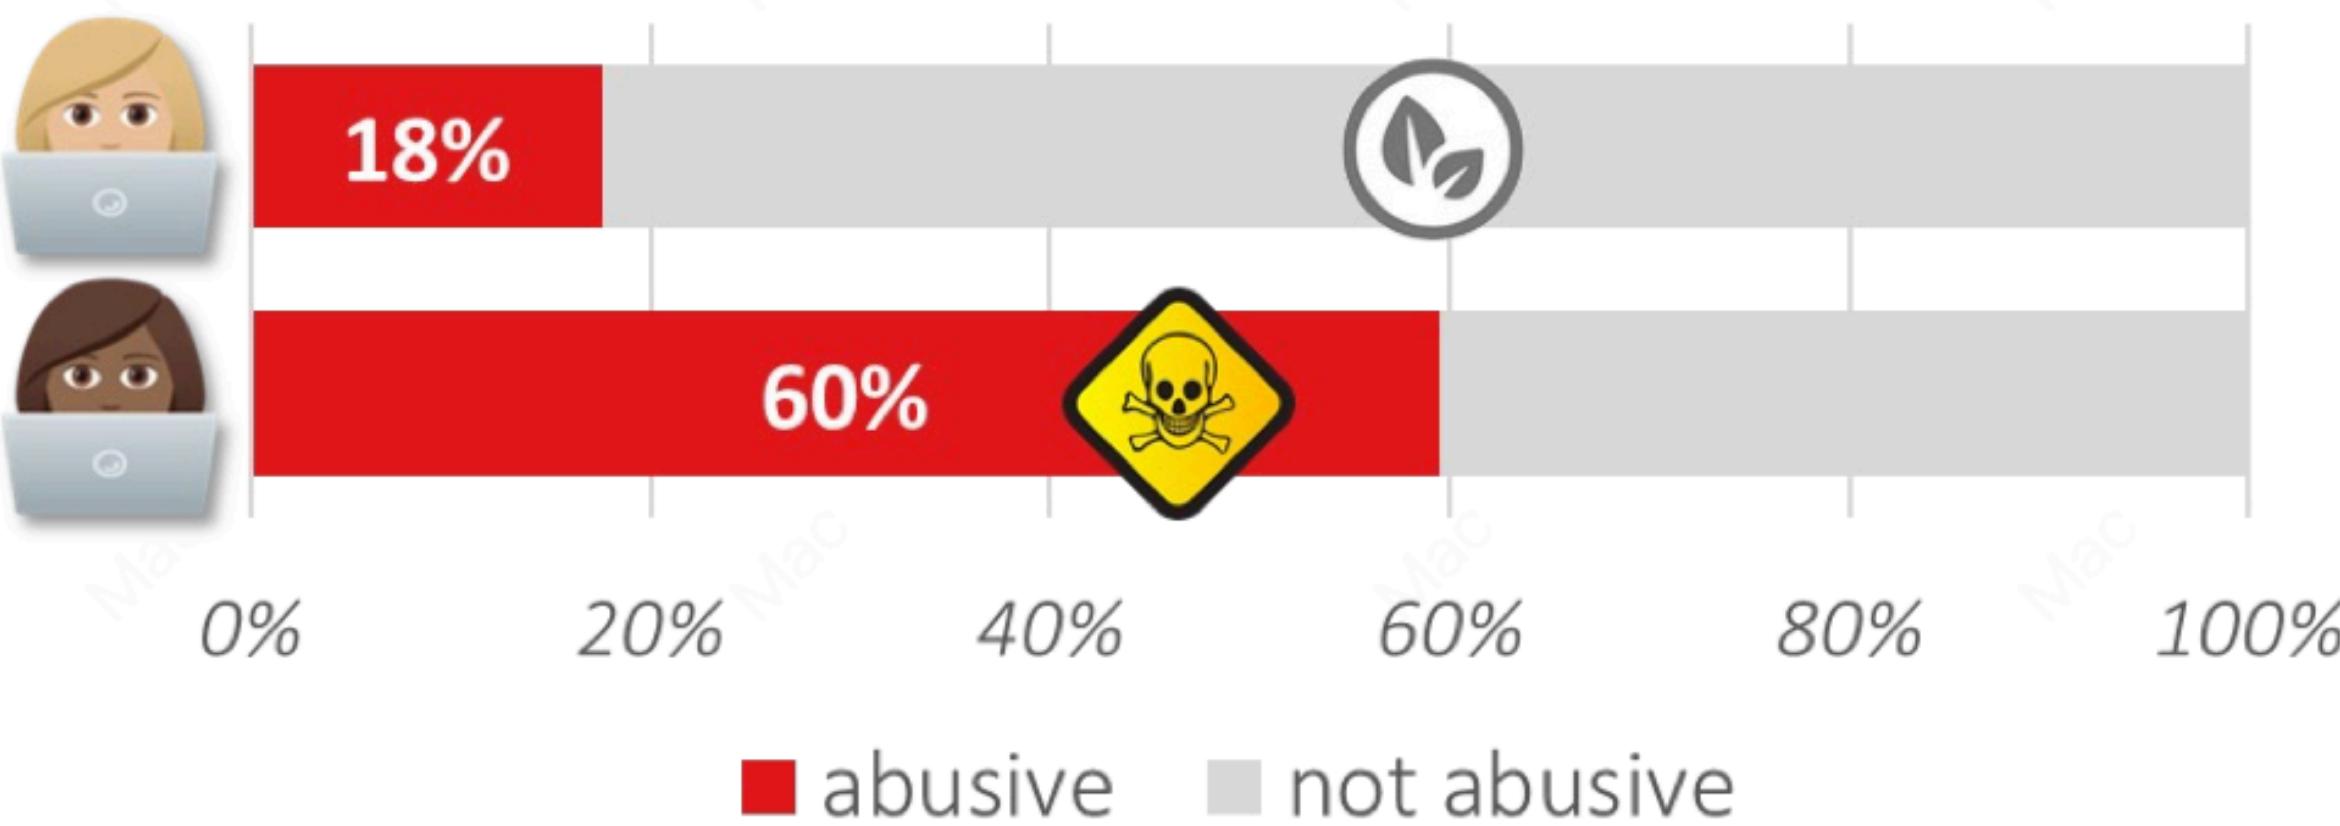
Goal: find and flag hateful or toxic content online, to make the internet less toxic

Racial biases in popular datasets



TWT-HATEBASE
(Davidson et al., 2017)

Both datasets have biases w.r.t. AAE tweets



TWT-BOOTSTRAP
(Founta et al., 2018)

Why did these biases occur? Why didn't NLP system designers think about these issues beforehand

The world itself is biased

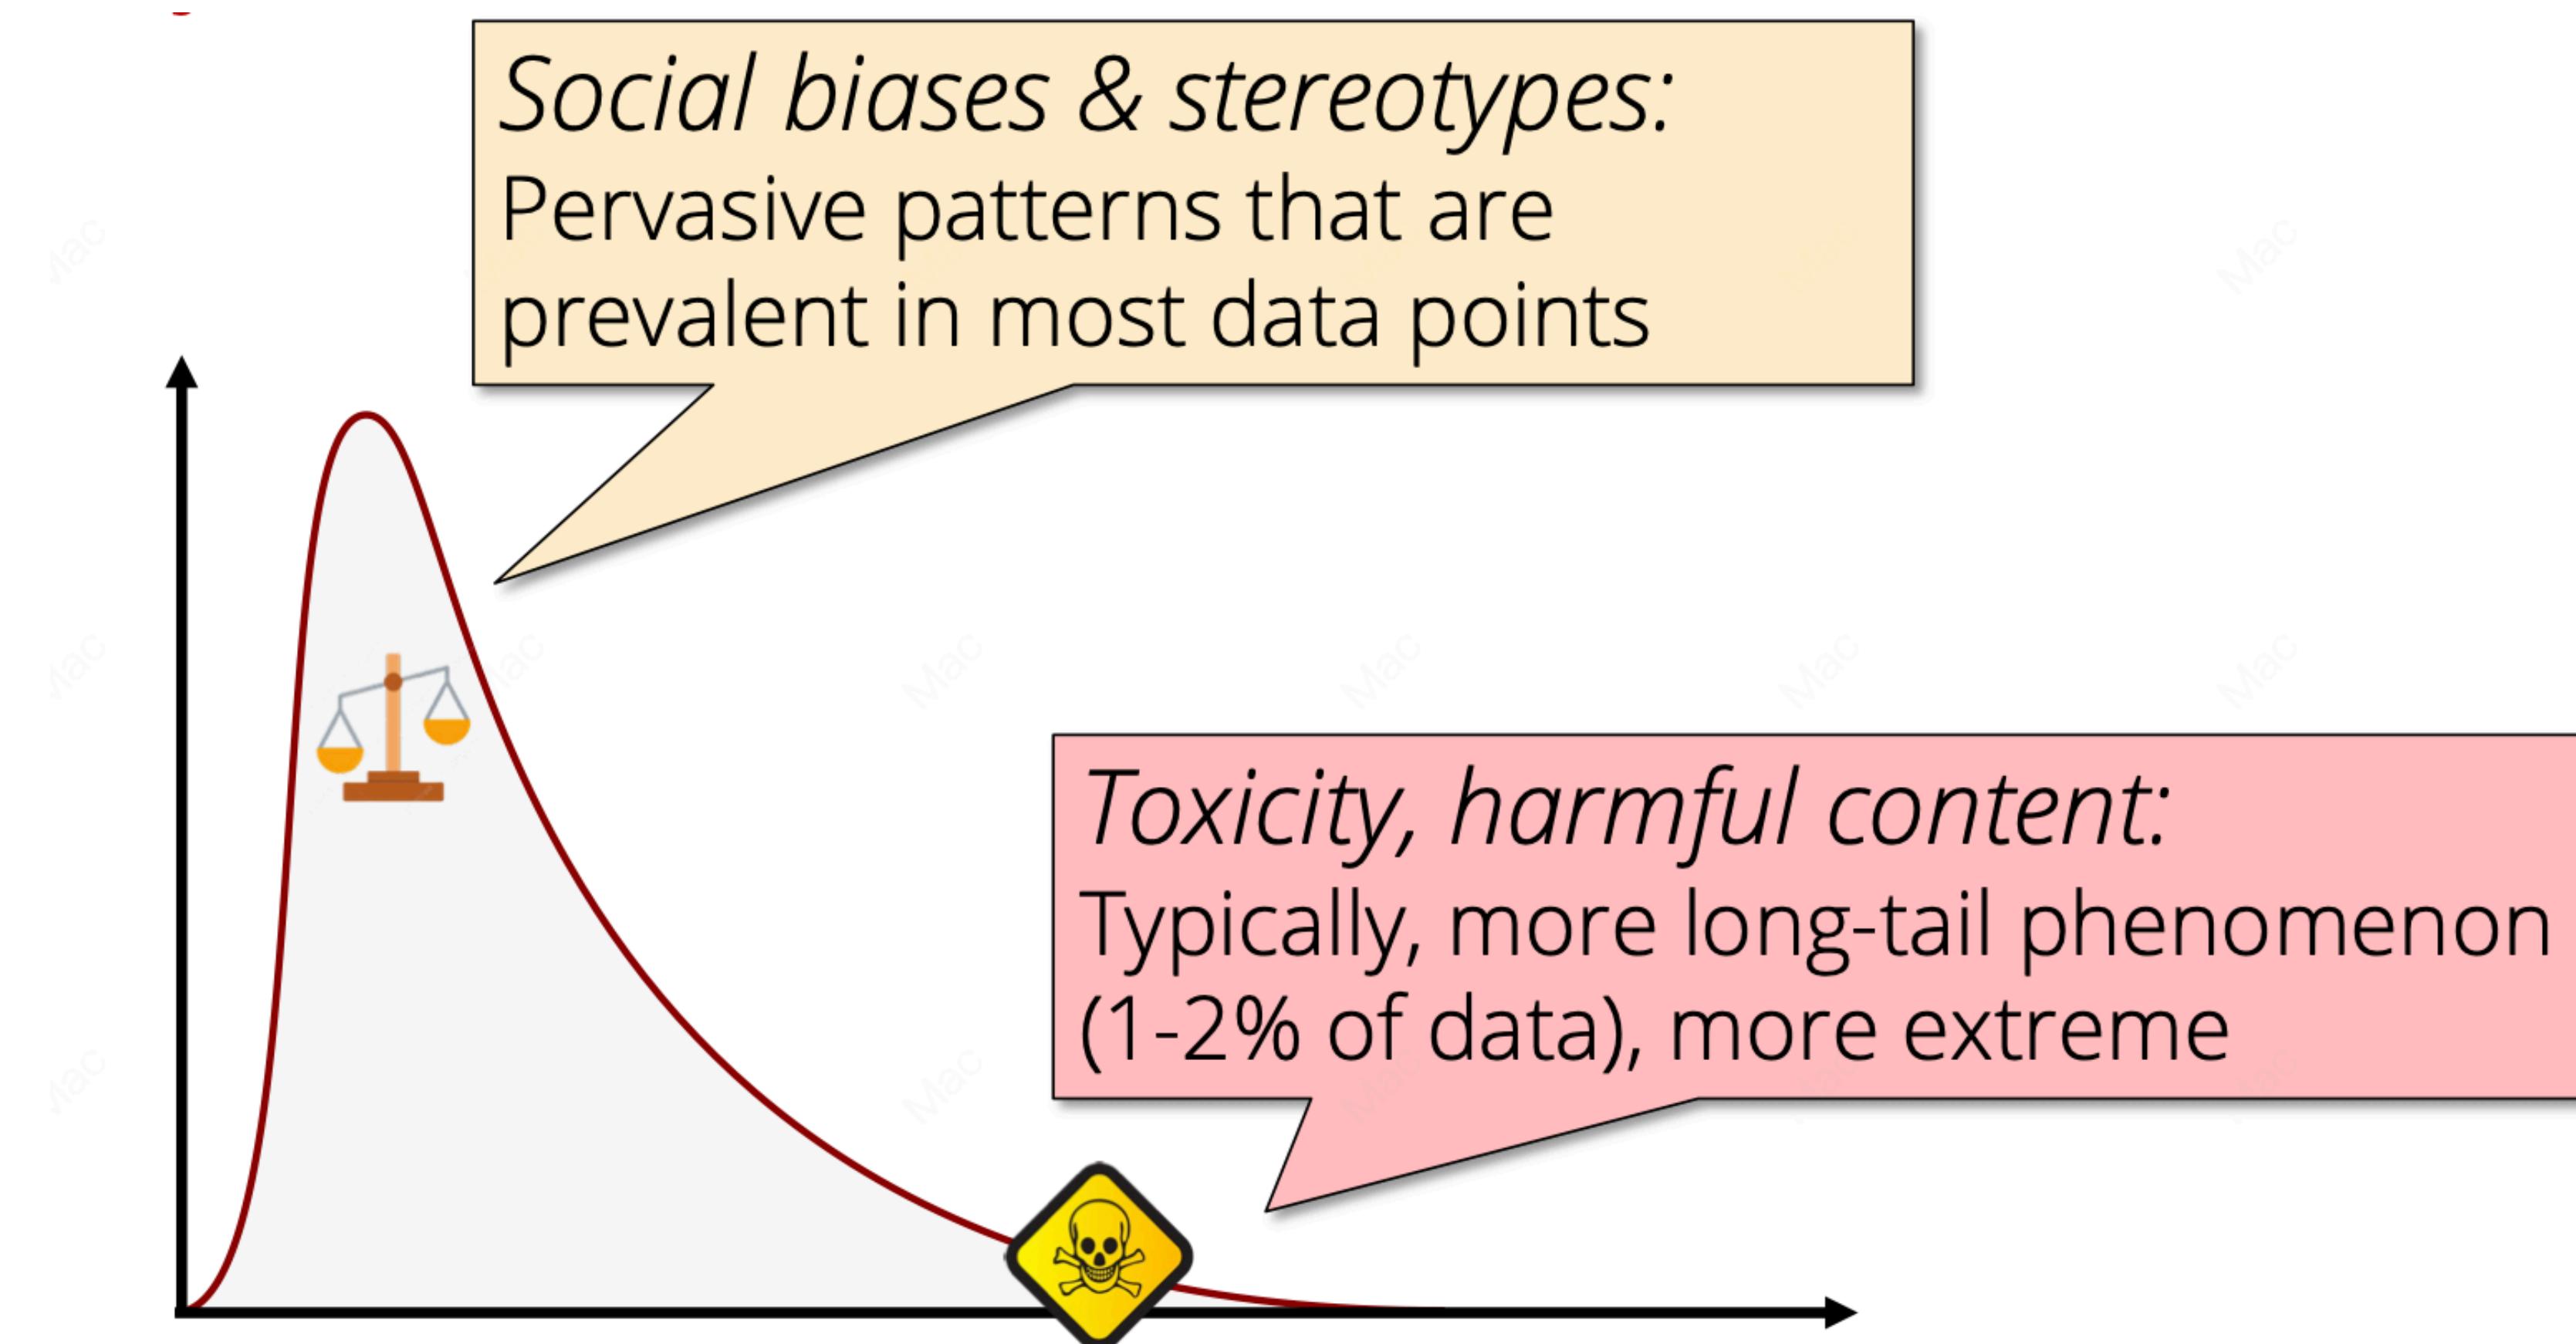
System designers have our own biases because of their *positionality*, i.e., set of perspectives that we hold due to our lived experiences and identity.

Positionality affects all our choices (e.g., assuming 1-1 mapping between languages and gendered pronouns, assuming toxicity looks the same in different dialects)



Harmful Content & Toxicity

Biases vs. Toxicity



Toxicity in LLMs

- Gehman et al (2020) introduced concept of *neural toxic degeneration* in LLMs
- Out of a 100 generations sampled from models, at least one toxic sentence
 - 65-70% toxicity from GPT2, GPT3
 - 85% toxicity from GPT1
- Model size affects toxicity: larger models have more toxicity [Touvron et al 2023]

Gehman et a. RealToxicityPrompts: Evaluating Neural Toxic Degeneration in Language Models. 2020

Touvron et a. LLaMA: Open and Efficient Foundation Language Models. 2023

**Why are these models learning so
much undesirable content?**

Problems with Self-Supervised Pretraining

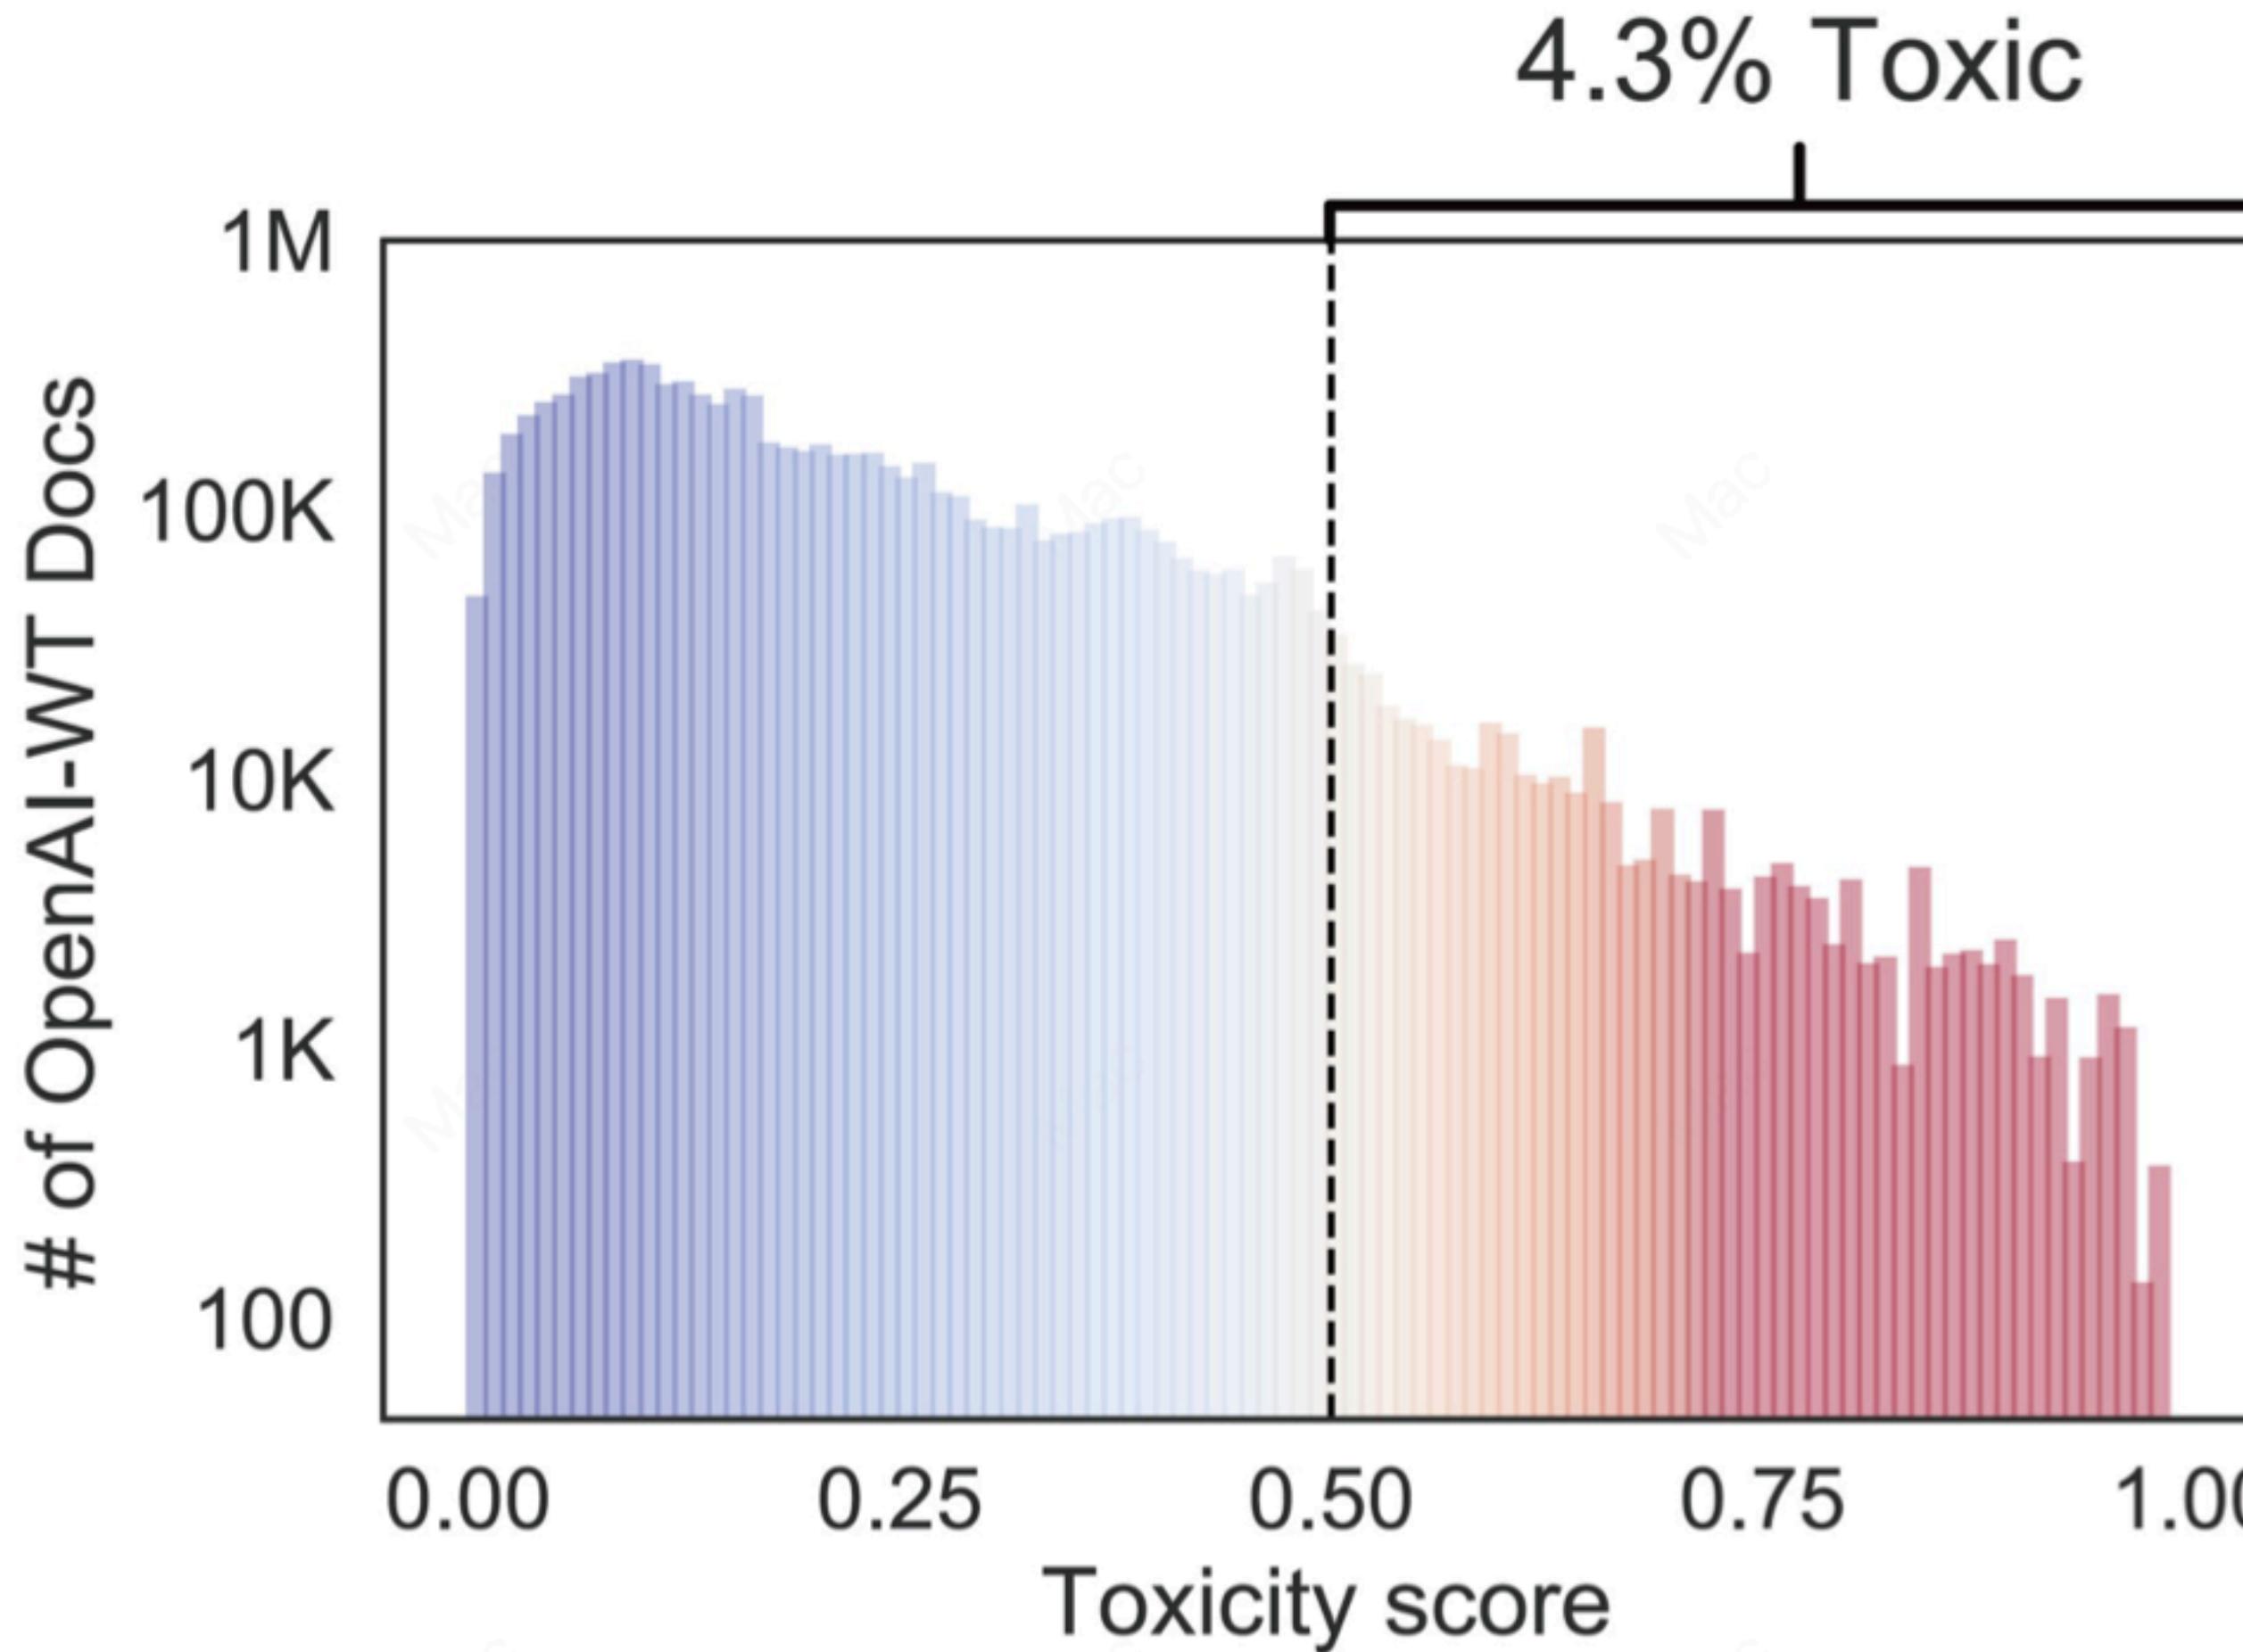
"Feeding AI systems on the world's beauty, ugliness, and cruelty, but expecting it to reflect only the beauty is a fantasy"



Prof. Ruha Benjamin, PhD

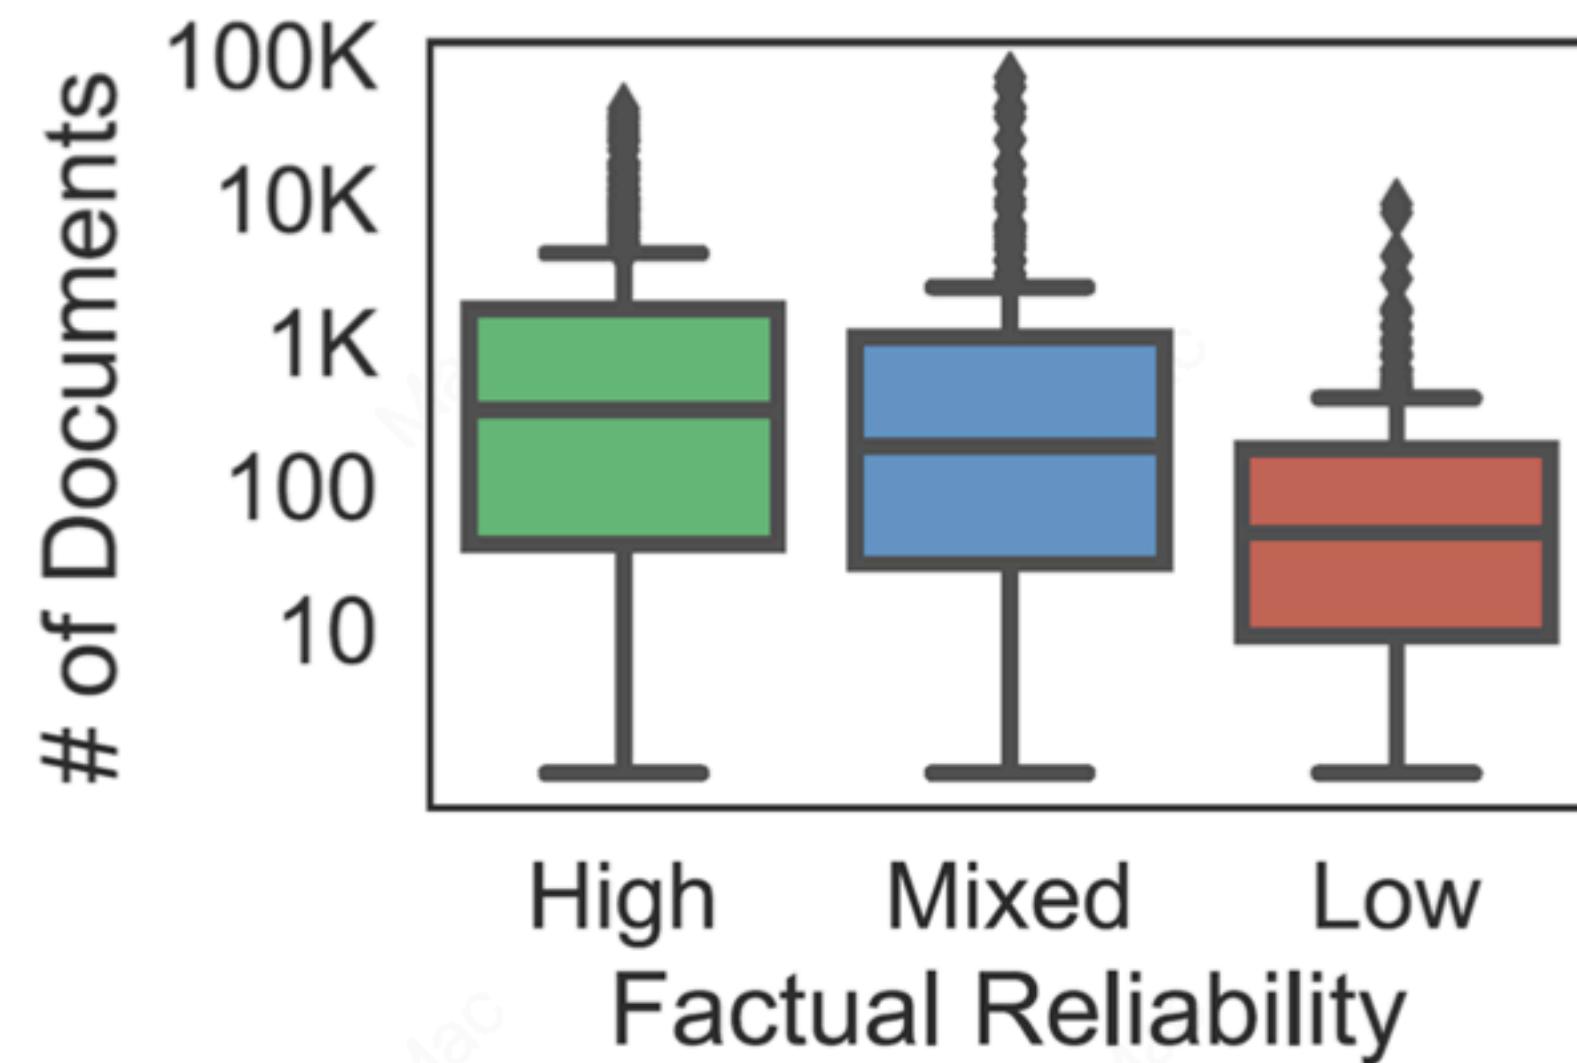
- Recipe: scrape as much pretraining data as you can to train your LM
- Consequence: LM ends up learning toxicity, biases, extremism, hate speech...

Toxicity in GPT-2's pretraining data



Fake news in GPT-2's Pretraining data

- Also looked at sources of documents in training data
- Cross-referencing sources of documents with known factual reliability categorization
 - >272K (3.4%) docs from low/mixed reliability sources
- Examining source where document is shared
 - >200K (3%) docs linked from banned/quarantined subreddits, which typically are more toxic docs
- Important to examine training data
 - Can only do that if publicly released!



Thank You!